

# Building a Single-Cycle RISC-V Processor



# The von Neumann Model

---

- Almost all modern computers are based on the von Neumann model (John von Neumann, 1945)

# The von Neumann Model


---

- Almost all modern computers are based on the von Neumann model (John von Neumann, 1945)
- Components:

# The von Neumann Model

---

- Almost all modern computers are based on the von Neumann model (John von Neumann, 1945)
- Components:




- **Main memory** holds programs and their data

# The von Neumann Model

---

- Almost all modern computers are based on the von Neumann model (John von Neumann, 1945)
- Components:




- **Main memory** holds programs and their data
- **Central processing unit** accesses and processes memory values

# The von Neumann Model

---

- Almost all modern computers are based on the von Neumann model (John von Neumann, 1945)
- Components:

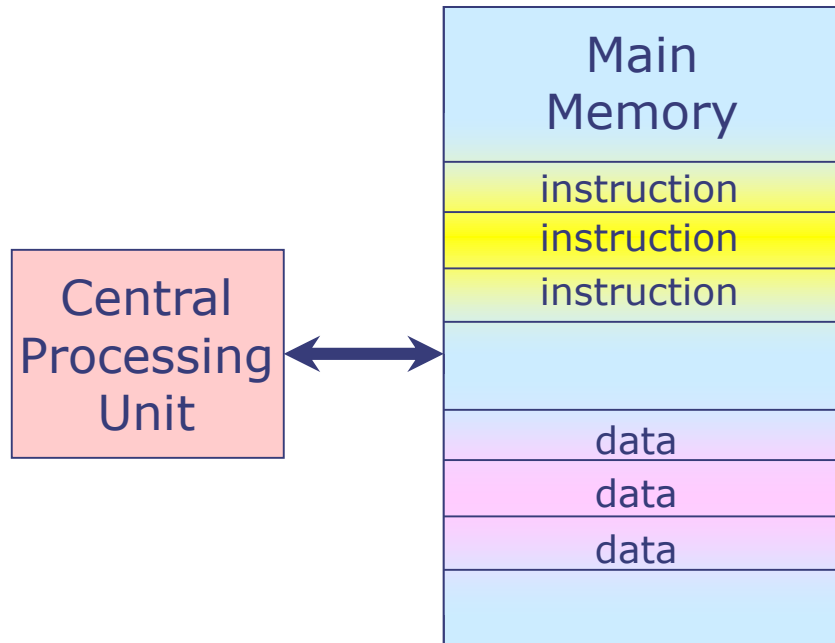


- **Main memory** holds programs and their data
- **Central processing unit** accesses and processes memory values
- **Input/output devices** to communicate with the outside world

# Key Idea: Stored-Program Computer

---

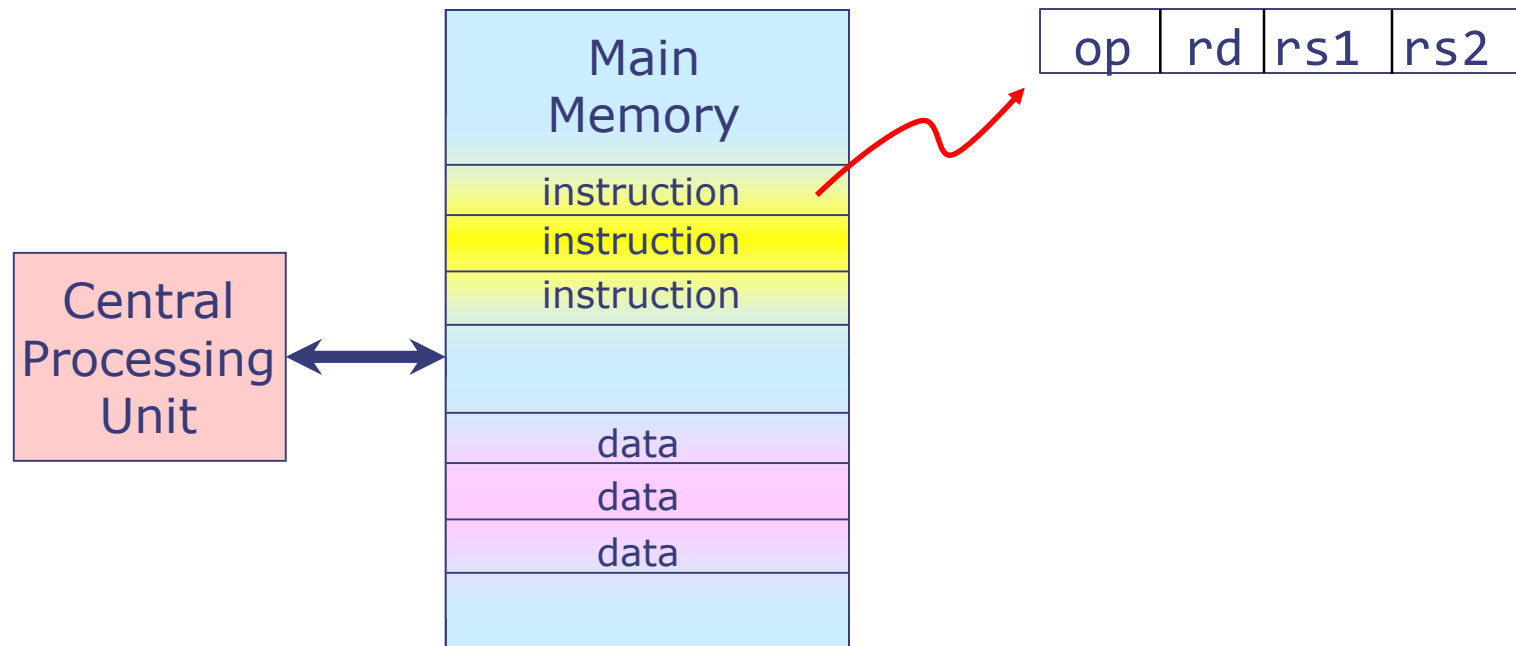
- Express program as a sequence of **coded instructions**
- Memory holds both data and instructions
- CPU fetches, interprets, and executes successive instructions of the program



# Key Idea: Stored-Program Computer

---

- Express program as a sequence of **coded instructions**
- Memory holds both data and instructions
- CPU fetches, interprets, and executes successive instructions of the program

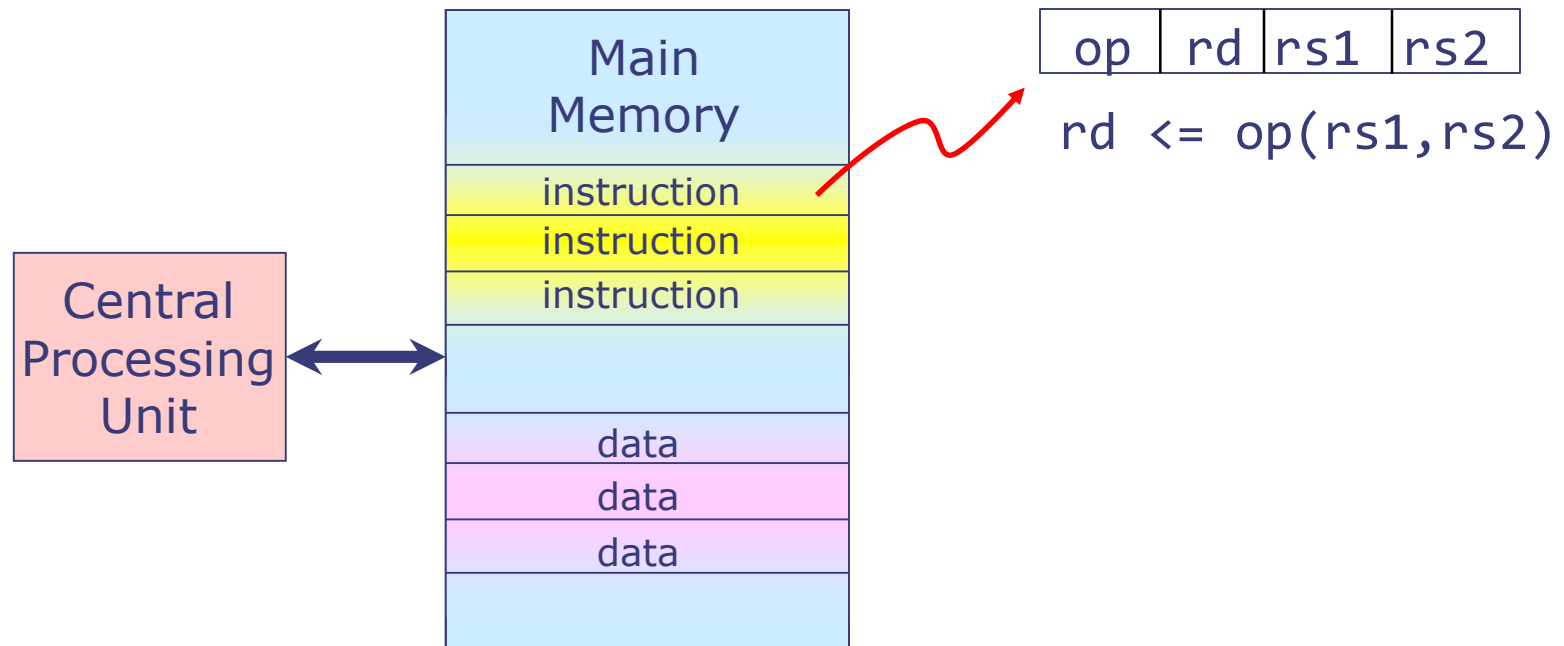




# Key Idea: Stored-Program Computer

---

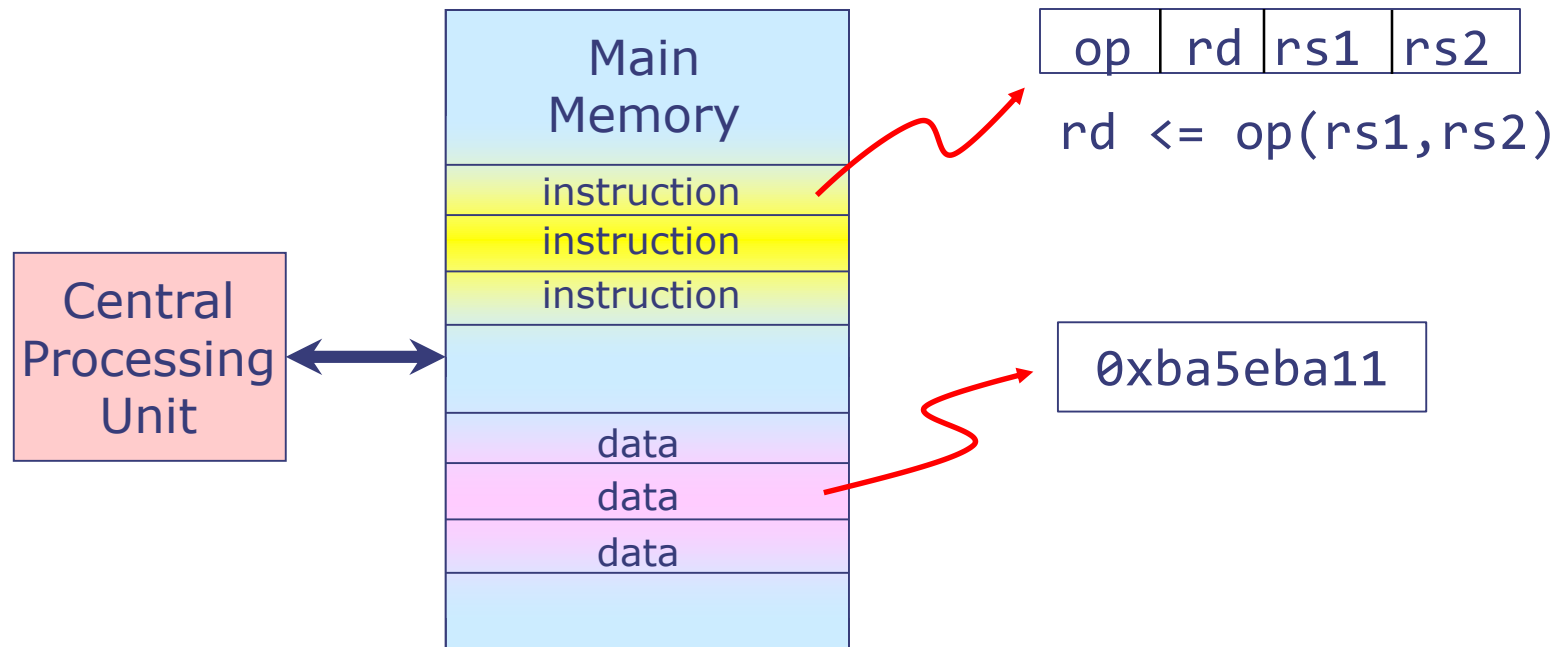
- Express program as a sequence of **coded instructions**
- Memory holds both data and instructions
- CPU fetches, interprets, and executes successive instructions of the program



# Key Idea: Stored-Program Computer

---

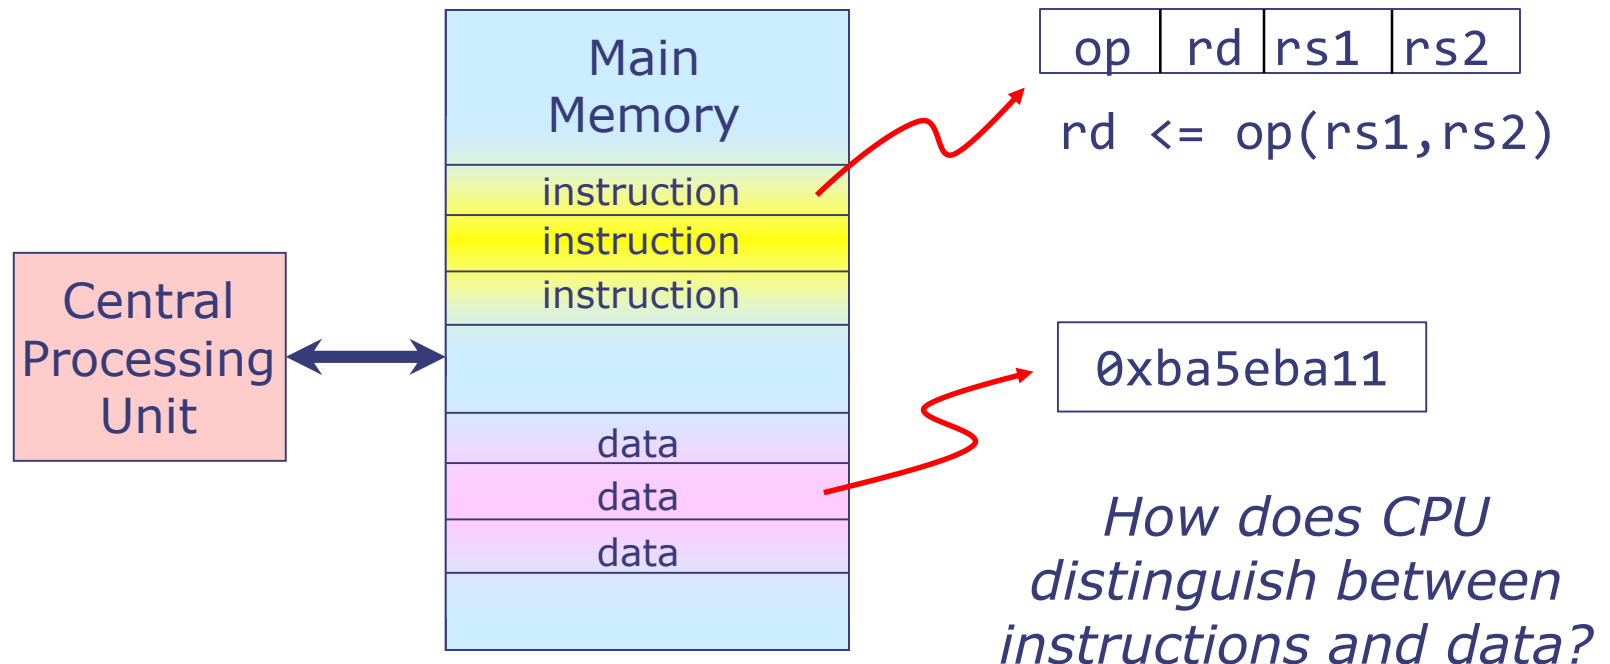
- Express program as a sequence of **coded instructions**
- Memory holds both data and instructions
- CPU fetches, interprets, and executes successive instructions of the program



# Key Idea: Stored-Program Computer


---

- Express program as a sequence of **coded instructions**
- Memory holds both data and instructions
- CPU fetches, interprets, and executes successive instructions of the program




# Anatomy of a von Neumann Computer

---

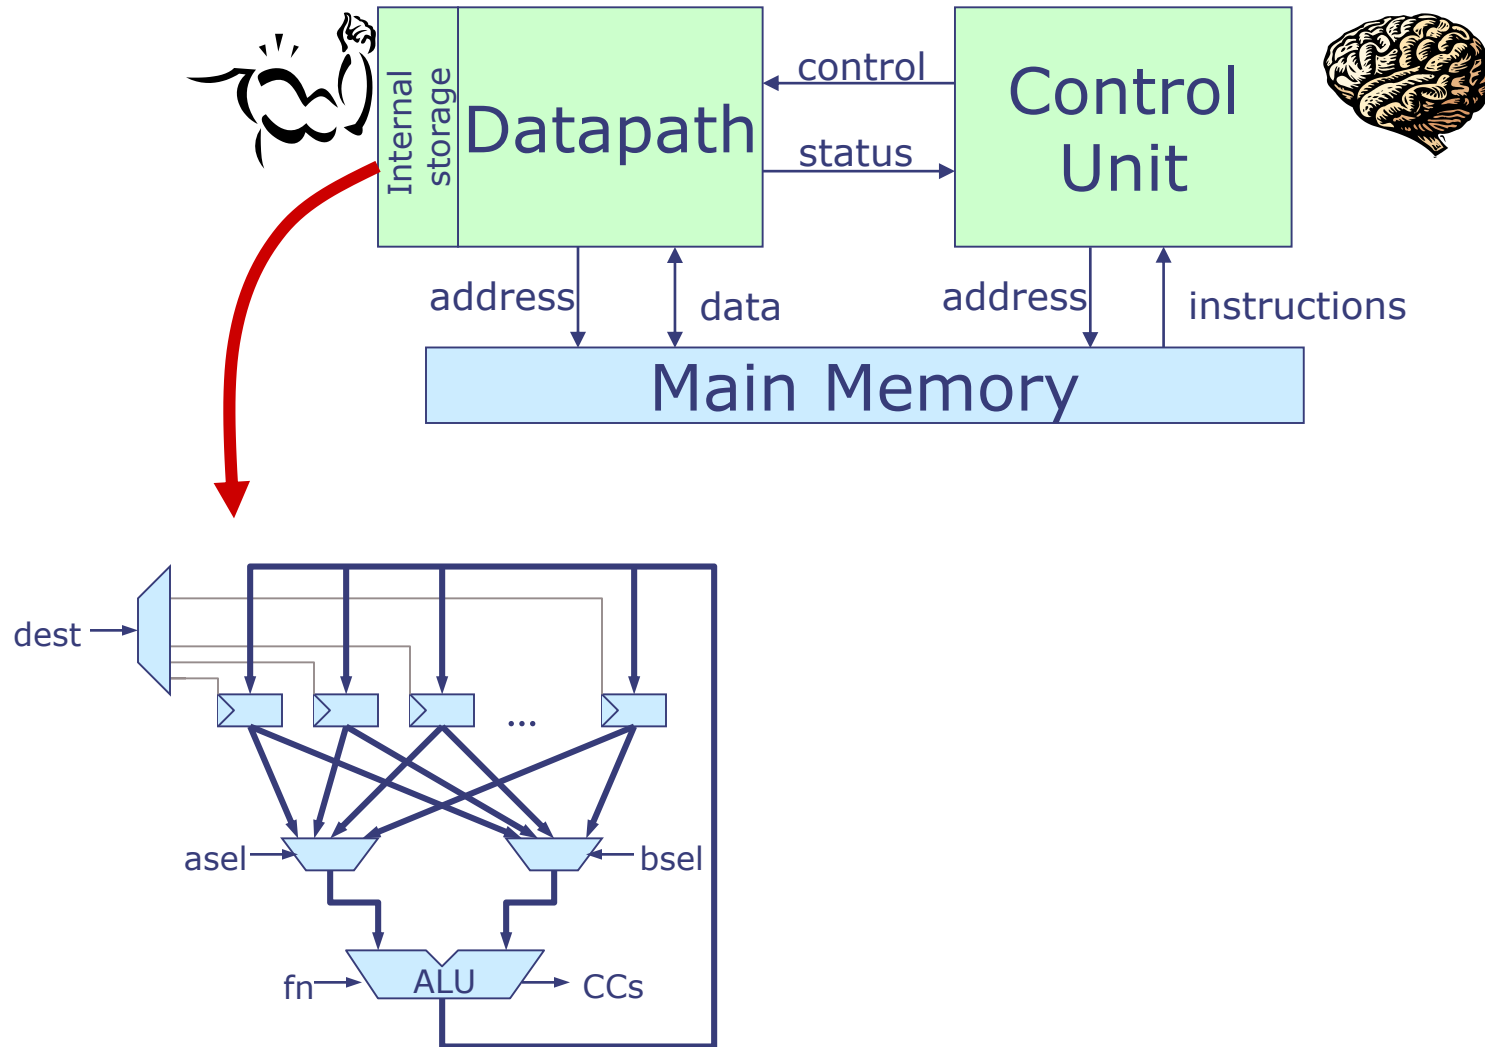


# Anatomy of a von Neumann Computer

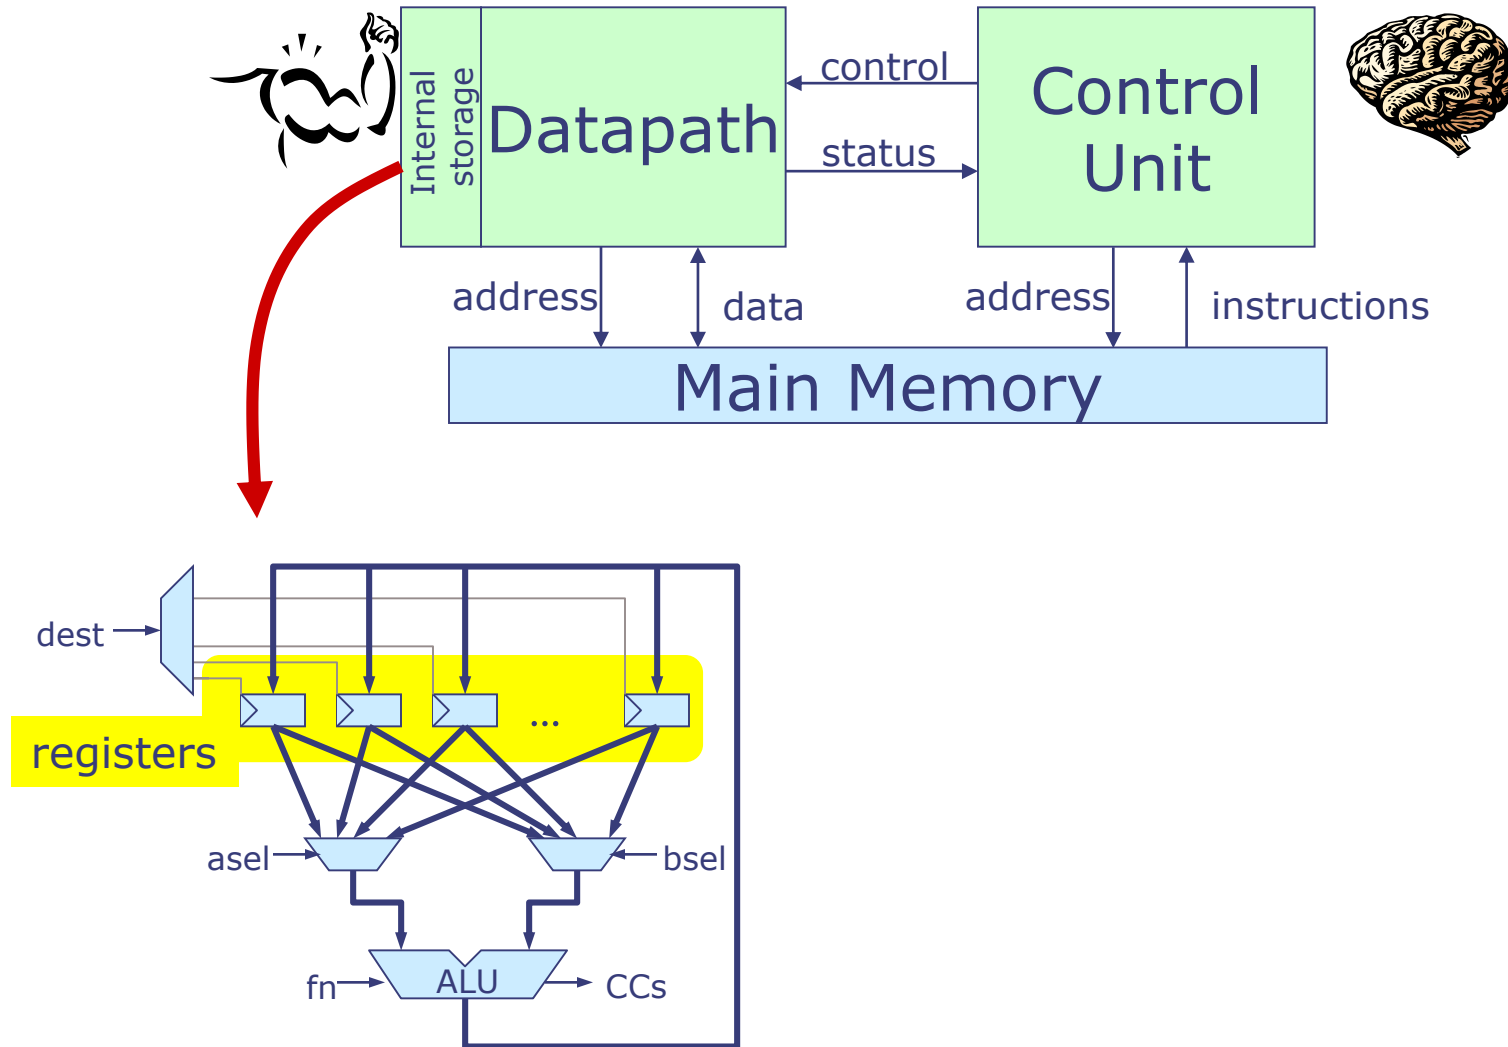
---



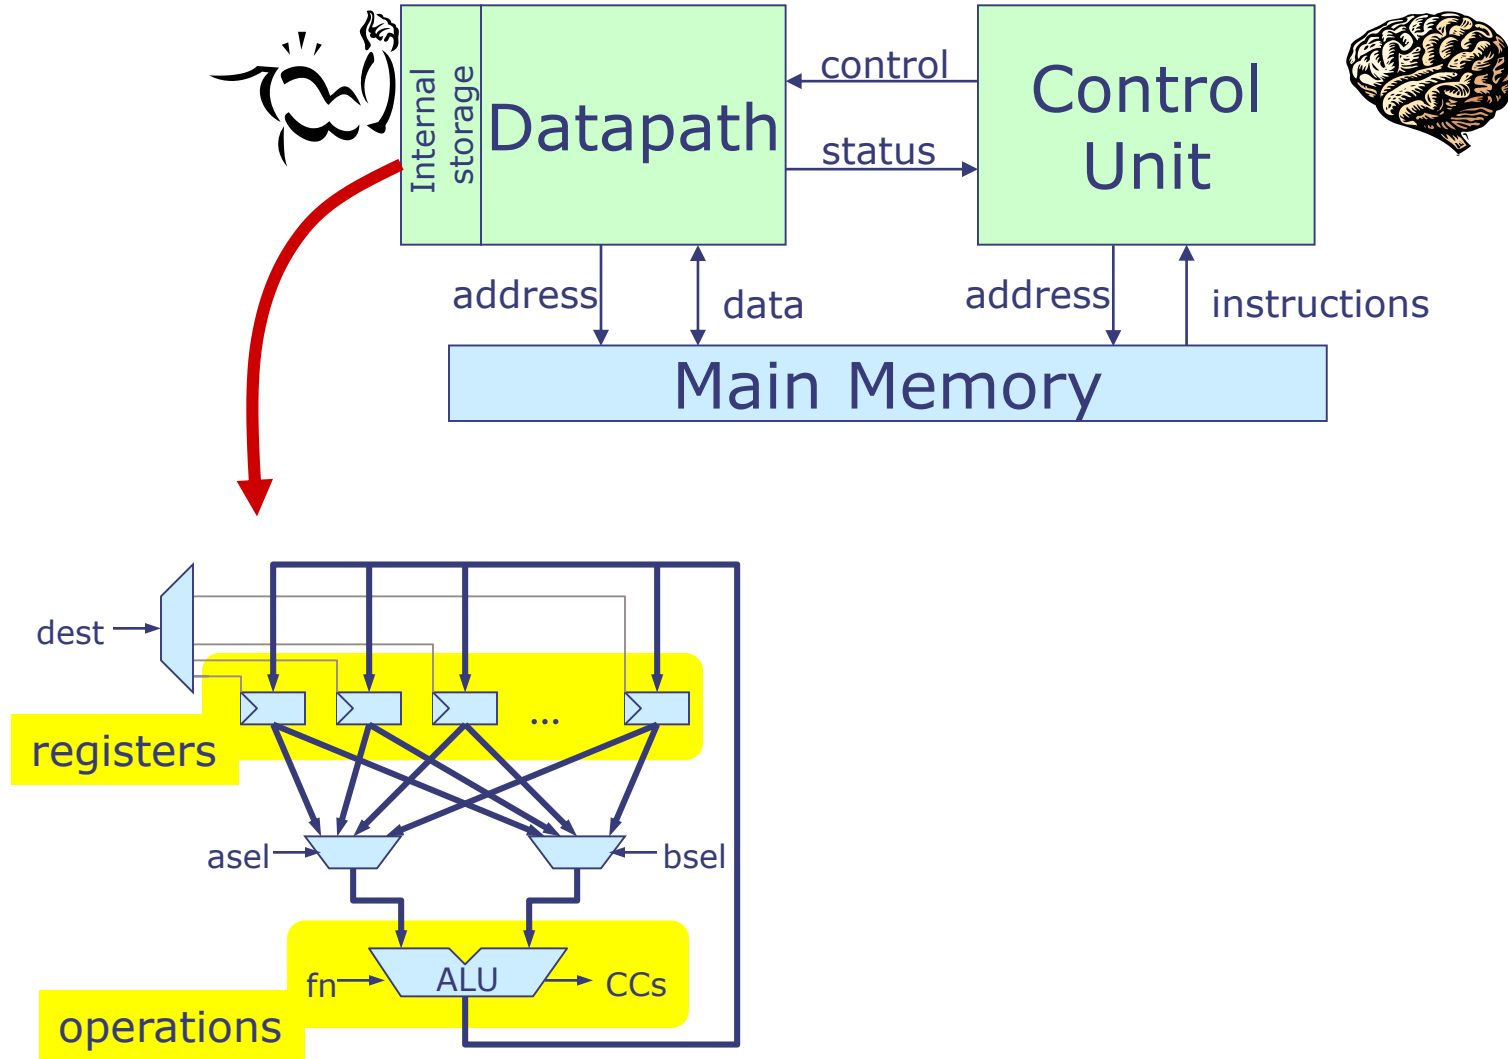
# Anatomy of a von Neumann Computer



# Anatomy of a von Neumann Computer

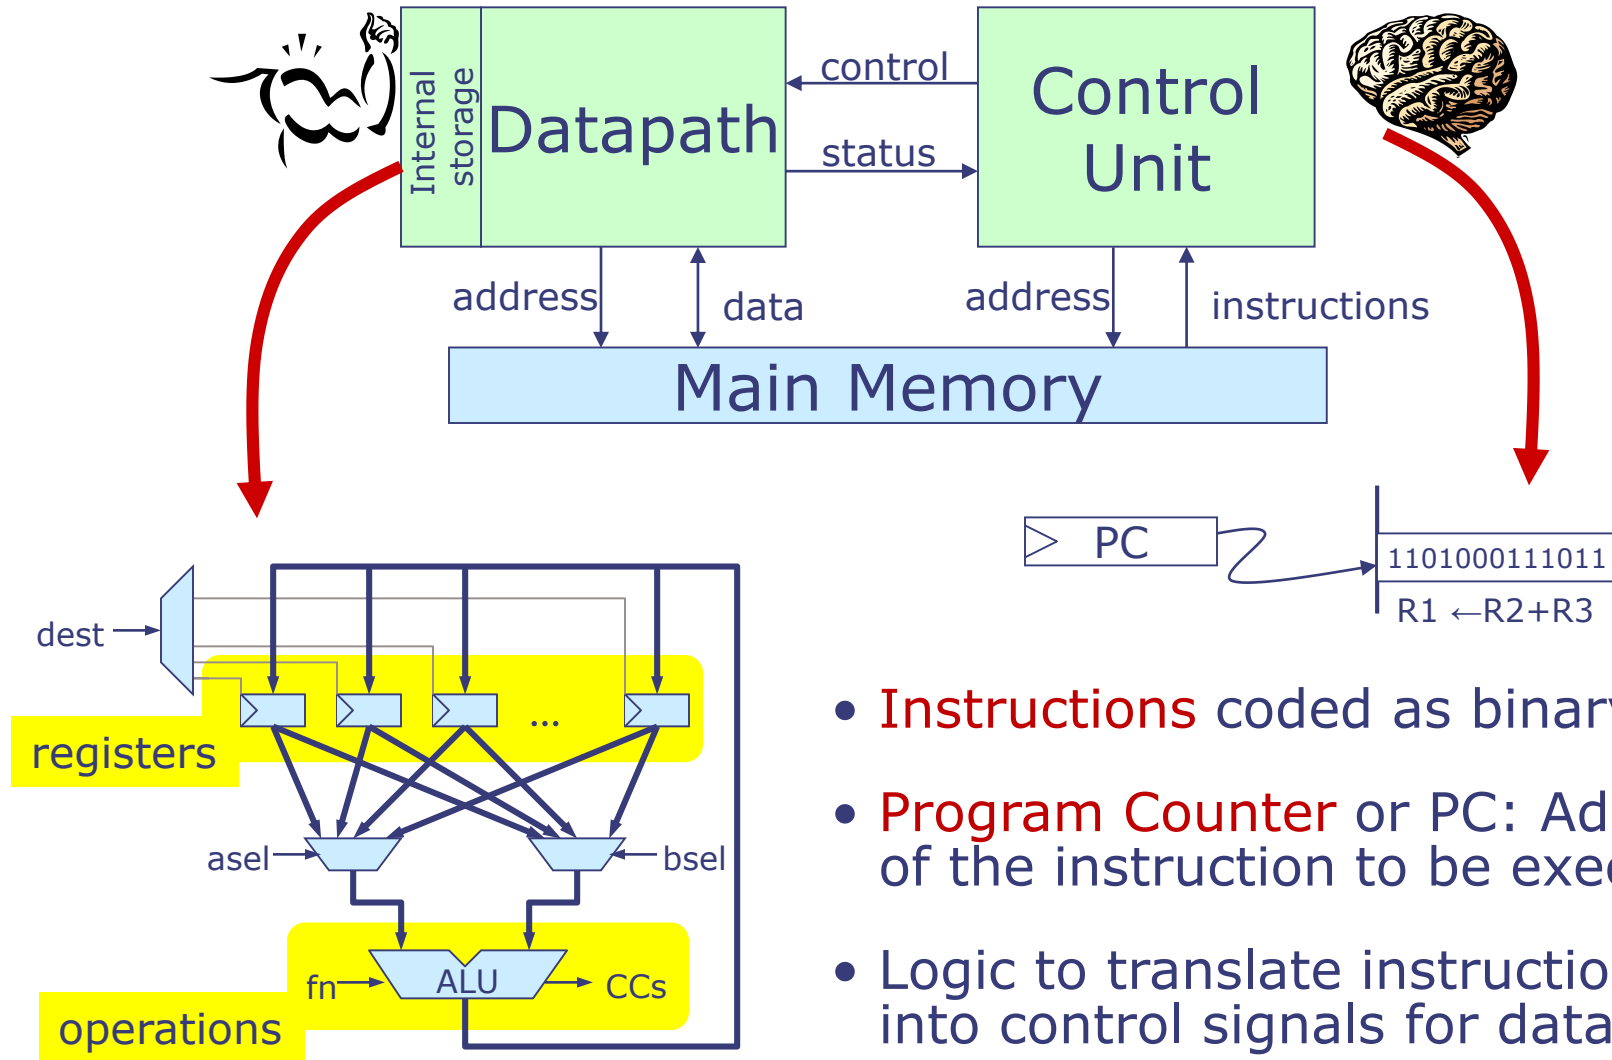


# Anatomy of a von Neumann Computer





# Anatomy of a von Neumann Computer



- **Instructions** coded as binary data
- **Program Counter** or PC: Address of the instruction to be executed
- Logic to translate instructions into control signals for datapath

# Instructions

---

- Instructions are the fundamental unit of work

# Instructions

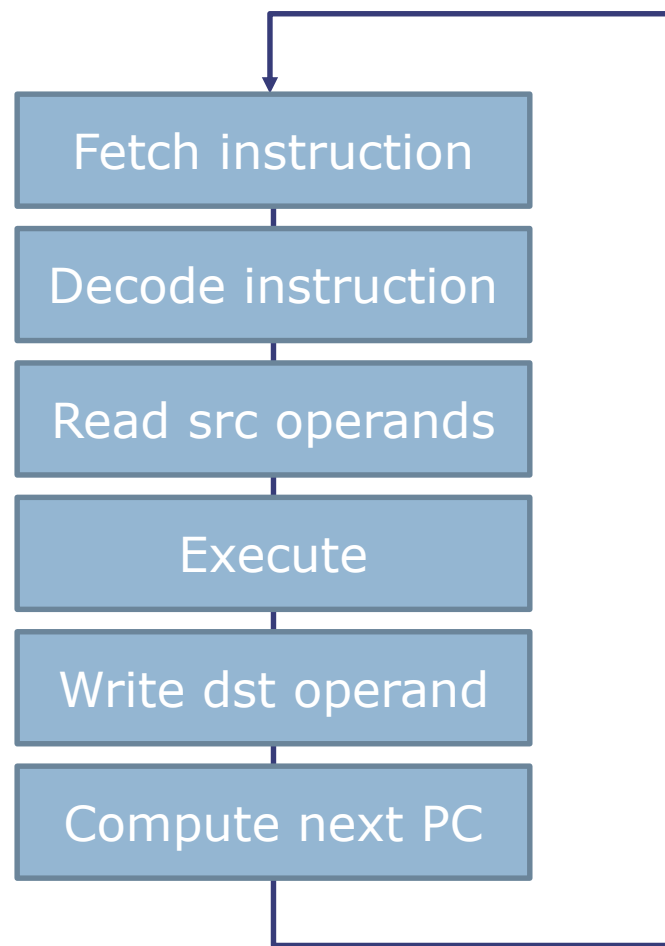
---

- Instructions are the fundamental unit of work
- Each instruction specifies:
  - An operation or **opcode** to be performed
  - **Source operands** and **destination** for the result

# Instructions

---

- Instructions are the fundamental unit of work
- Each instruction specifies:
  - An operation or **opcode** to be performed
  - **Source operands** and **destination** for the result
- In a von Neumann machine, instructions are executed sequentially
  - CPU logically implements this loop:
  - By default, the next PC is current PC + size of current instruction unless the instruction says otherwise



# Processor Performance

---

- “Iron Law” of performance:

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \cdot \frac{\text{Cycles}}{\text{Instruction}} \cdot \frac{\text{Time}}{\text{Cycle}} \quad \text{Perf} = \frac{1}{\text{Time}}$$

# Processor Performance

---

- “Iron Law” of performance:

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \cdot \frac{\text{Cycles}}{\text{Instruction}} \cdot \frac{\text{Time}}{\text{Cycle}} \quad \text{Perf} = \frac{1}{\text{Time}}$$

- Options to reduce execution time:
  - Executed instructions ↓ (work/instruction ↑)

# Processor Performance

---

- “Iron Law” of performance:

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \cdot \frac{\text{Cycles}}{\text{Instruction}} \cdot \frac{\text{Time}}{\text{Cycle}} \quad \text{Perf} = \frac{1}{\text{Time}}$$

- Options to reduce execution time:
  - Executed instructions ↓ (work/instruction ↑)
  - Cycles per instruction (CPI) ↓

# Processor Performance

---

- “Iron Law” of performance:

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \cdot \frac{\text{Cycles}}{\text{Instruction}} \cdot \frac{\text{Time}}{\text{Cycle}} \quad \text{Perf} = \frac{1}{\text{Time}}$$

- Options to reduce execution time:
  - Executed instructions ↓ (work/instruction ↑)
  - Cycles per instruction (CPI) ↓
  - Cycle time ↓ (frequency ↑)



# Processor Performance

---

- “Iron Law” of performance:

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \cdot \frac{\text{Cycles}}{\text{Instruction}} \cdot \frac{\text{Time}}{\text{Cycle}} \quad \text{Perf} = \frac{1}{\text{Time}}$$

- Options to reduce execution time:
  - Executed instructions ↓ (work/instruction ↑)
  - Cycles per instruction (CPI) ↓
  - Cycle time ↓ (frequency ↑)
- Today: Simple **single-cycle processor**, executes one instruction from start to end each clock cycle

# Processor Performance

---

- “Iron Law” of performance:

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \cdot \frac{\text{Cycles}}{\text{Instruction}} \cdot \frac{\text{Time}}{\text{Cycle}} \quad \text{Perf} = \frac{1}{\text{Time}}$$

- Options to reduce execution time:
  - Executed instructions ↓ (work/instruction ↑)
  - Cycles per instruction (CPI) ↓
  - Cycle time ↓ (frequency ↑)
- Today: Simple **single-cycle processor**, executes one instruction from start to end each clock cycle
  - CPI = 1, but low frequency
  - Later: Pipelining to increase frequency

# Approach: Incremental Featurism

---

We'll implement datapaths for each instruction class individually, and merge them (using MUXes)

# Approach: Incremental Featurism

---

We'll implement datapaths for each instruction class individually, and merge them (using MUXes)

## Steps:

1. ALU instructions
2. Load & store instructions
3. Branch & jump instructions

# Approach: Incremental Featurism

---

We'll implement datapaths for each instruction class individually, and merge them (using MUXes)

## Steps:

1. ALU instructions
2. Load & store instructions
3. Branch & jump instructions

Component  
Repertoire:



Registers

# Approach: Incremental Featurism

---

We'll implement datapaths for each instruction class individually, and merge them (using MUXes)

## Steps:

1. ALU instructions
2. Load & store instructions
3. Branch & jump instructions

Component  
Repertoire:



Registers



Muxes

# Approach: Incremental Featurism

---

We'll implement datapaths for each instruction class individually, and merge them (using MUXes)

## Steps:

1. ALU instructions
2. Load & store instructions
3. Branch & jump instructions

Component  
Repertoire:



Registers



Muxes



"Black box" ALU

# Approach: Incremental Featurism

---

We'll implement datapaths for each instruction class individually, and merge them (using MUXes)

## Steps:

1. ALU instructions
2. Load & store instructions
3. Branch & jump instructions

## Component Repertoire:



Registers




Muxes



"Black box" ALU

Register File





# Approach: Incremental Featurism

---

We'll implement datapaths for each instruction class individually, and merge them (using MUXes)

## Steps:

1. ALU instructions
2. Load & store instructions
3. Branch & jump instructions

## Component Repertoire:



Registers





Muxes



"Black box" ALU


Register File




Memories

# Multi-Ported Register File


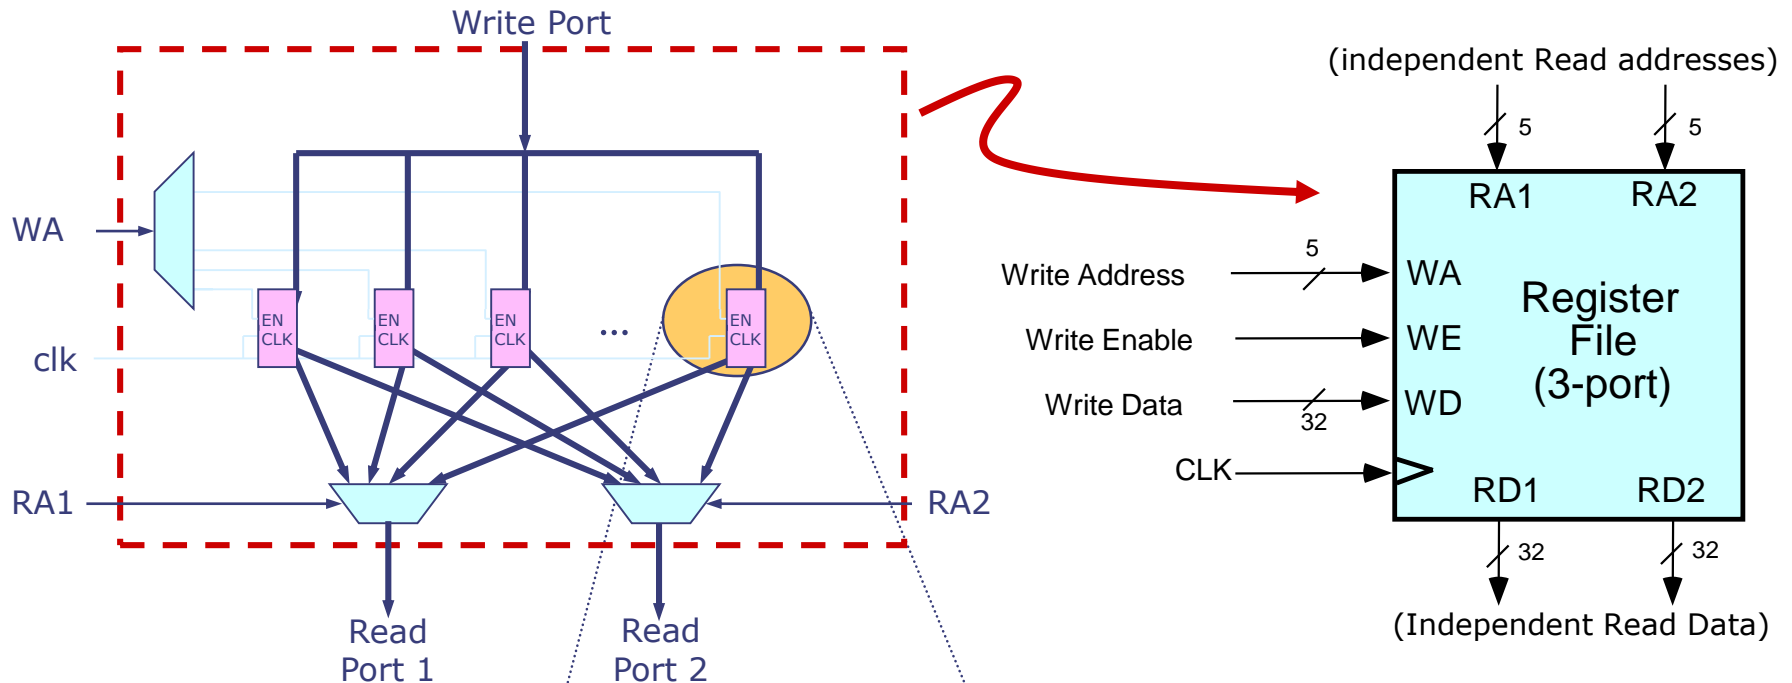
---



# Multi-Ported Register File



# Multi-Ported Register File

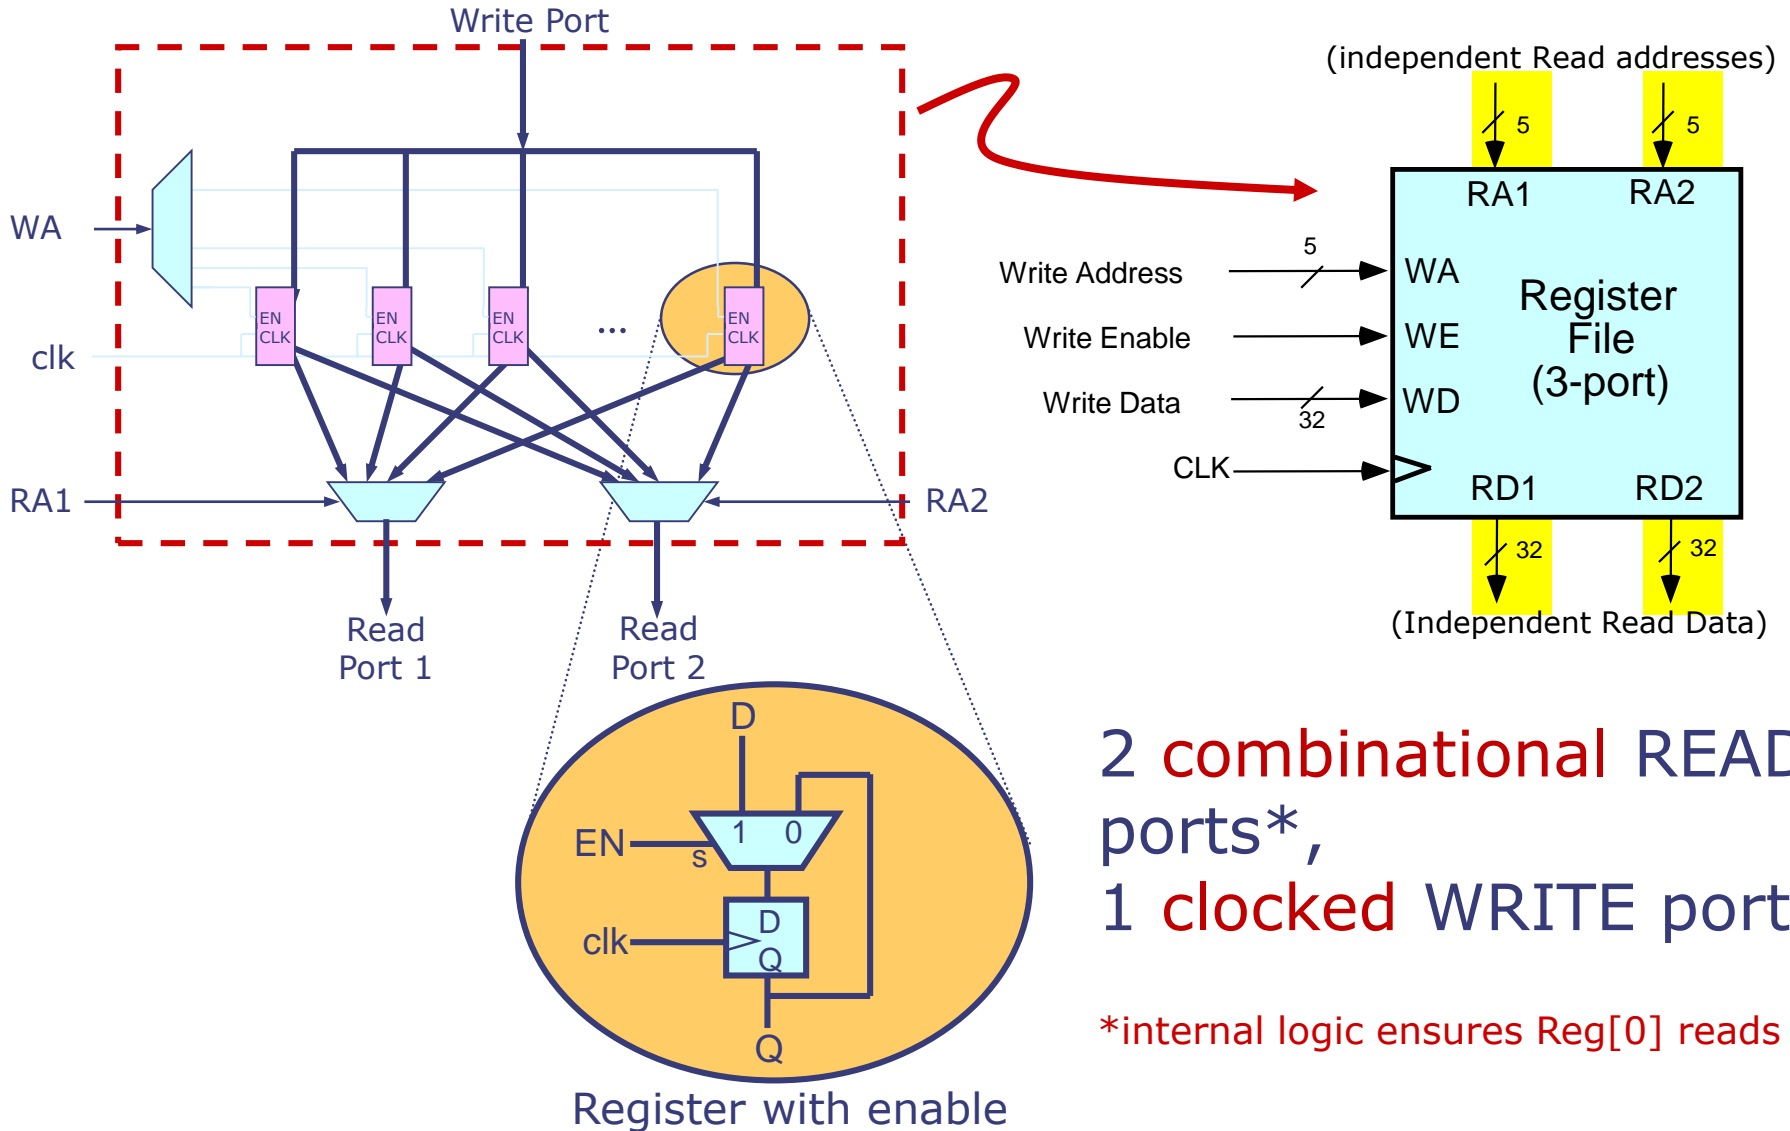


Register with enable

2 **combinational** READ ports\*,  
1 **clocked** WRITE port

\*internal logic ensures Reg[0] reads as 0

# Multi-Ported Register File


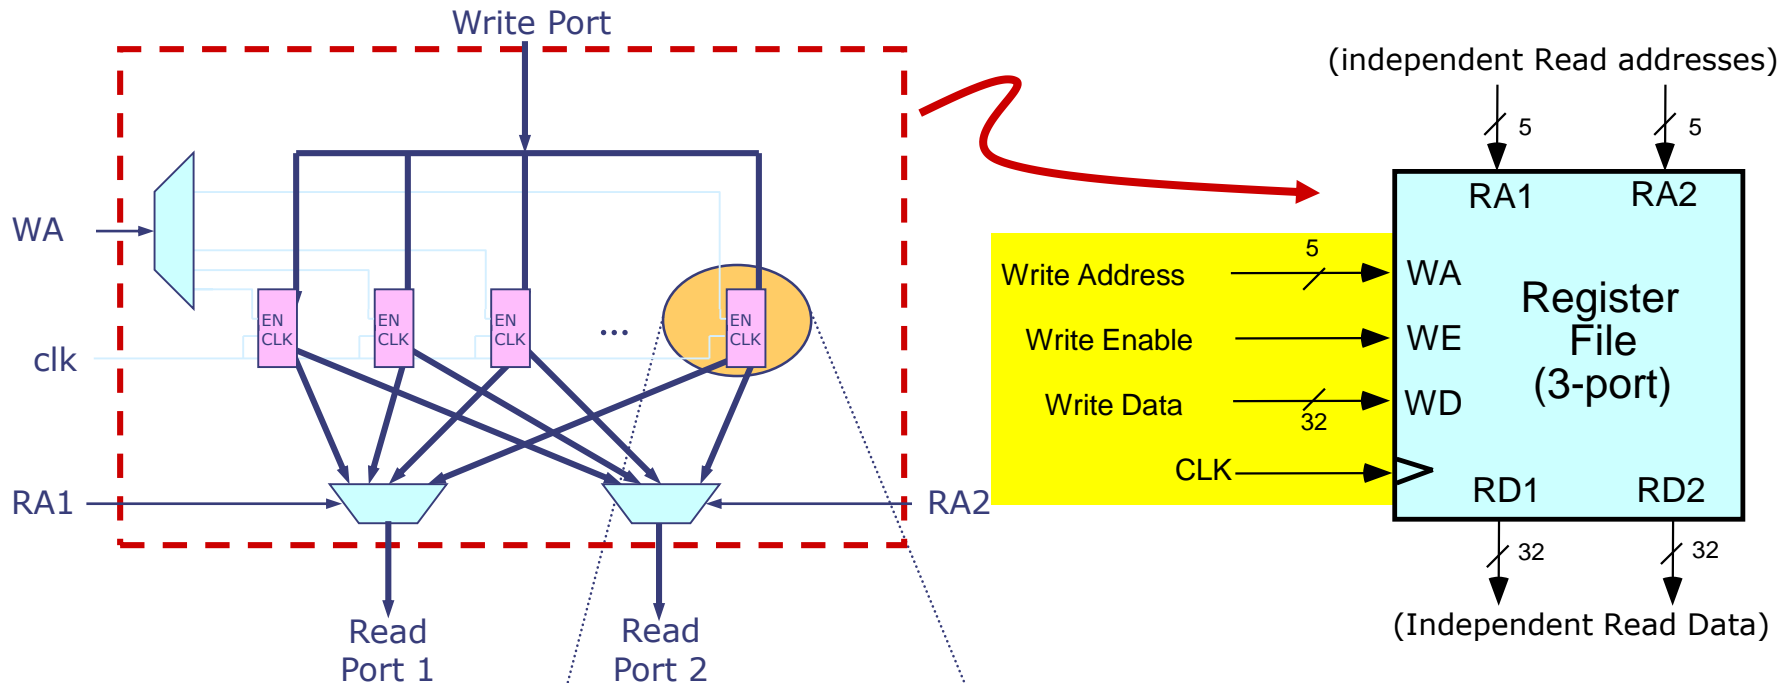


**2 combinational READ ports\*,**  
**1 clocked WRITE port**

\*internal logic ensures Reg[0] reads as 0

Register with enable

# Multi-Ported Register File




Register with enable

2 **combinational** READ ports\*,  
1 **clocked** WRITE port

\*internal logic ensures Reg[0] reads as 0

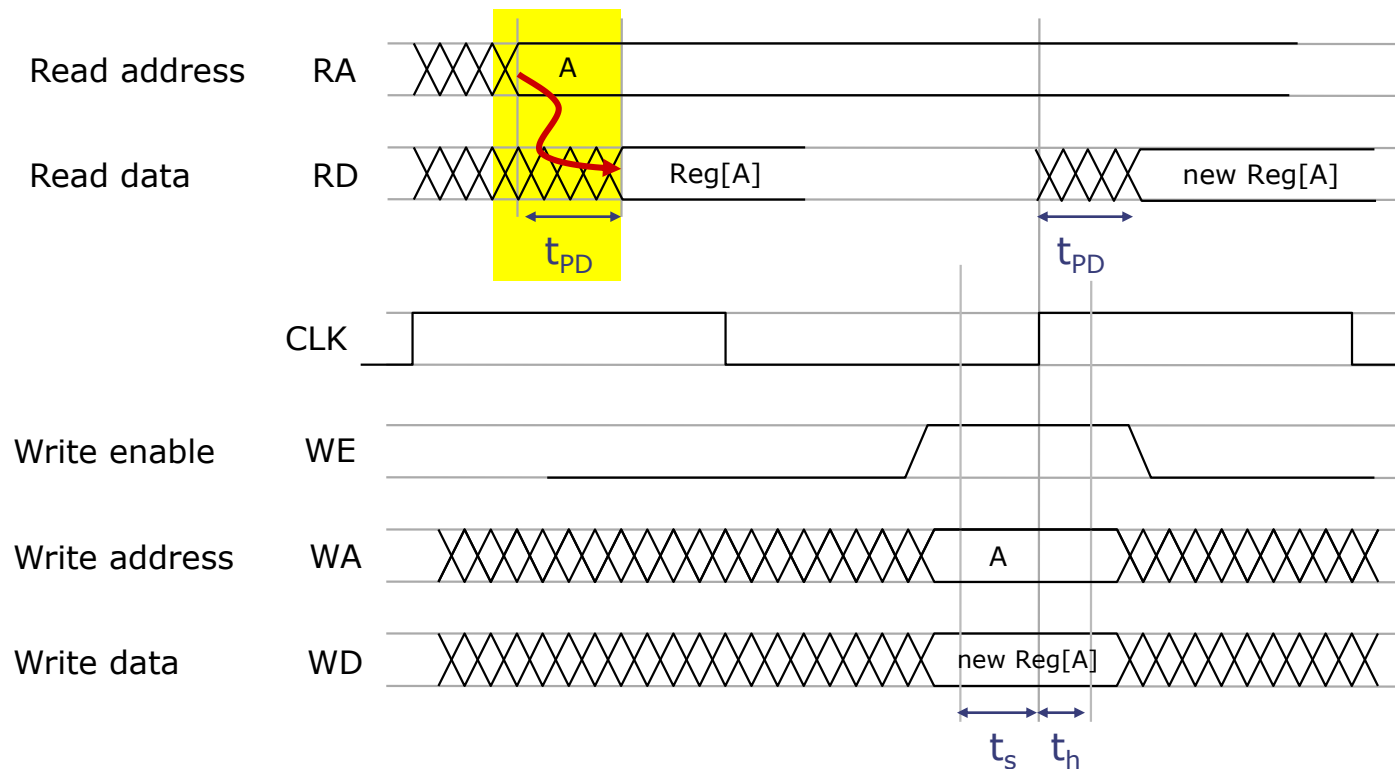
# Register File Timing

2 combinational READ ports, 1 clocked WRITE port



# Register File Timing

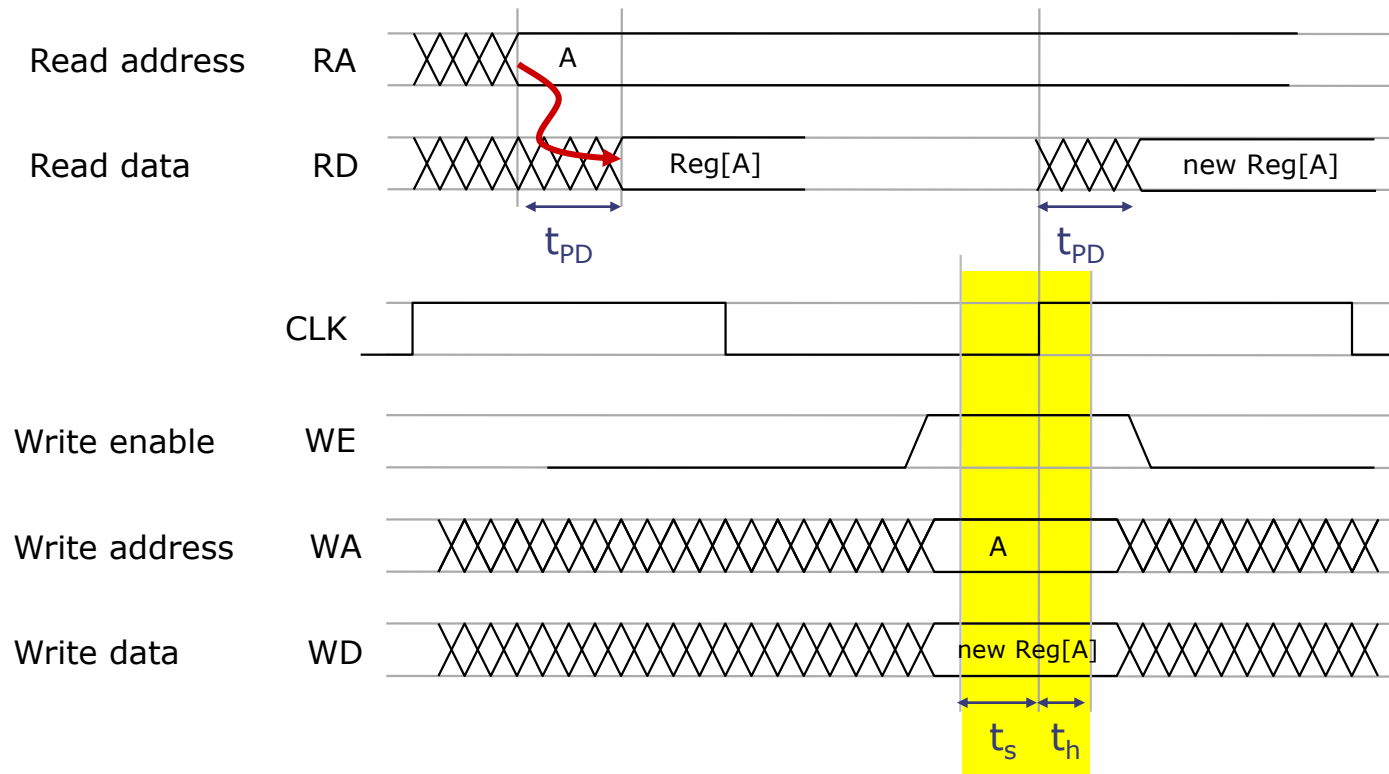
2 combinational READ ports, 1 clocked WRITE port





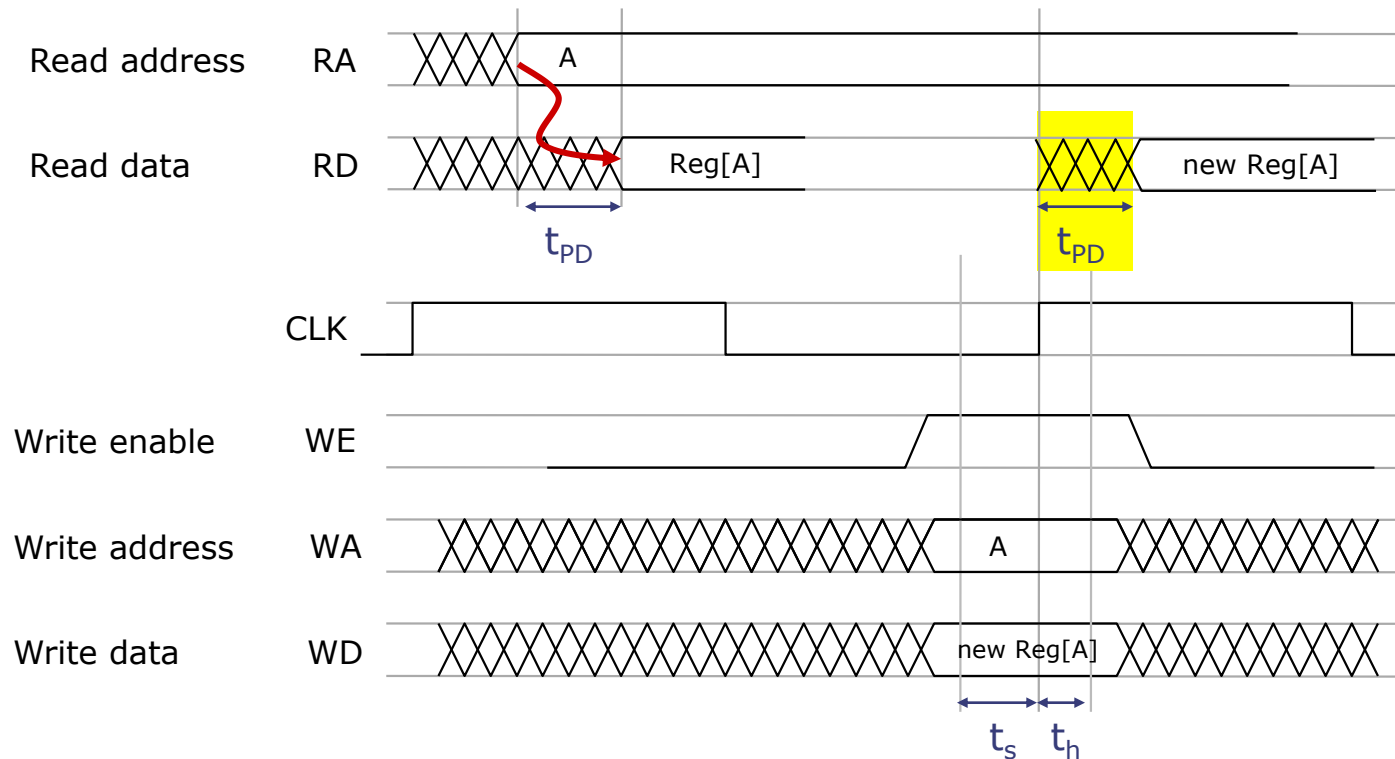
# Register File Timing

2 combinational READ ports, 1 clocked WRITE port



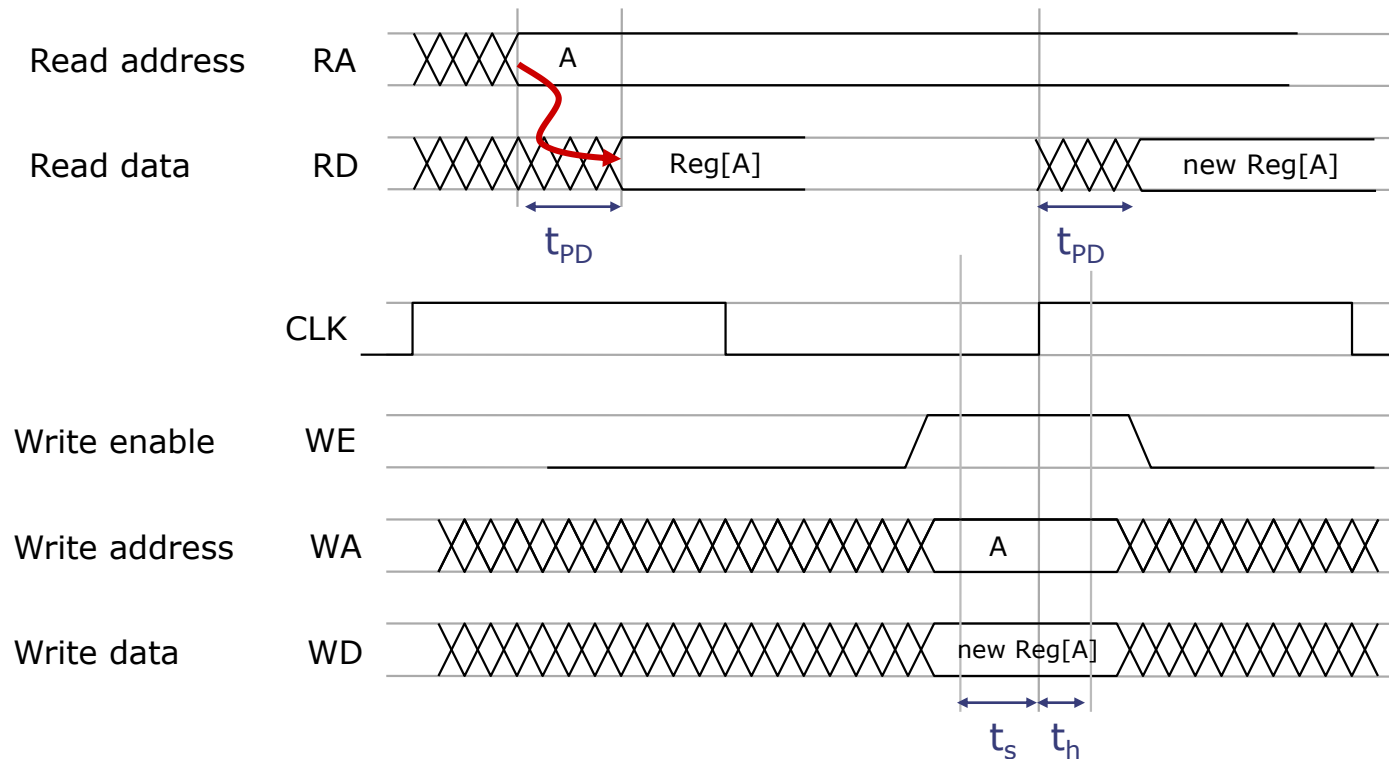
# Register File Timing

2 combinational READ ports, 1 clocked WRITE port



# Register File Timing

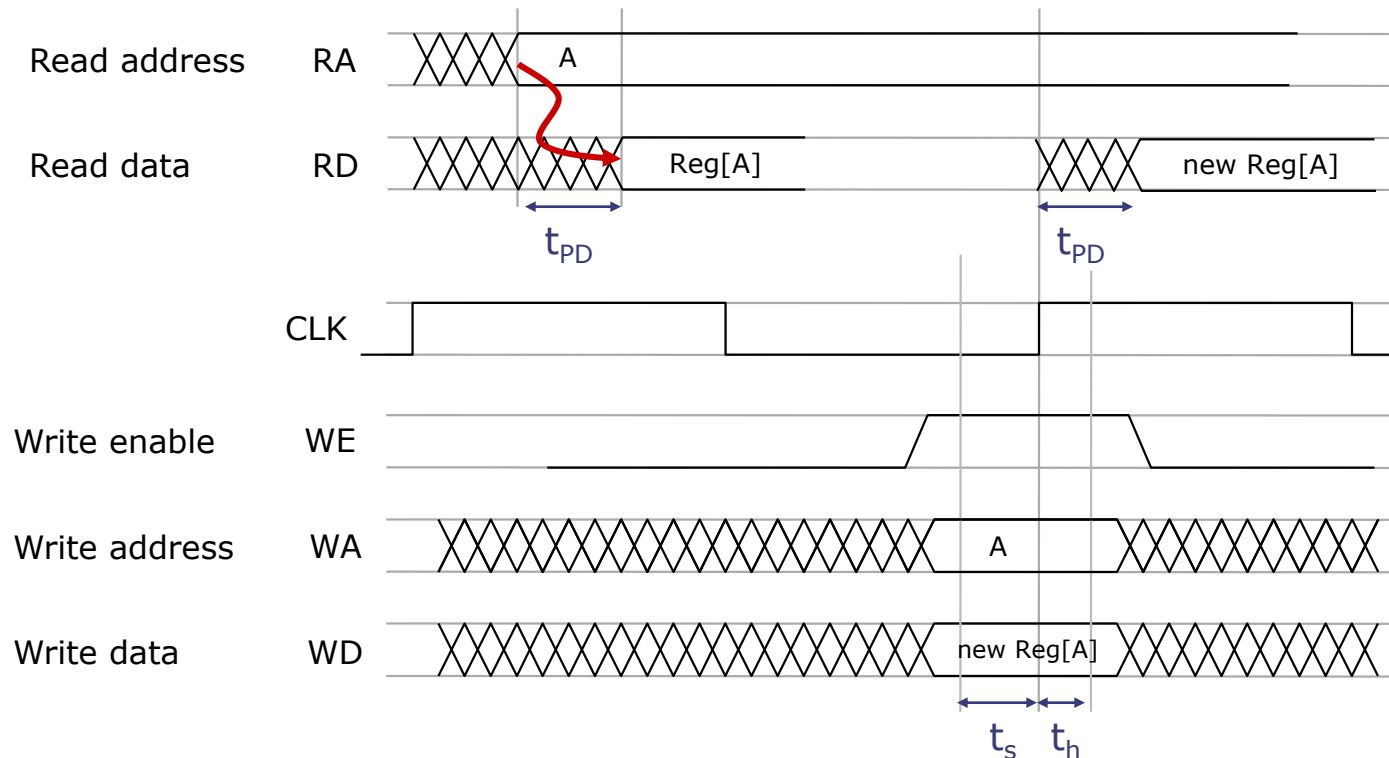
2 combinational READ ports, 1 clocked WRITE port



What if  $WA=RA$ ?

# Register File Timing

2 combinational READ ports, 1 clocked WRITE port



What if  $WA=RA1$ ?

RD1 reads "old" value of  $Reg[RA1]$   
until next clock edge!

# Memory Timing

---

- For now (lab 6), we will assume that our memories behave just like our register file in terms of timing

# Memory Timing

---

- For now (lab 6), we will assume that our memories behave just like our register file in terms of timing
  - Loads are combinational: Data is returned in the same clock cycle as load request

# Memory Timing

---

- For now (lab 6), we will assume that our memories behave just like our register file in terms of timing
  - Loads are combinational: Data is returned in the same clock cycle as load request
  - Stores are clocked

# Memory Timing

---

- For now (lab 6), we will assume that our memories behave just like our register file in terms of timing
  - Loads are combinational: Data is returned in the same clock cycle as load request
  - Stores are clocked
  - In lab 6, you will see the memory module referred to as a **magic memory**, since it's not realistic



# Memory Timing

---

- For now (lab 6), we will assume that our memories behave just like our register file in terms of timing
  - Loads are combinational: Data is returned in the same clock cycle as load request
  - Stores are clocked
  - In lab 6, you will see the memory module referred to as a **magic memory**, since it's not realistic
- Next week we will learn about the implementation of these memories

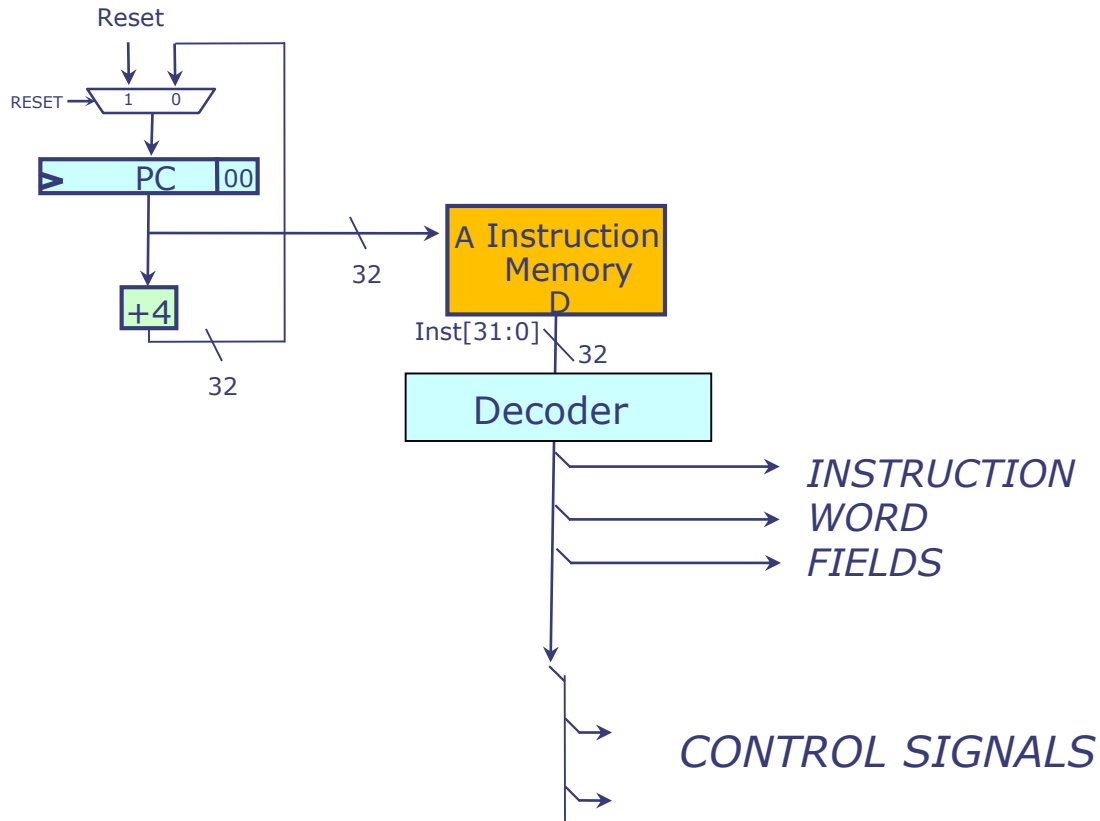
# Memory Timing

---

- For now (lab 6), we will assume that our memories behave just like our register file in terms of timing
  - Loads are combinational: Data is returned in the same clock cycle as load request
  - Stores are clocked
  - In lab 6, you will see the memory module referred to as a **magic memory**, since it's not realistic
- Next week we will learn about the implementation of these memories
- In lab 7 and the design project, we will use realistic memories for our processor

# Instruction Fetch/Decode

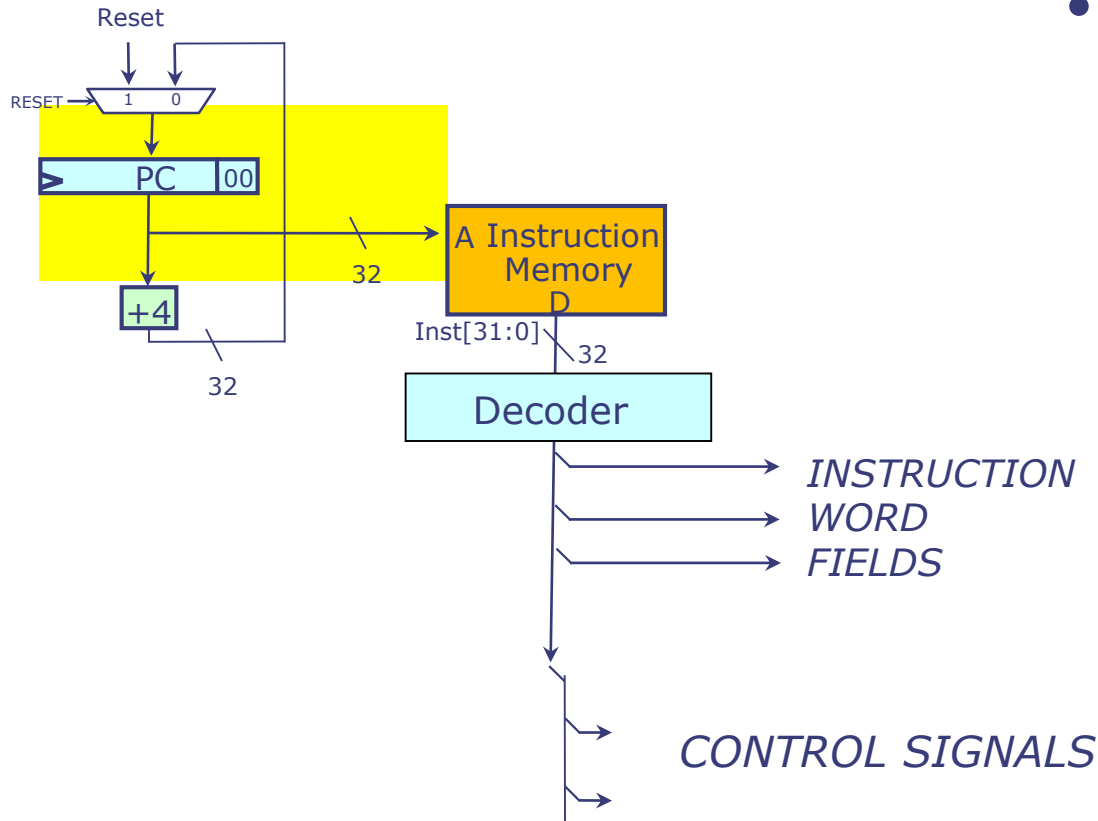
Use Program Counter (PC) to fetch the next instruction:



# Instruction Fetch/Decode

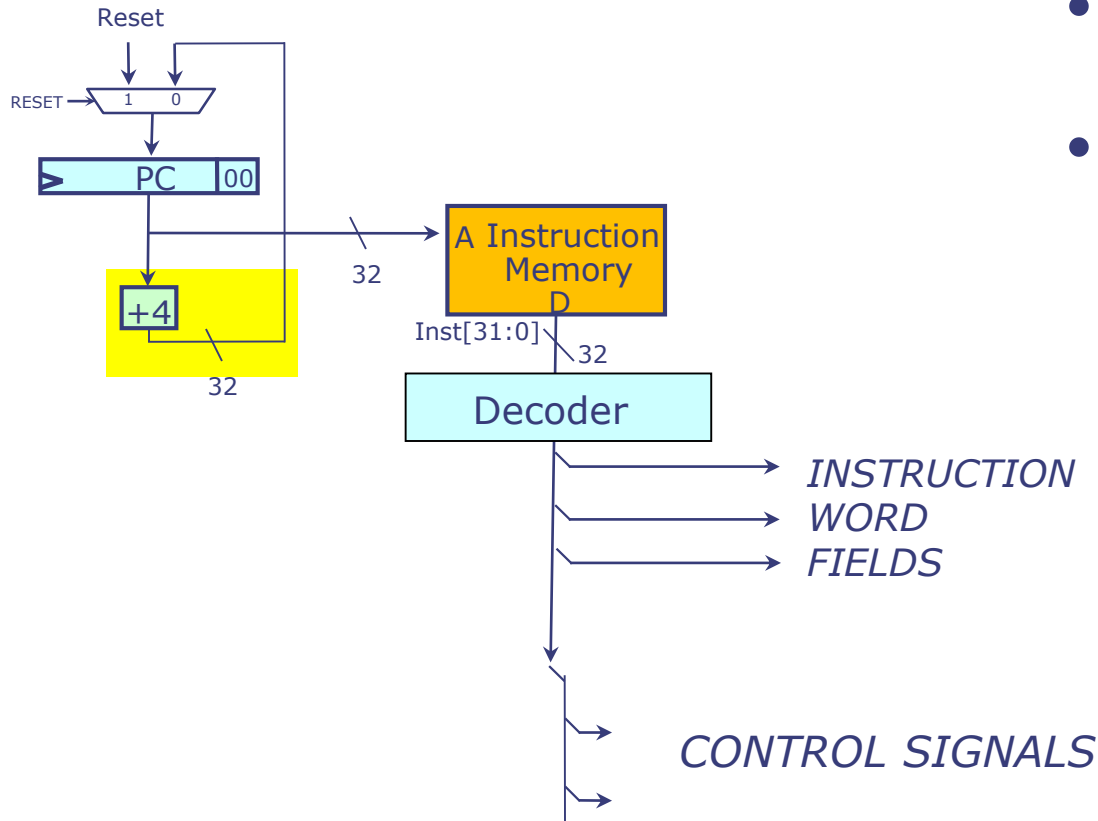
Use Program Counter (PC) to fetch the next instruction:

- Use PC as memory address



# Instruction Fetch/Decode

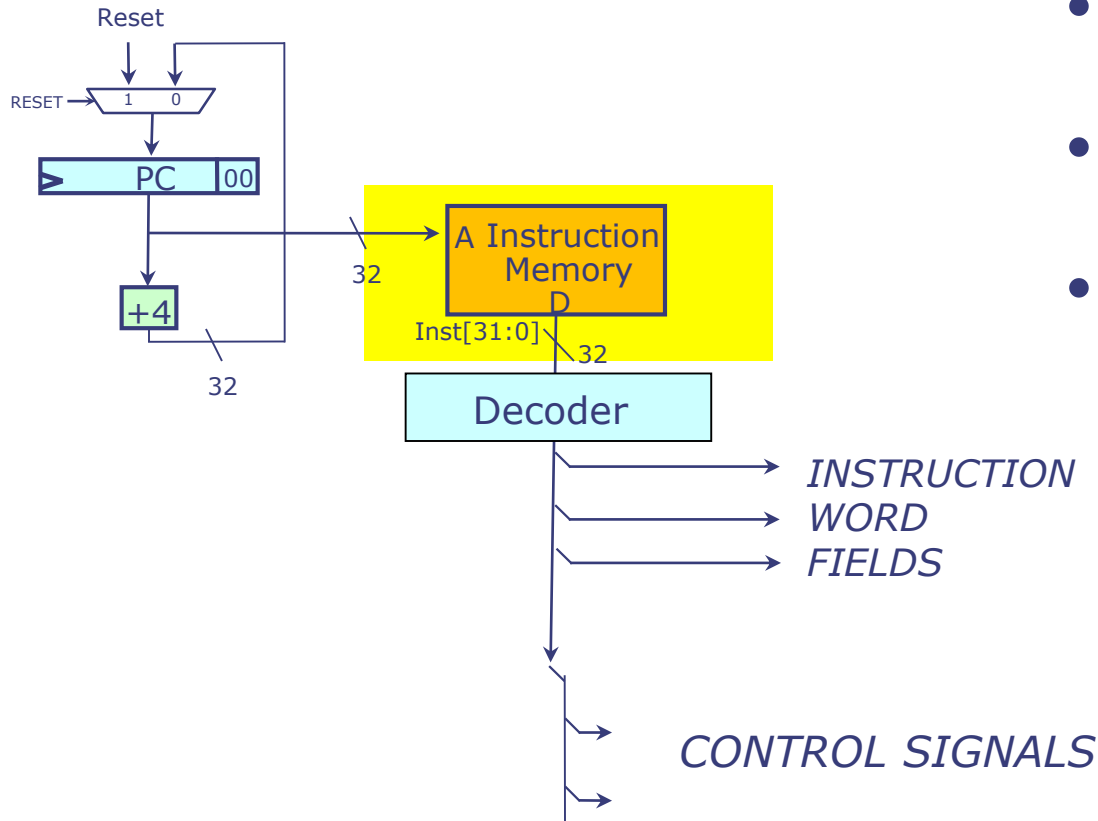
Use Program Counter (PC) to fetch the next instruction:



- Use PC as memory address
- Add 4 to PC, load new value at end of cycle

# Instruction Fetch/Decode

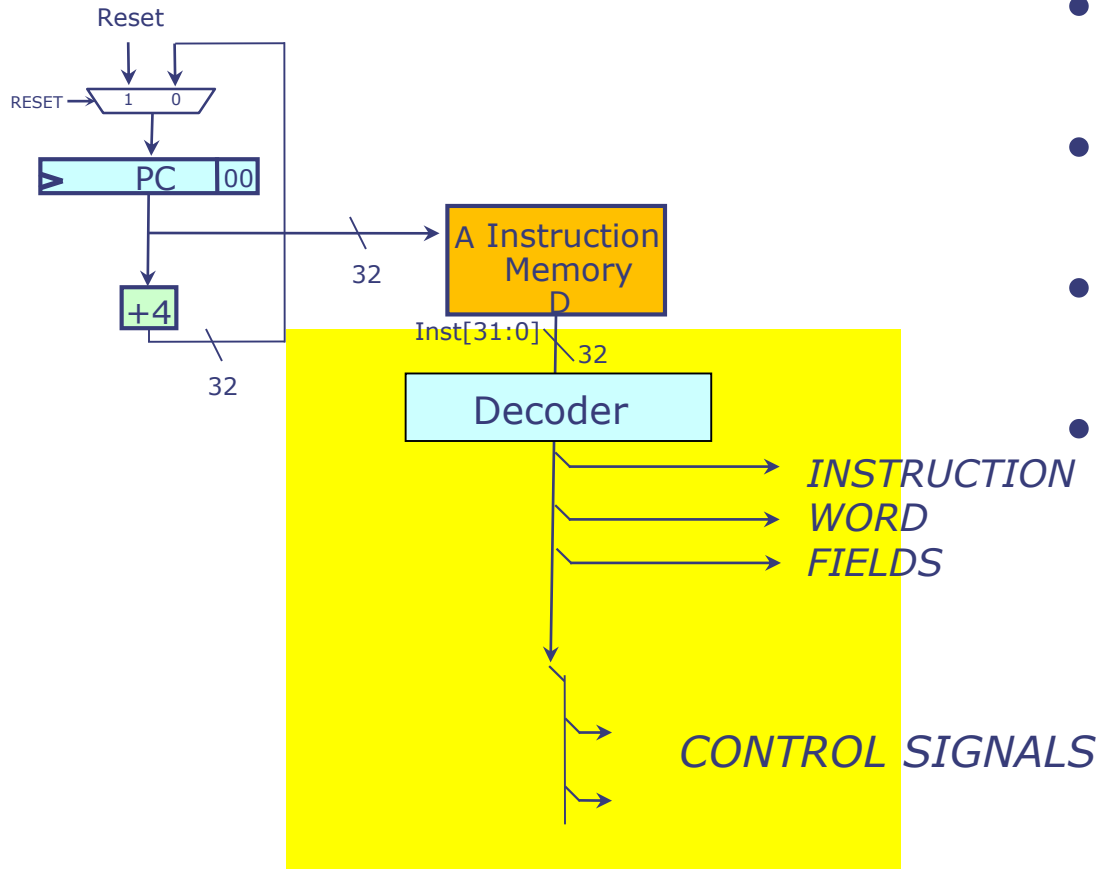
Use Program Counter (PC) to fetch the next instruction:



- Use PC as memory address
- Add 4 to PC, load new value at end of cycle
- Fetch instruction from memory

# Instruction Fetch/Decode

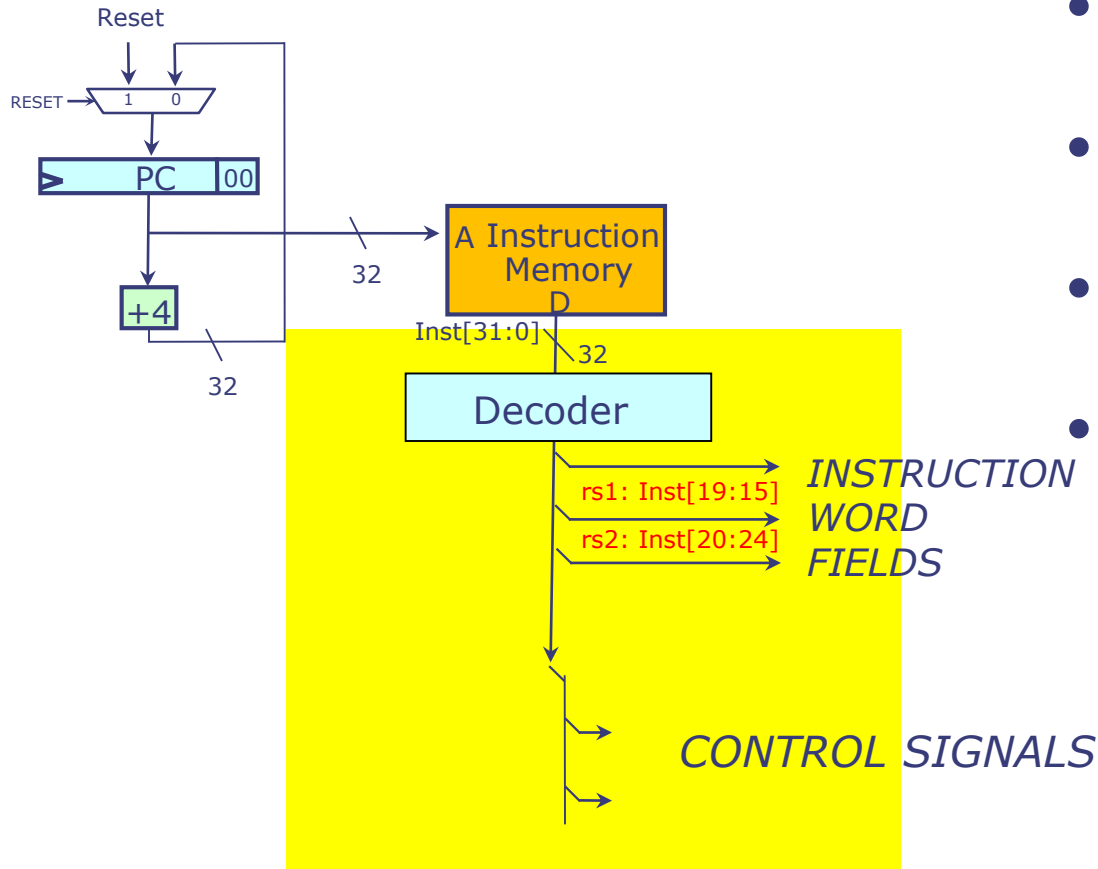
Use Program Counter (PC) to fetch the next instruction:



- Use PC as memory address
- Add 4 to PC, load new value at end of cycle
- Fetch instruction from memory
- Decode instruction:

# Instruction Fetch/Decode

Use Program Counter (PC) to fetch the next instruction:

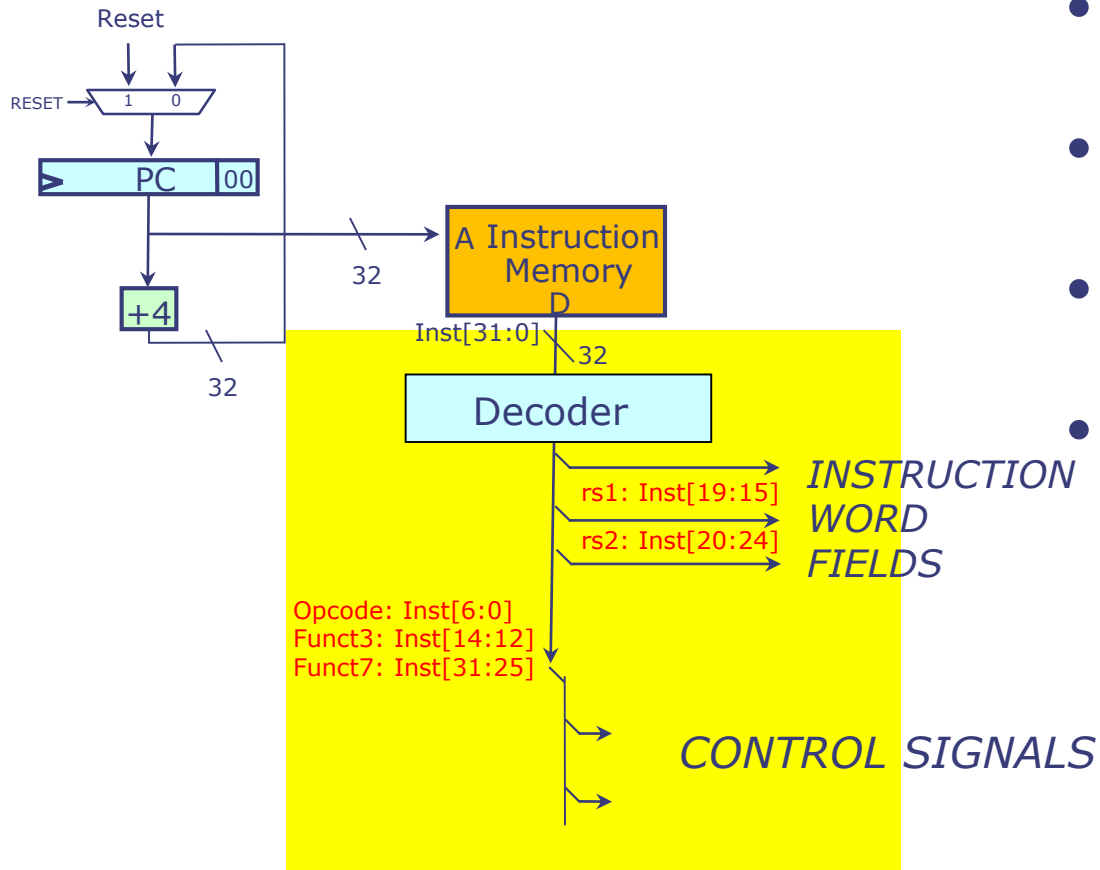


- Use PC as memory address
- Add 4 to PC, load new value at end of cycle
- Fetch instruction from memory
- Decode instruction:
  - Use some instruction fields directly (register indexes, immediate values)



# Instruction Fetch/Decode

Use Program Counter (PC) to fetch the next instruction:



- Use PC as memory address
- Add 4 to PC, load new value at end of cycle
- Fetch instruction from memory
- Decode instruction:
  - Use some instruction fields directly (register indexes, immediate values)
  - Use opcode, funct3, and funct7 bits to generate control signals

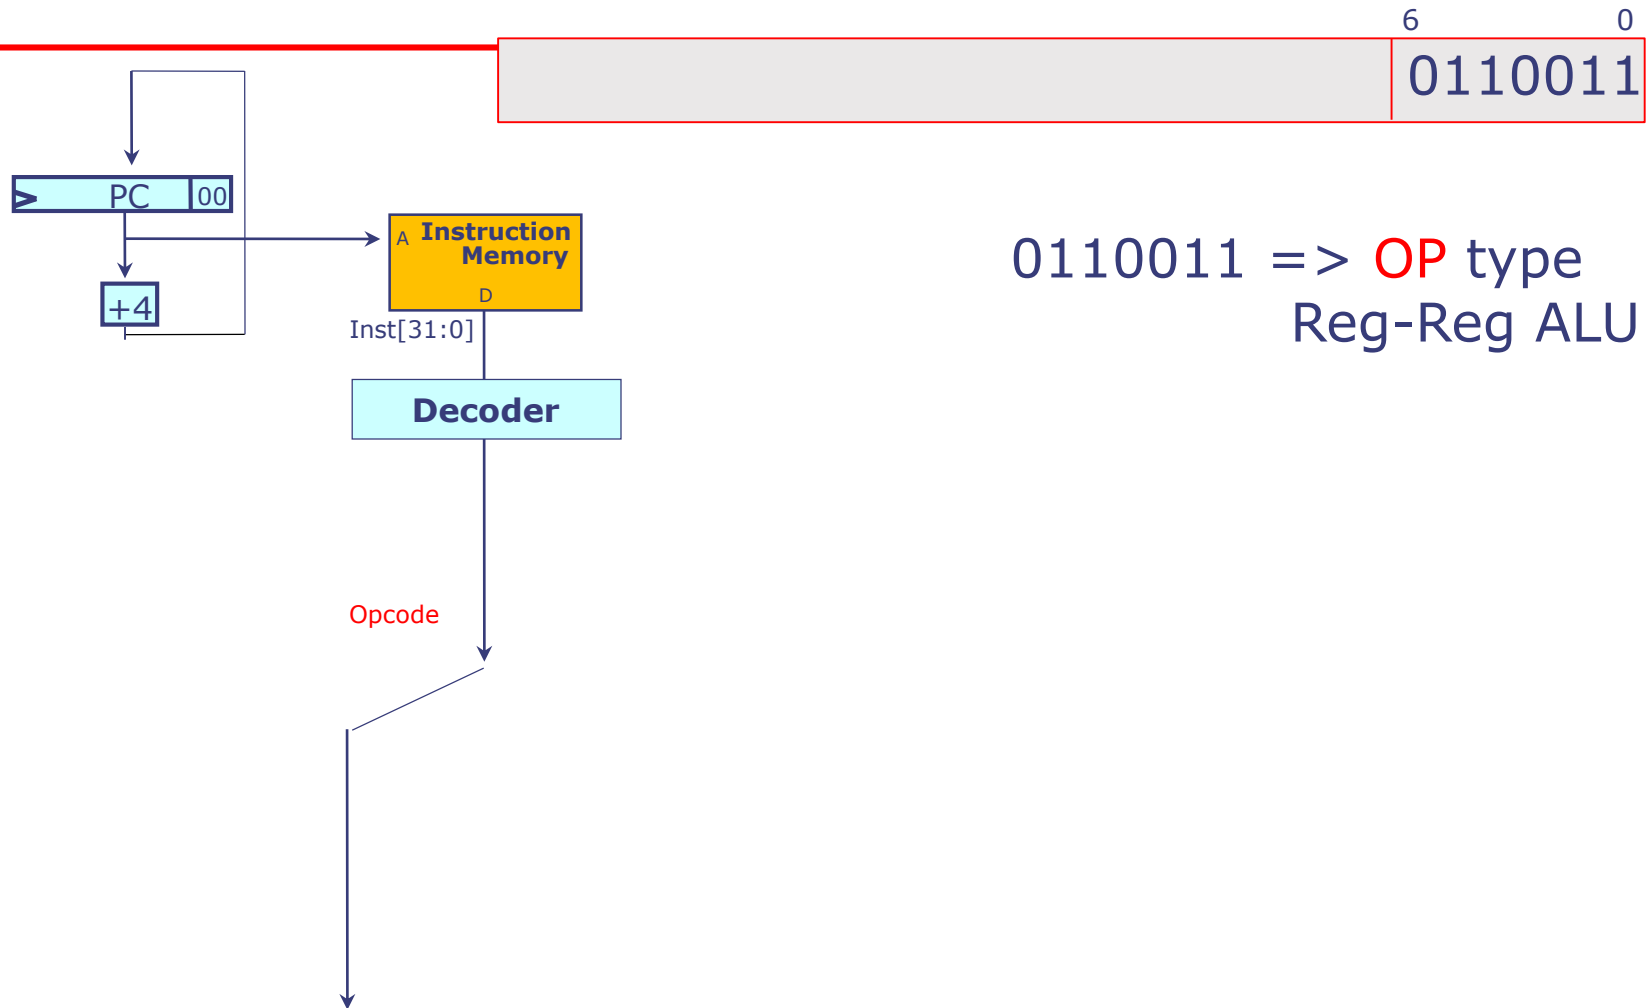
# ALU Instructions

Differ only in the ALU op to be performed

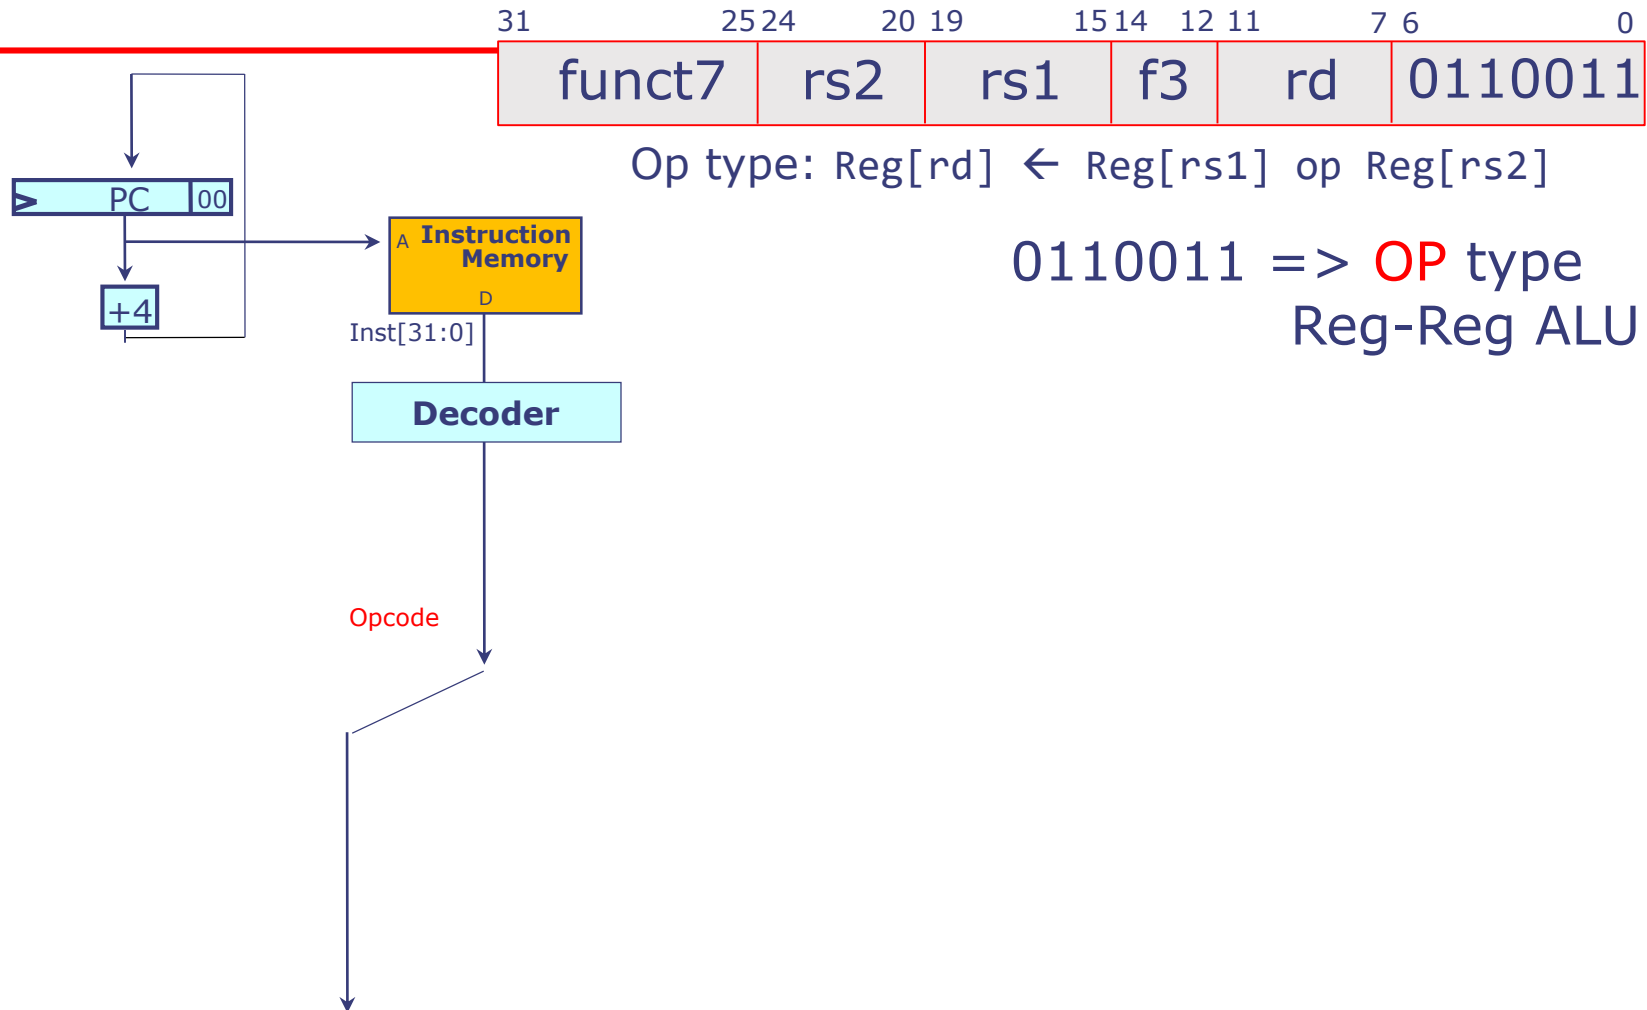
Instruction	Description	Execution
ADD rd, rs1, rs2	Add	$\text{reg}[\text{rd}] \leq \text{reg}[\text{rs1}] + \text{reg}[\text{rs2}]$
SUB rd, rs1, rs2	Sub	$\text{reg}[\text{rd}] \leq \text{reg}[\text{rs1}] - \text{reg}[\text{rs2}]$
SLL rd, rs1, rs2	Shift Left Logical	$\text{reg}[\text{rd}] \leq \text{reg}[\text{rs1}] \ll \text{reg}[\text{rs2}]$
SLT rd, rs1, rs2	Set if < (Signed)	$\text{reg}[\text{rd}] \leq (\text{reg}[\text{rs1}] <_s \text{reg}[\text{rs2}]) ? 1 : 0$
SLTU rd, rs1, rs2	Set if < (Unsigned)	$\text{reg}[\text{rd}] \leq (\text{reg}[\text{rs1}] <_u \text{reg}[\text{rs2}]) ? 1 : 0$
XOR rd, rs1, rs2	Xor	$\text{reg}[\text{rd}] \leq \text{reg}[\text{rs1}] \wedge \text{reg}[\text{rs2}]$
SRL rd, rs1, rs2	Shift Right Logical	$\text{reg}[\text{rd}] \leq \text{reg}[\text{rs1}] \gg_u \text{reg}[\text{rs2}]$
SRA rd, rs1, rs2	Shift Right Arithmetic	$\text{reg}[\text{rd}] \leq \text{reg}[\text{rs1}] \gg_s \text{reg}[\text{rs2}]$
OR rd, rs1, rs2	Or	$\text{reg}[\text{rd}] \leq \text{reg}[\text{rs1}]   \text{reg}[\text{rs2}]$
AND rd, rs1, rs2	And	$\text{reg}[\text{rd}] \leq \text{reg}[\text{rs1}] \& \text{reg}[\text{rs2}]$

These instructions are grouped in a category called OP with fields (AluFunc, rd, rs1, rs2)

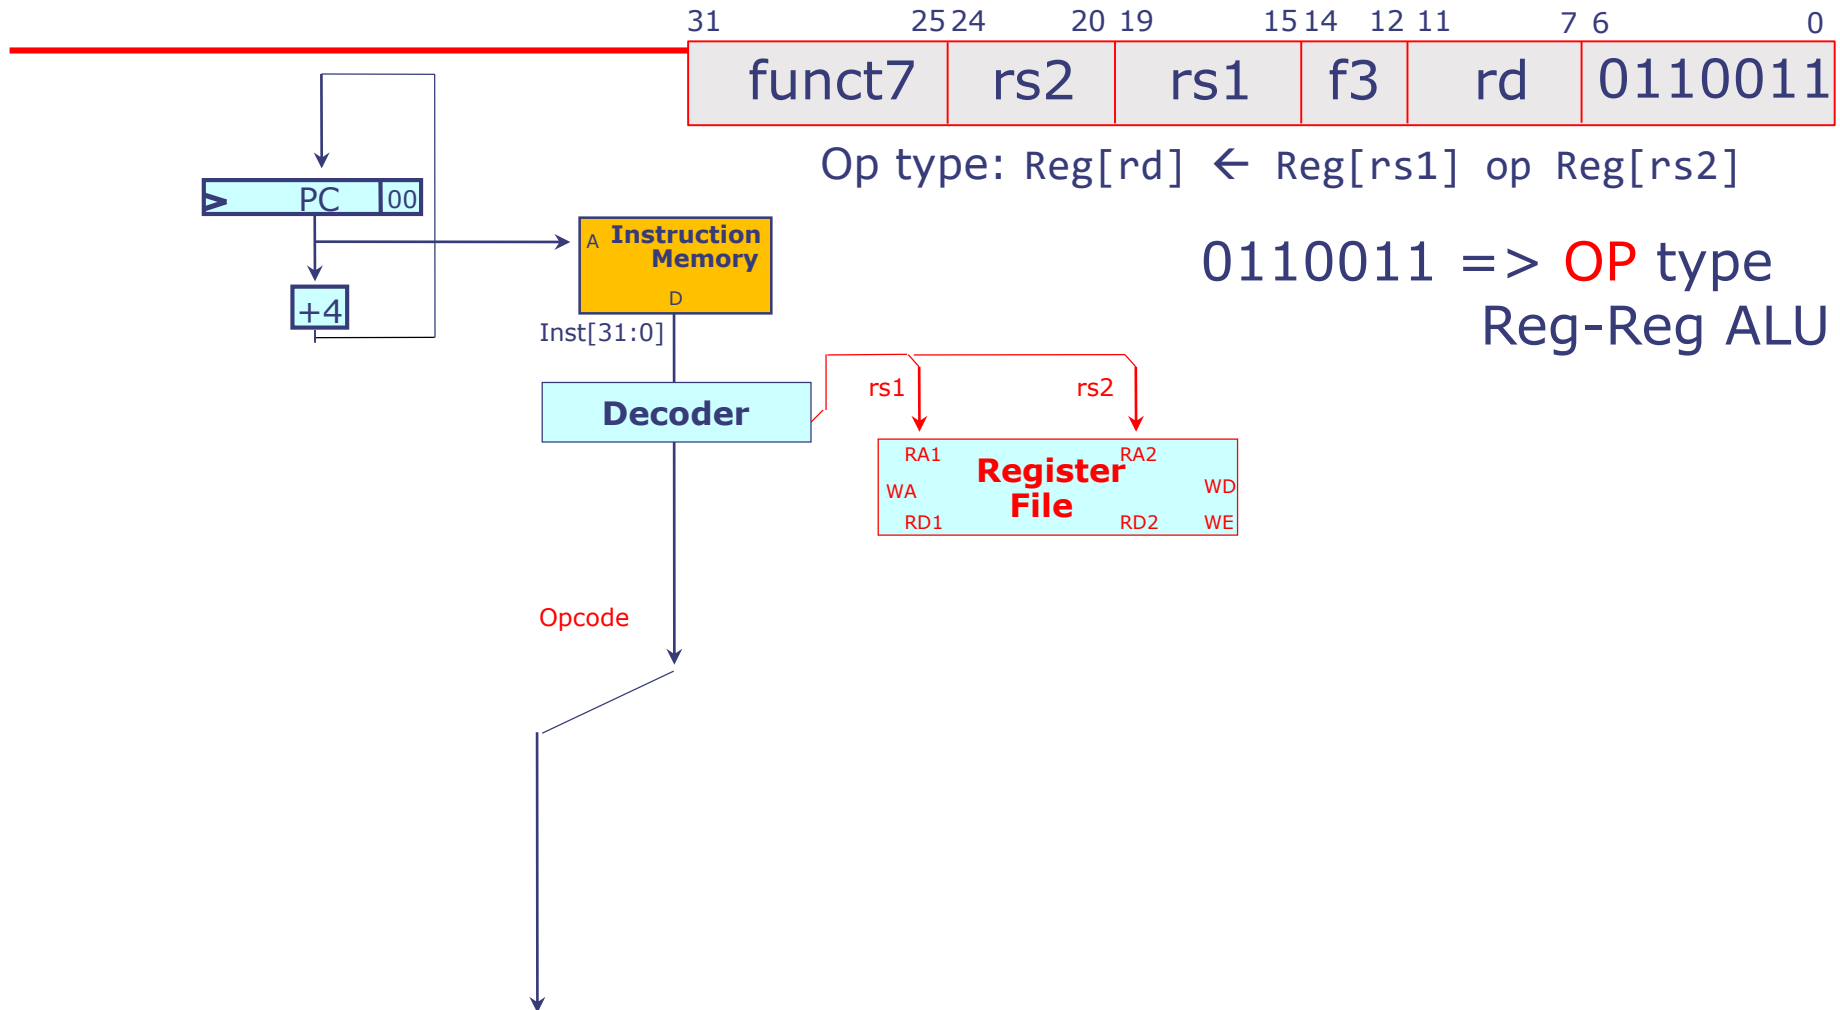
# Register-Register ALU Datapath



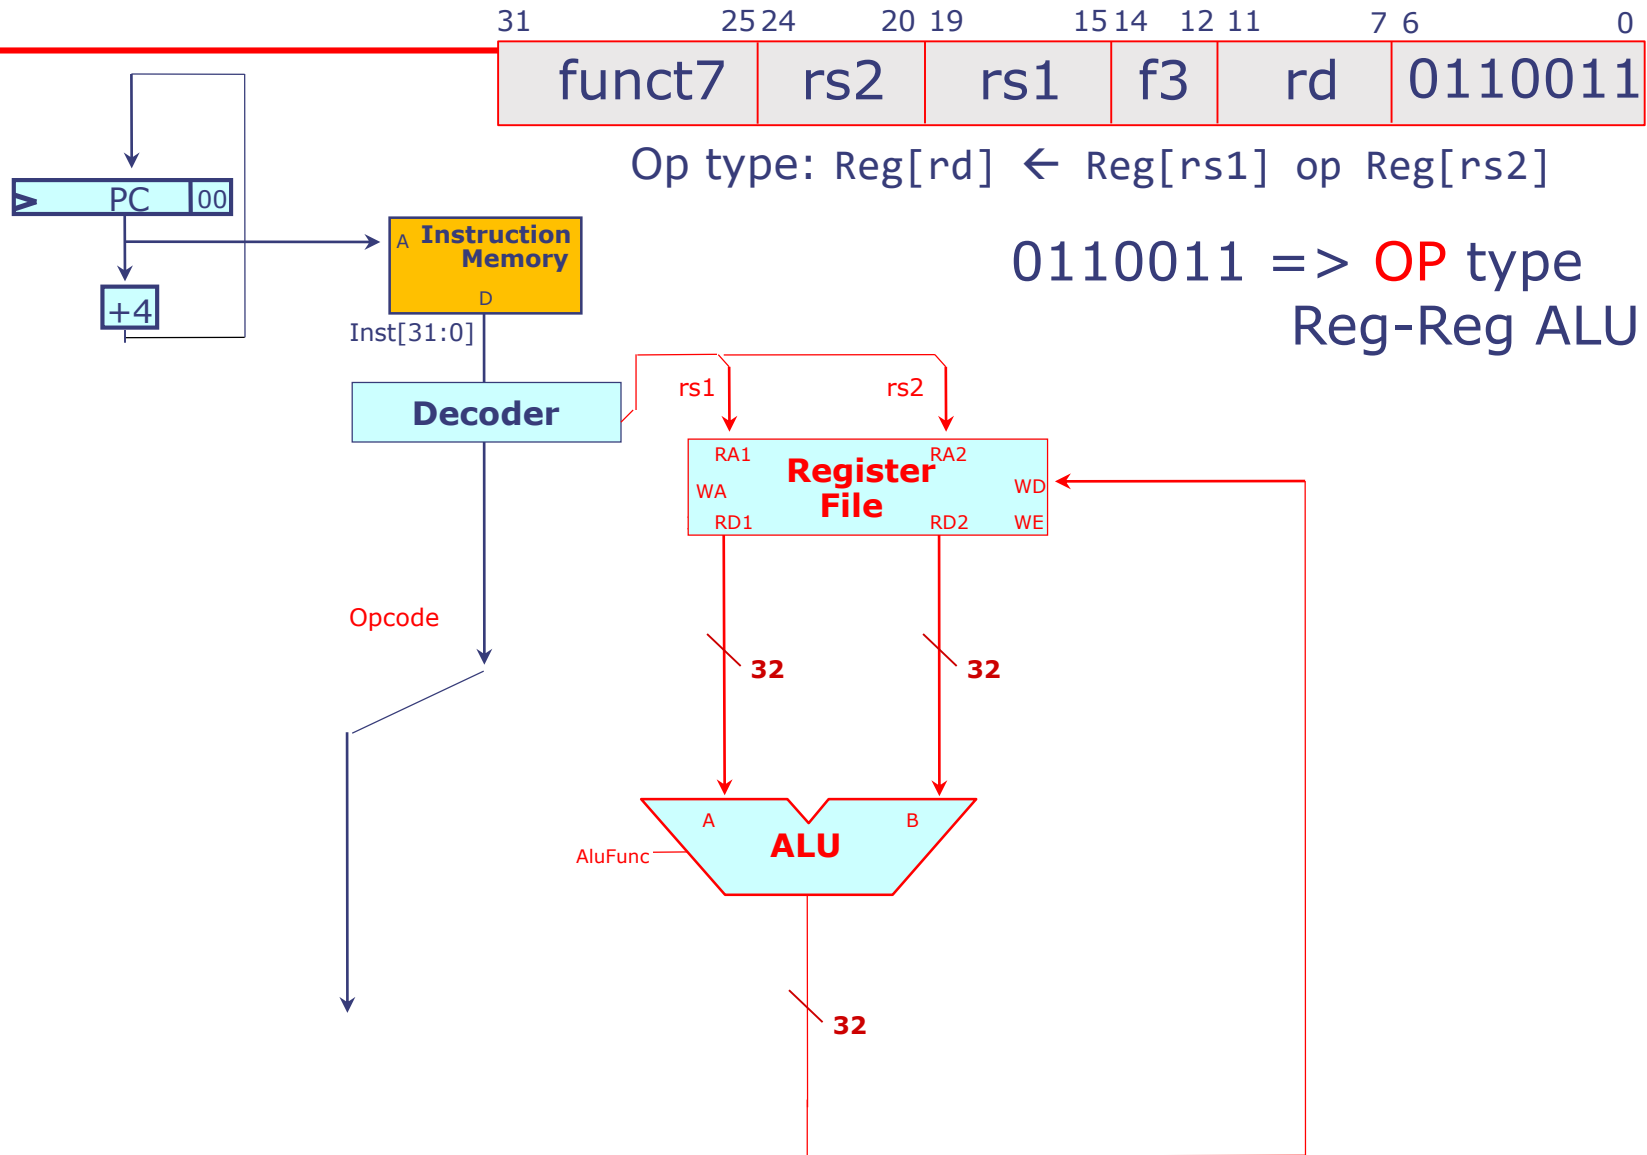
# Register-Register ALU Datapath



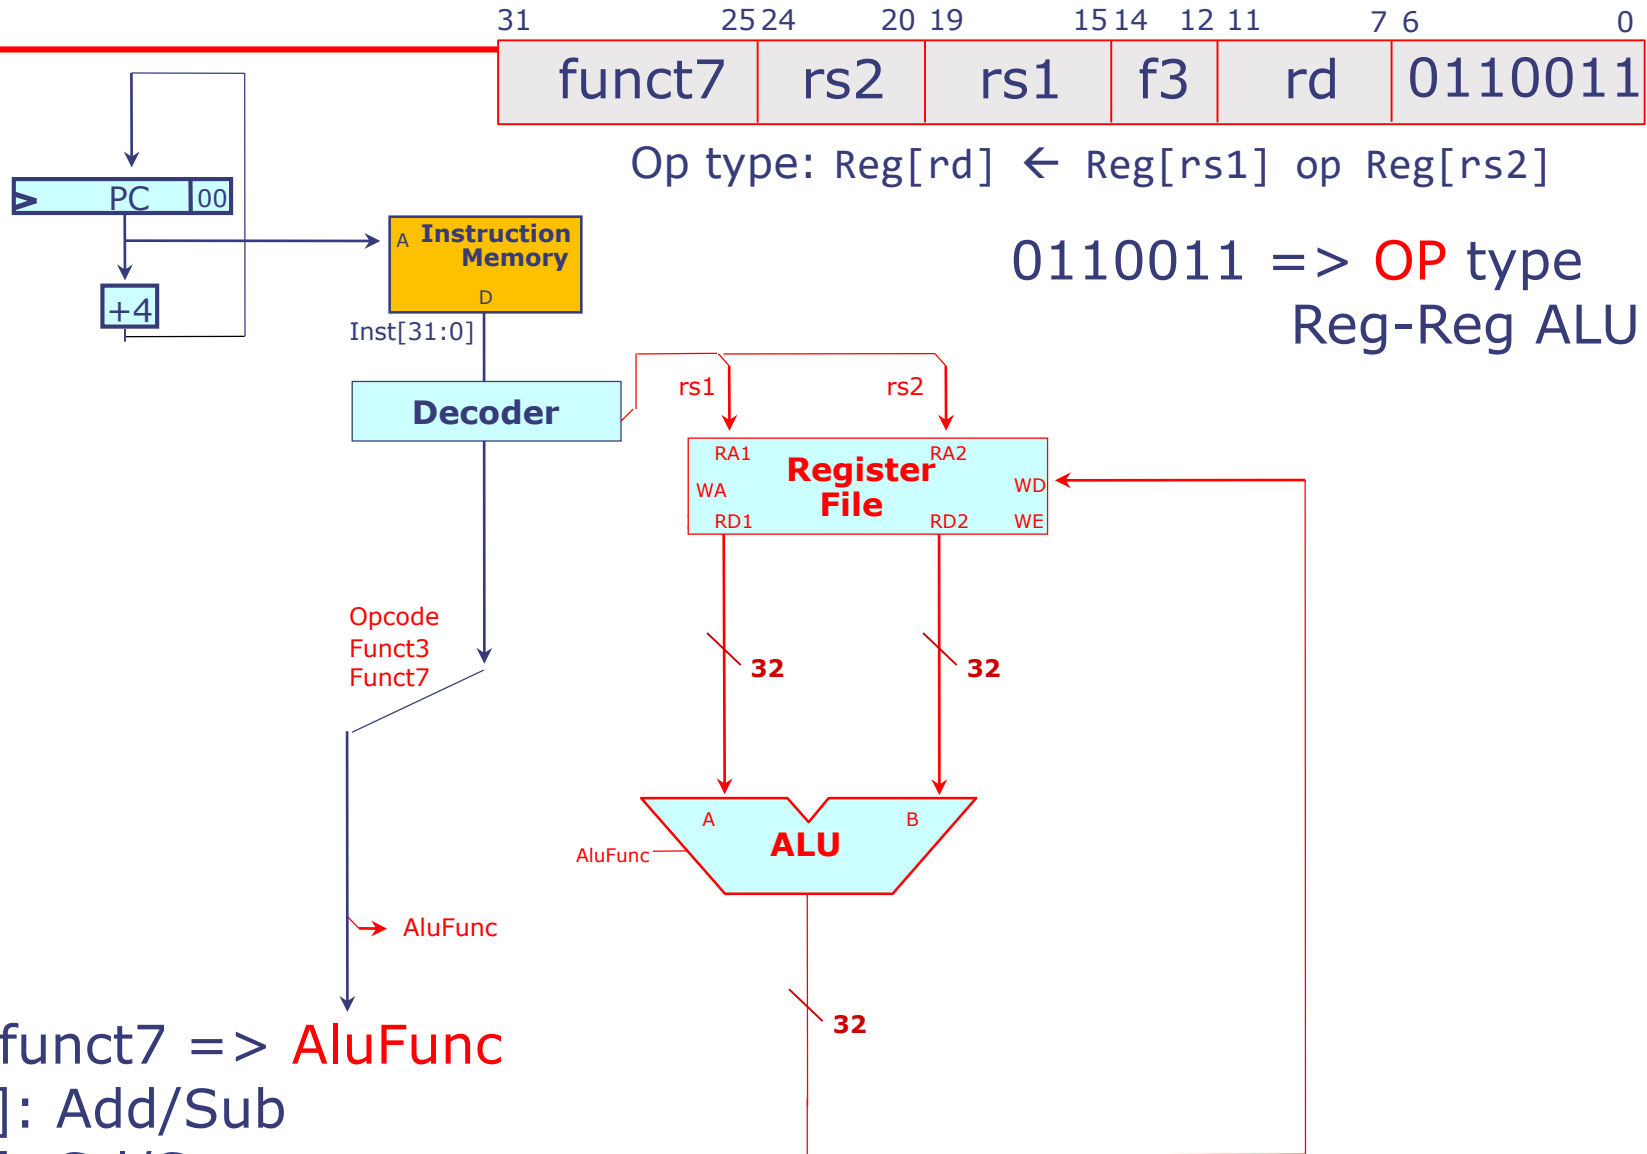
# Register-Register ALU Datapath



# Register-Register ALU Datapath



# Register-Register ALU Datapath

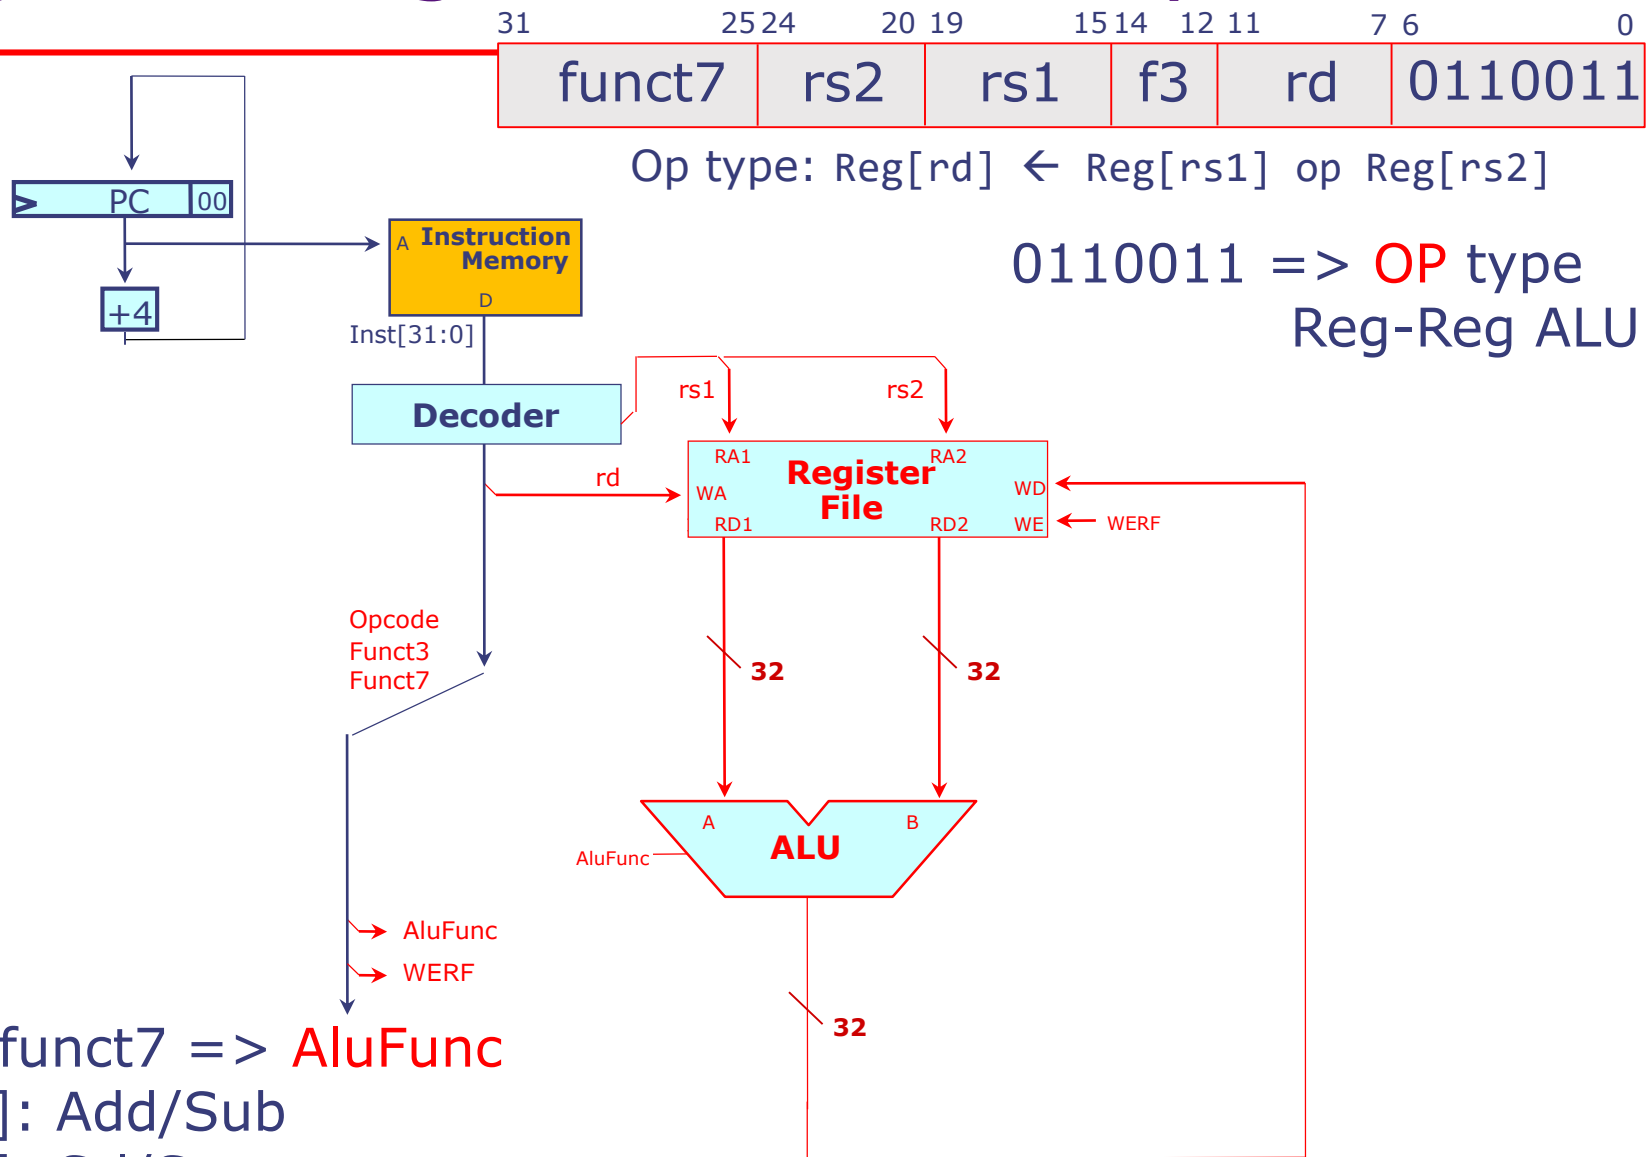


funct3, funct7 => **AluFunc**

Inst[30]: Add/Sub

Inst[30]: Srl/Sra

# Register-Register ALU Datapath



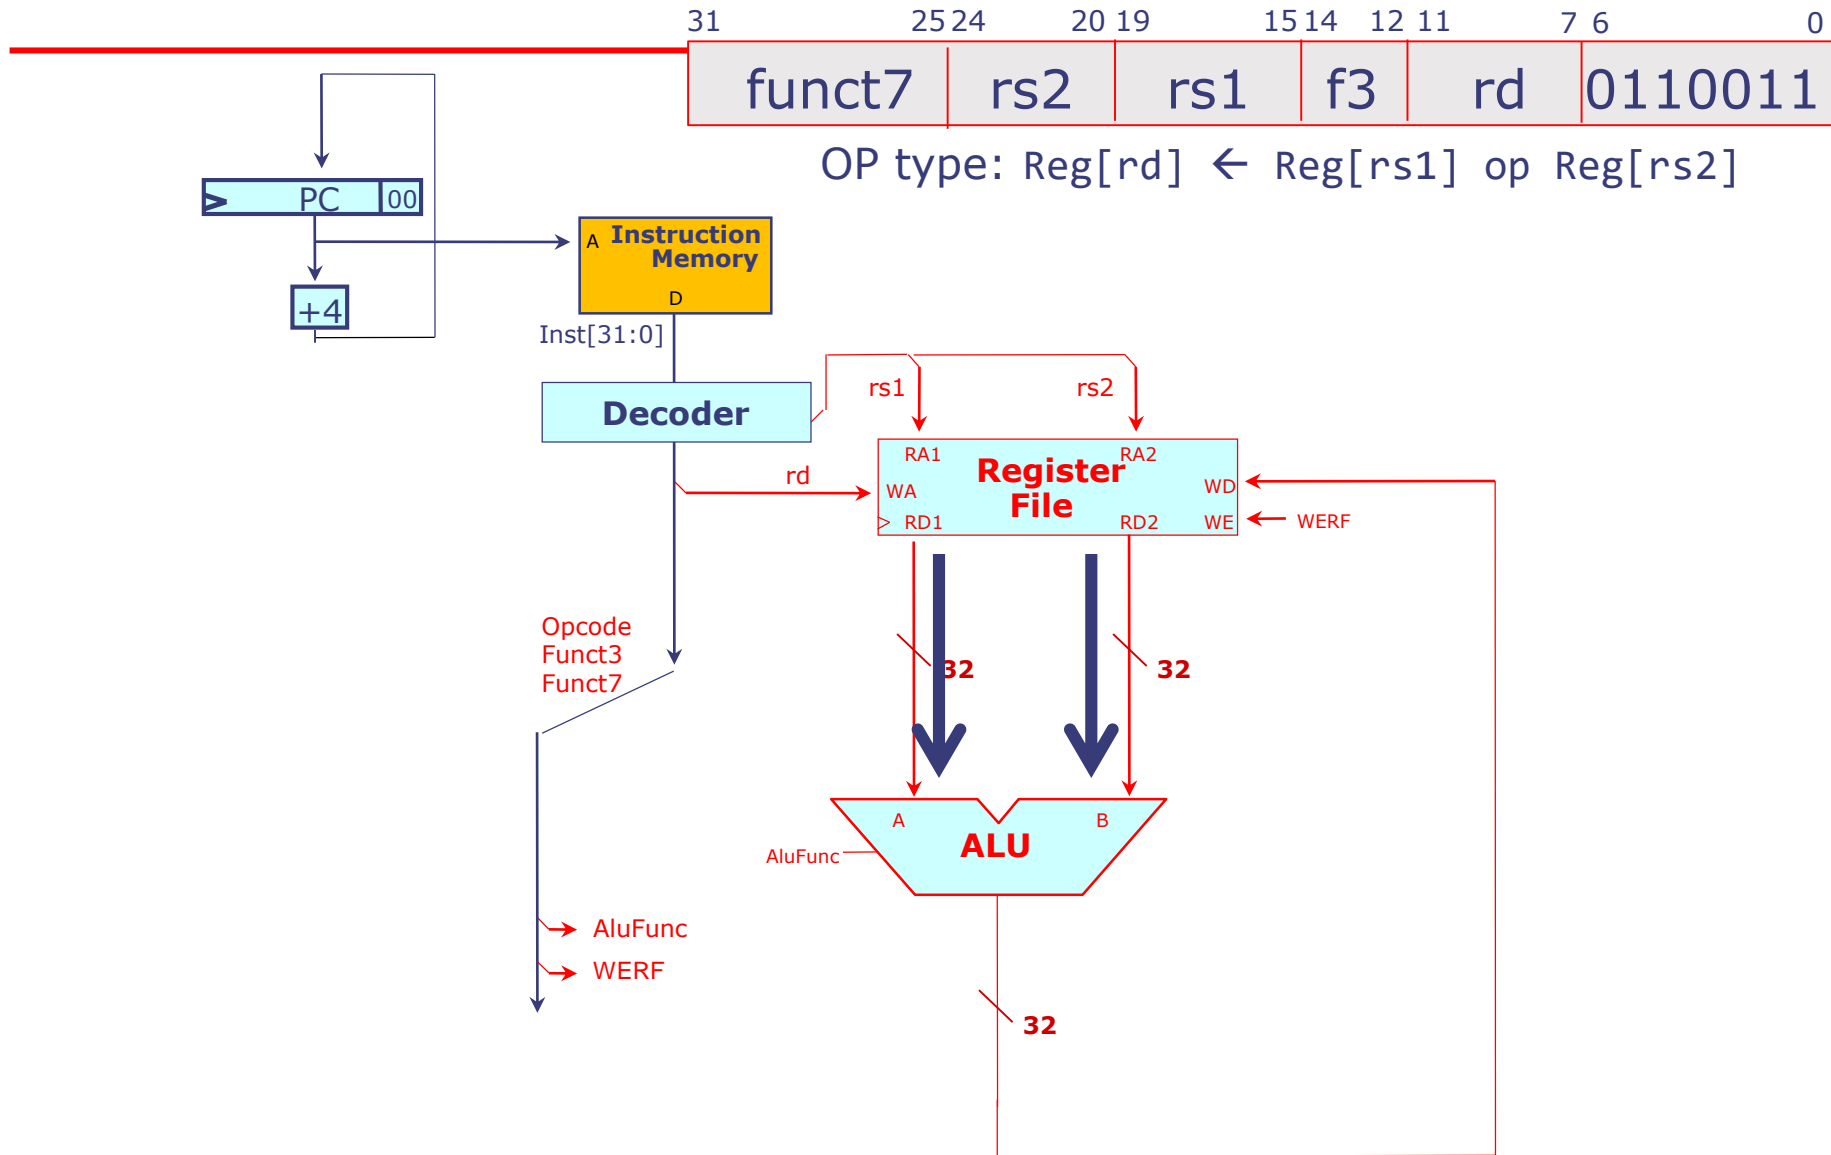
funct3, funct7 => **AluFunc**

Inst[30]: Add/Sub

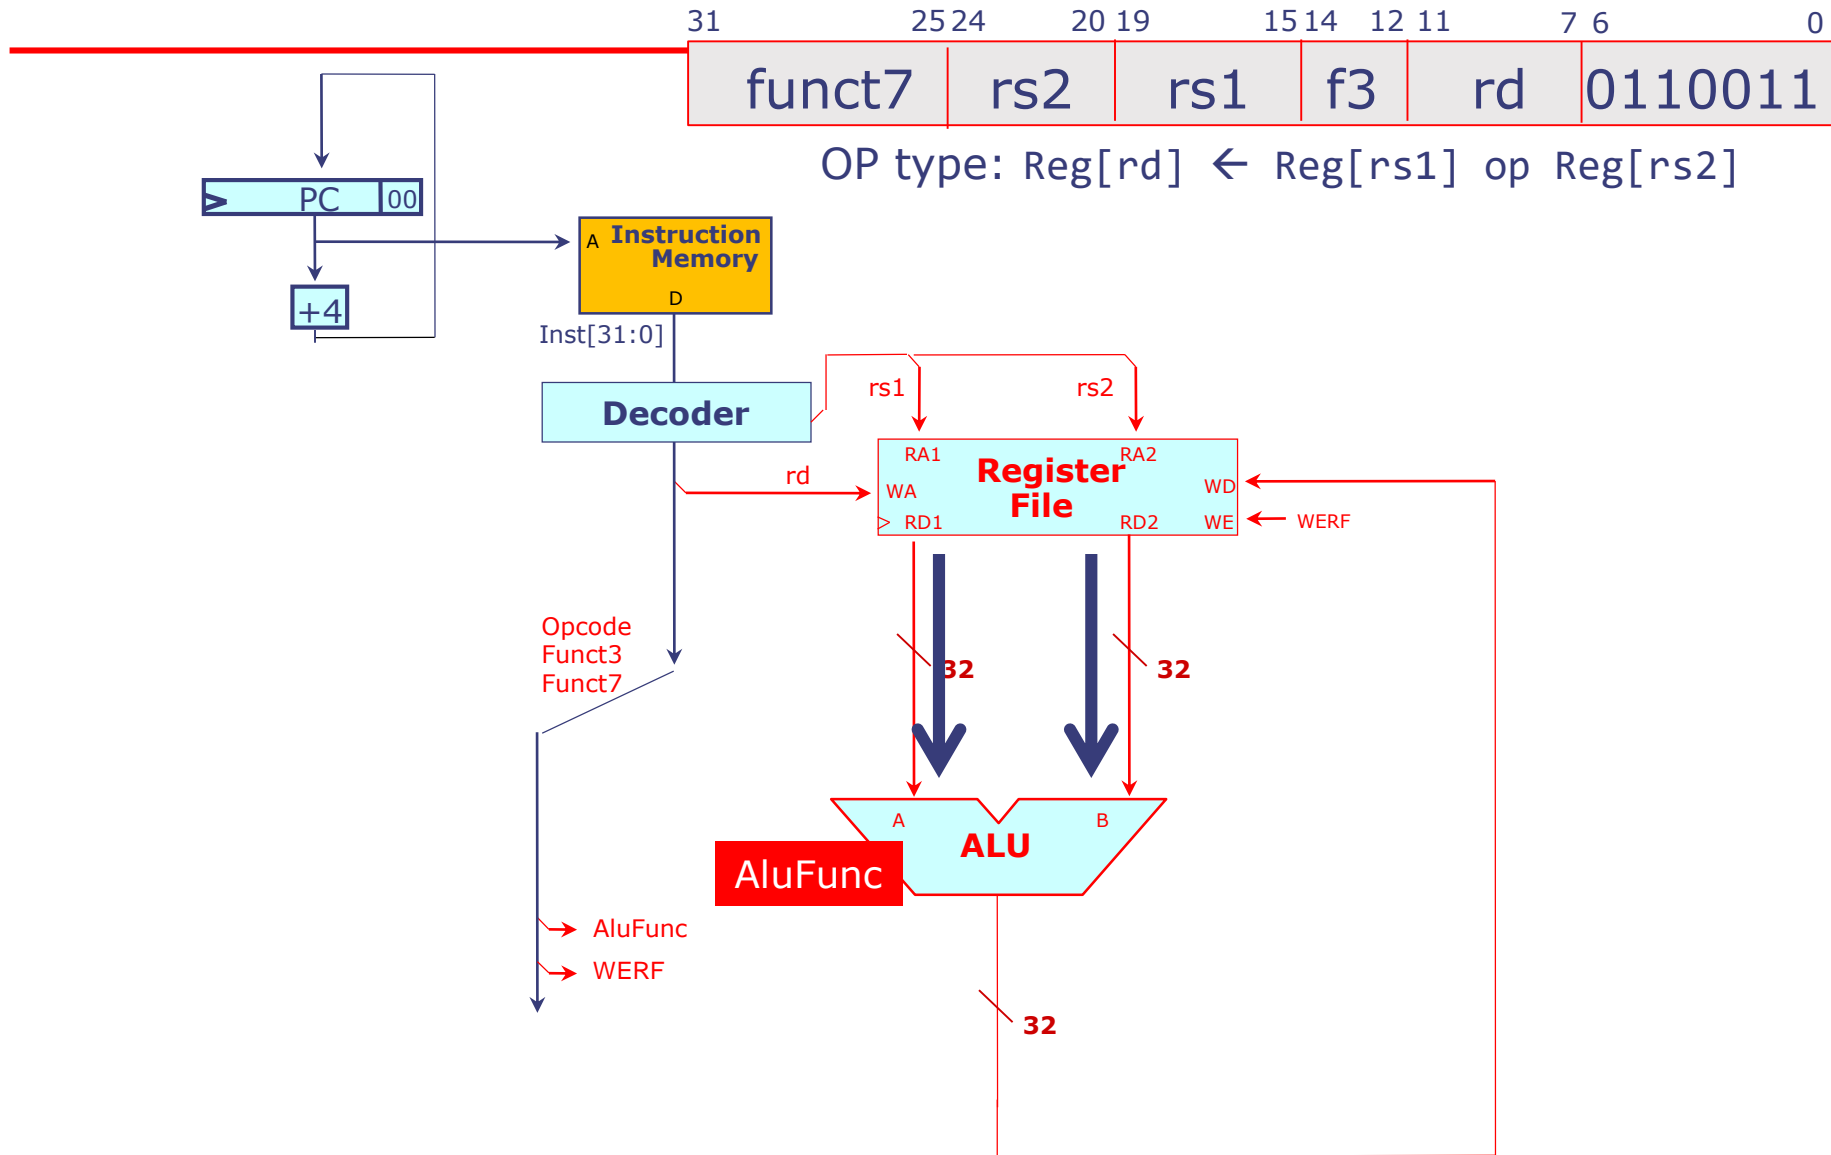
Inst[30]: Srl/Sra



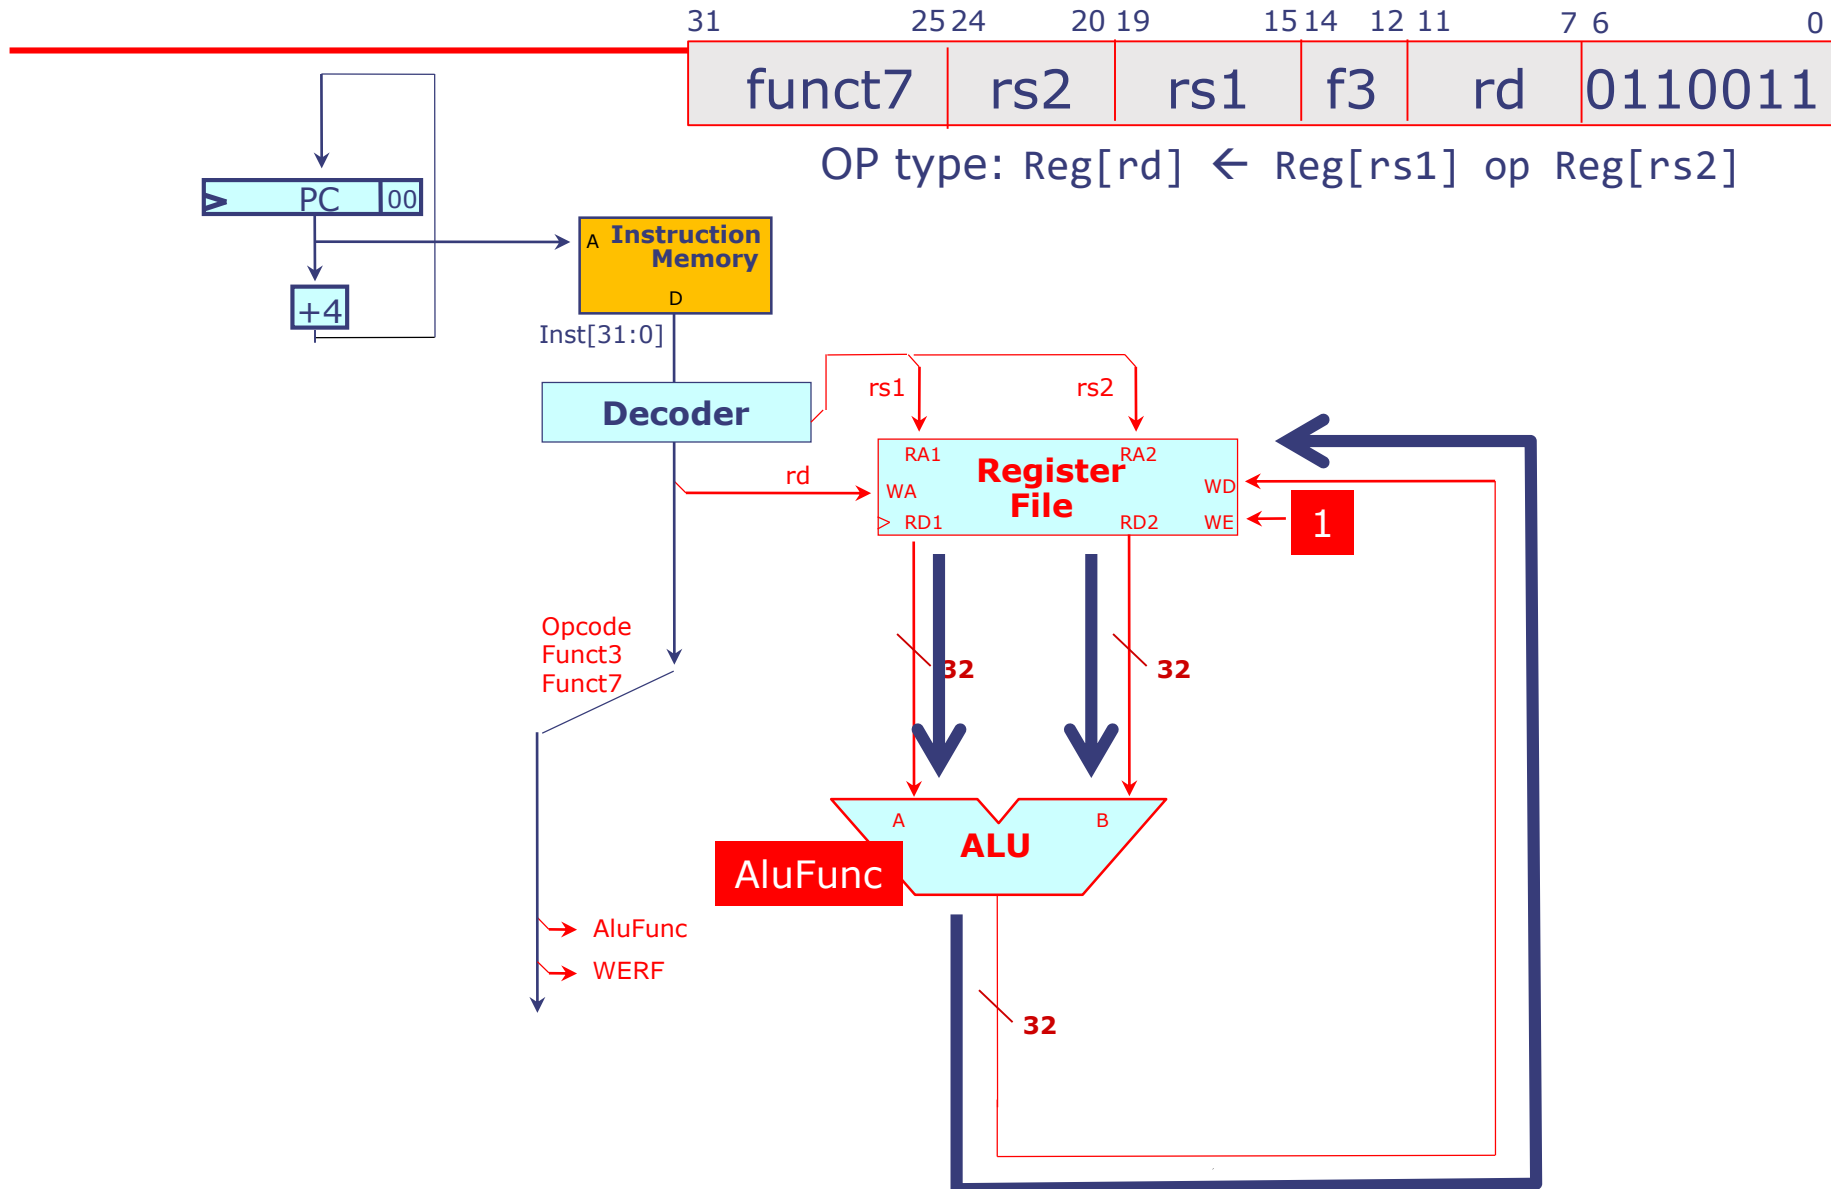
# Register-Register ALU Datapath



# Register-Register ALU Datapath



# Register-Register ALU Datapath



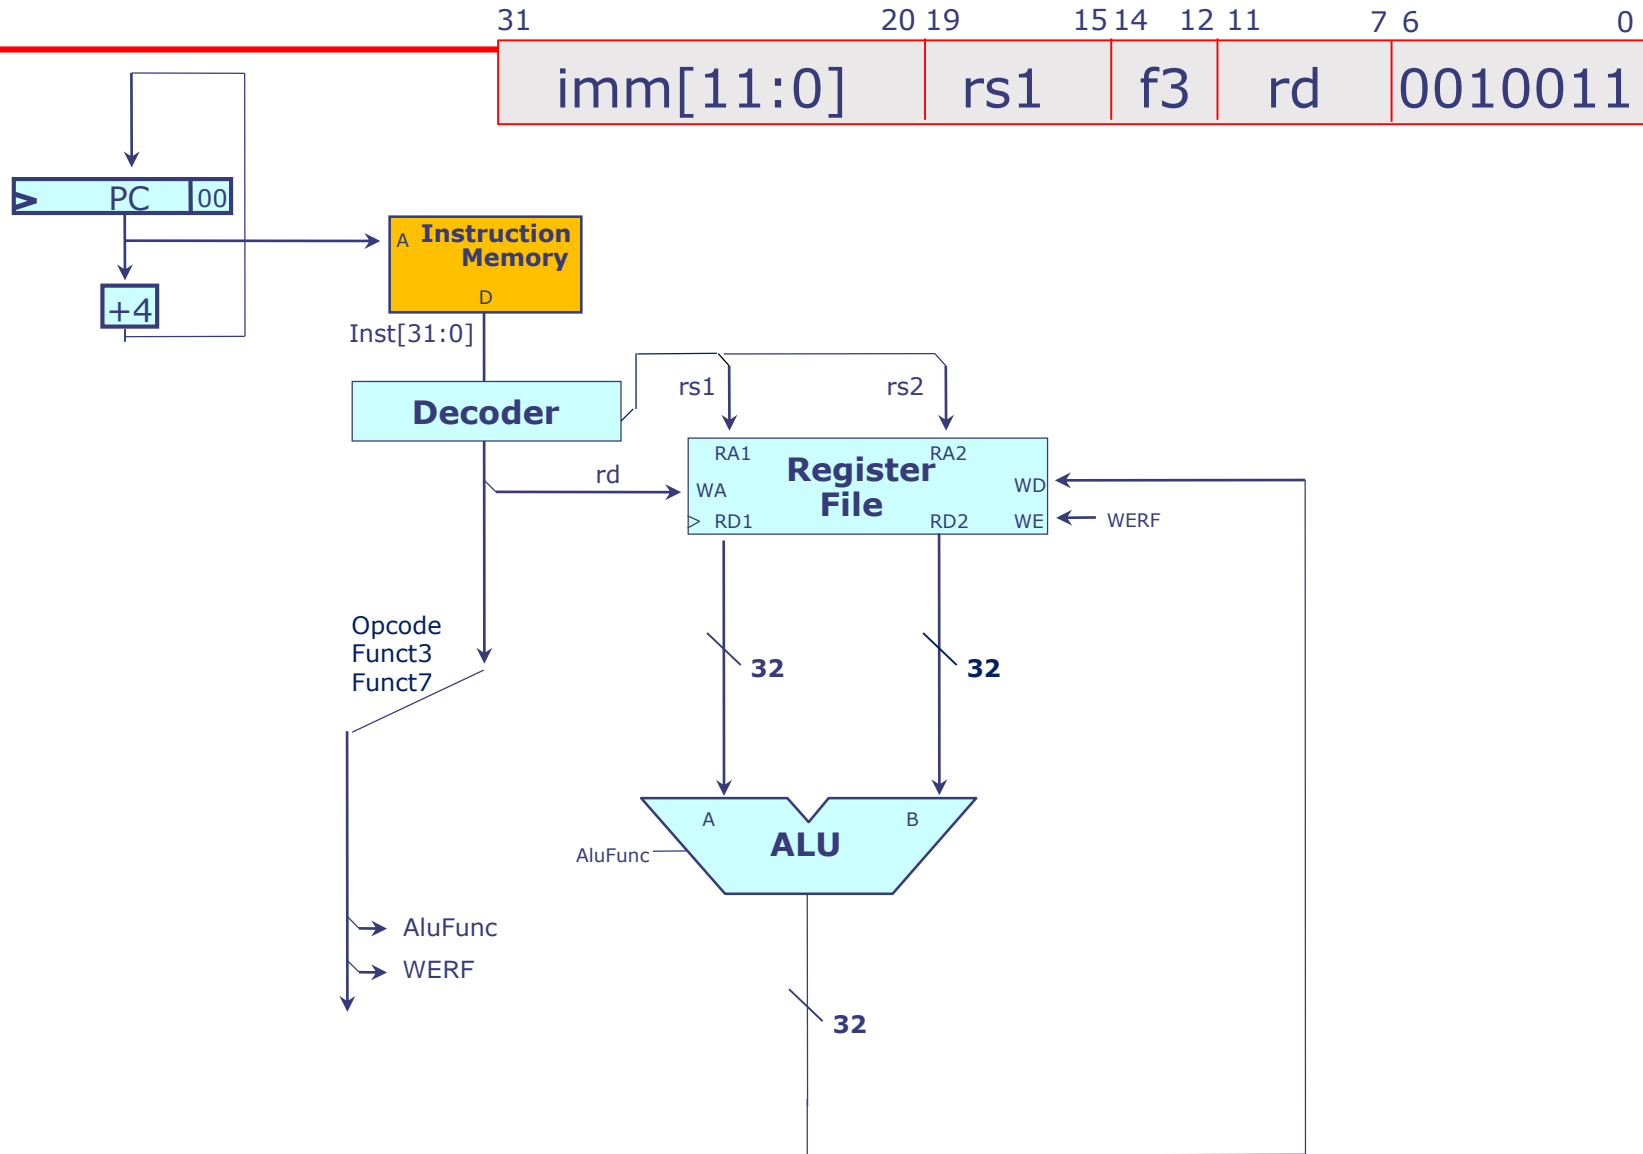
# ALU Instructions

with one Immediate operand

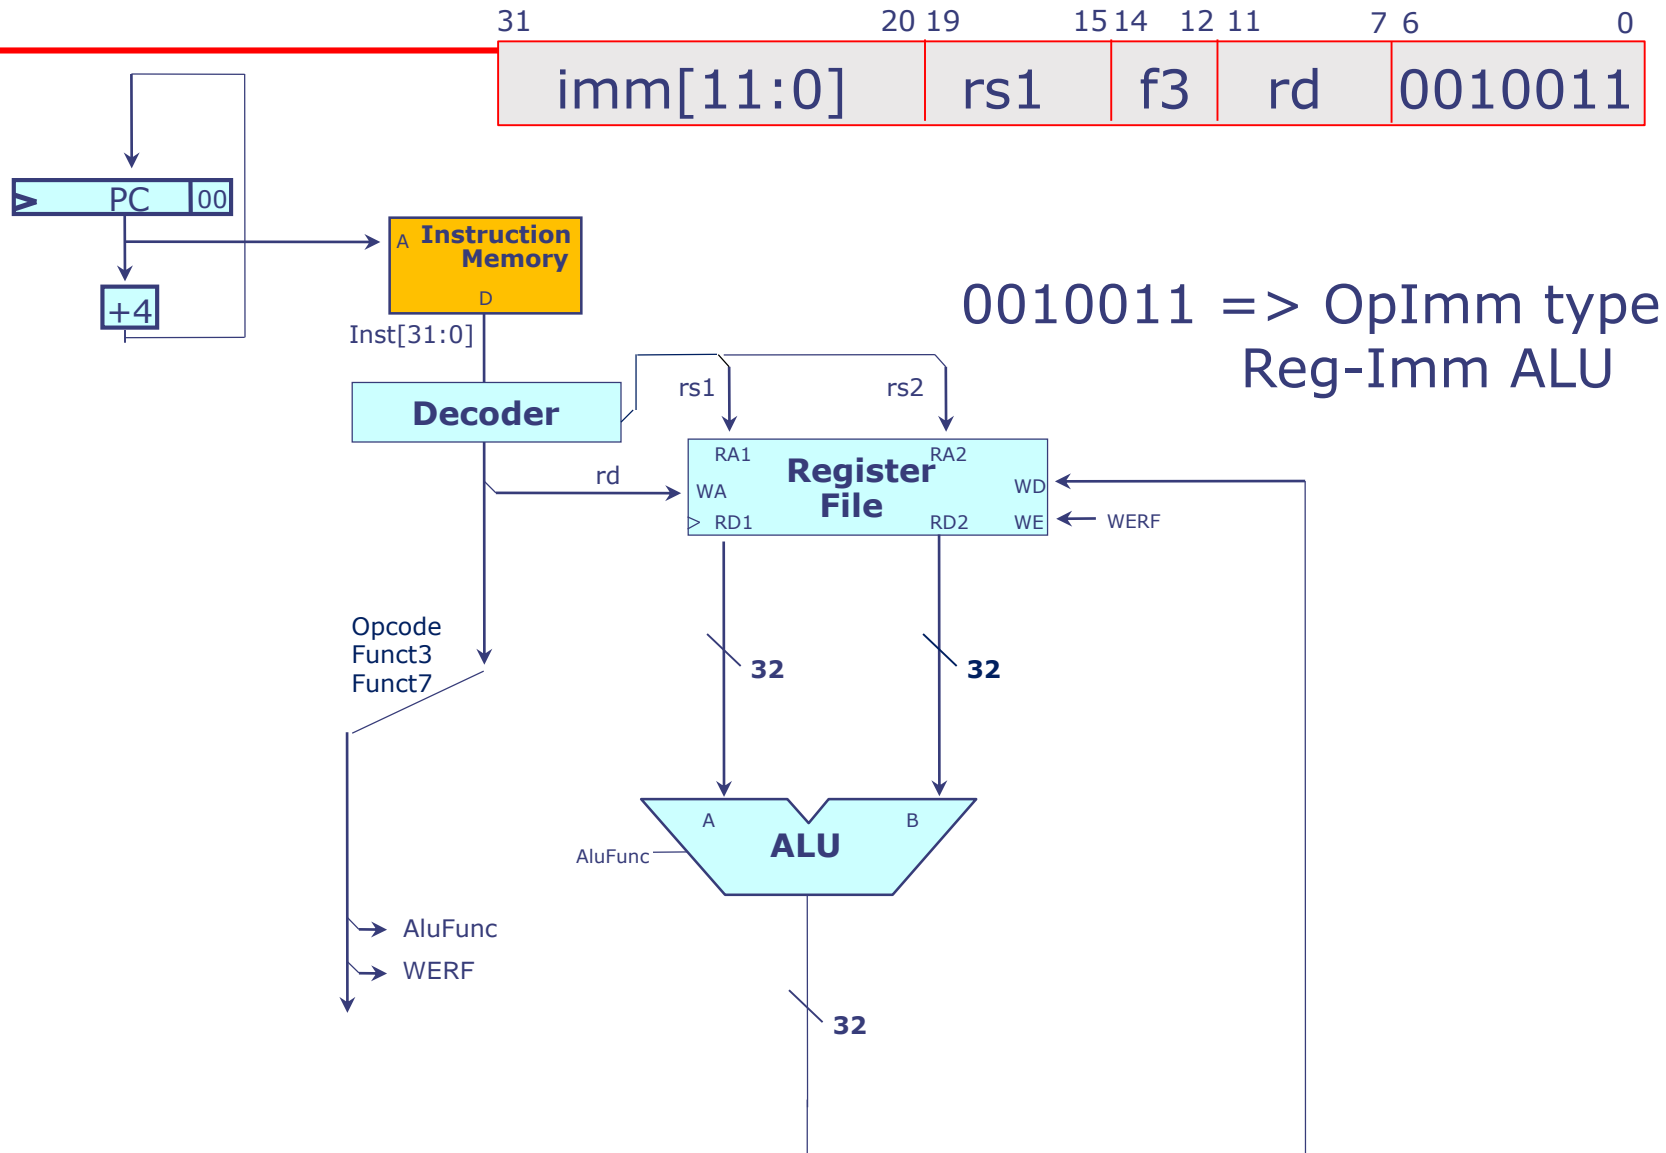
Instruction	Description	Execution
ADDI rd, rs1, const	Add Immediate	$\text{reg}[\text{rd}] \leq \text{reg}[\text{rs1}] + \text{const}$
SLTI rd, rs1, const	Set if < Immediate (Signed)	$\text{reg}[\text{rd}] \leq (\text{reg}[\text{rs1}] <_s \text{const}) ? 1 : 0$
SLTIU rd, rs1, const	Set if < Immediate (Unsigned)	$\text{reg}[\text{rd}] \leq (\text{reg}[\text{rs1}] <_u \text{const}) ? 1 : 0$
XORI rd, rs1, const	Xor Immediate	$\text{reg}[\text{rd}] \leq \text{reg}[\text{rs1}] \wedge \text{const}$
ORI rd, rs1, const	Or Immediate	$\text{reg}[\text{rd}] \leq \text{reg}[\text{rs1}] \mid \text{const}$
ANDI rd, rs1, const	And Immediate	$\text{reg}[\text{rd}] \leq \text{reg}[\text{rs1}] \& \text{const}$
SLLI rd, rs1, shamt	Shift Left Logical Immediate	$\text{reg}[\text{rd}] \leq \text{reg}[\text{rs1}] \ll \text{shamt}$
SRLI rd, rs1, shamt	Shift Right Logical Immediate	$\text{reg}[\text{rd}] \leq \text{reg}[\text{rs1}] \gg_u \text{shamt}$
SRAI rd, rs1, shamt	Shift Right Arithmetic Immediate	$\text{reg}[\text{rd}] \leq \text{reg}[\text{rs1}] \gg_s \text{shamt}$

These instructions are grouped in a category called OPIMM with fields (AluFunc, rd, rs1, immI32)

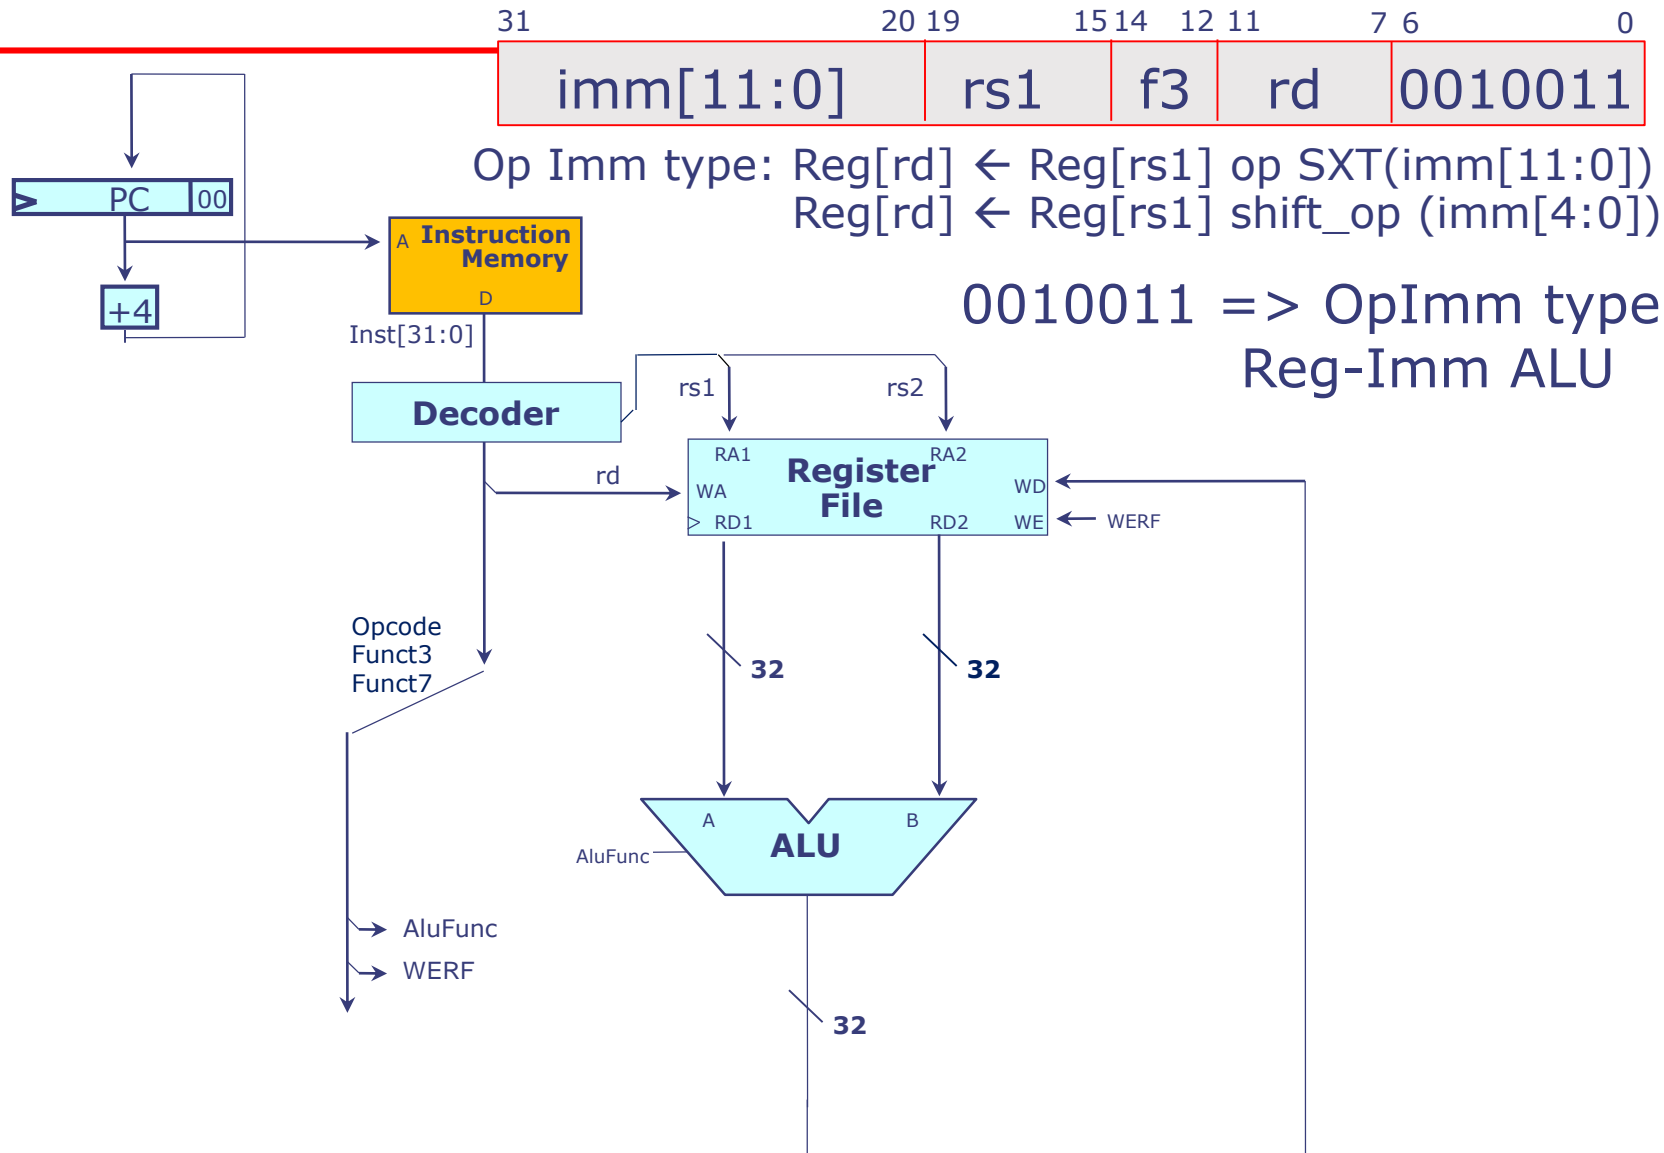
# Register-Immediate ALU Datapath



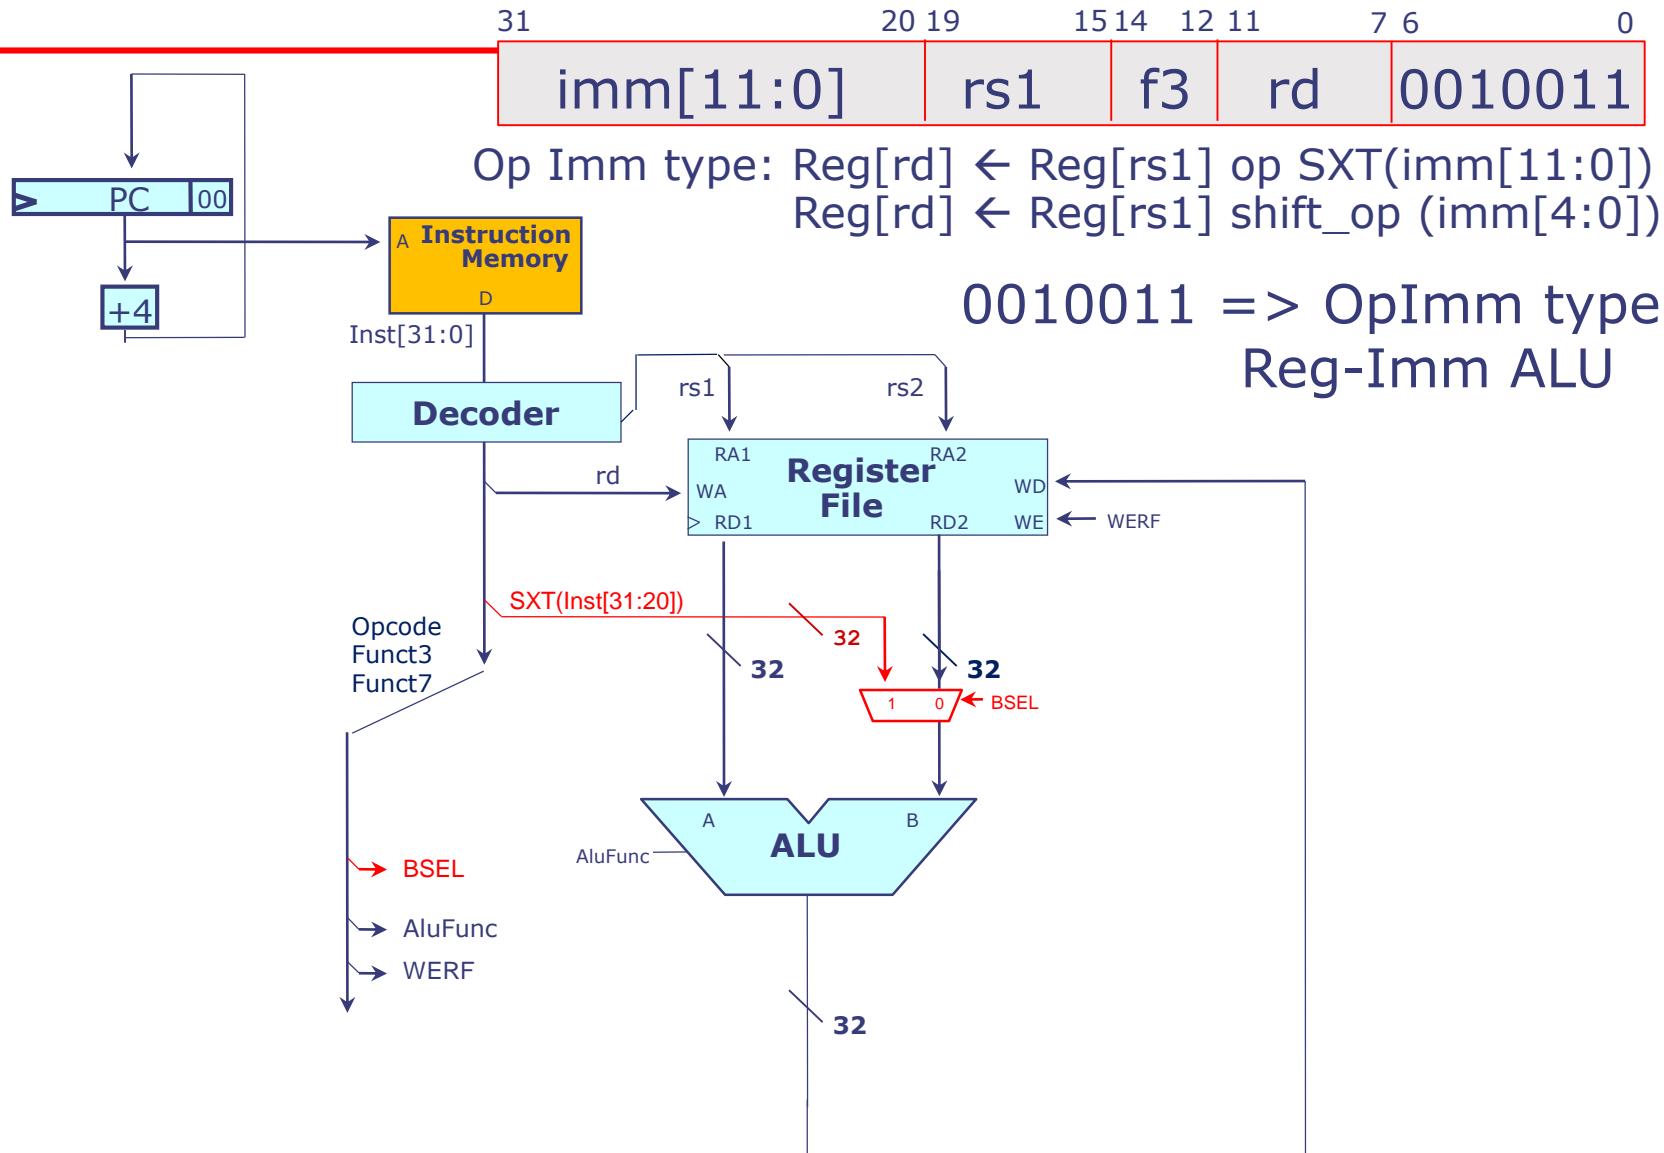
# Register-Immediate ALU Datapath



# Register-Immediate ALU Datapath

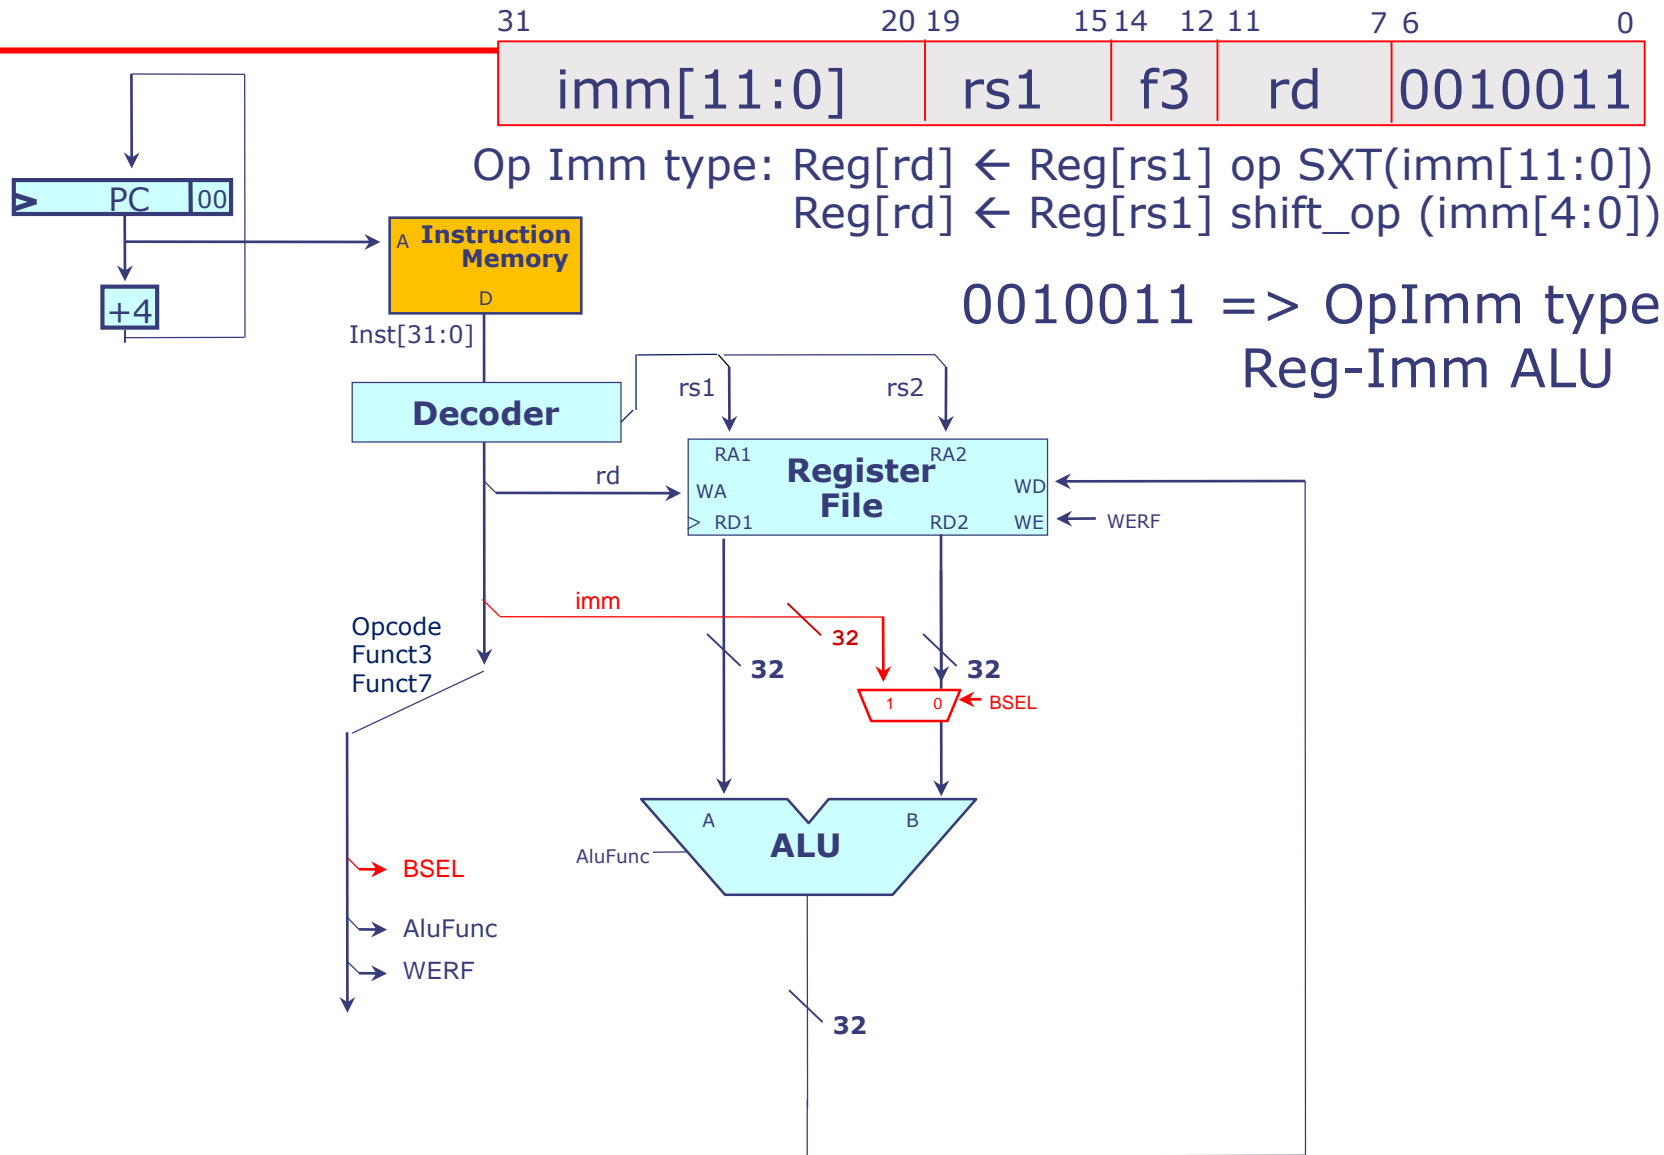


# Register-Immediate ALU Datapath

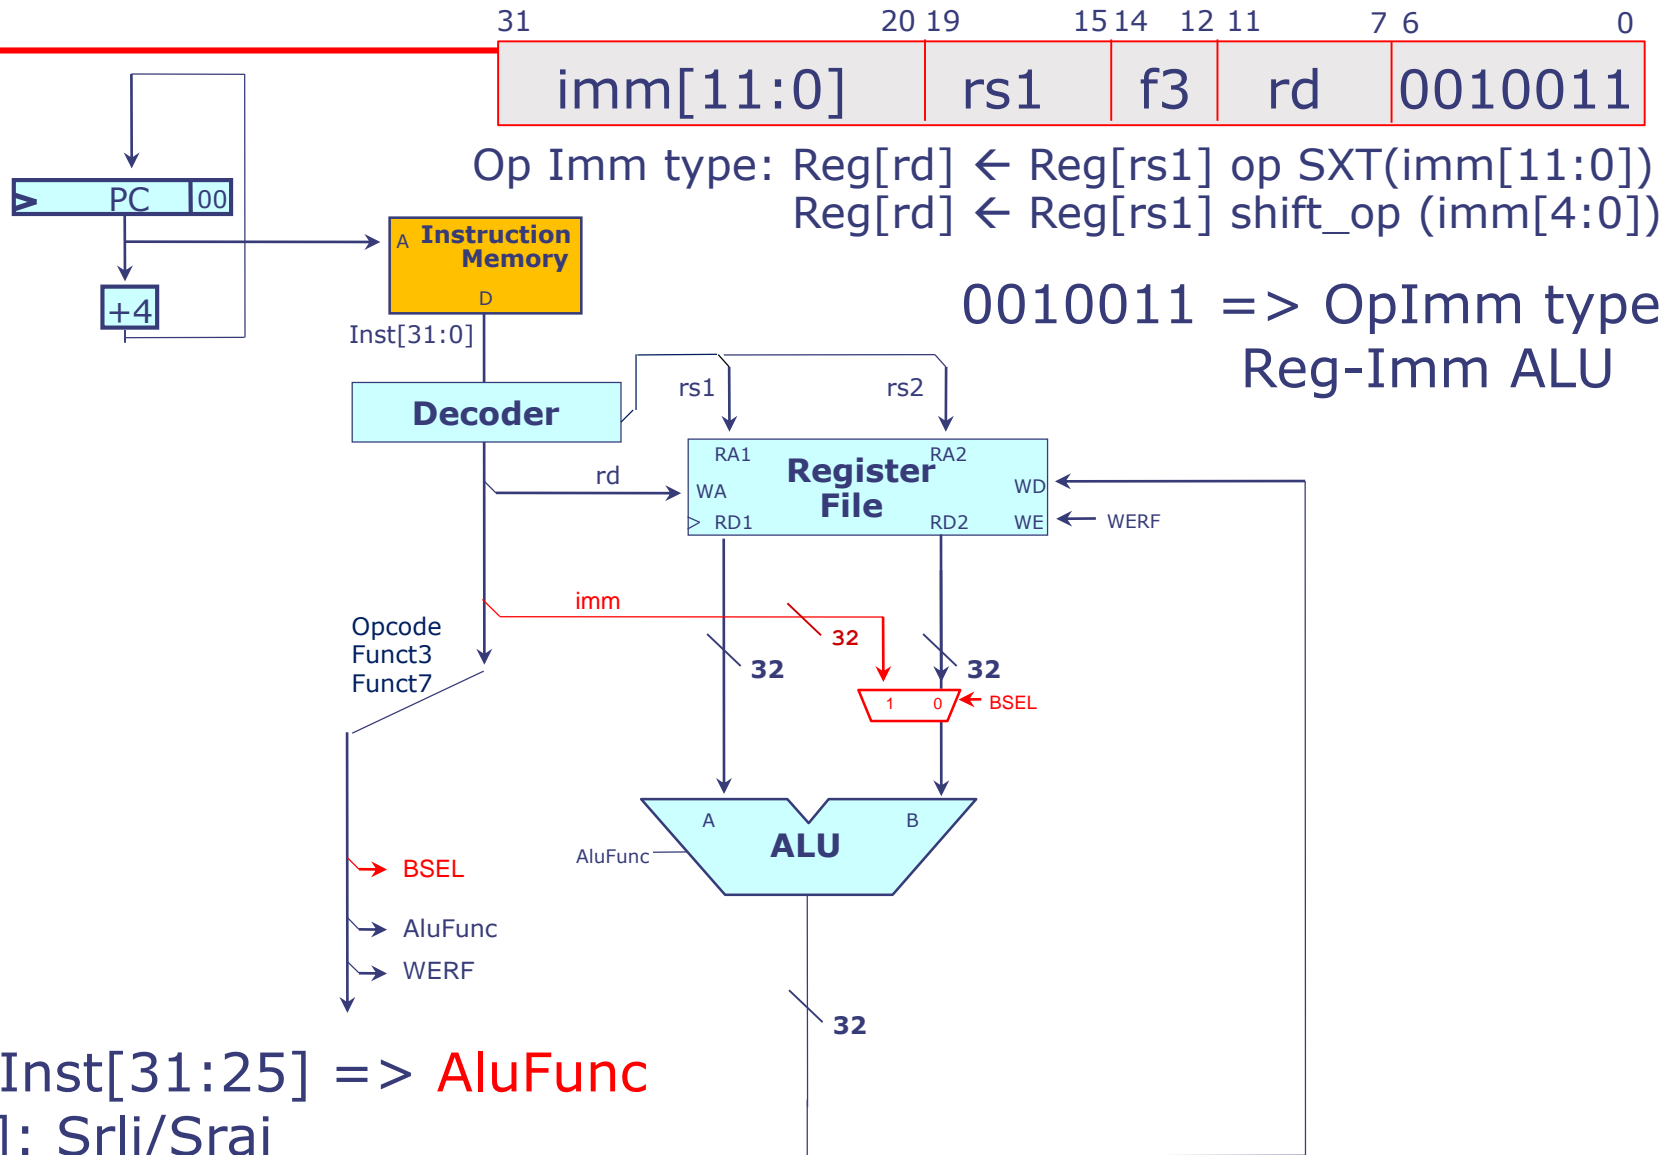




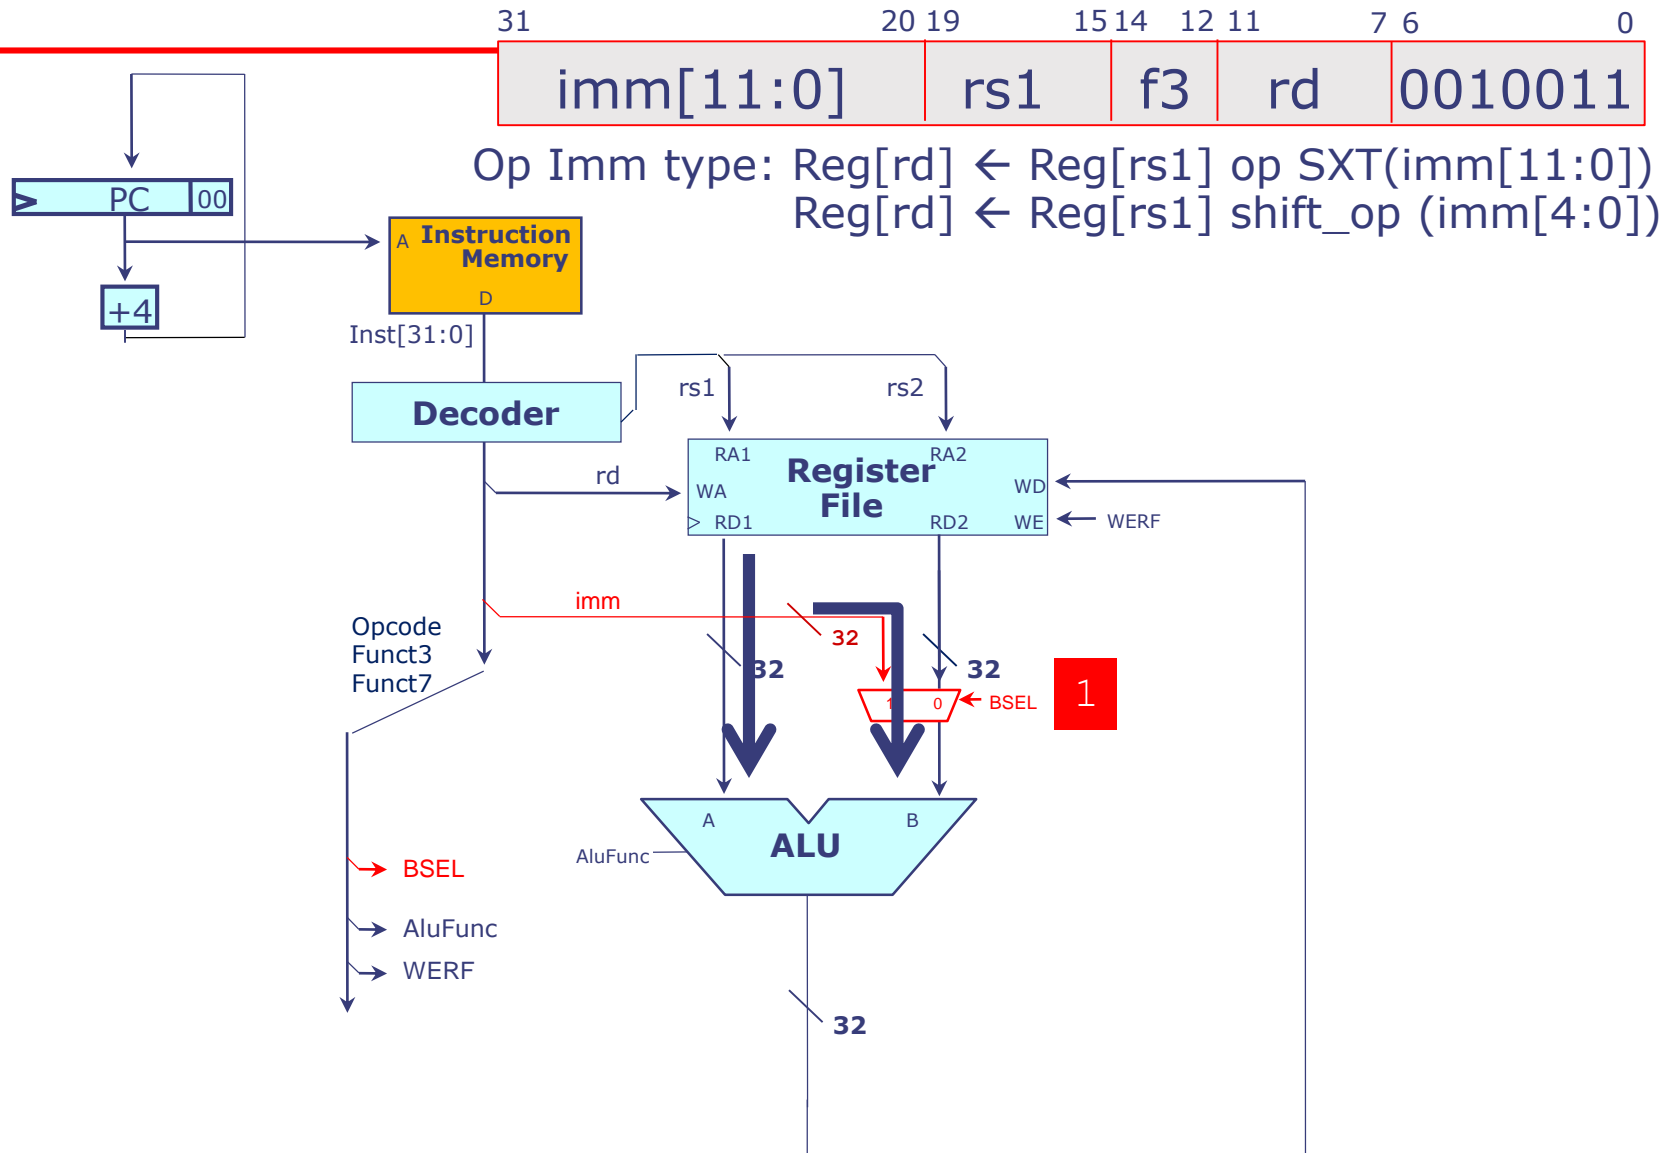
# Register-Immediate ALU Datapath



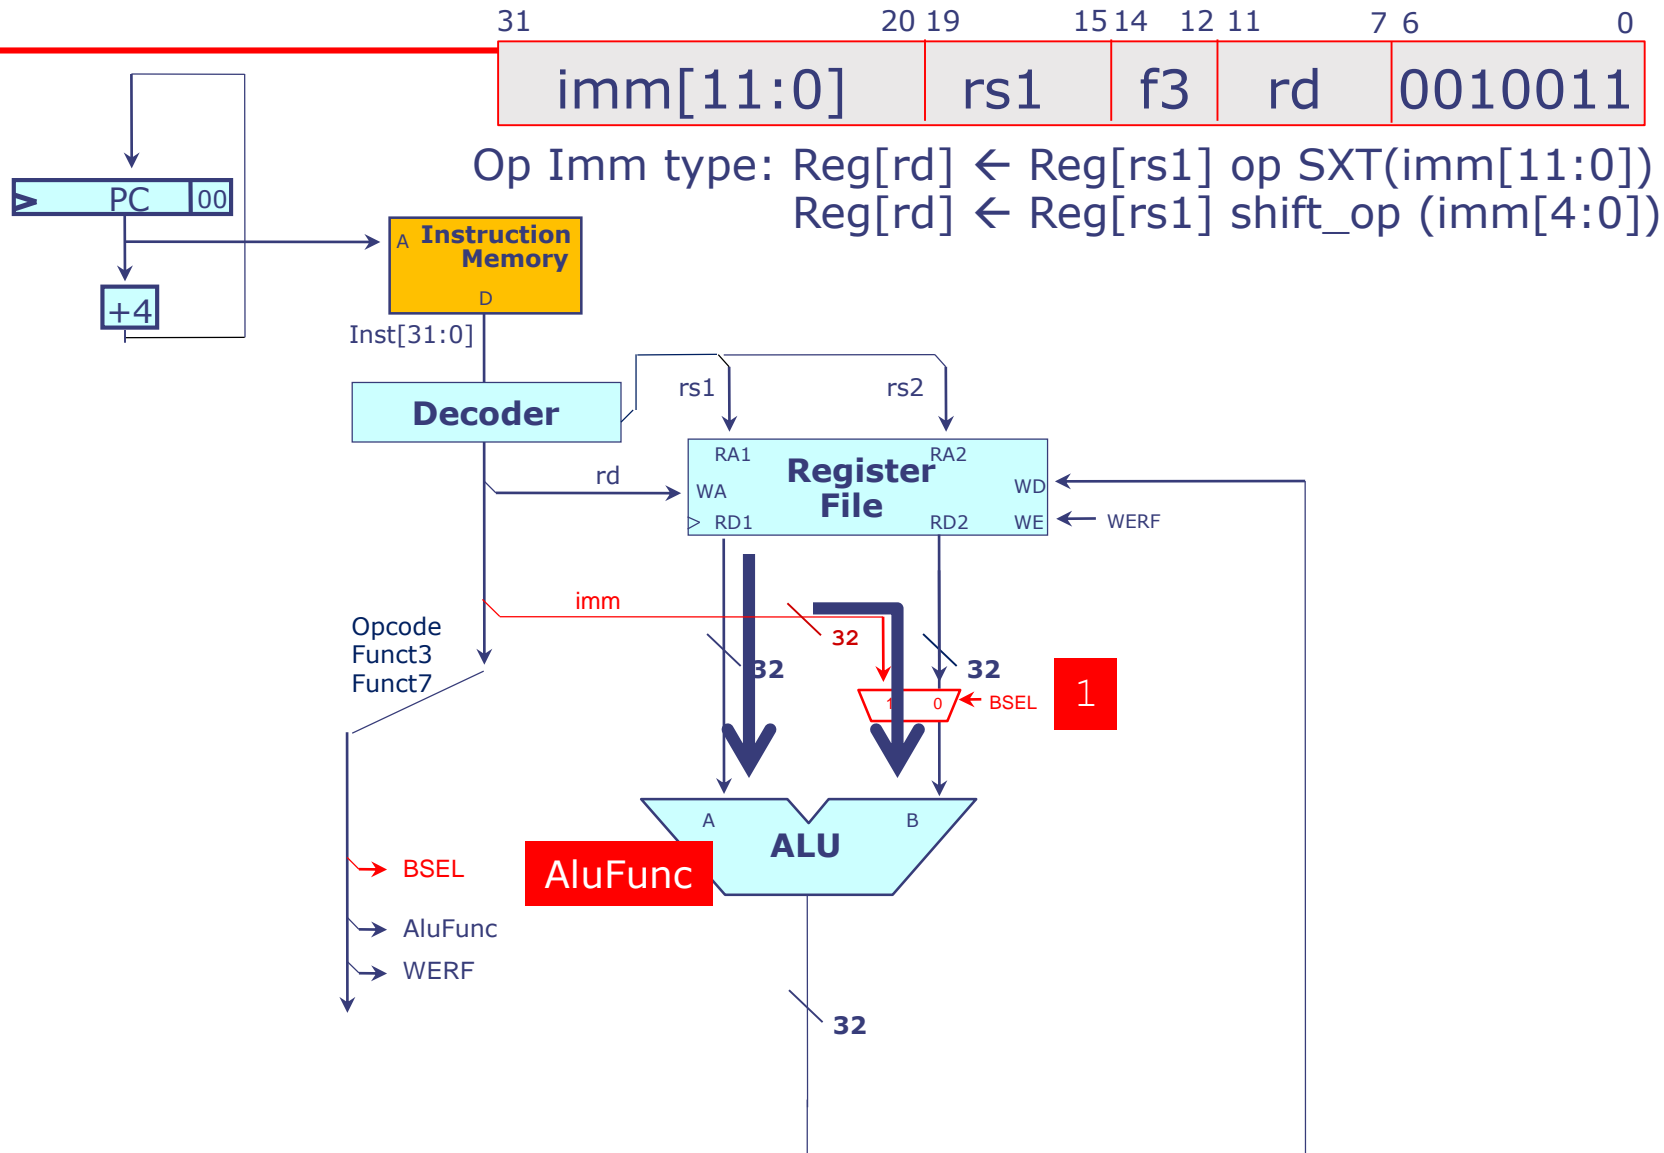
# Register-Immediate ALU Datapath



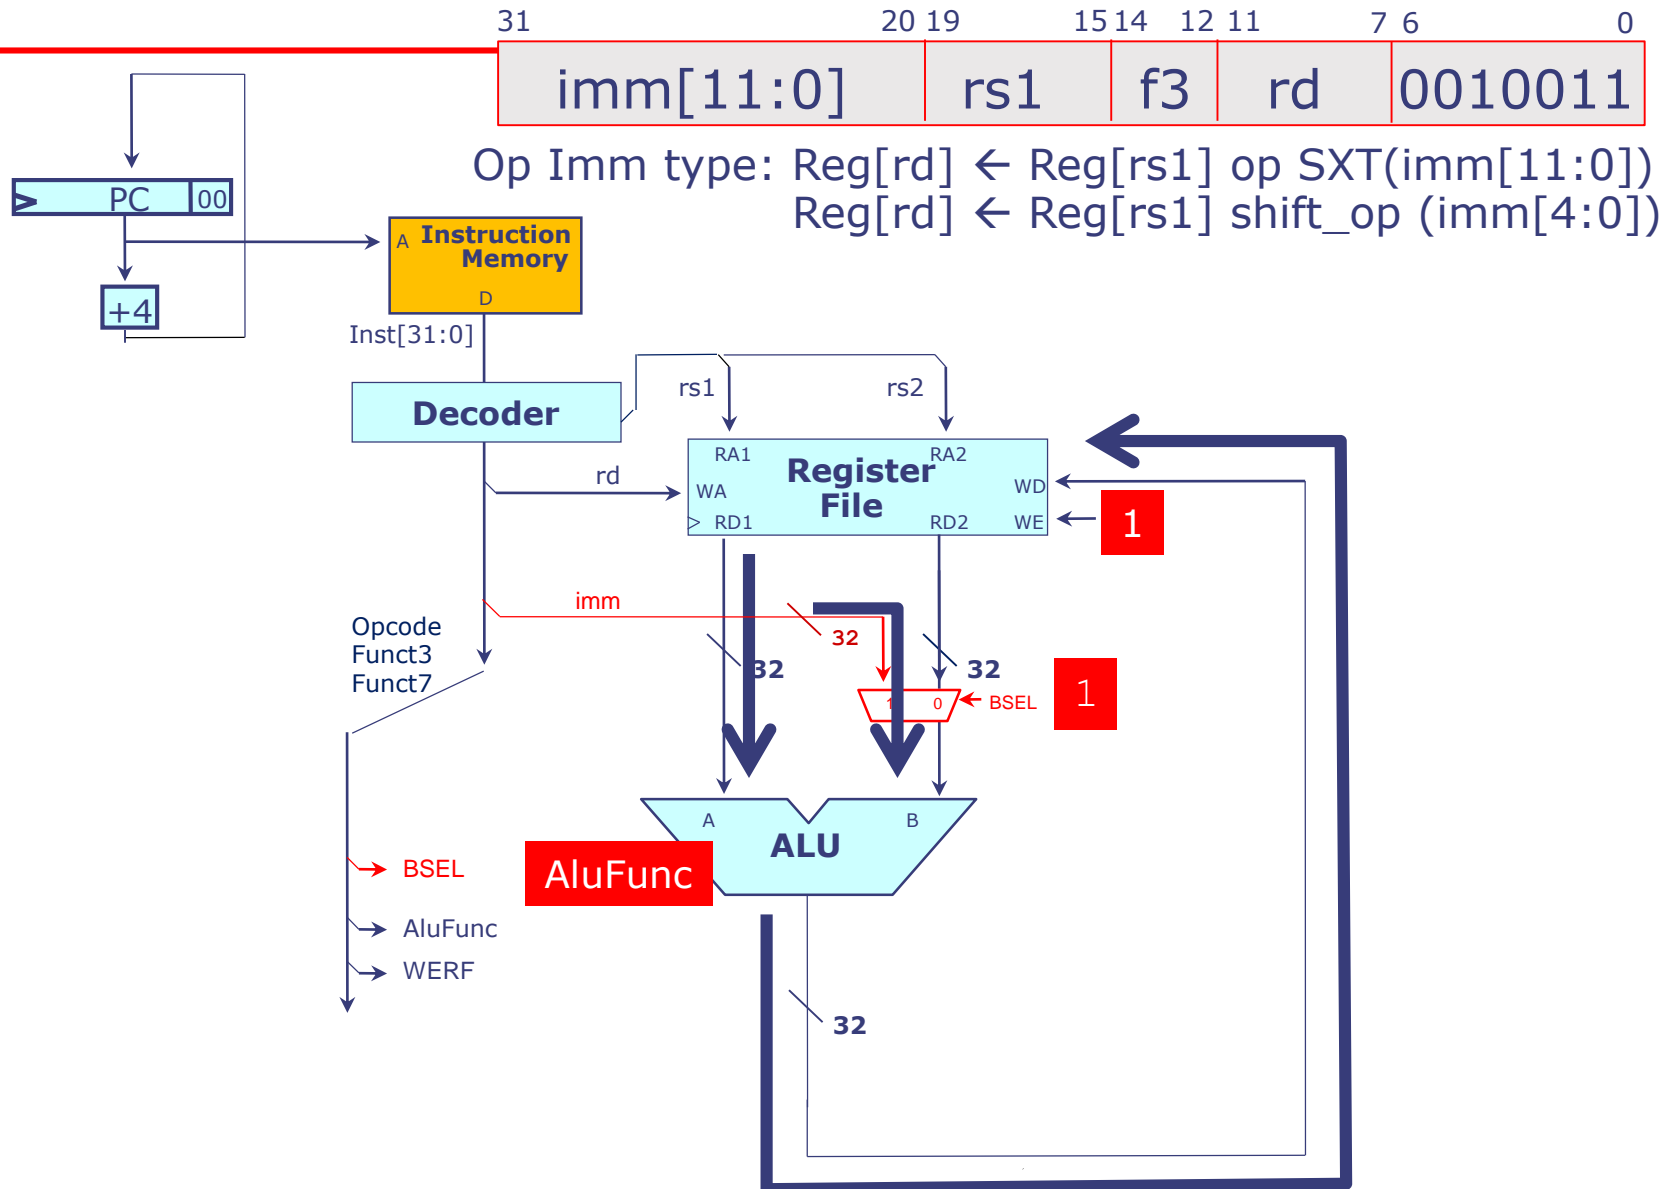
# Register-Immediate ALU Datapath



# Register-Immediate ALU Datapath



# Register-Immediate ALU Datapath



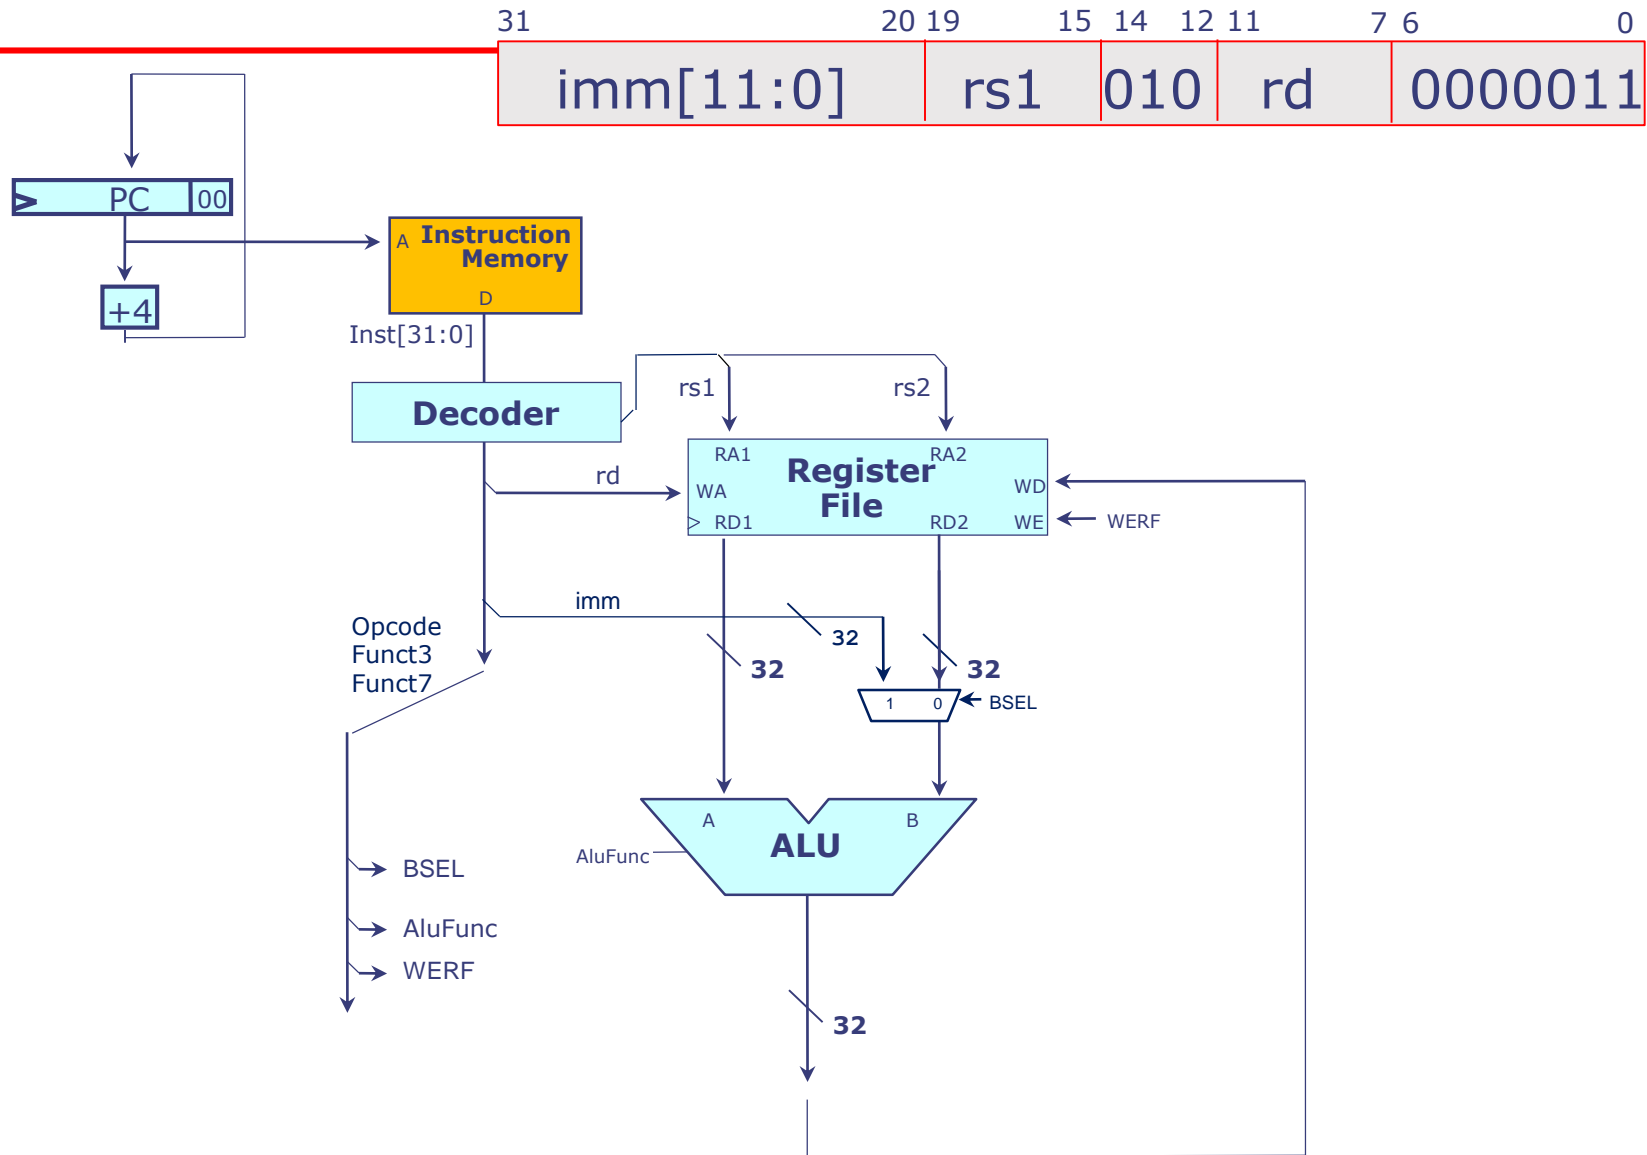
# Load and Store Instructions

---

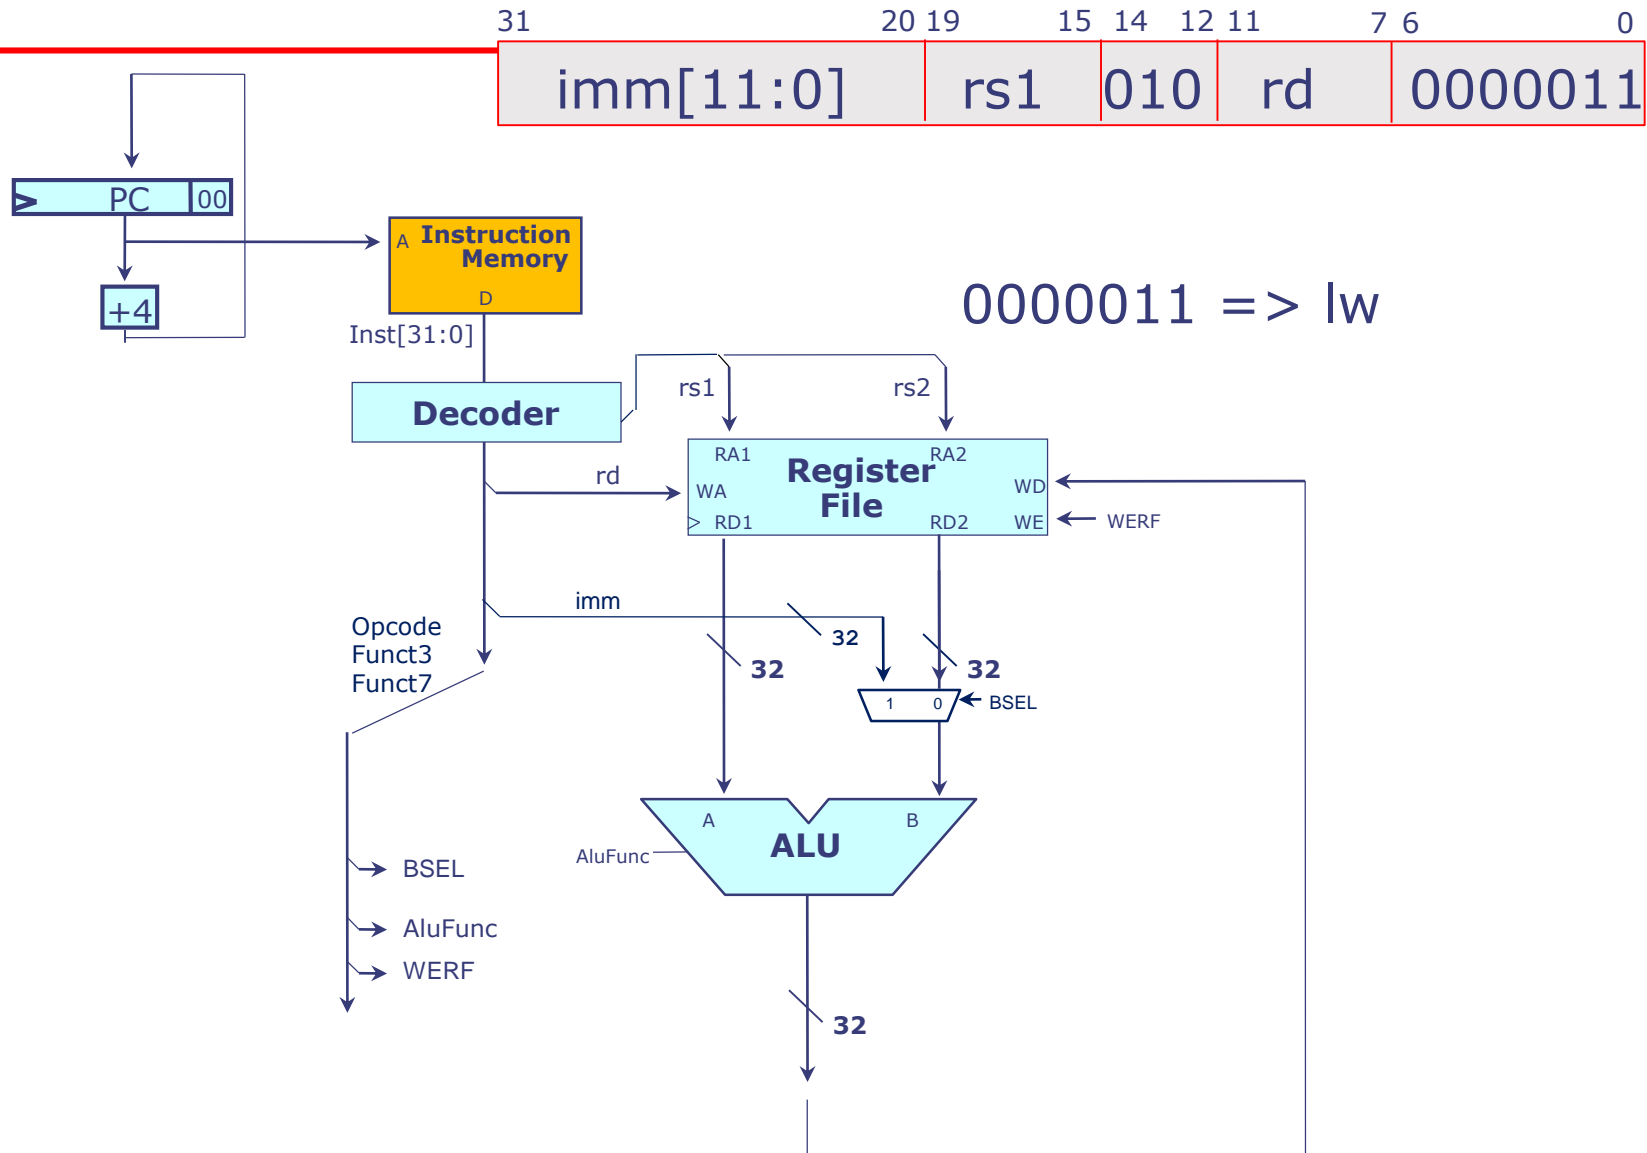
Instruction	Description	Execution
LW rd, offset(rs1)	Load Word	$\text{reg}[\text{rd}] \leftarrow \text{mem}[\text{reg}[\text{rs1}] + \text{offset}]$
SW rs2, offset(rs1)	Store Word	$\text{mem}[\text{reg}[\text{rs1}] + \text{offset}] \leftarrow \text{reg}[\text{rs2}]$

LW and SW need to access memory for execution, so they need to compute an effective memory address

# Load Instruction

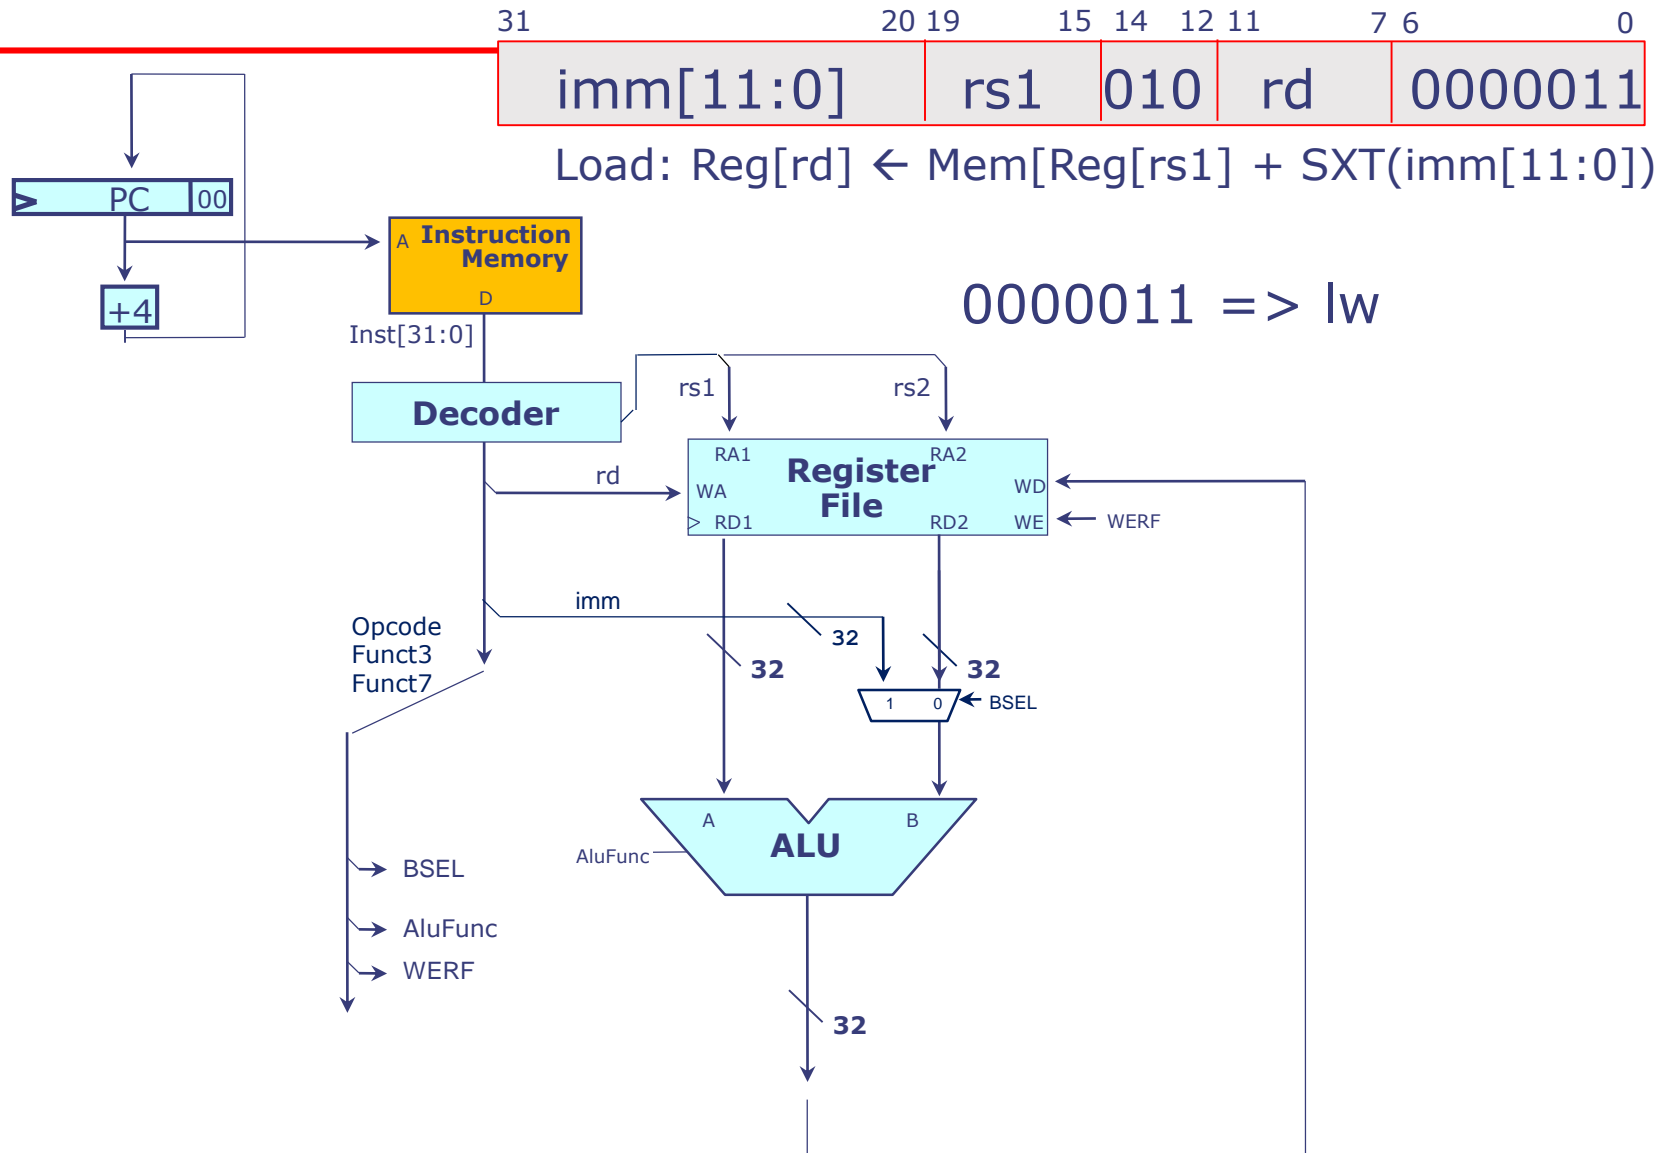


# Load Instruction

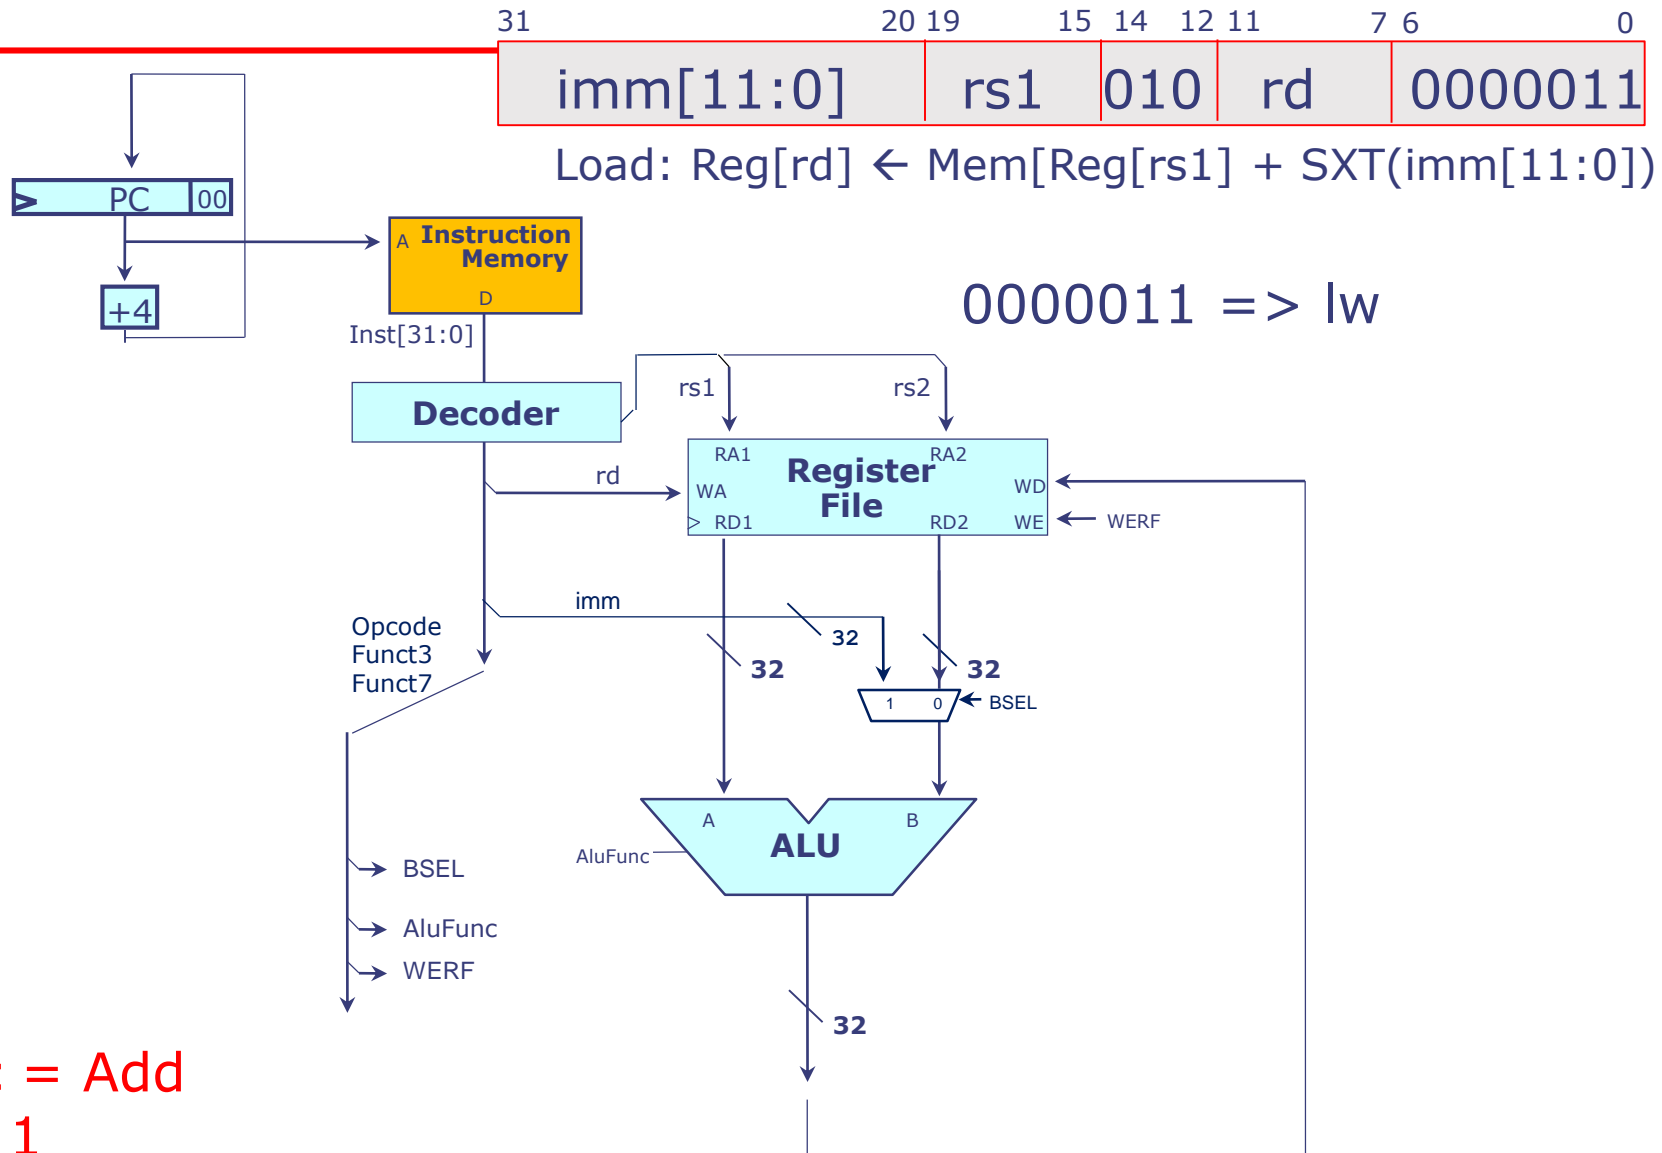




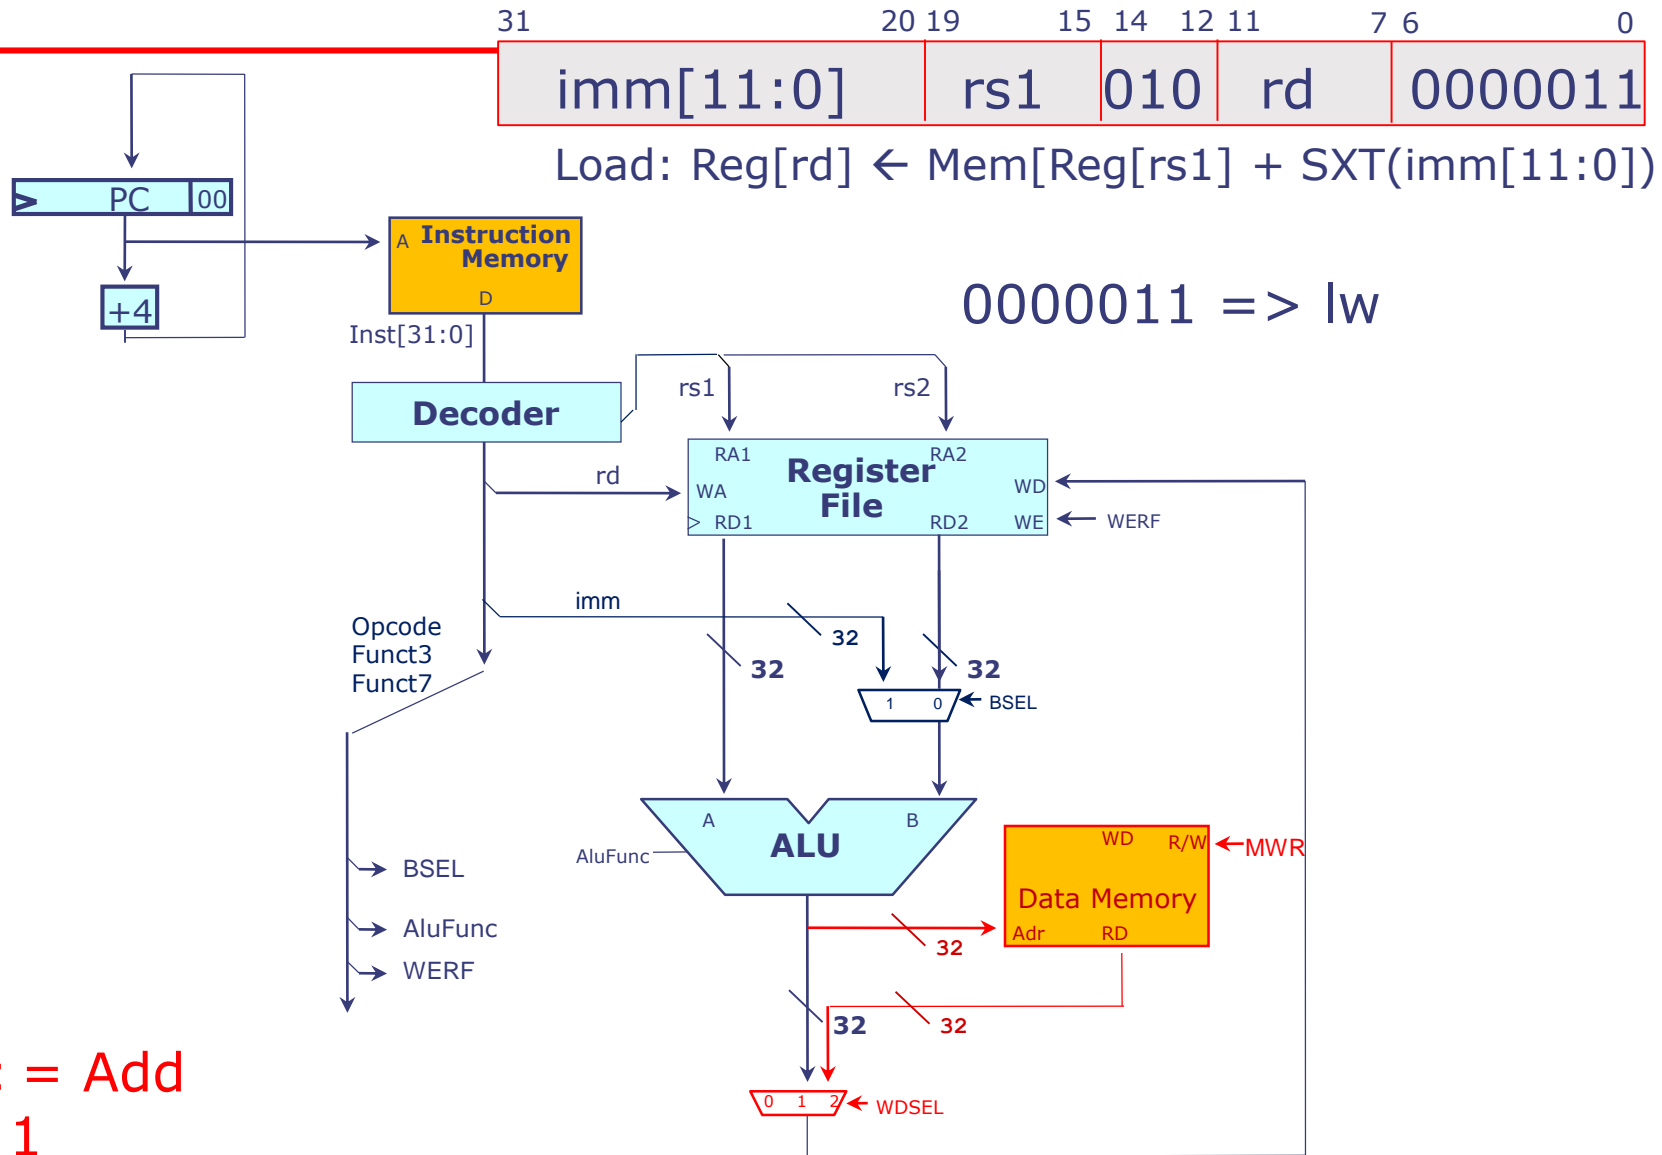
# Load Instruction



# Load Instruction

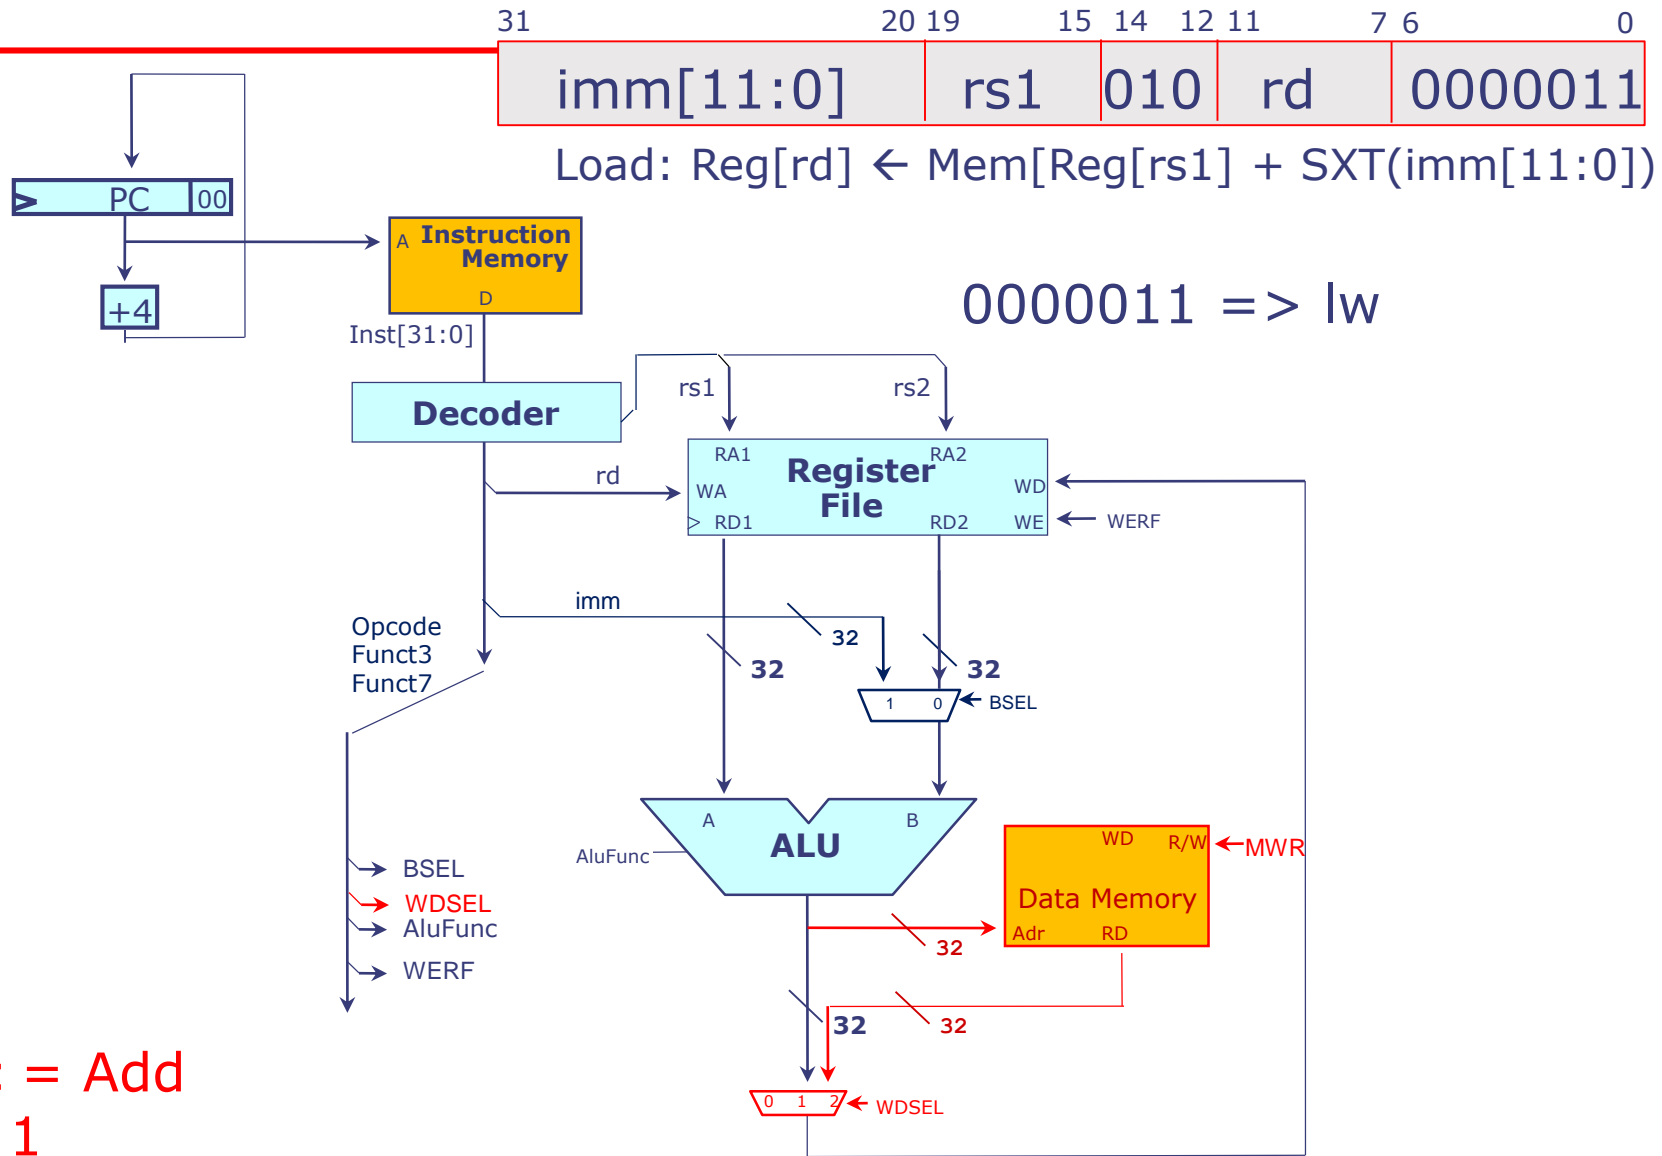


# Load Instruction



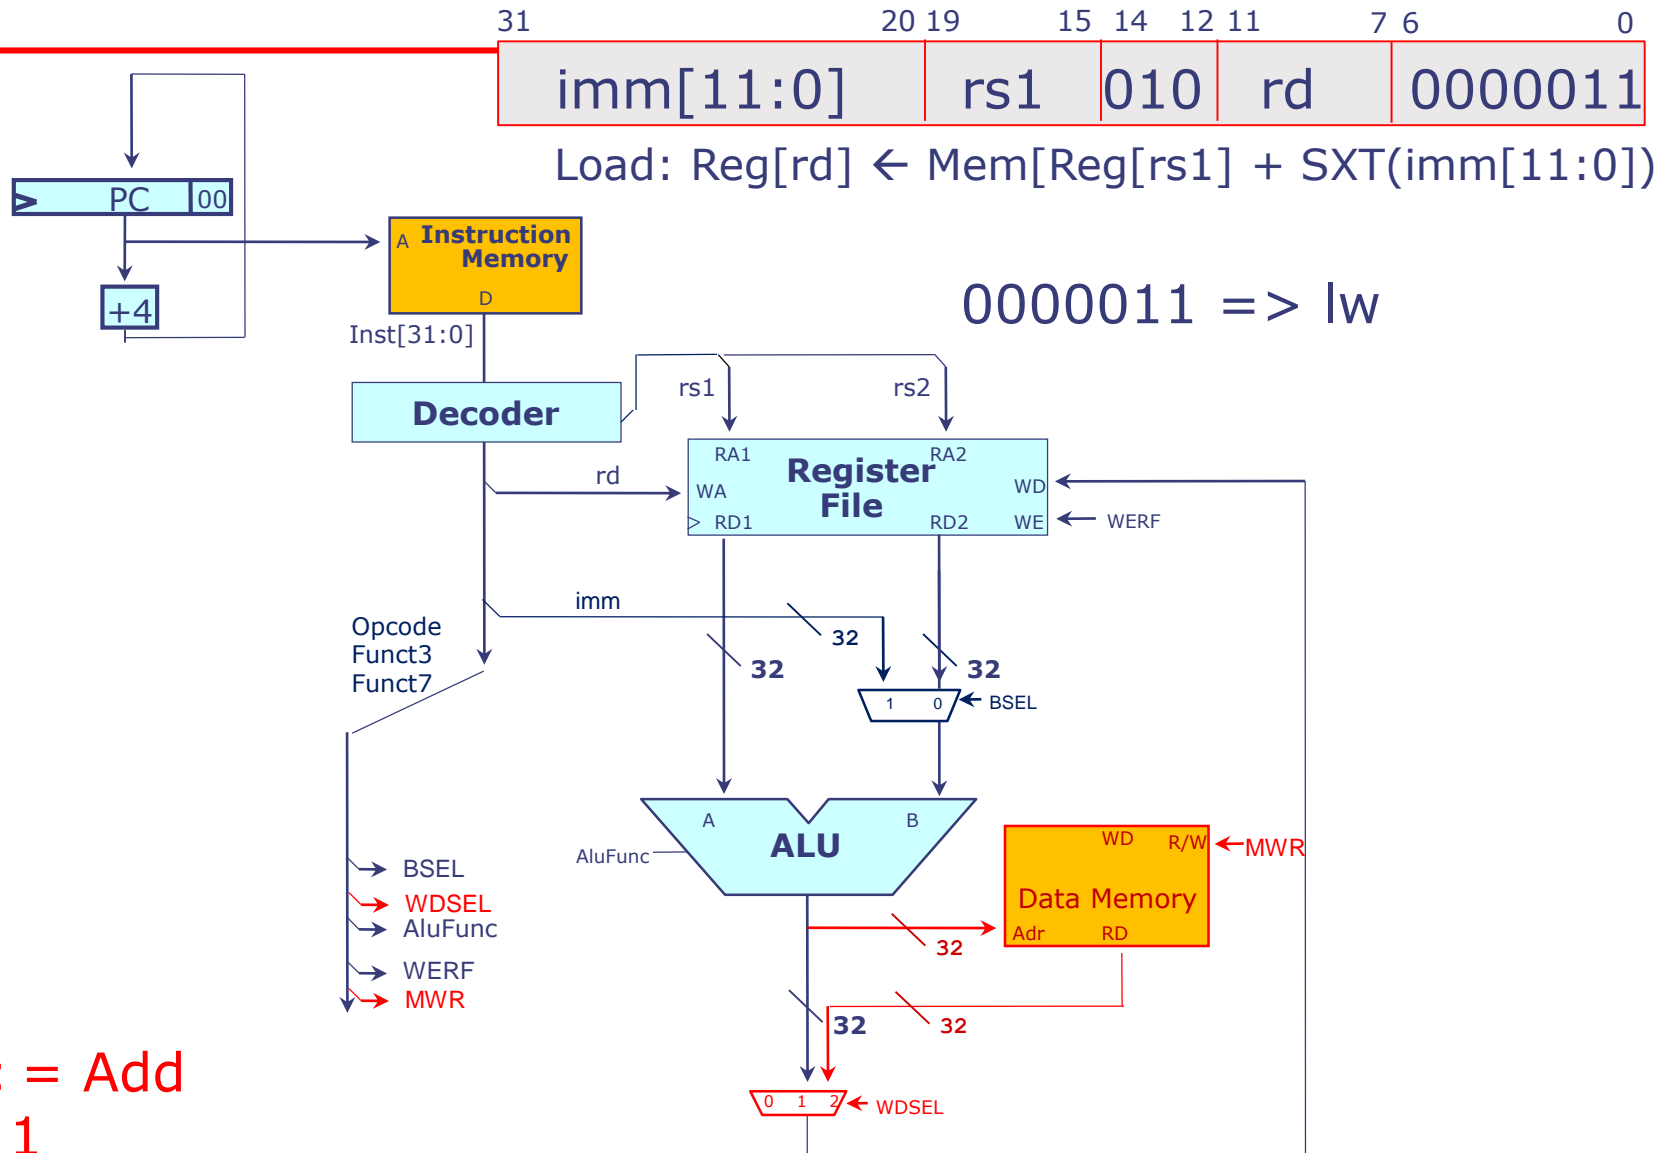
AluFunc = Add  
BSEL = 1

# Load Instruction

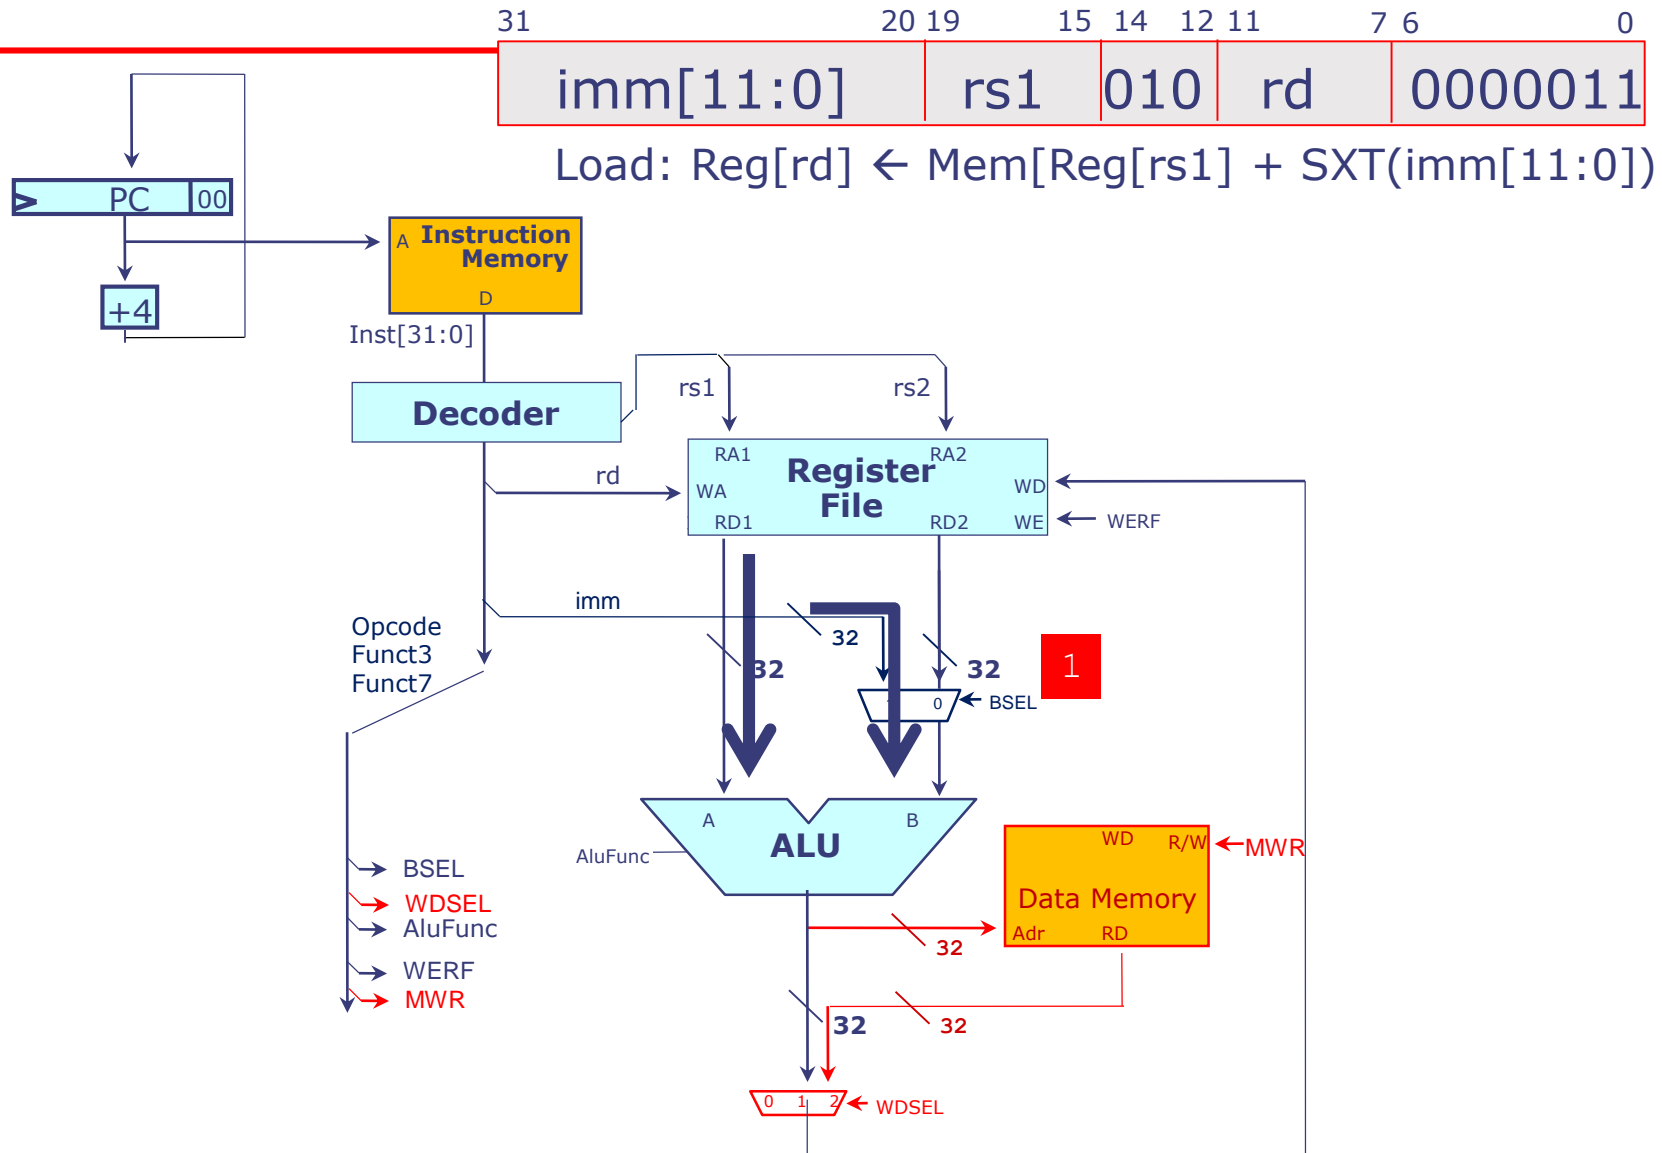


AluFunc = Add  
BSEL = 1

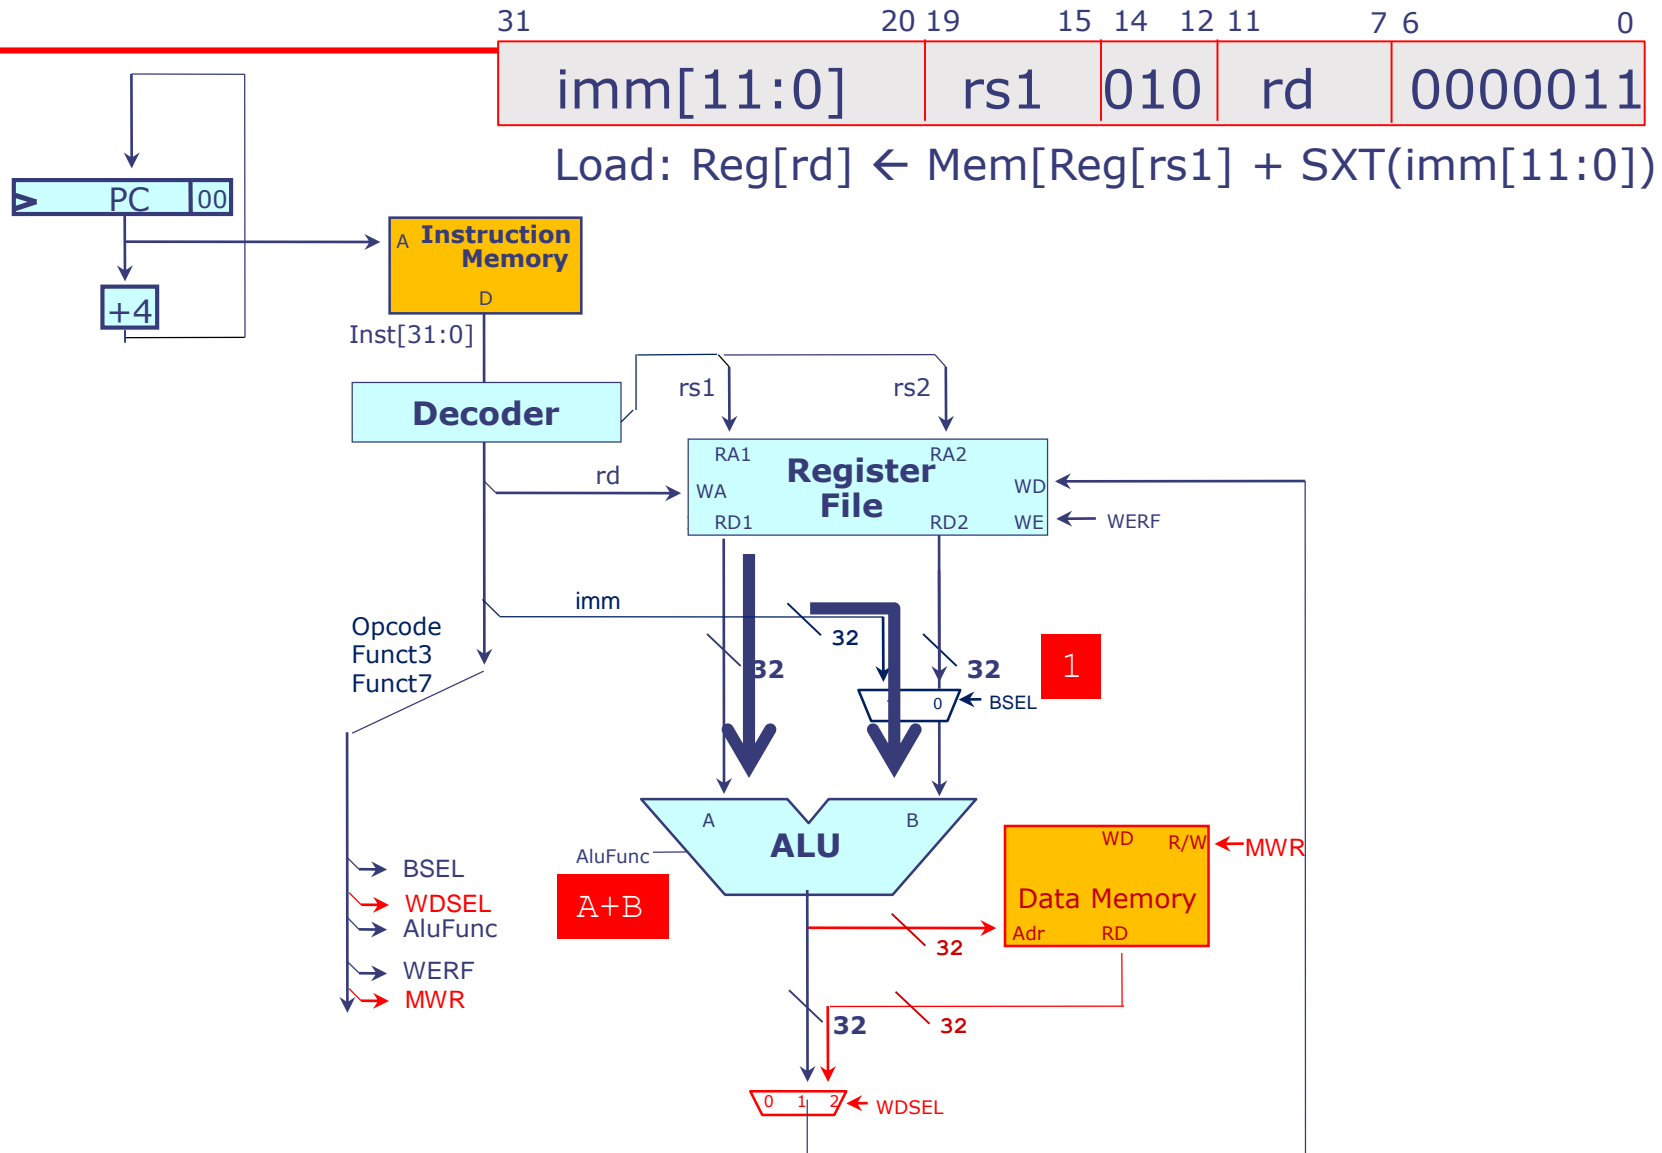
# Load Instruction



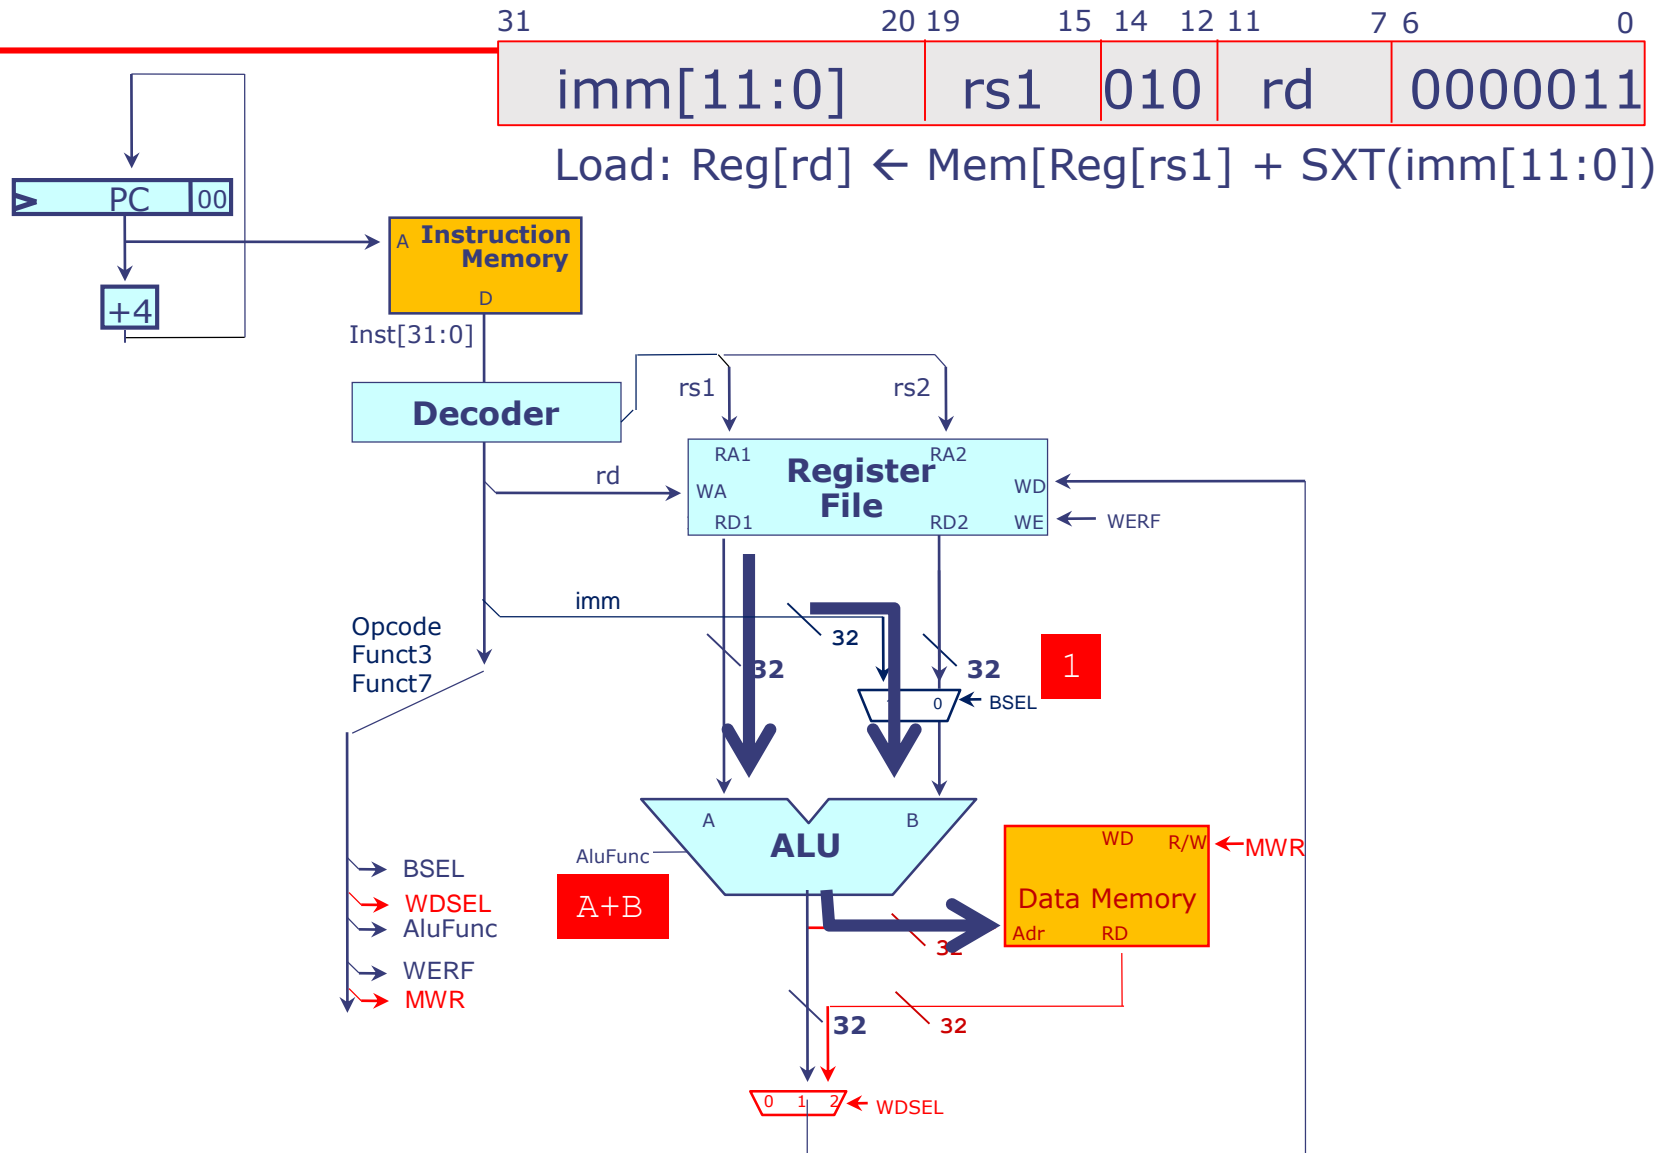
# Load Instruction



# Load Instruction

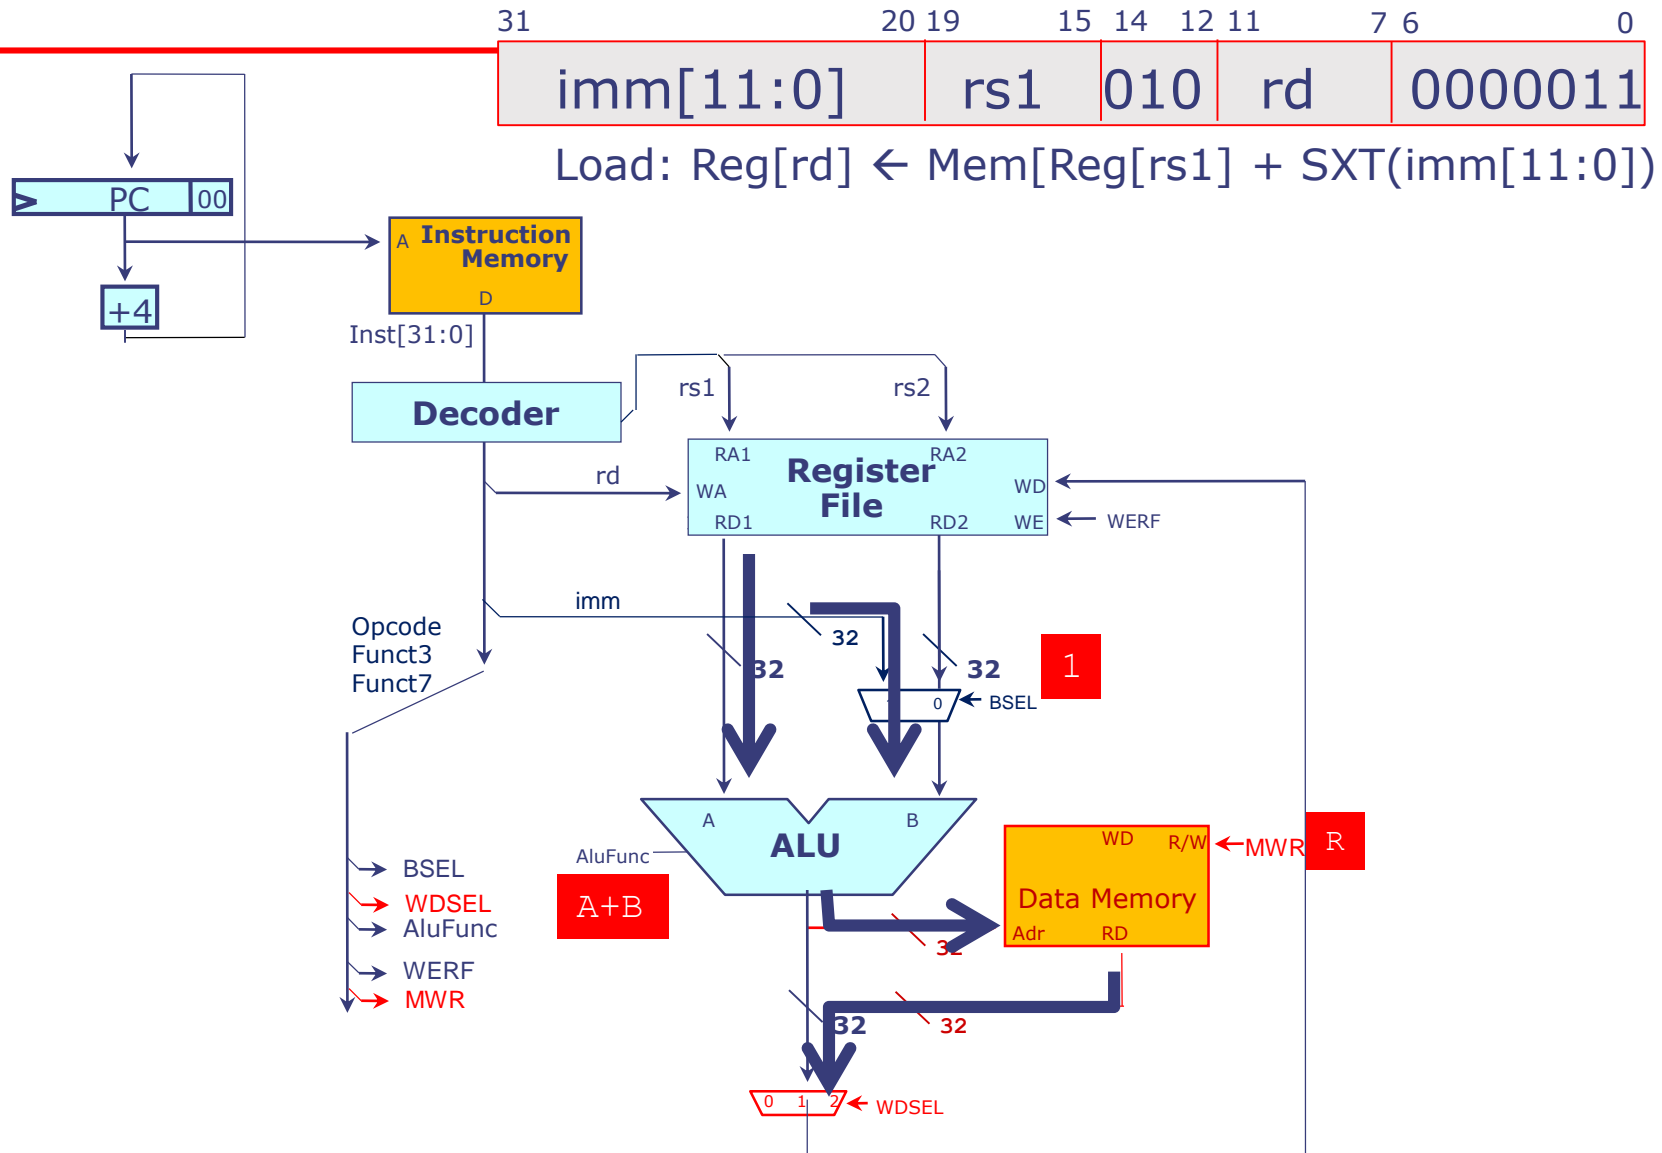


# Load Instruction

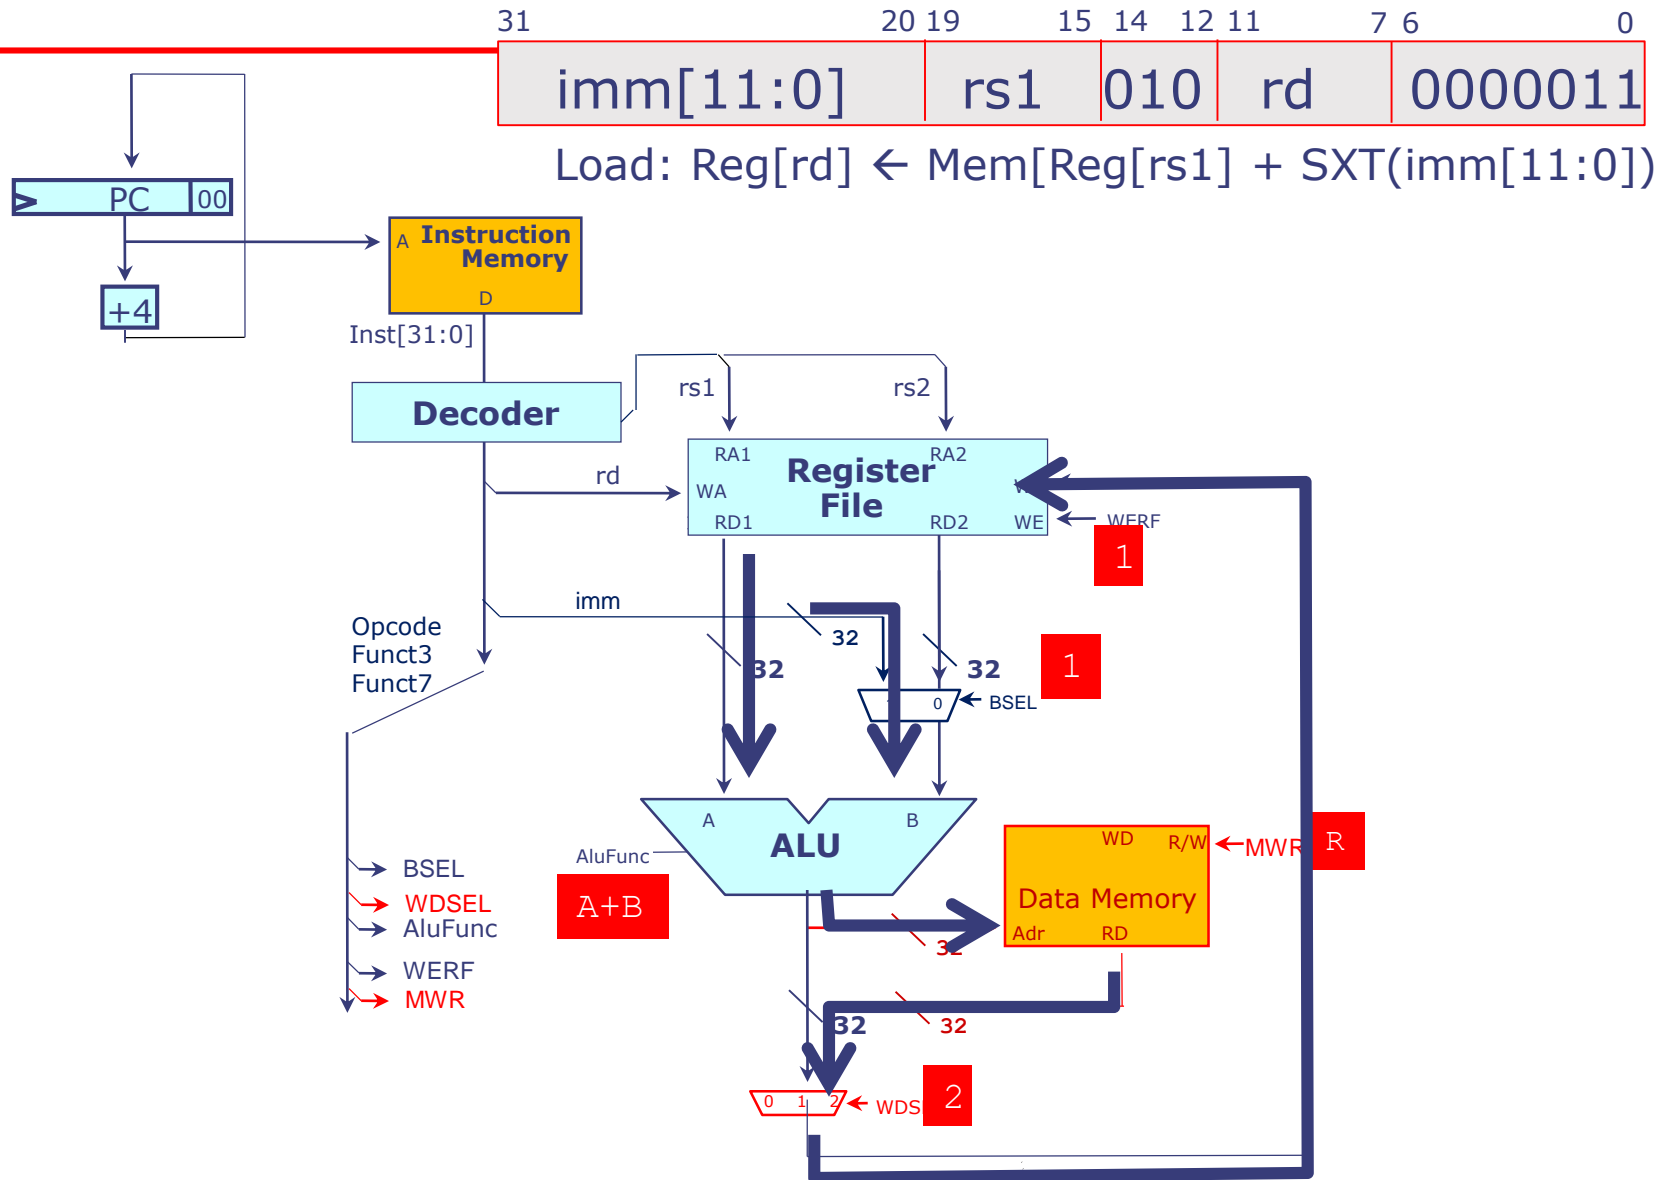




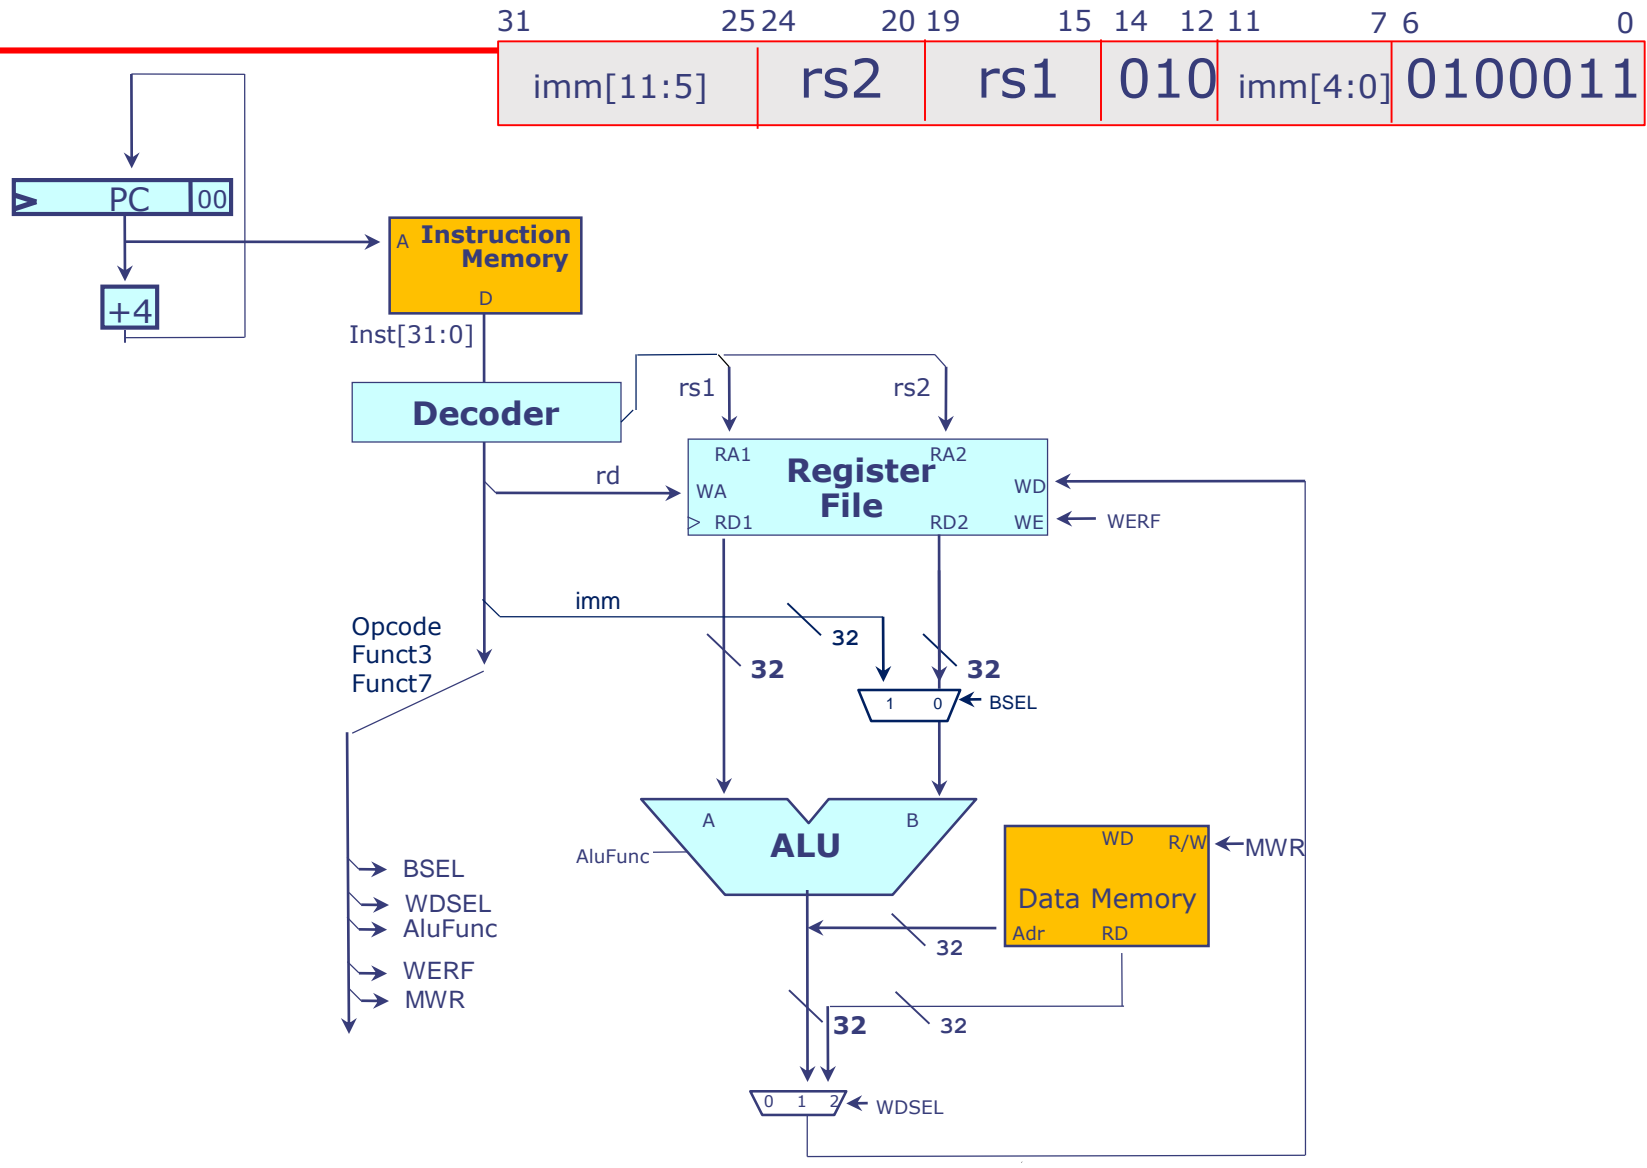
# Load Instruction



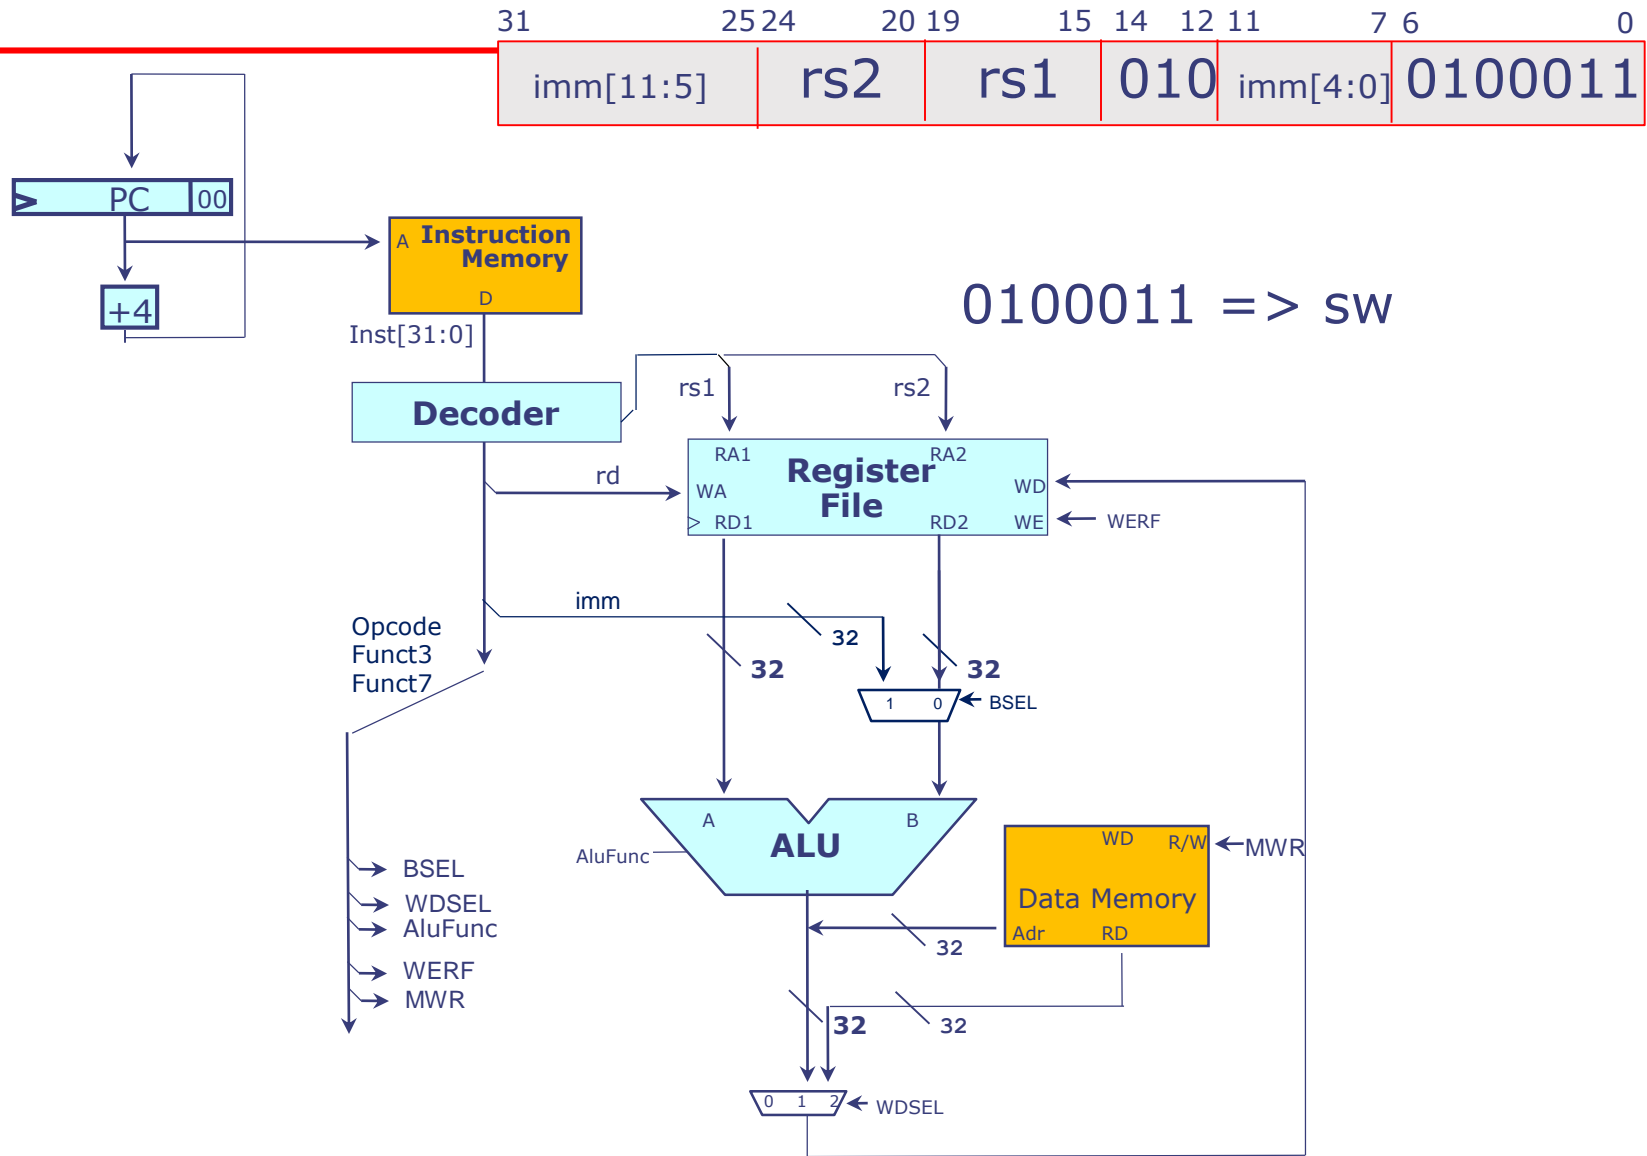
# Load Instruction



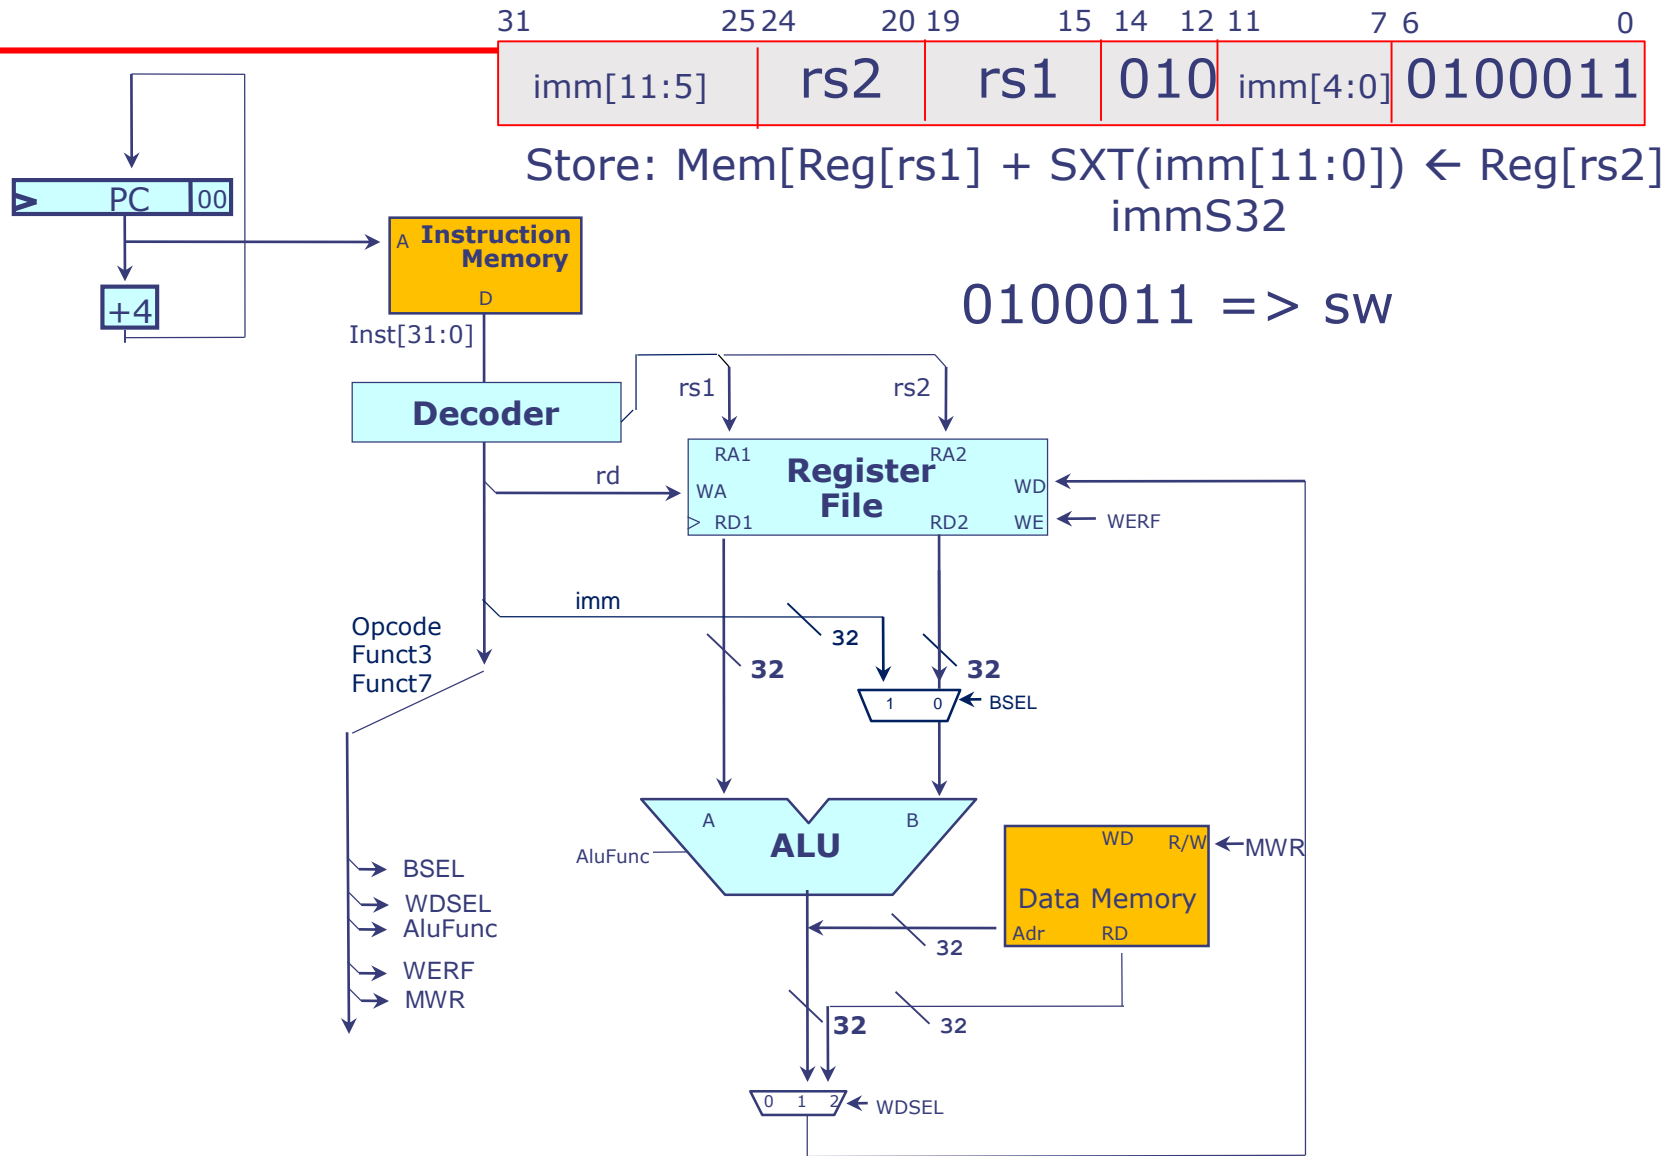
# Store Instruction



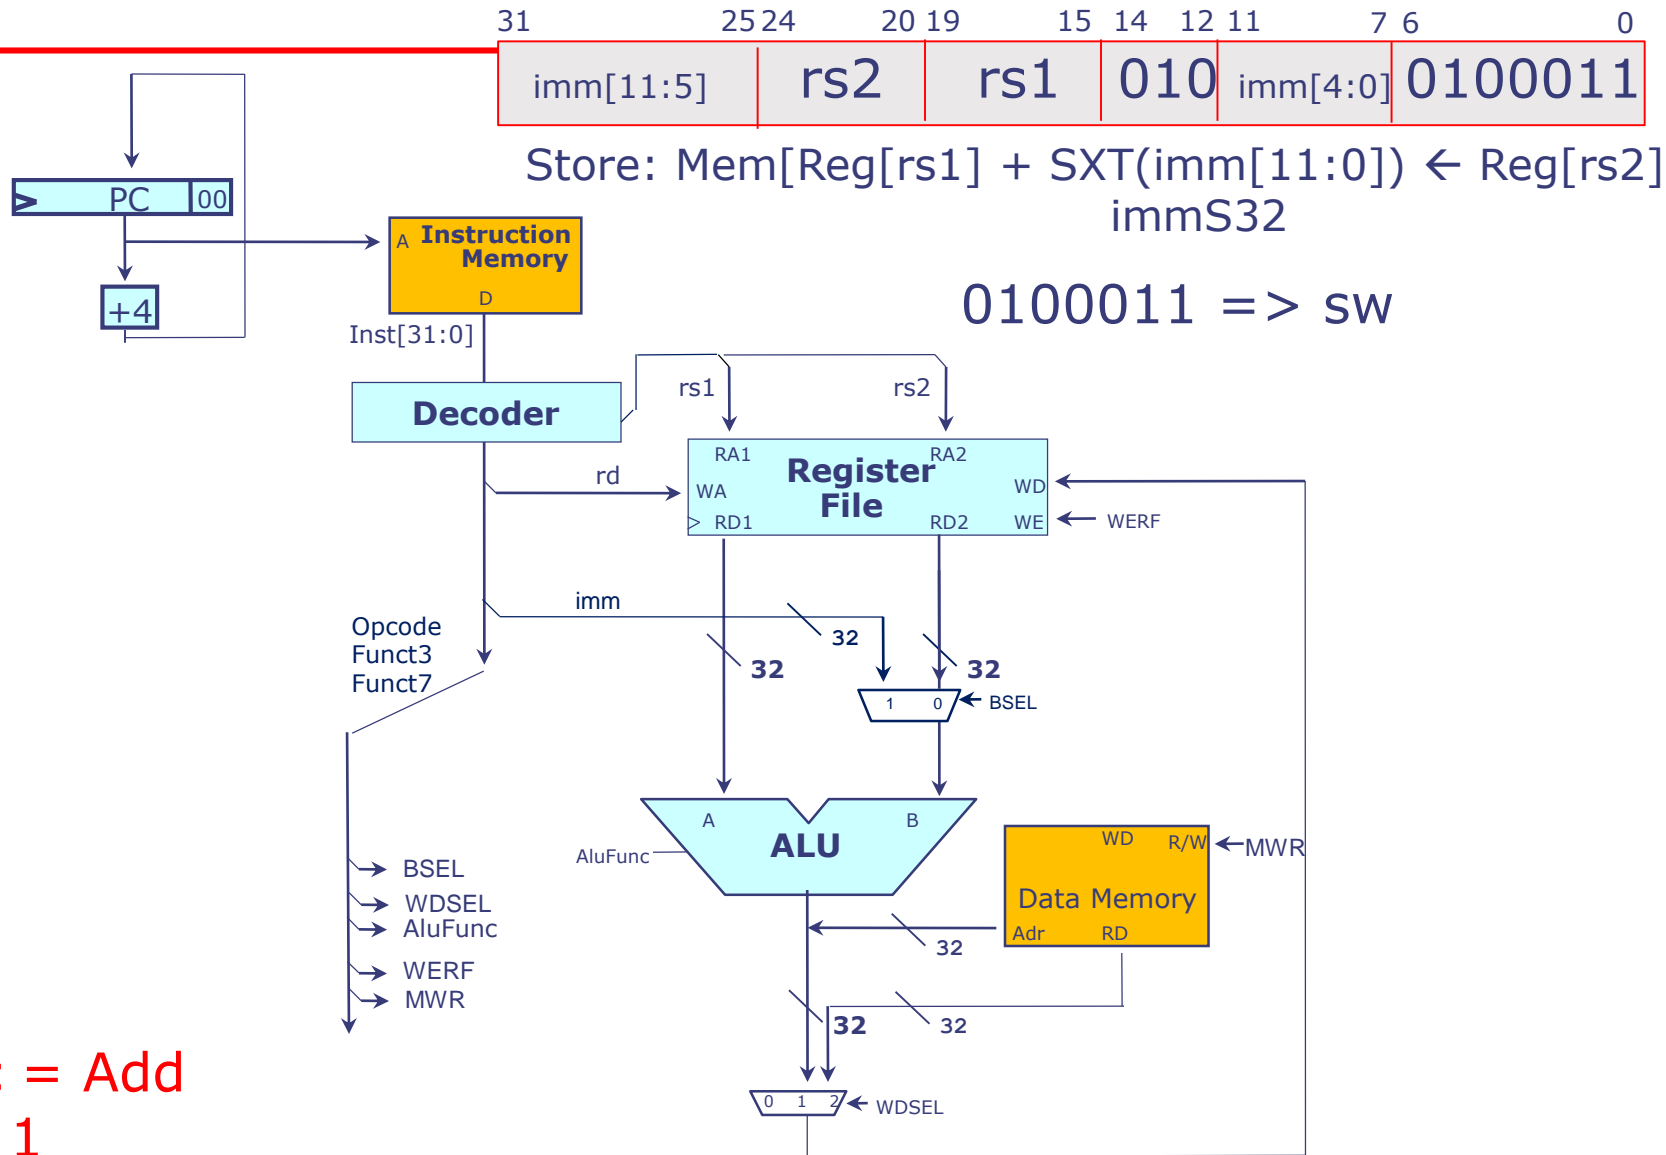
# Store Instruction



# Store Instruction

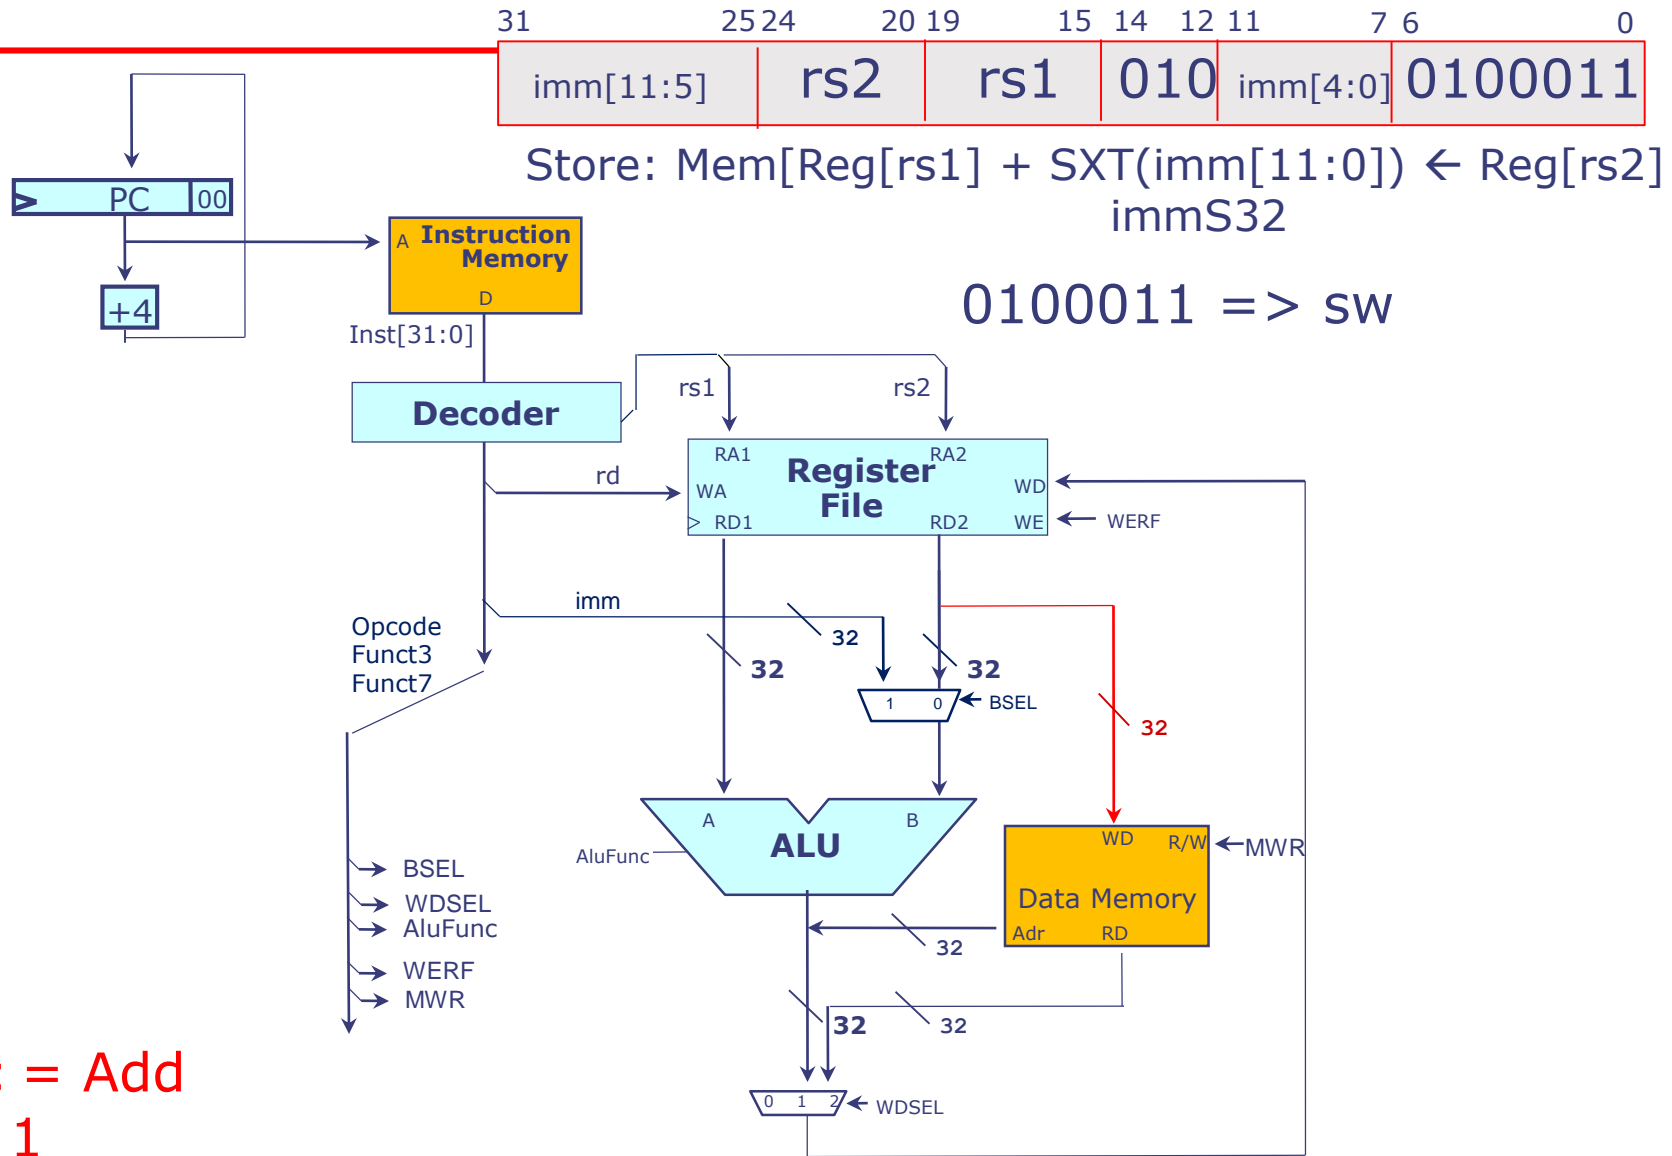


# Store Instruction



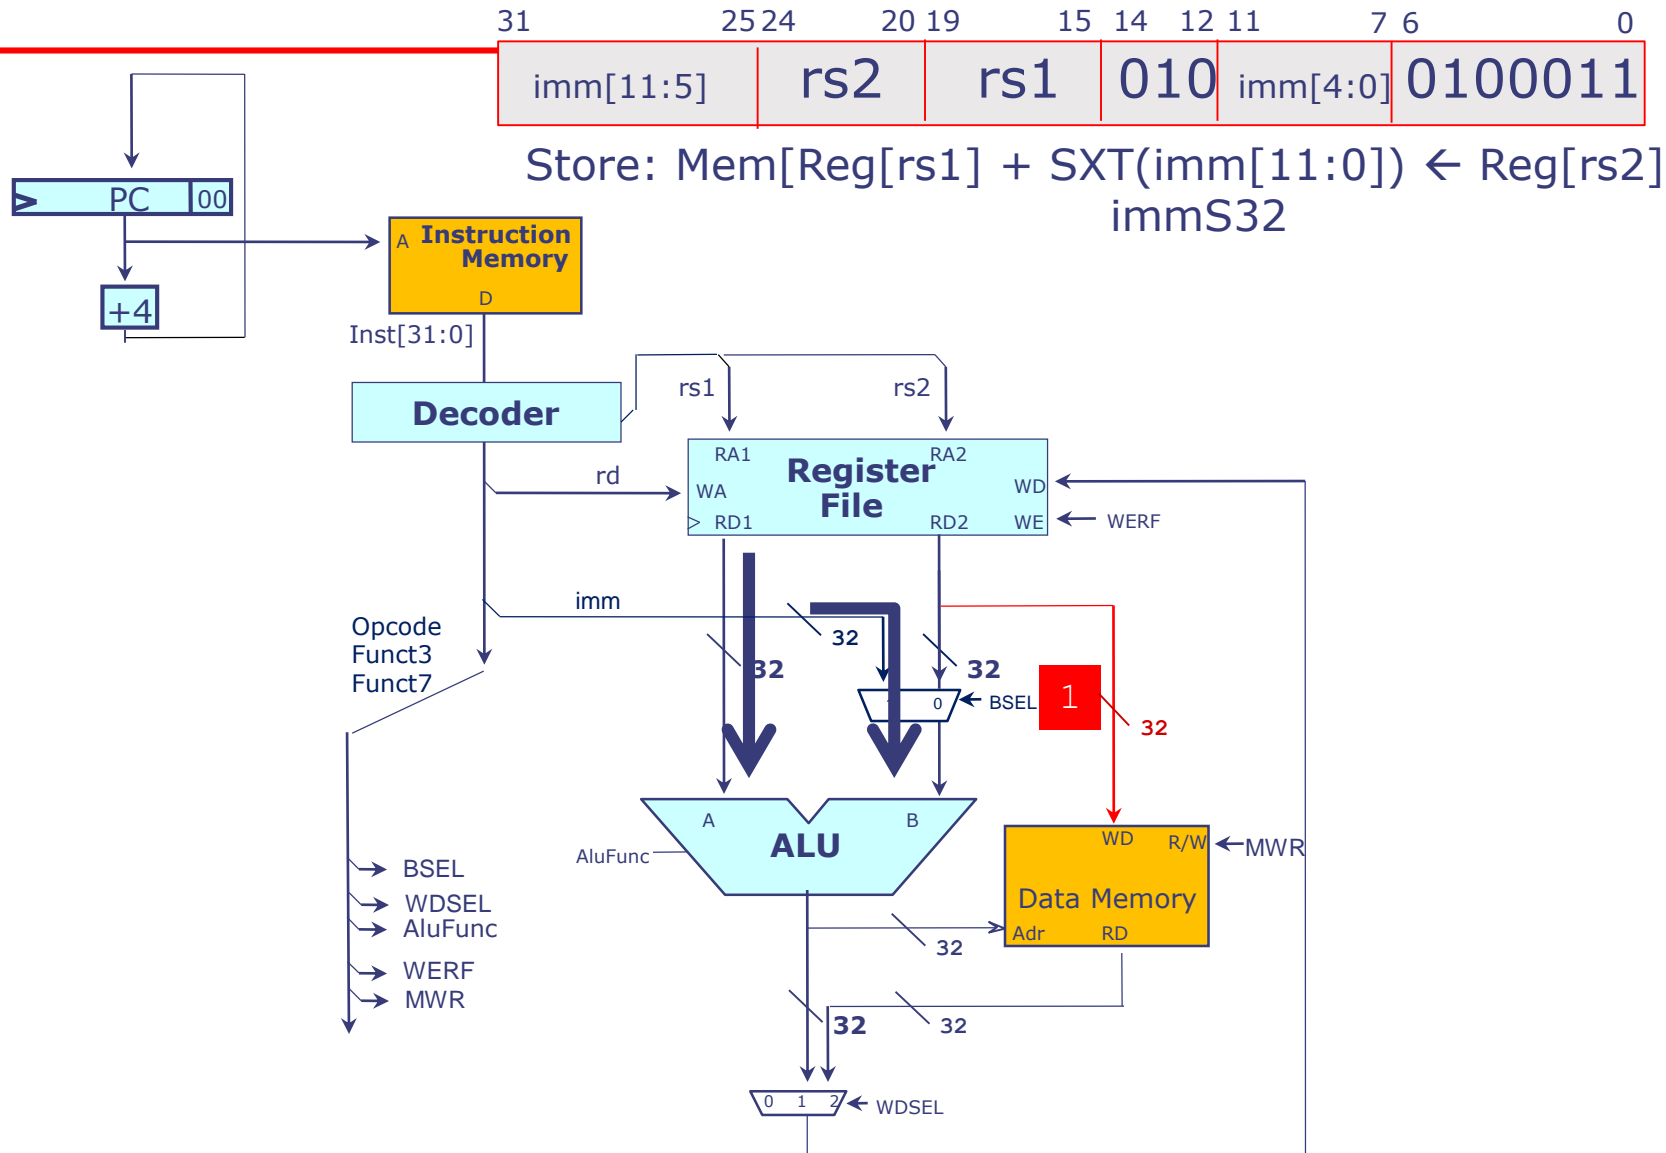
AluFunc = Add  
BSEL = 1

# Store Instruction



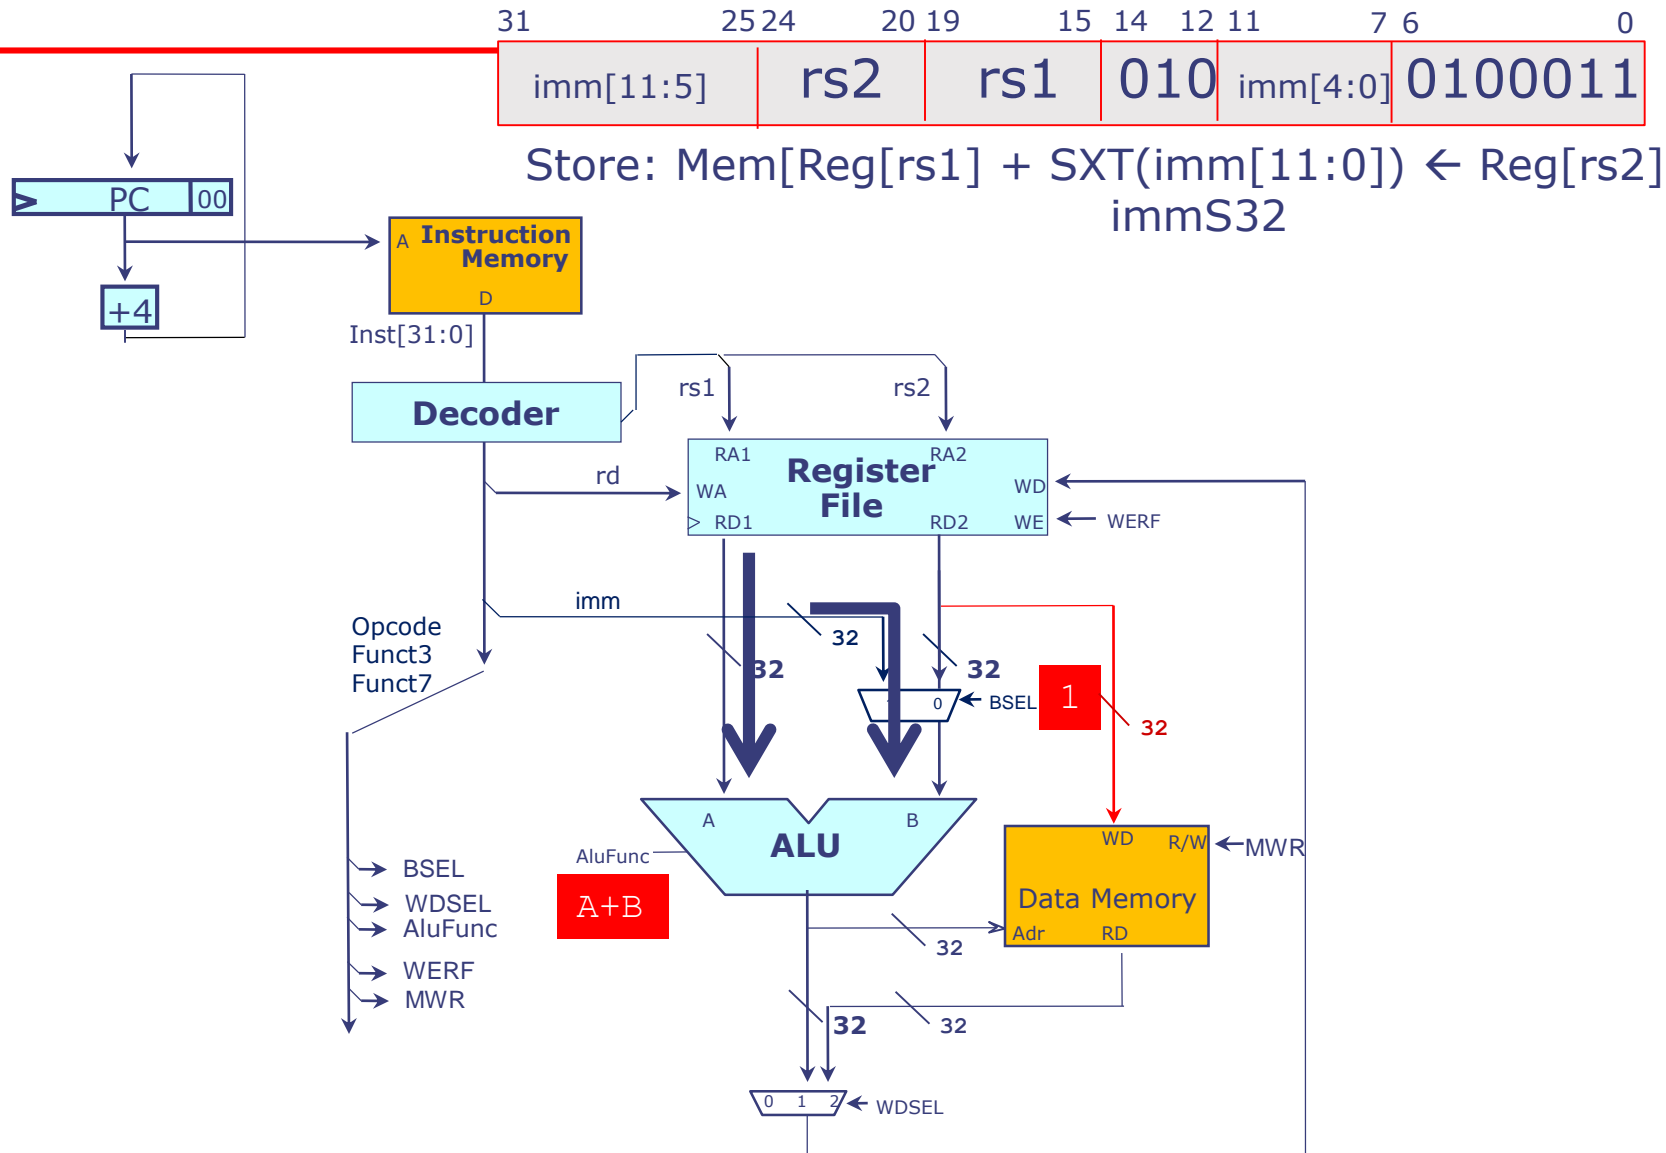
AluFunc = Add  
BSEL = 1

# Store Instruction

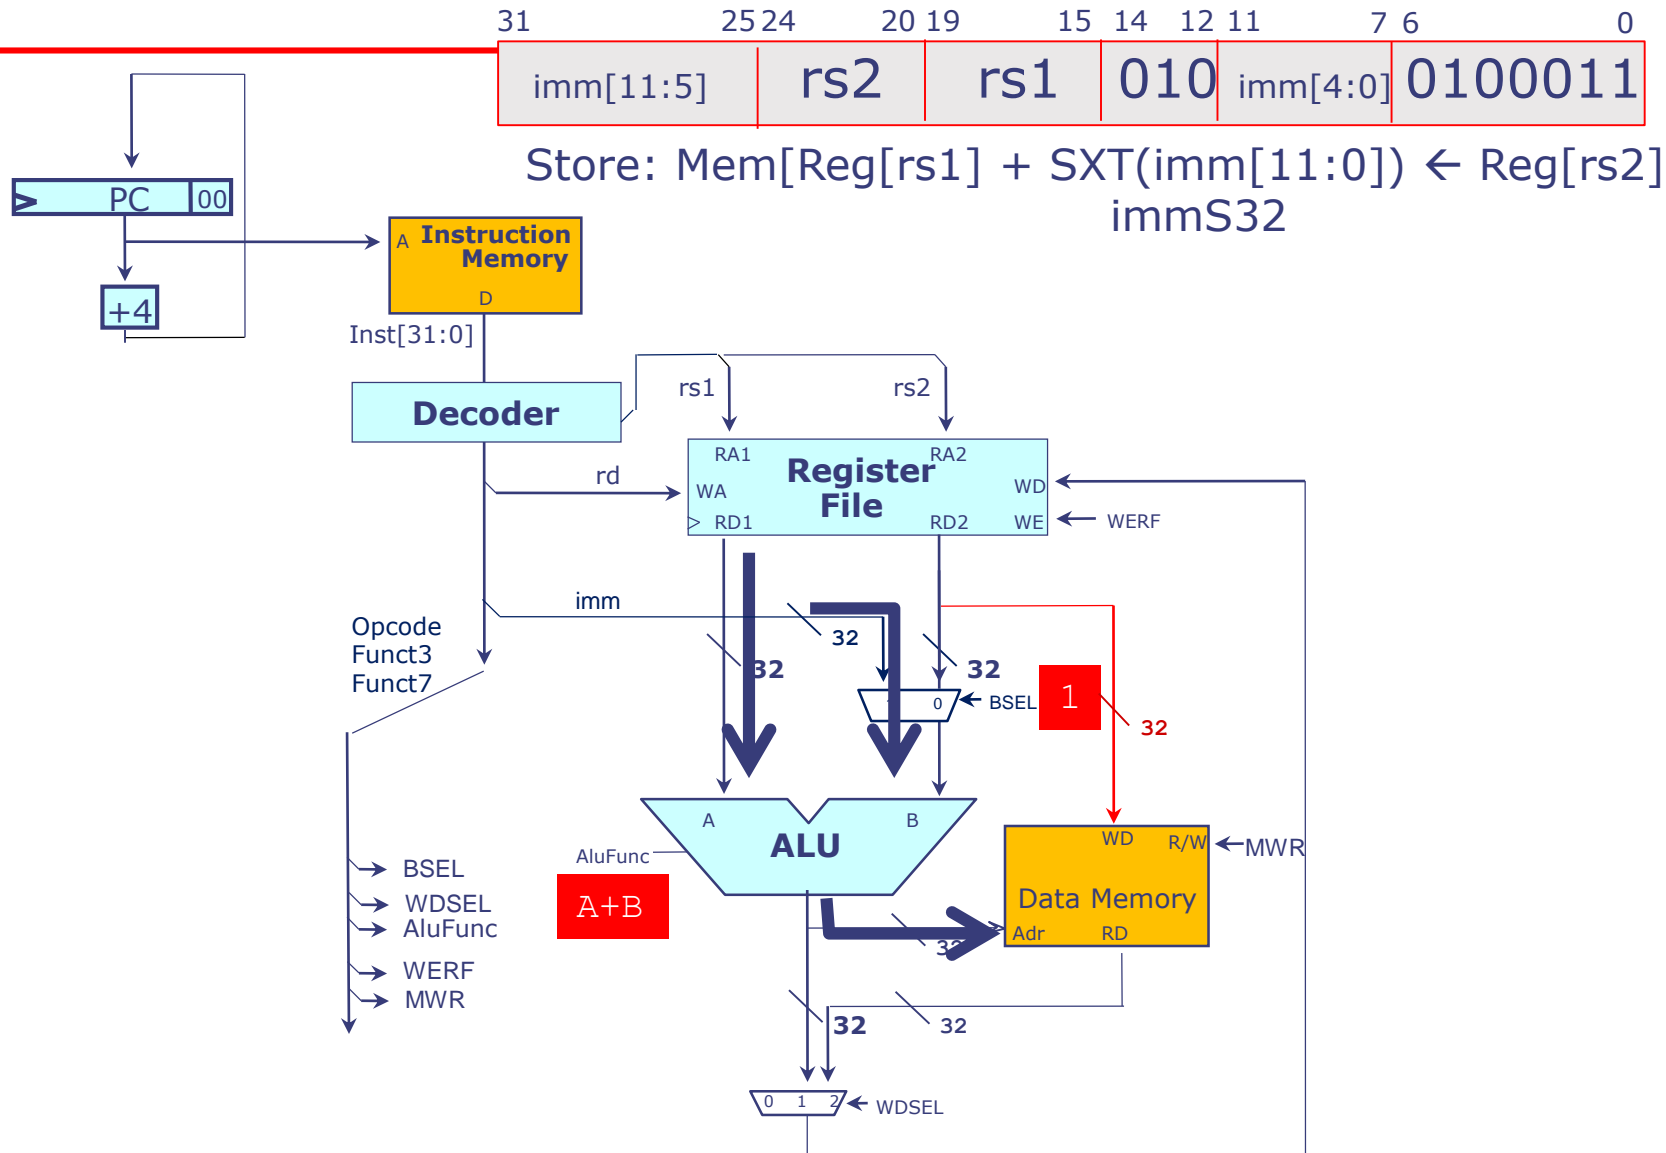




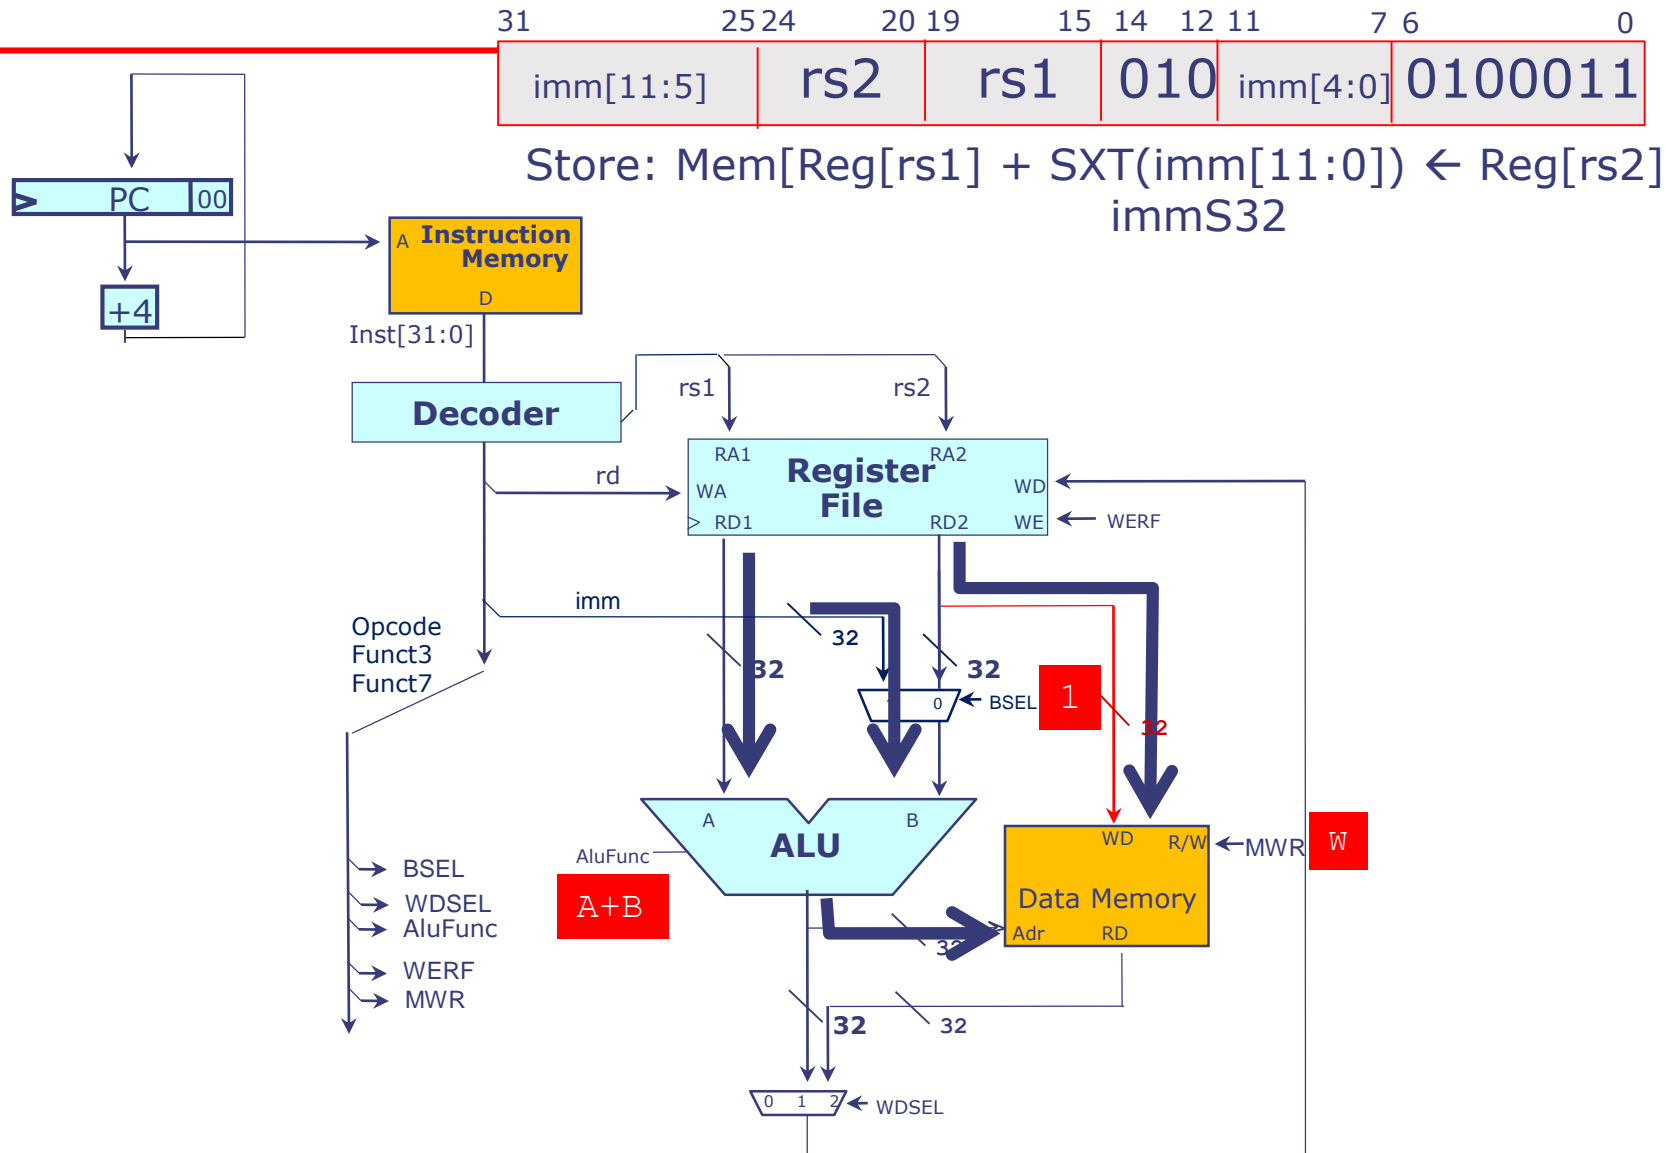
# Store Instruction



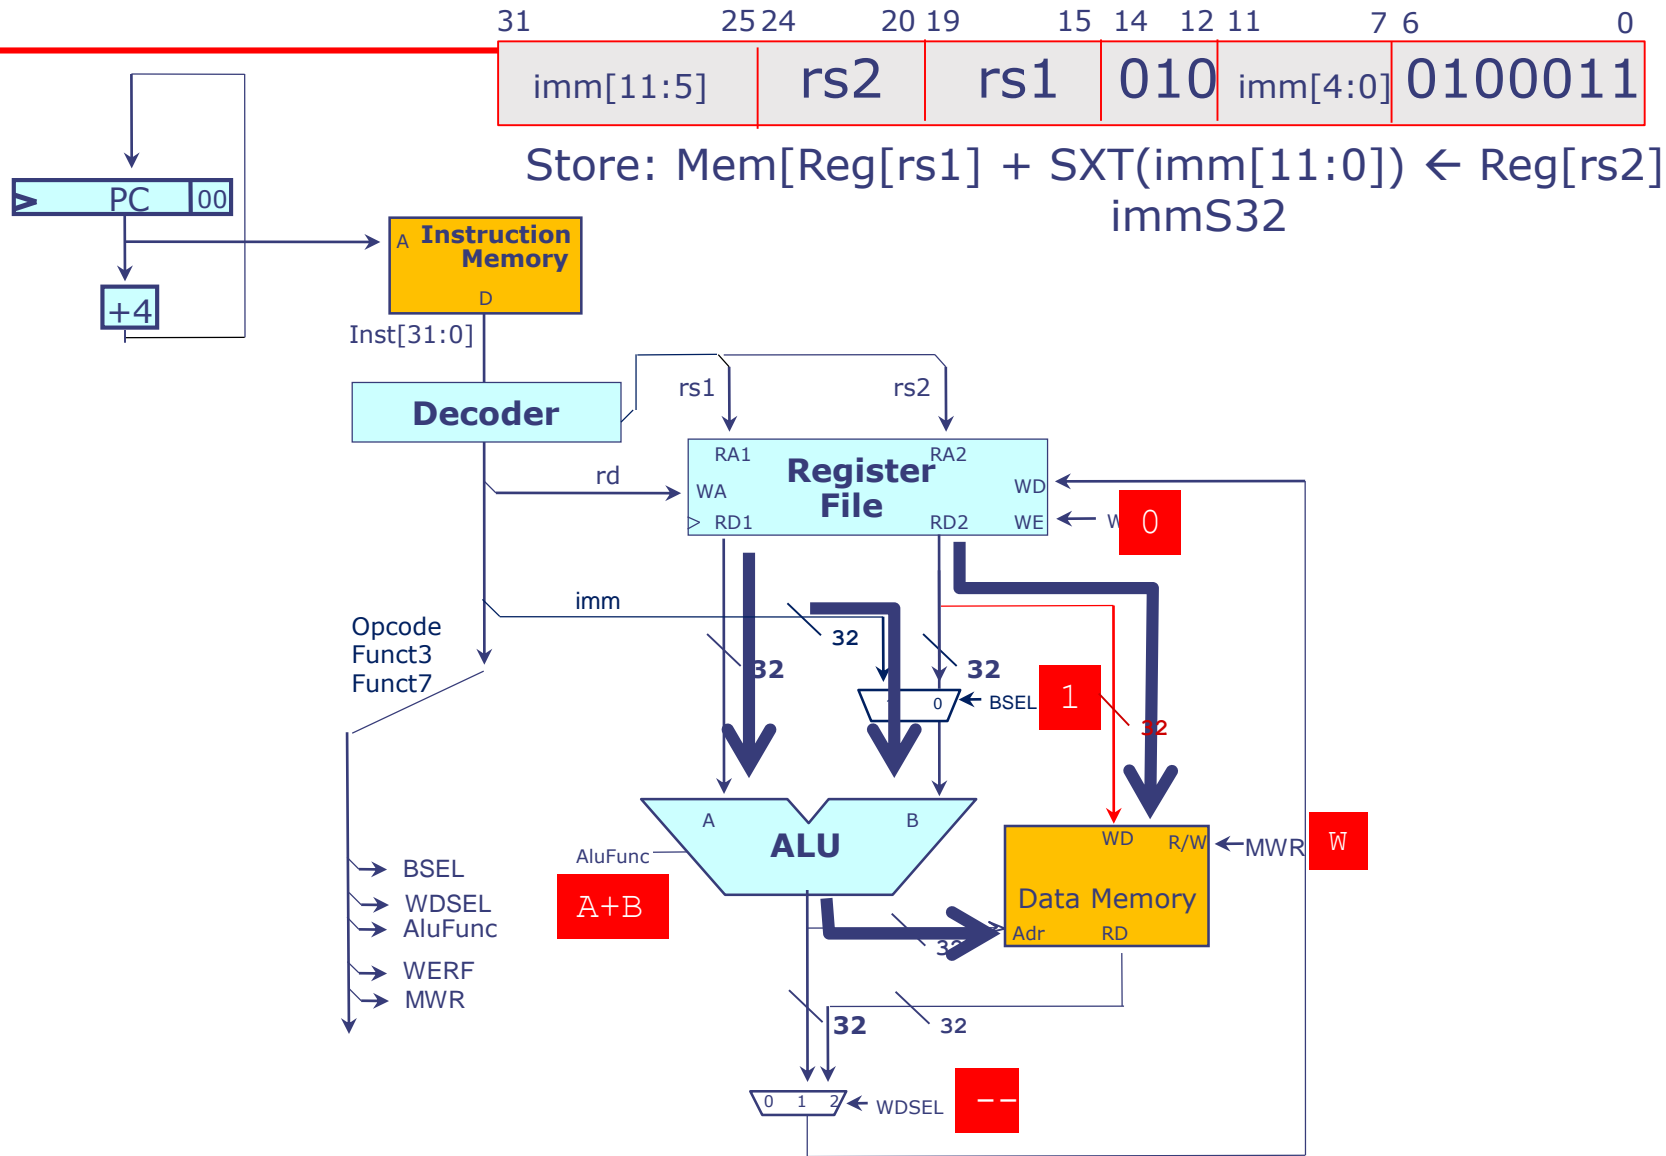
# Store Instruction



# Store Instruction



# Store Instruction



# Branch Instructions

differ only in the aluBr operation they perform

---


Instruction	Description	Execution
BEQ rs1, rs2, label	Branch =	$pc \leq (\text{reg}[\text{rs1}] == \text{reg}[\text{rs2}]) ? \text{label} : pc + 4$
BNE rs1, rs2, label	Branch !=	$pc \leq (\text{reg}[\text{rs1}] != \text{reg}[\text{rs2}]) ? \text{label} : pc + 4$
BLT rs1, rs2, label	Branch < (Signed)	$pc \leq (\text{reg}[\text{rs1}] <_s \text{reg}[\text{rs2}]) ? \text{label} : pc + 4$
BGE rs1, rs2, label	Branch $\geq$ (Signed)	$pc \leq (\text{reg}[\text{rs1}] \geq_s \text{reg}[\text{rs2}]) ? \text{label} : pc + 4$
BLTU rs1, rs2, label	Branch < (Unsigned)	$pc \leq (\text{reg}[\text{rs1}] <_u \text{reg}[\text{rs2}]) ? \text{label} : pc + 4$
BGEU rs1, rs2, label	Branch $\geq$ (Unsigned)	$pc \leq (\text{reg}[\text{rs1}] \geq_u \text{reg}[\text{rs2}]) ? \text{label} : pc + 4$

These instructions are grouped in a category called **BRANCH** with fields (brFunc, rs1, rs2, immB32)

# ALU for Branch Comparisons

---



Like ALU, but  
returns a Bool



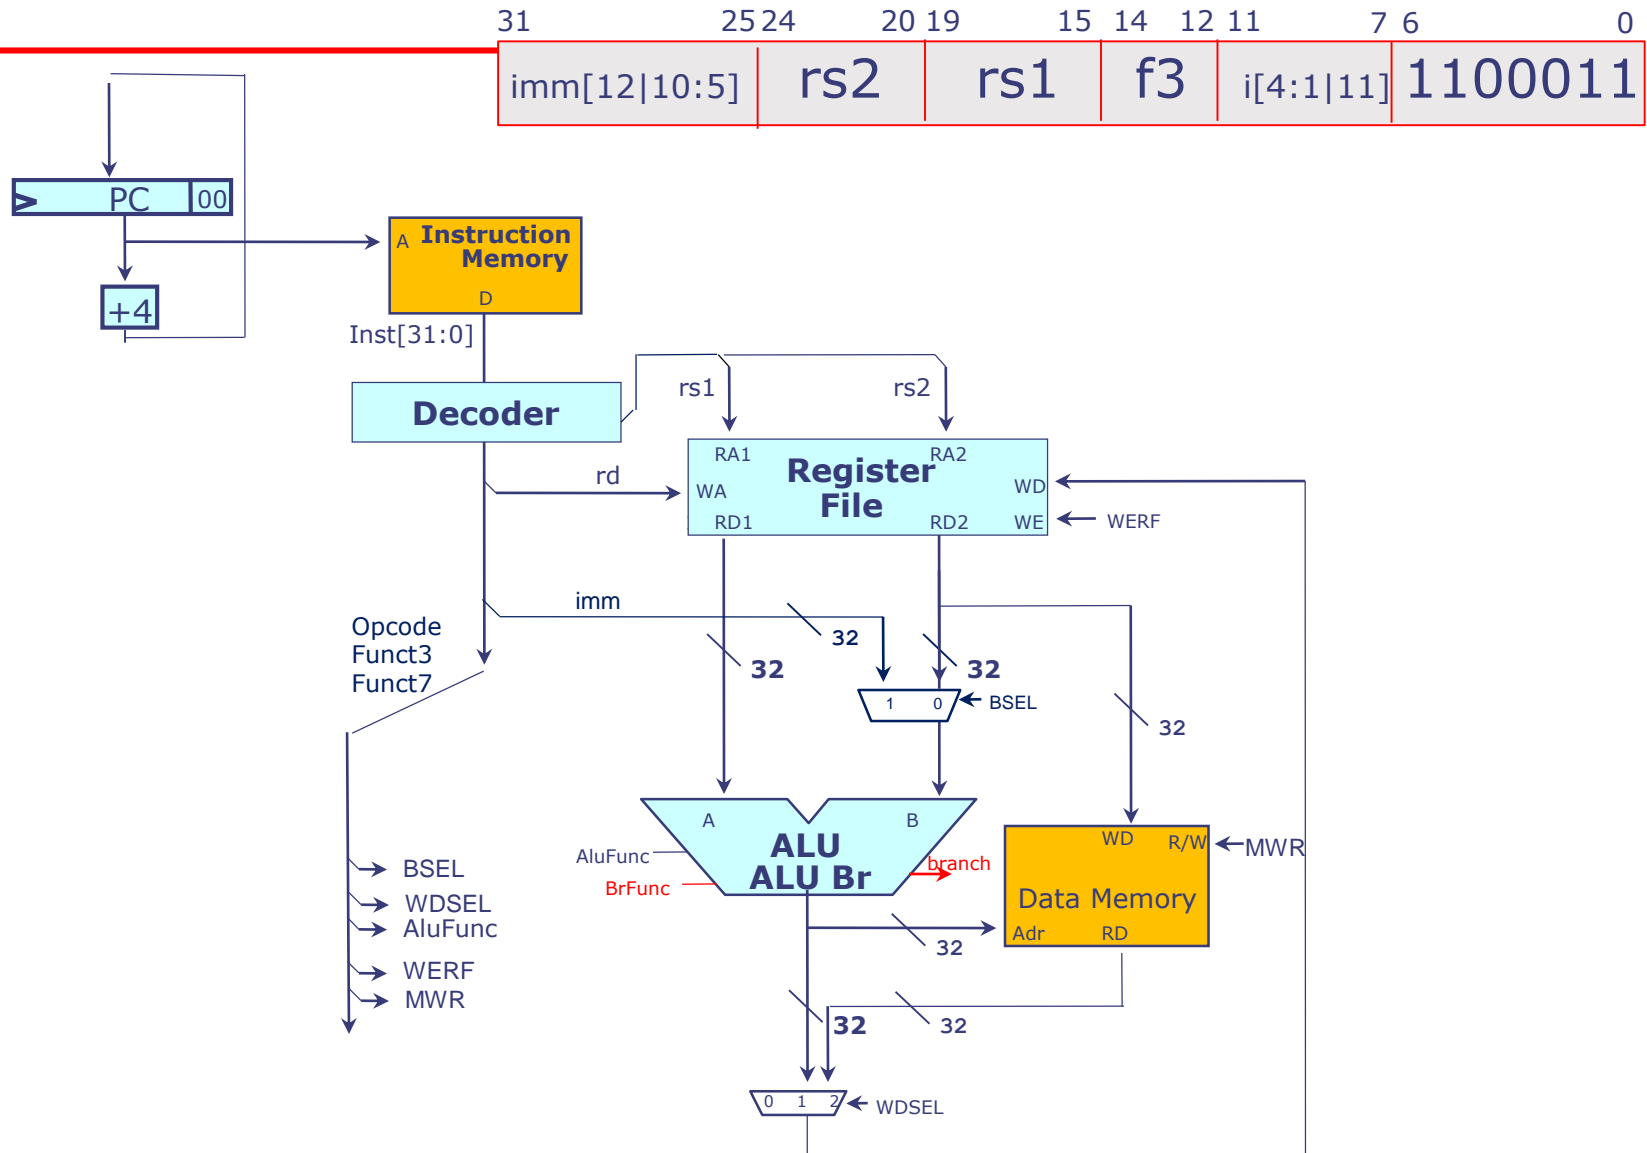
# ALU for Branch Comparisons

---

Like ALU, but  
returns a Bool

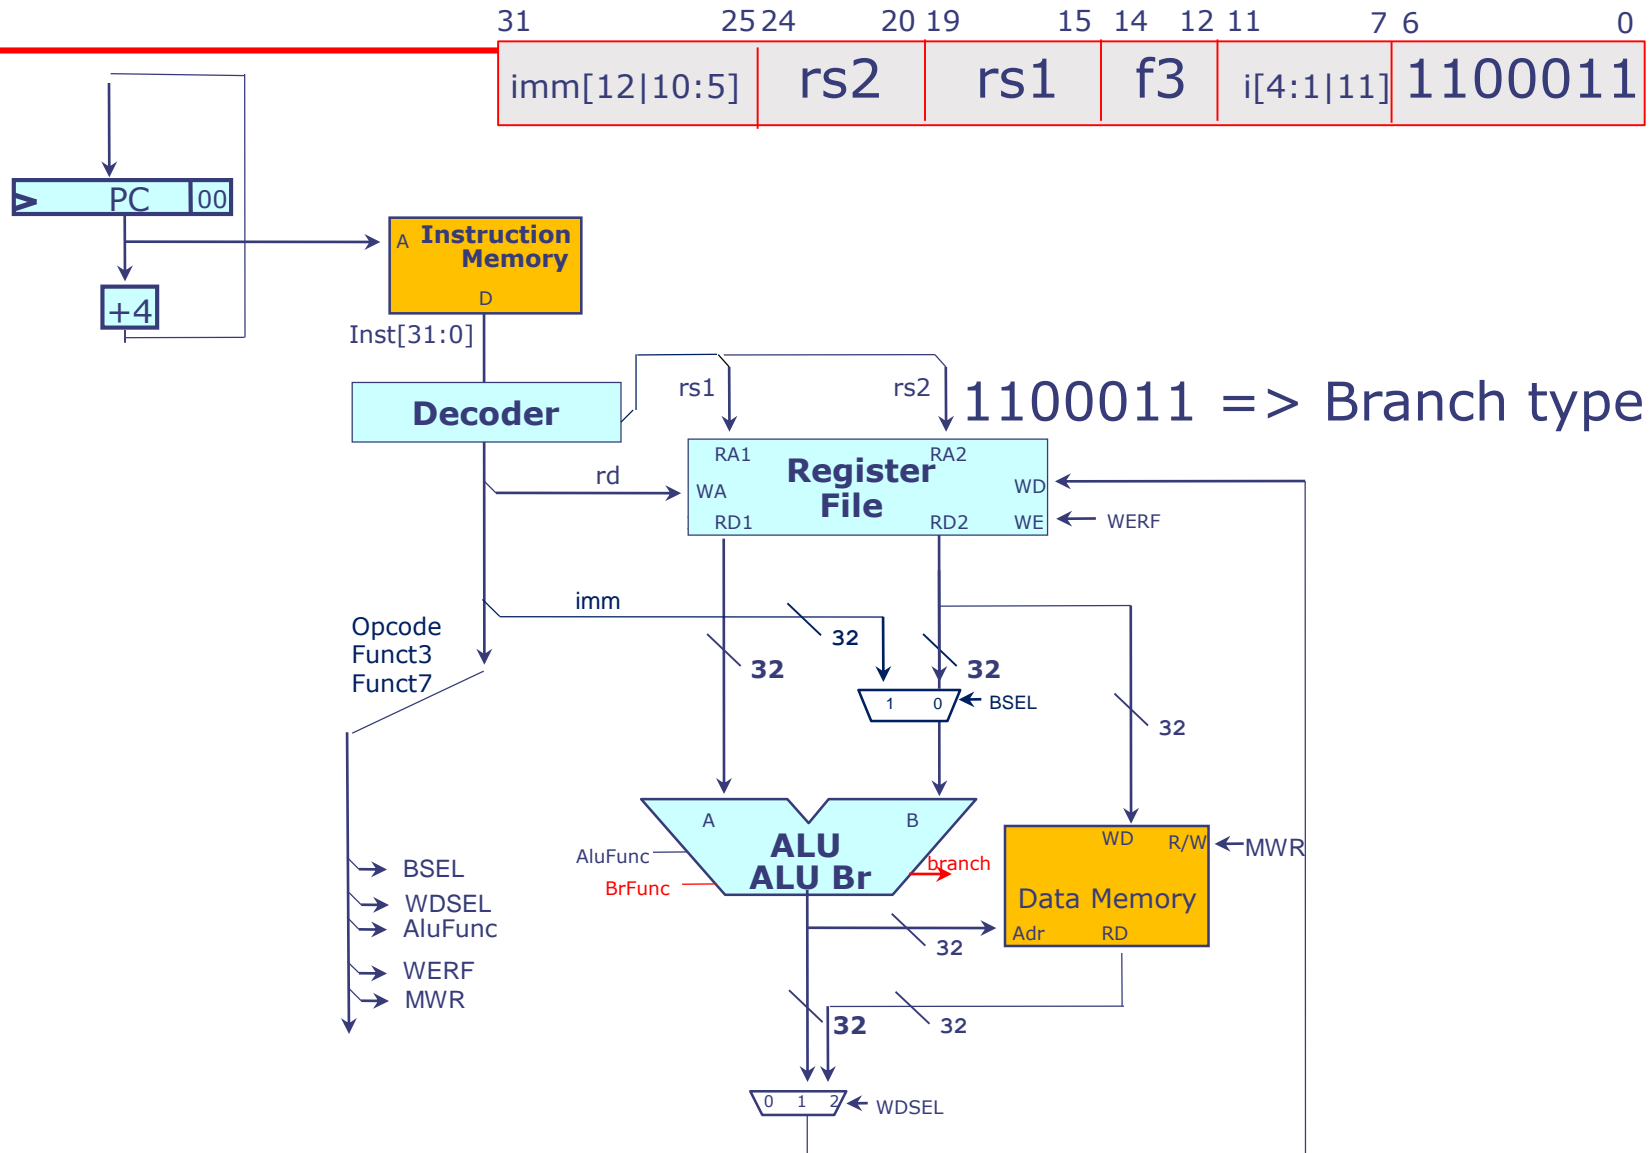


# Branch Instructions

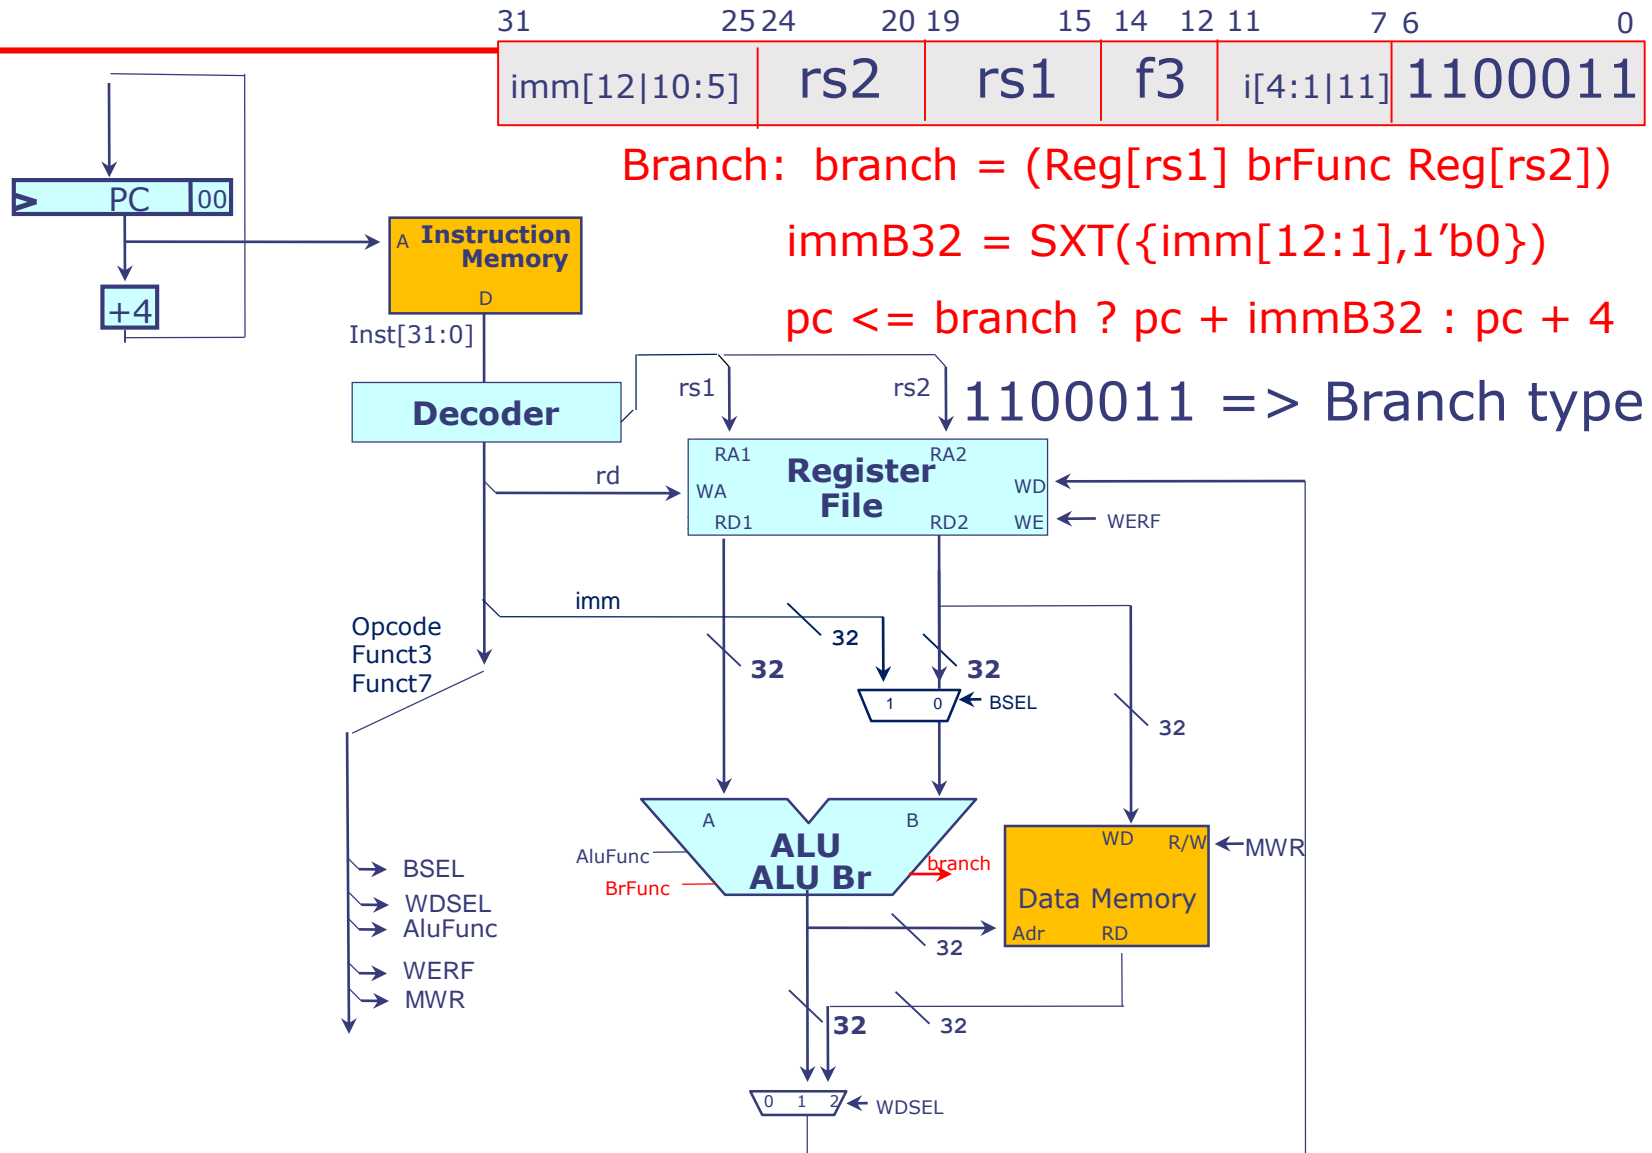




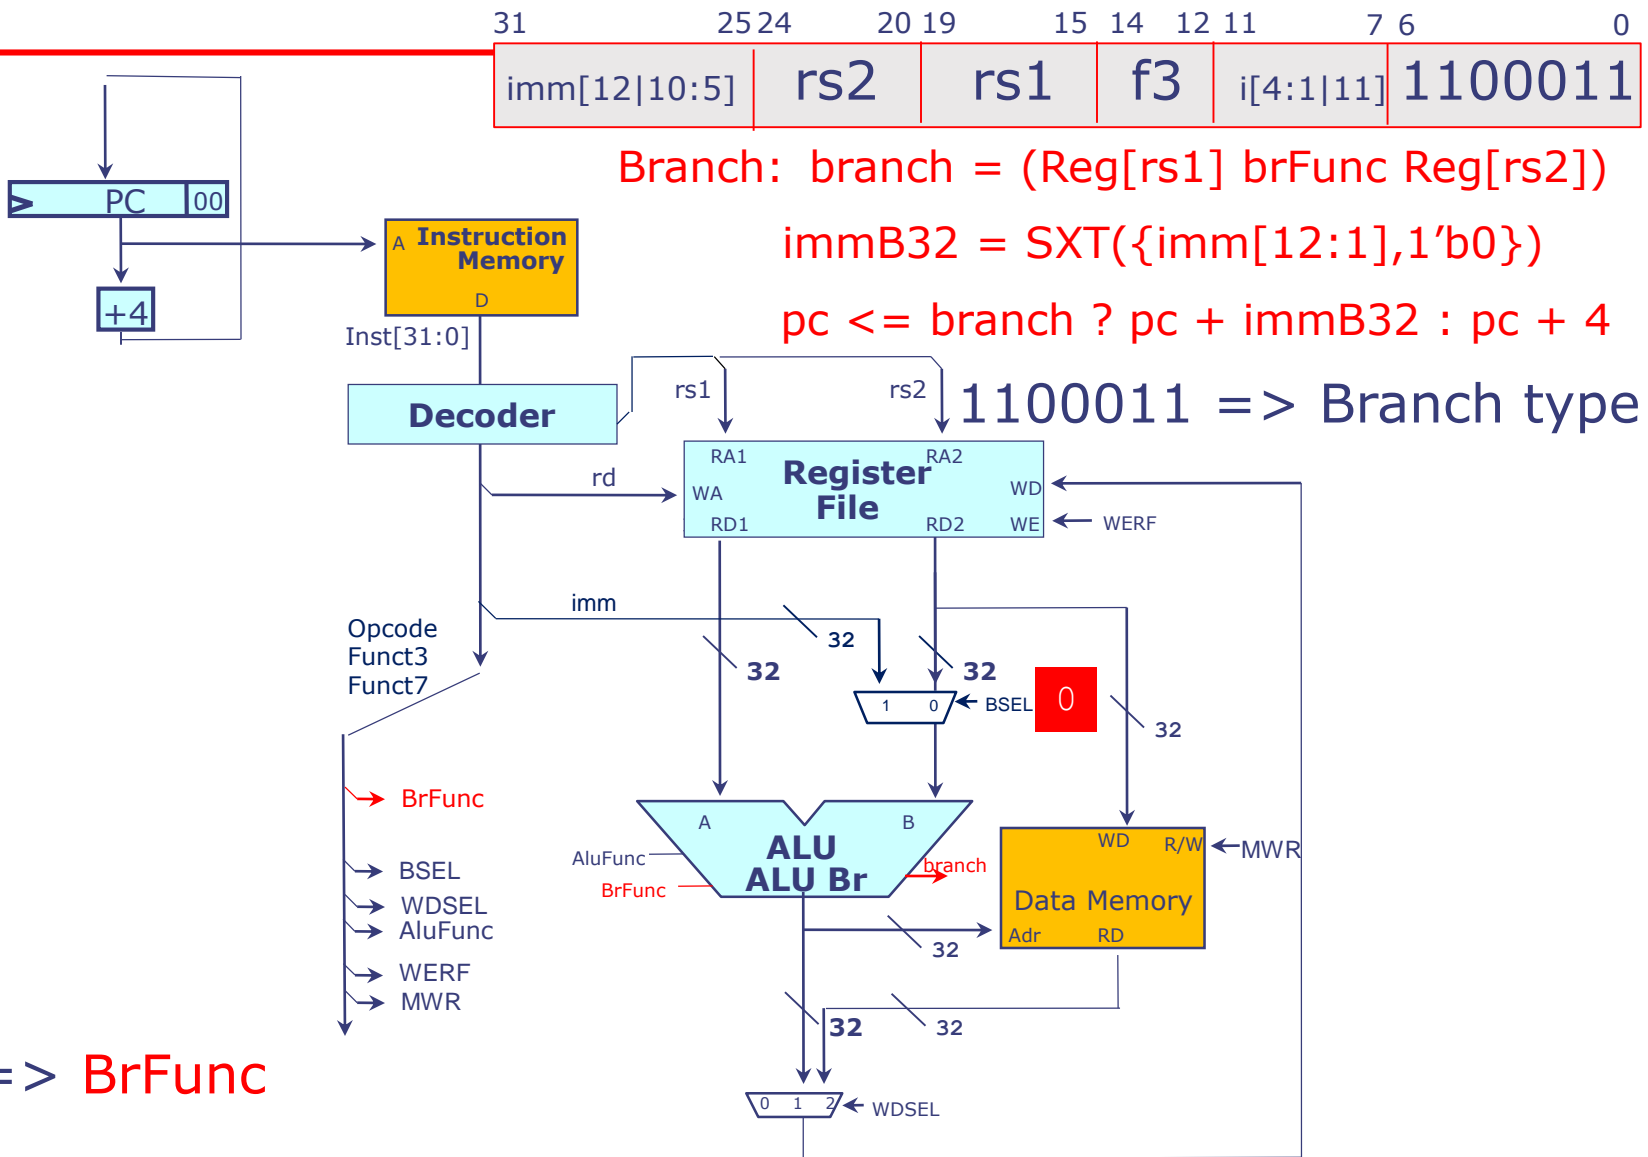
# Branch Instructions



# Branch Instructions

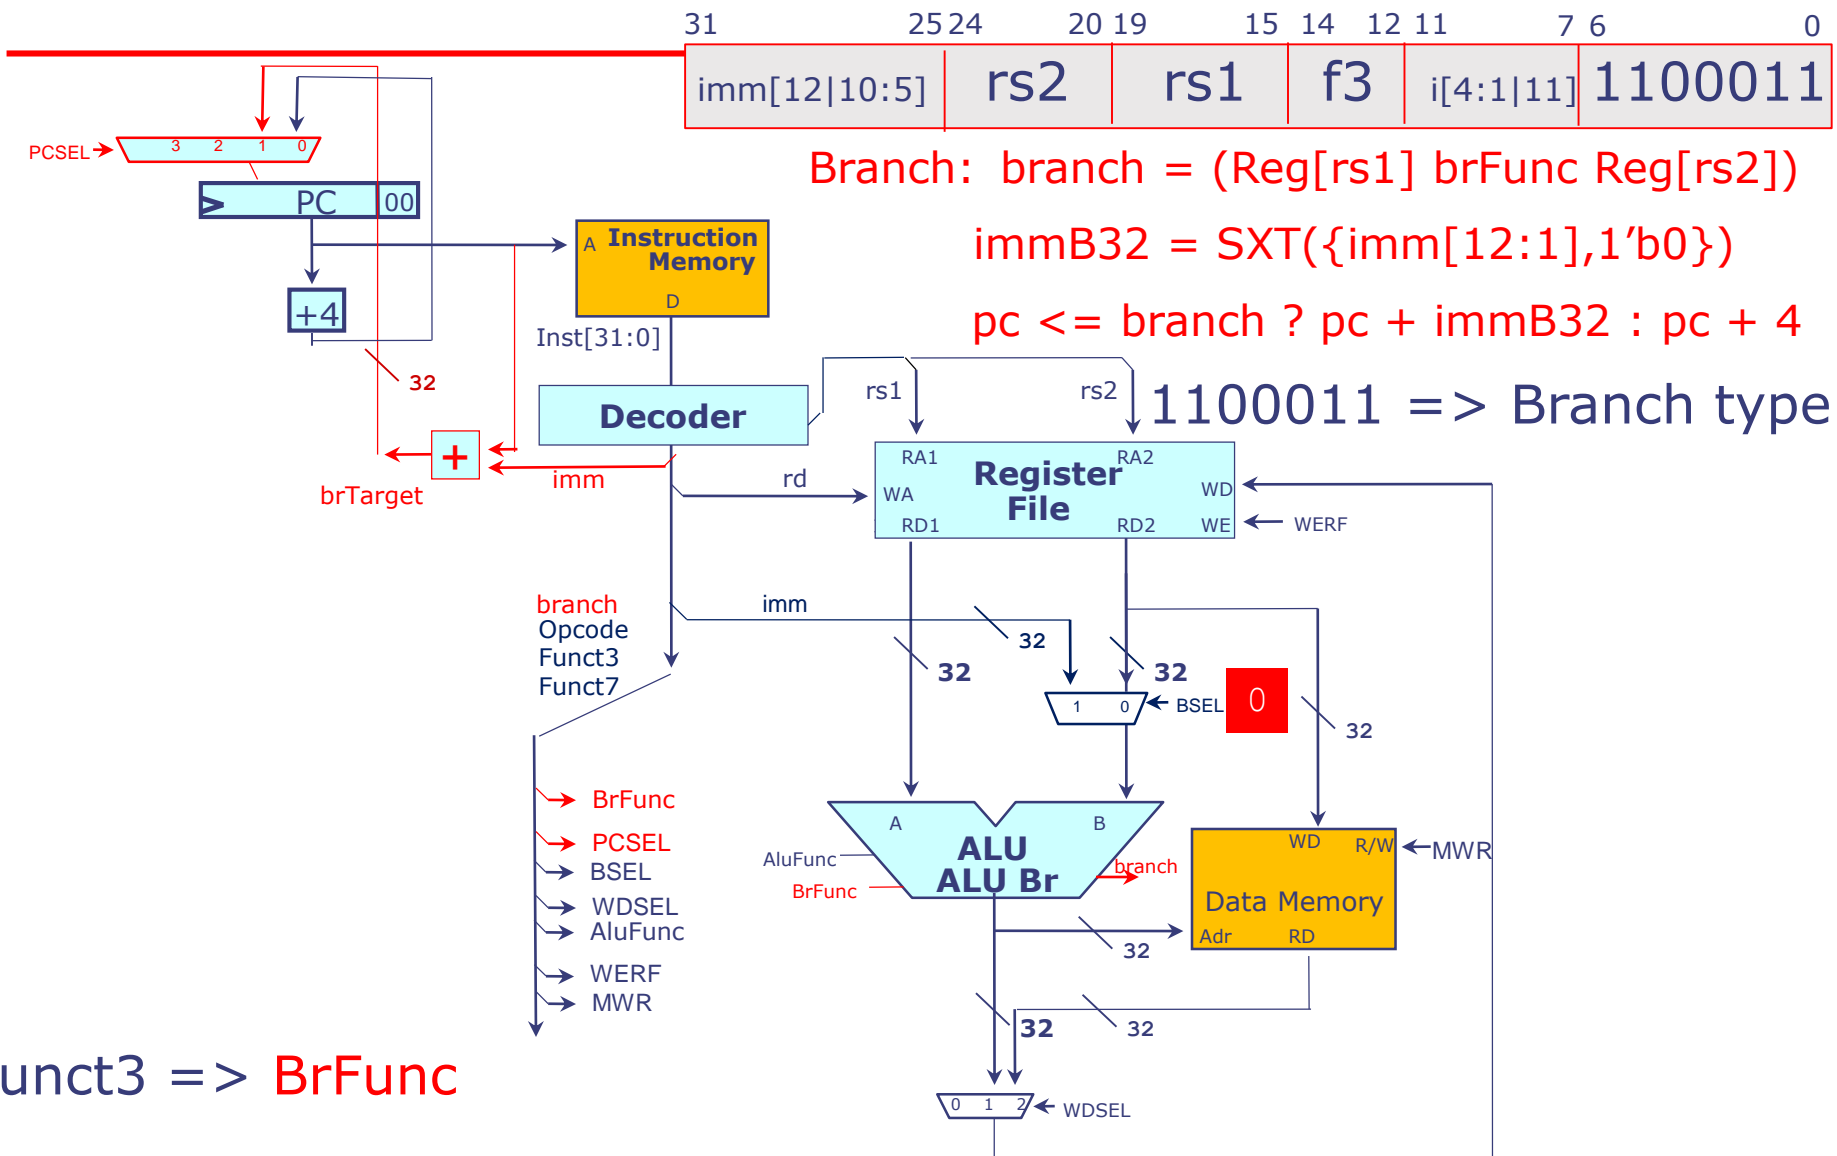


# Branch Instructions




funct3 => BrFunc

# Branch Instructions



funct3 => BrFunc

# Branch Instructions

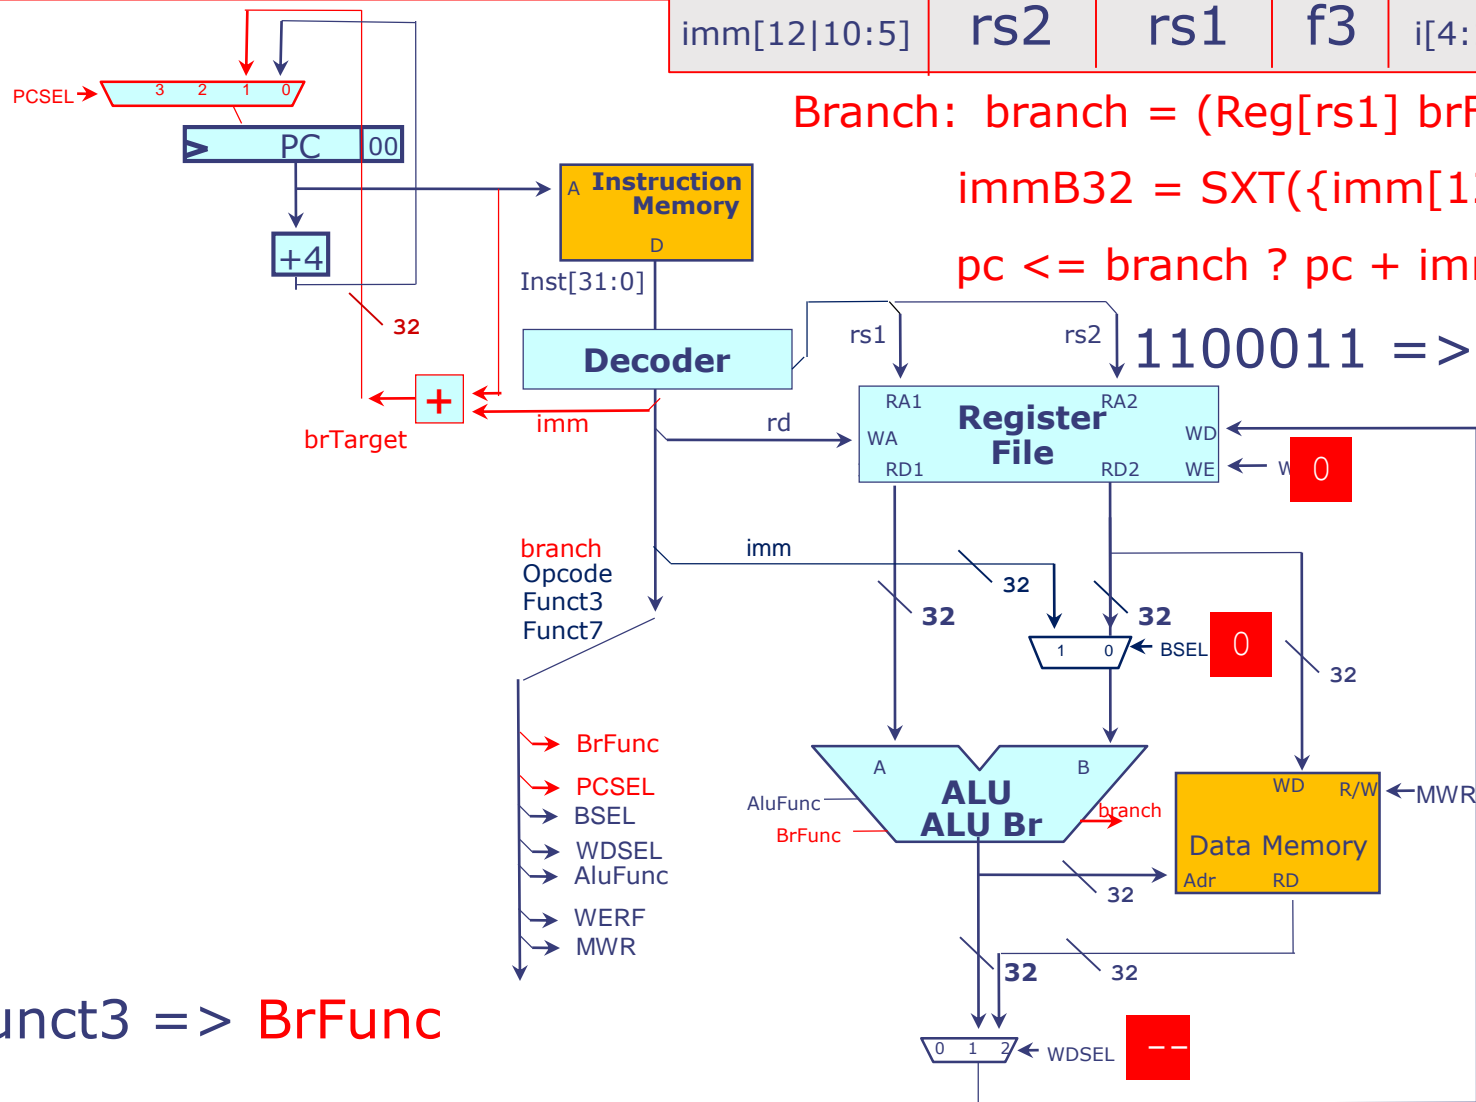


Branch:  $\text{branch} = (\text{Reg}[\text{rs1}] \text{ brFunc } \text{Reg}[\text{rs2}])$

$\text{immB32} = \text{SXT}(\{\text{imm}[12:1], 1'b0\})$

$\text{pc} \leq \text{branch} ? \text{pc} + \text{immB32} : \text{pc} + 4$

1100011 => Branch type



funct3 => BrFunc

# Remaining Instructions

---

Instruction	Description	Execution
JAL rd, label	Jump and Link	$\text{reg}[\text{rd}] \leq \text{pc} + 4$ $\text{pc} \leq \text{label}$
JALR rd, offset(rs1)	Jump and Link Register	$\text{reg}[\text{rd}] \leq \text{pc} + 4$ $\text{pc} \leq \{(\text{reg}[\text{rs1}] + \text{offset})[31:1], 1'b0\}$
LUI rd, luiConstant	Load Upper Immediate	$\text{reg}[\text{rd}] \leq \text{luiConstant} \ll 12$

Each of these instructions is in a category by itself and needs to extract different fields from the instruction.

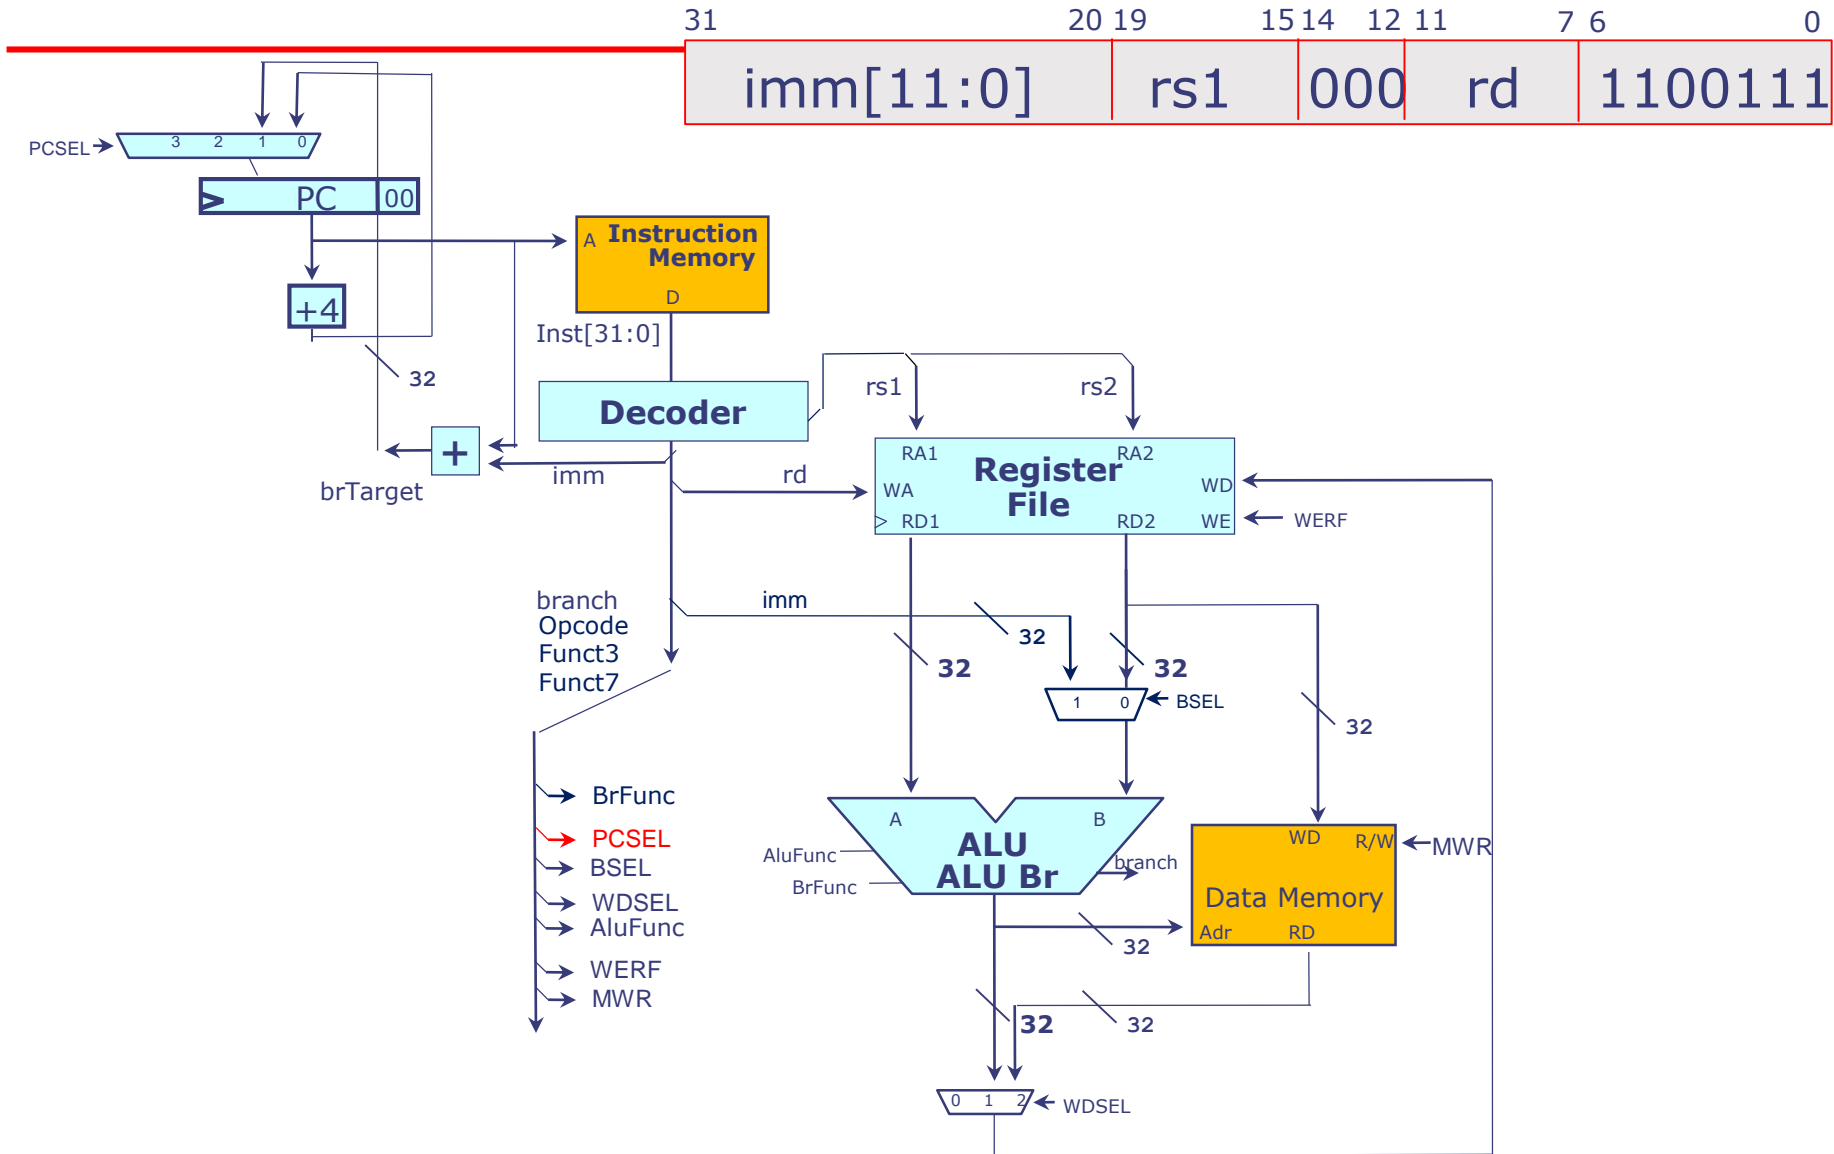
# Remaining Instructions

---

Instruction	Description	Execution
JAL rd, label	Jump and Link	$\text{reg}[\text{rd}] \leq \text{pc} + 4$ $\text{pc} \leq \text{label}$
JALR rd, offset(rs1)	Jump and Link Register	$\text{reg}[\text{rd}] \leq \text{pc} + 4$ $\text{pc} \leq \{(\text{reg}[\text{rs1}] + \text{offset})[31:1], 1'b0\}$
LUI rd, luiConstant	Load Upper Immediate	$\text{reg}[\text{rd}] \leq \text{luiConstant} \ll 12$

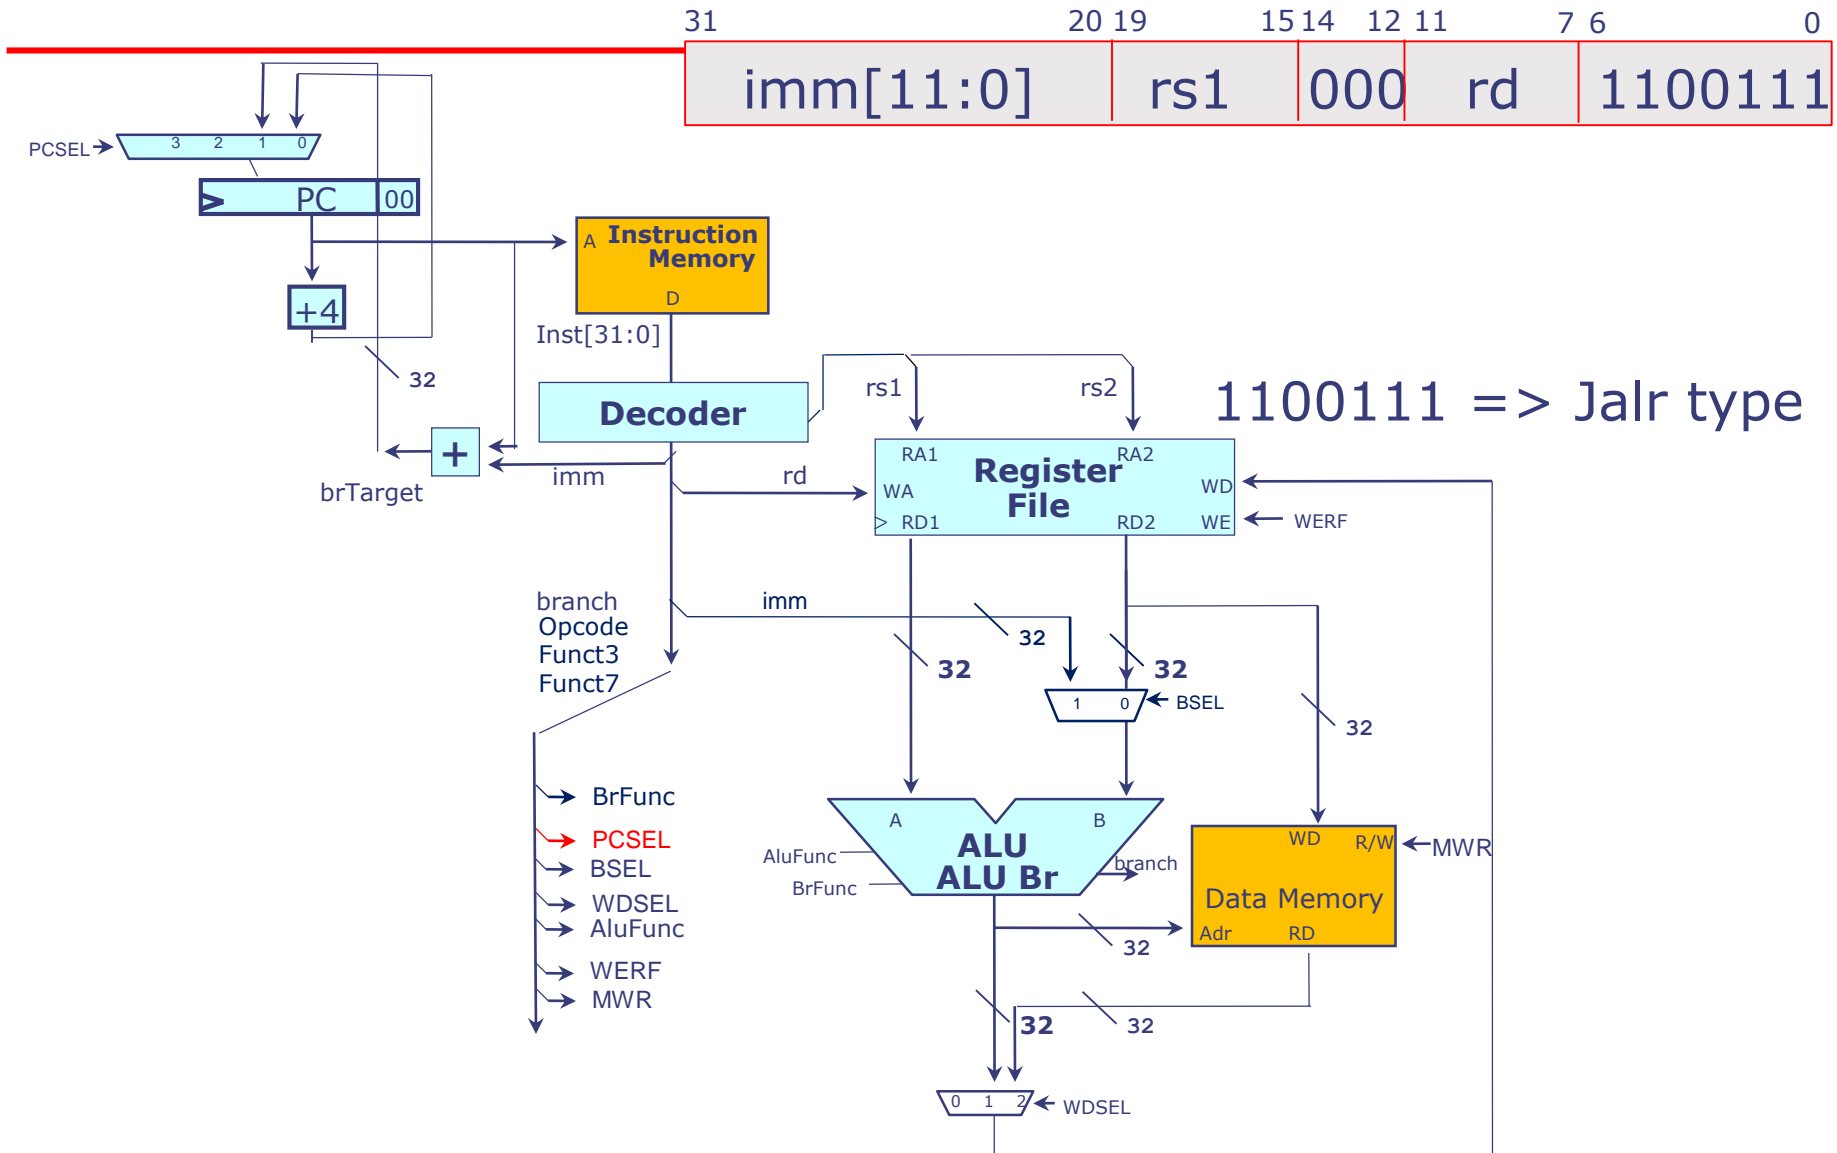
Each of these instructions is in a category by itself and needs to extract different fields from the instruction.  
jal and jalr update both pc and reg[rd].

# Jalr Instruction

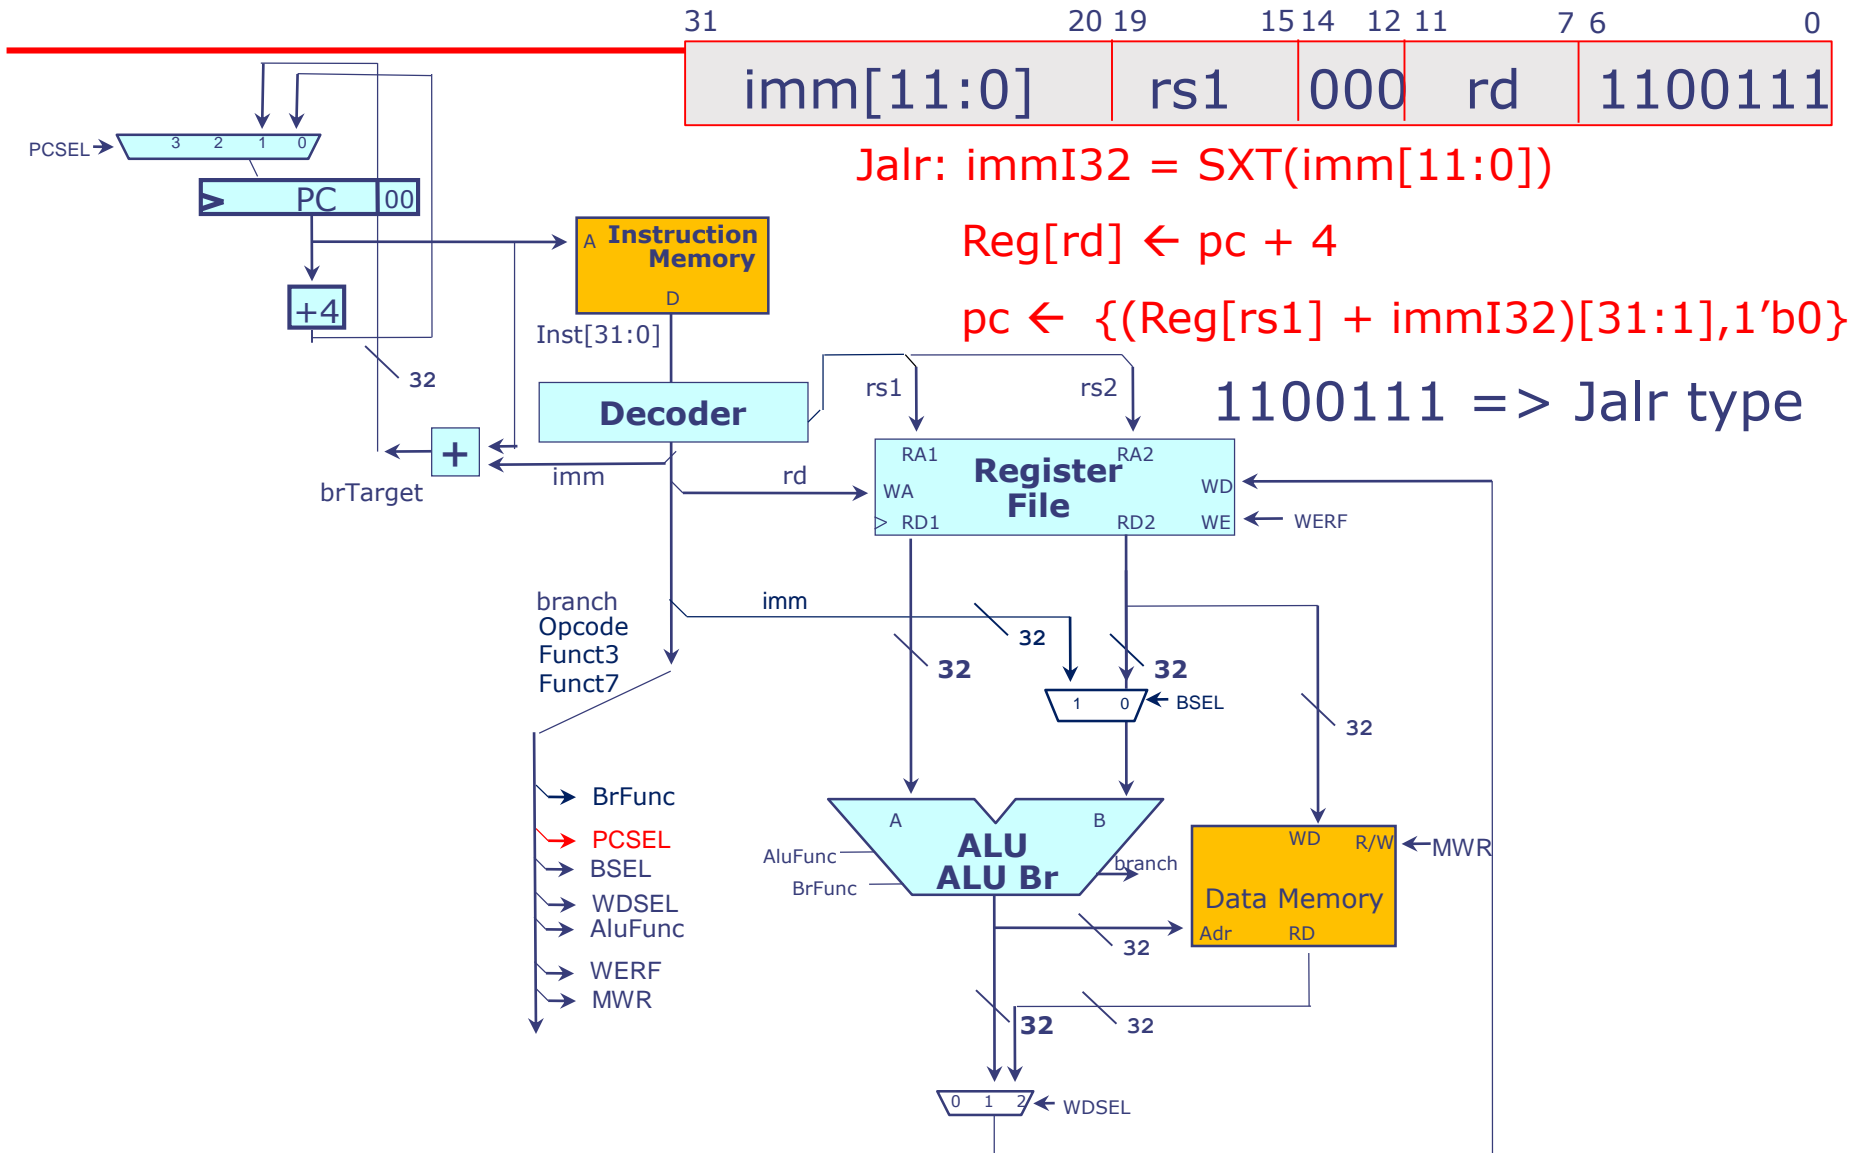




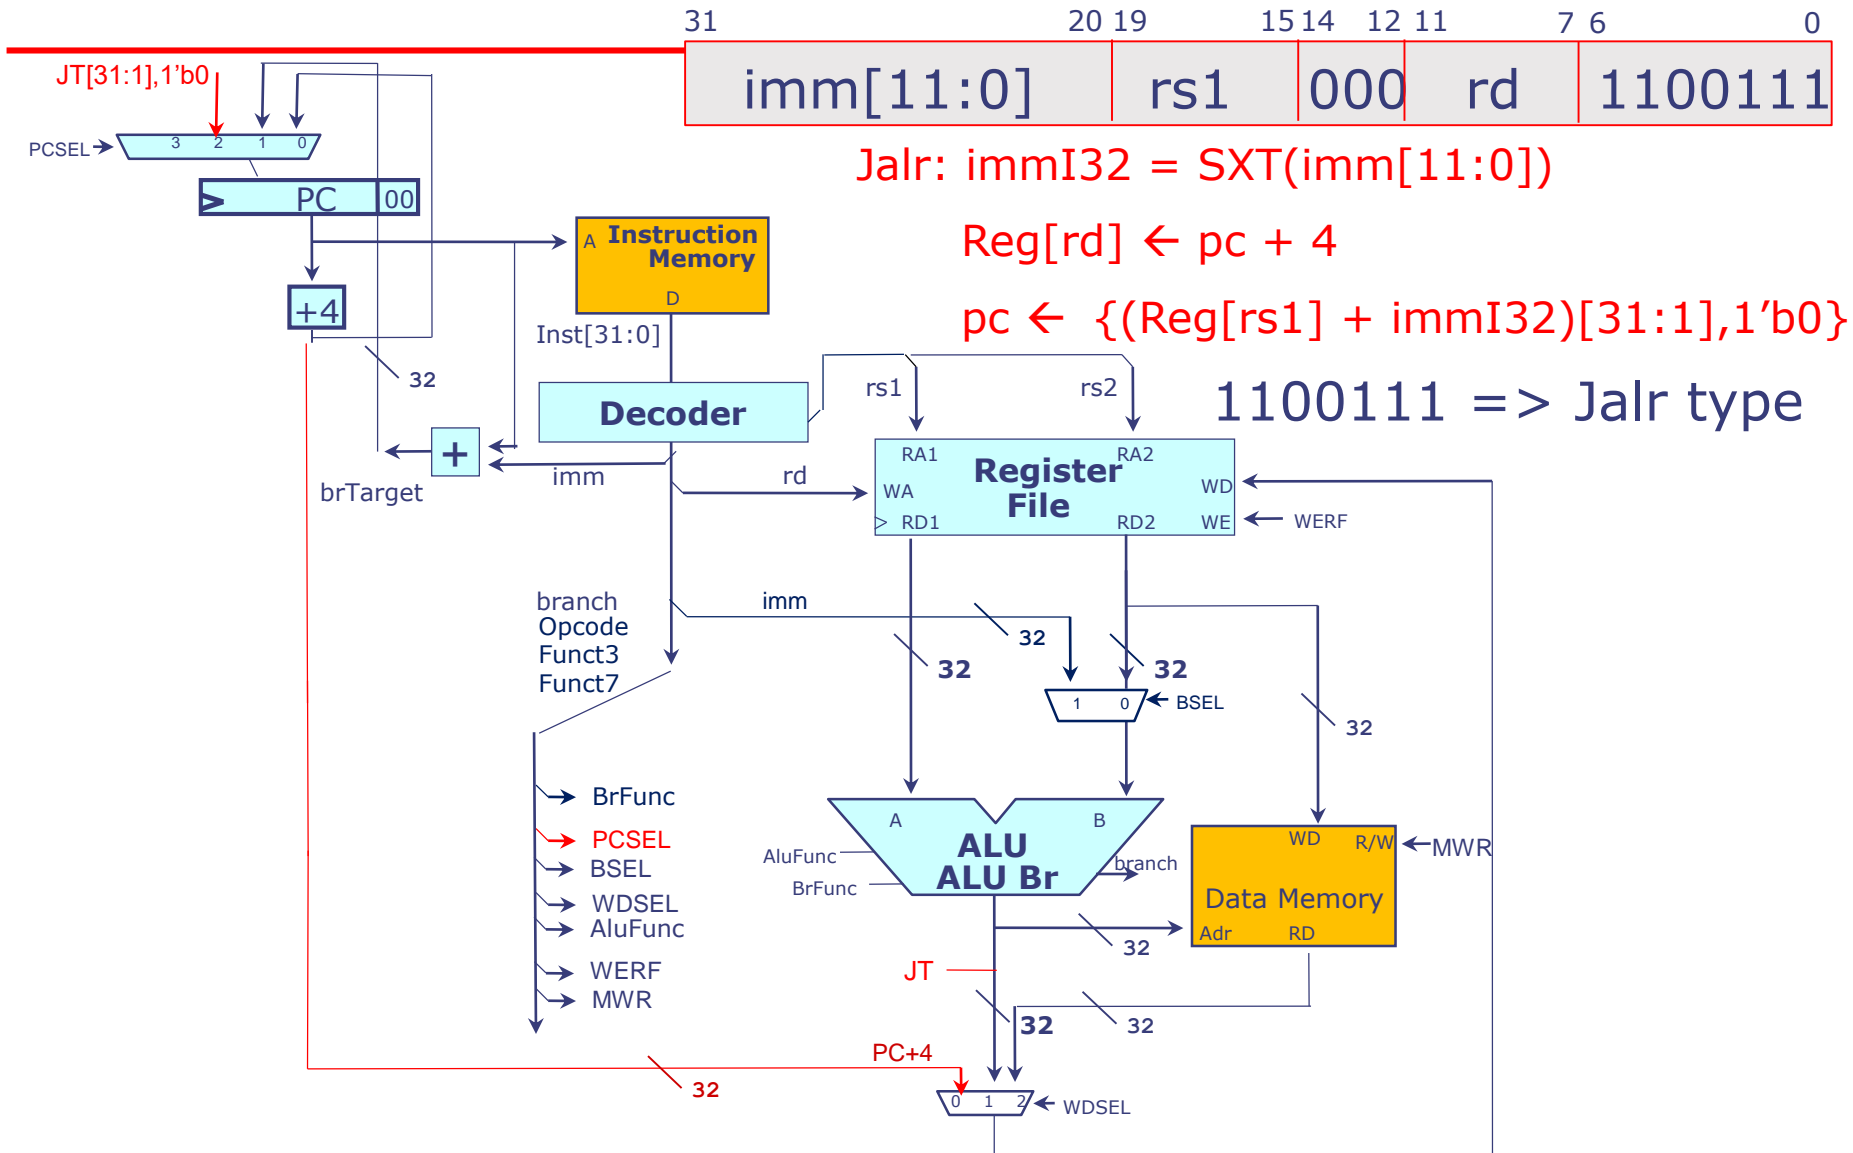
# Jalr Instruction



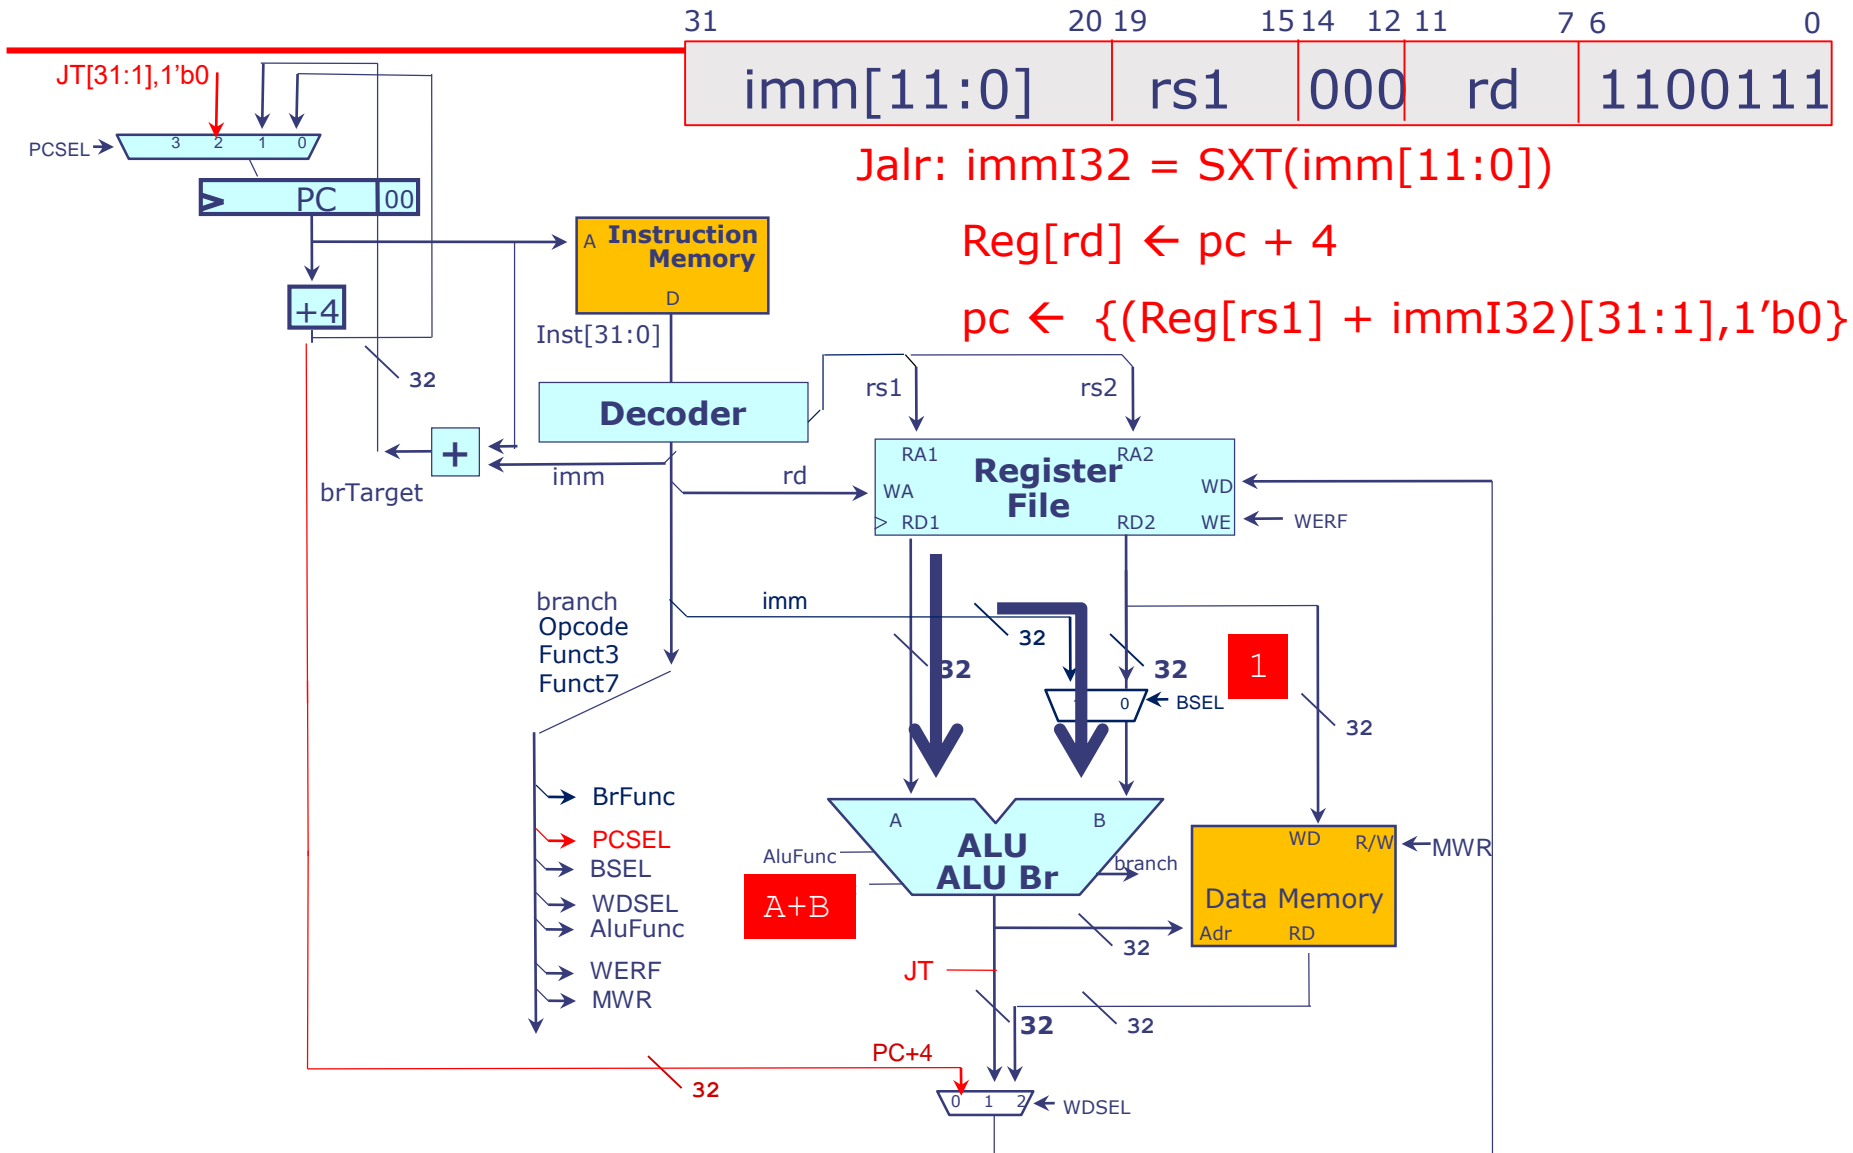
# Jalr Instruction



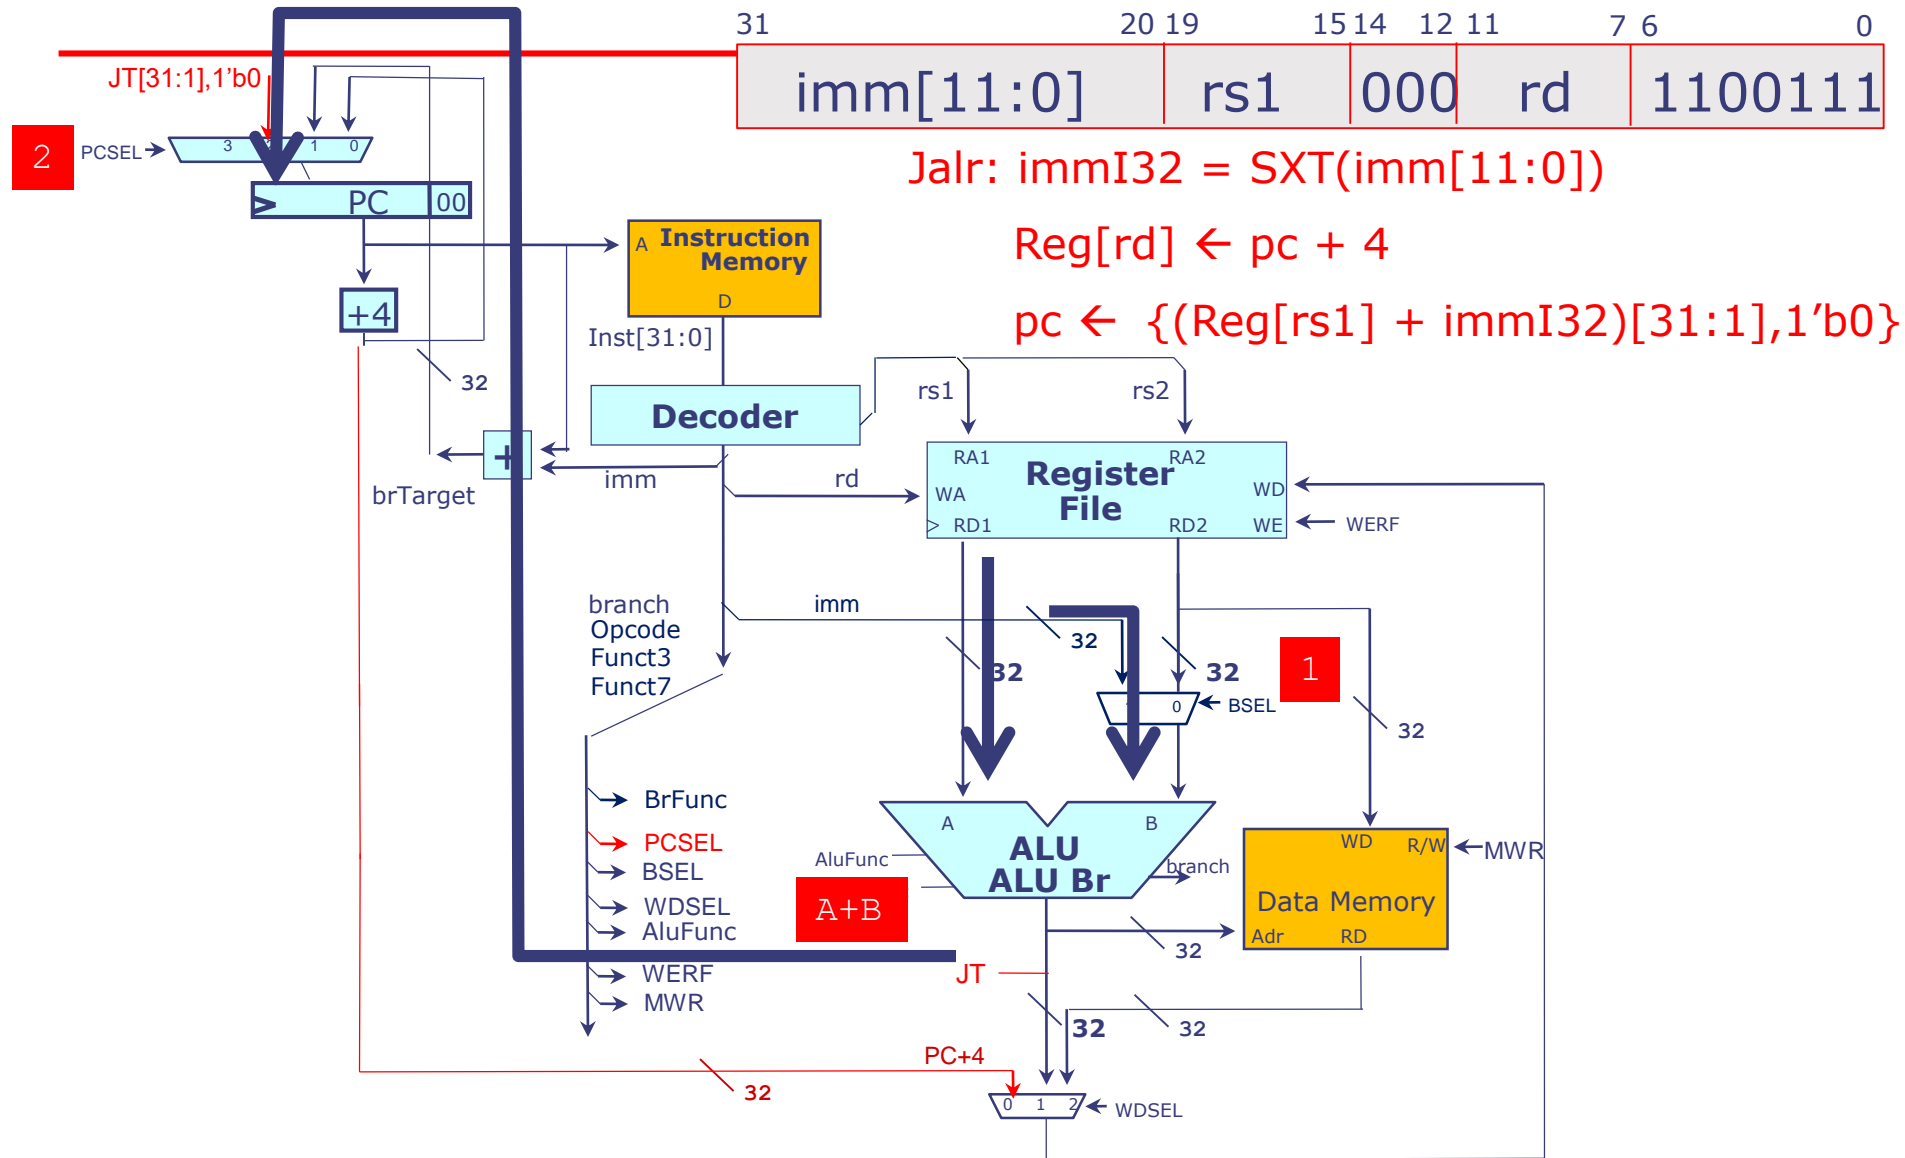
# Jalr Instruction



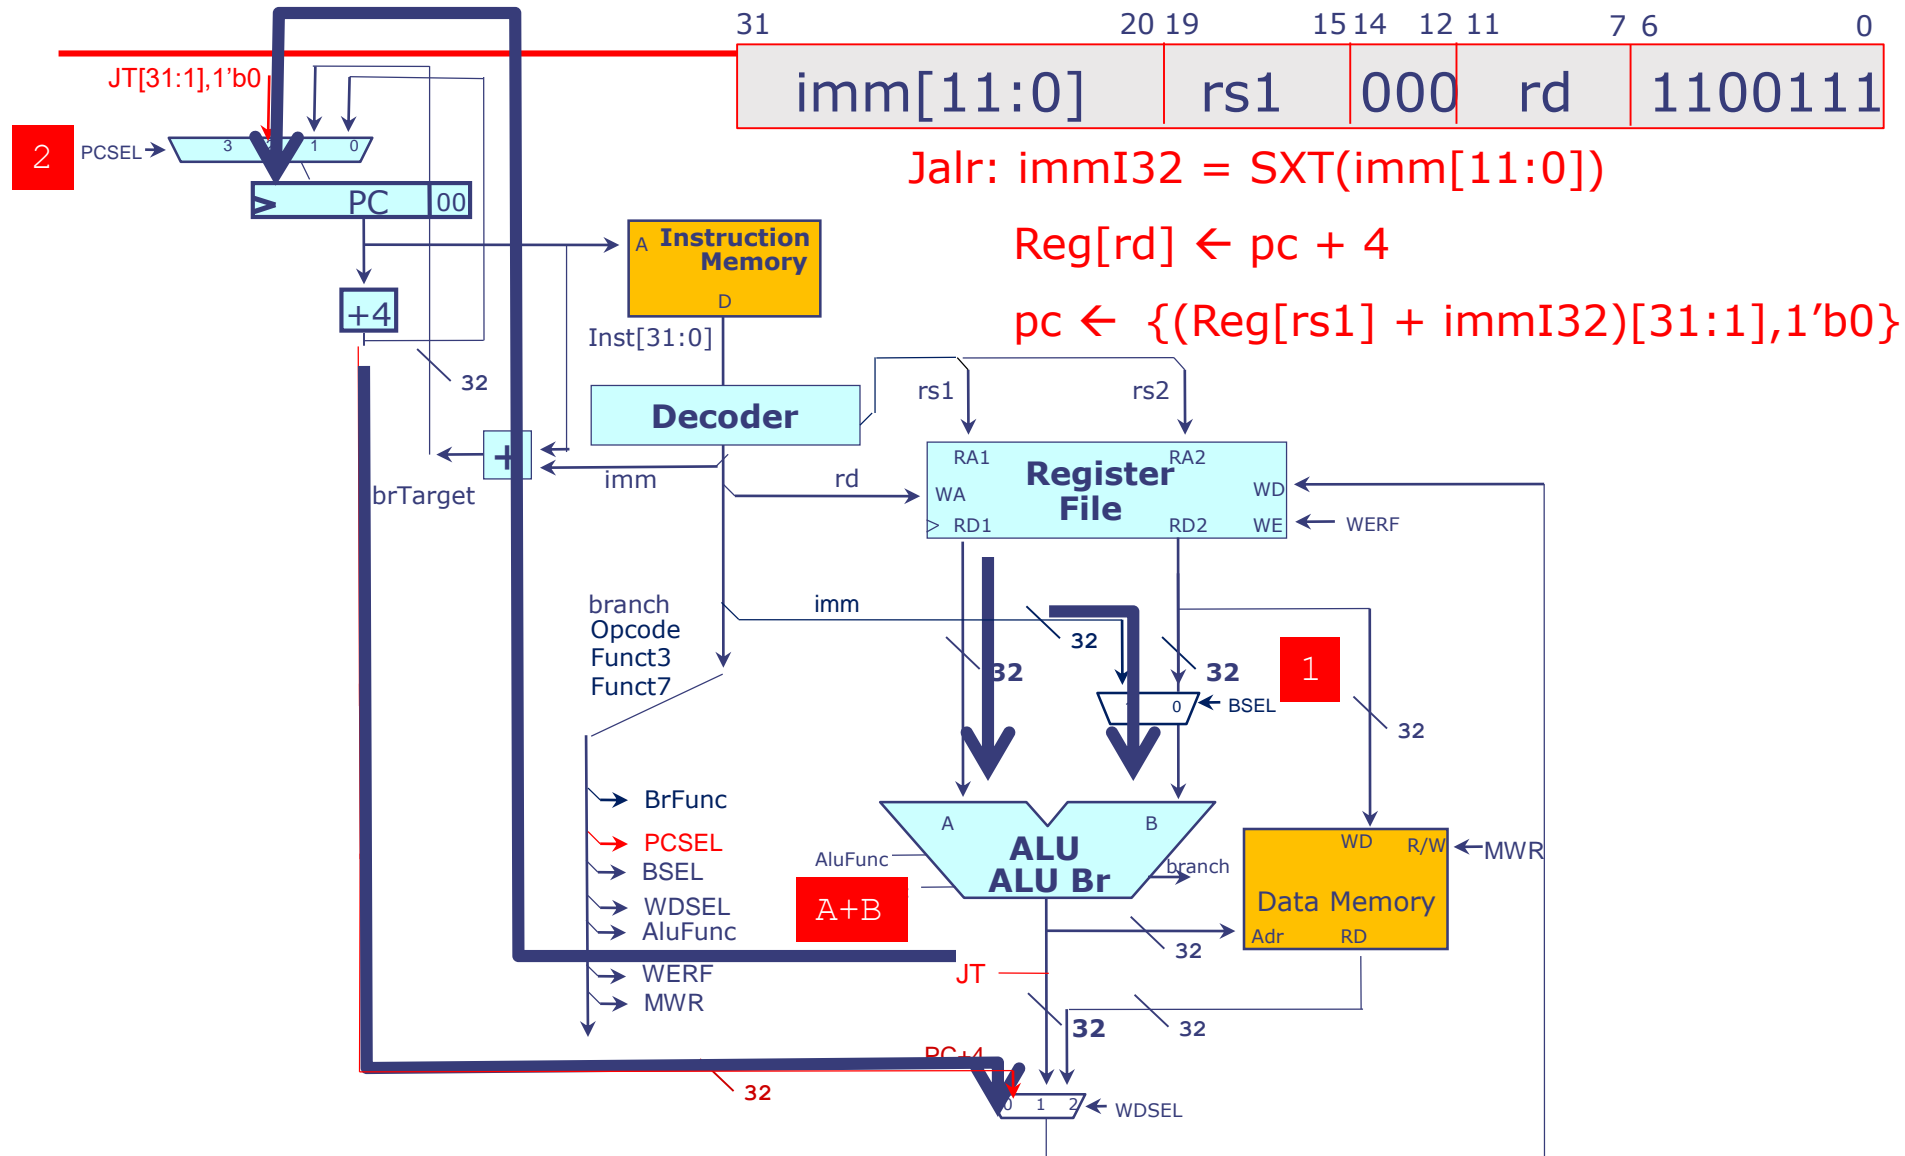
# Jalr Instruction



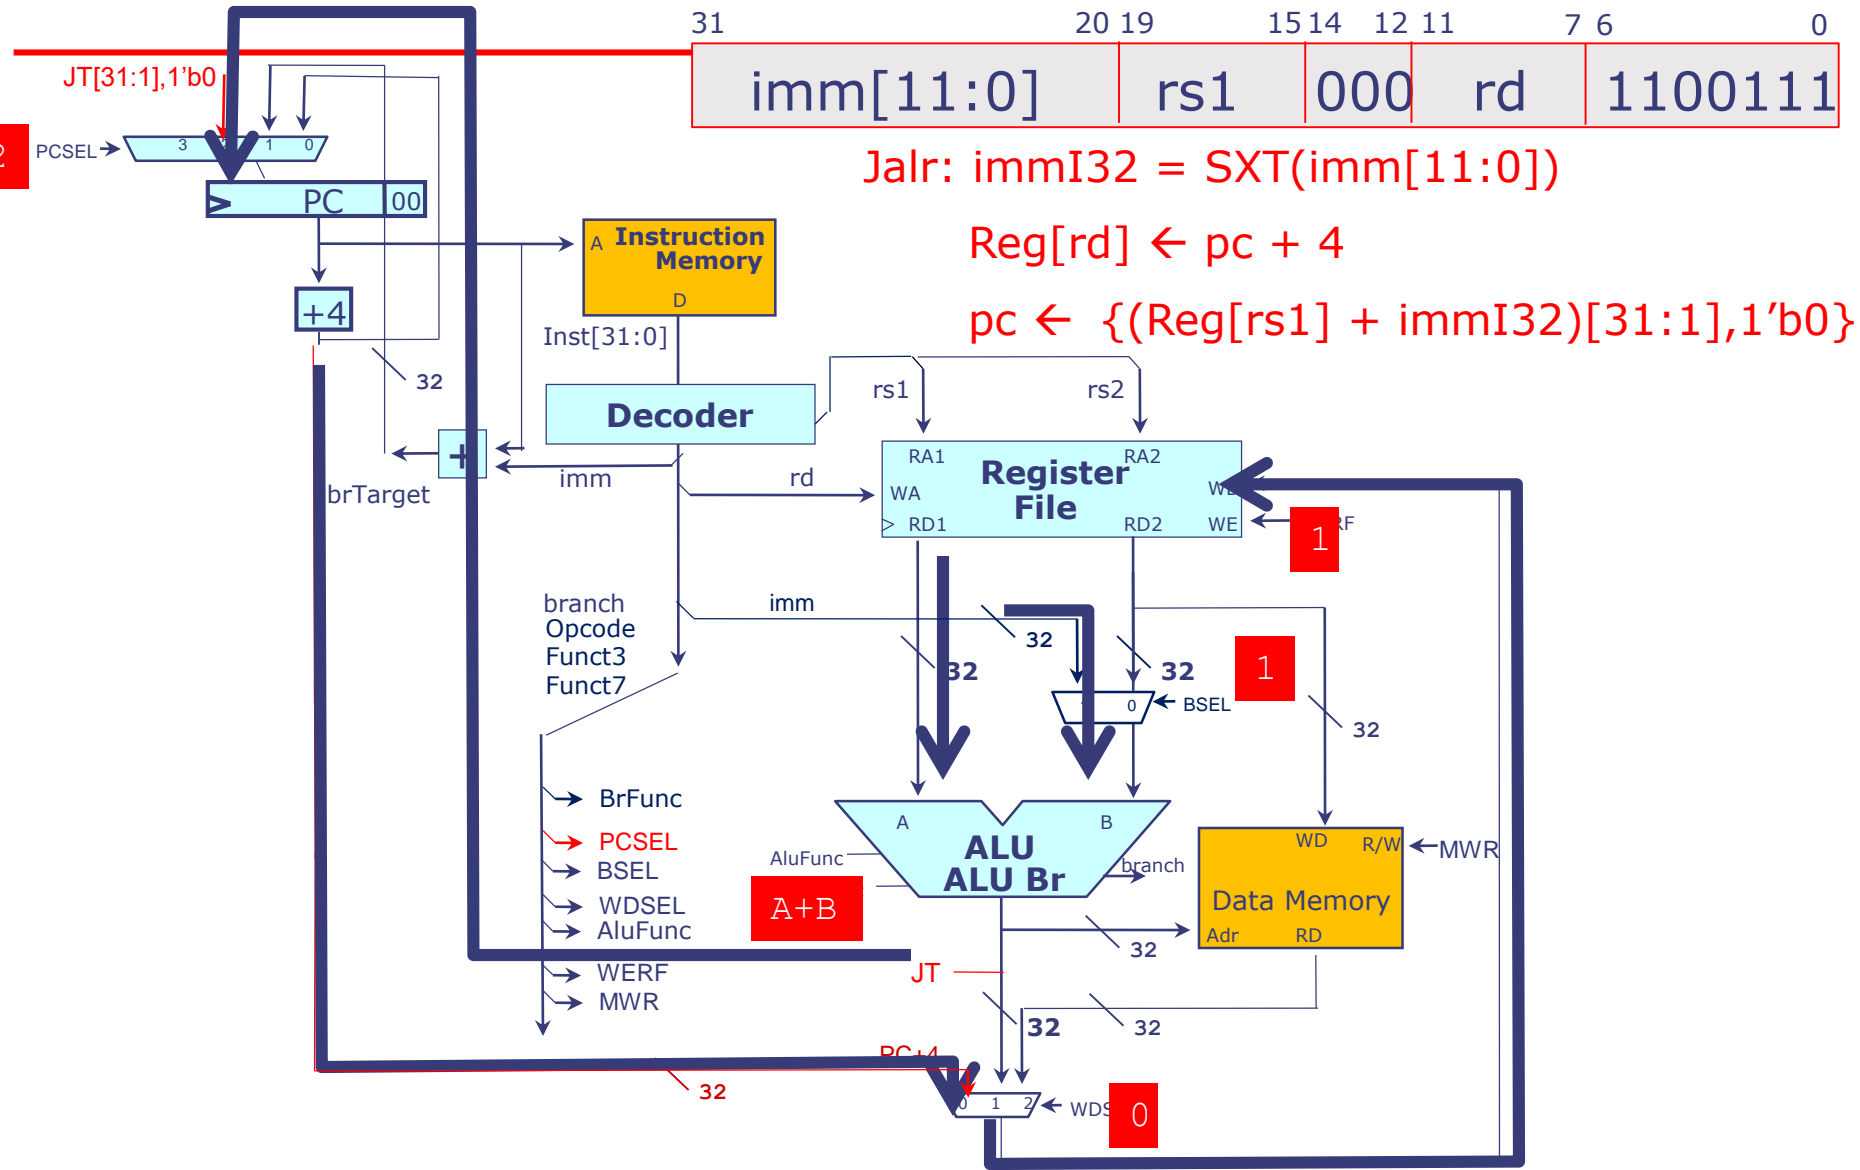
# Jalr Instruction



# Jalr Instruction



# Jalr Instruction



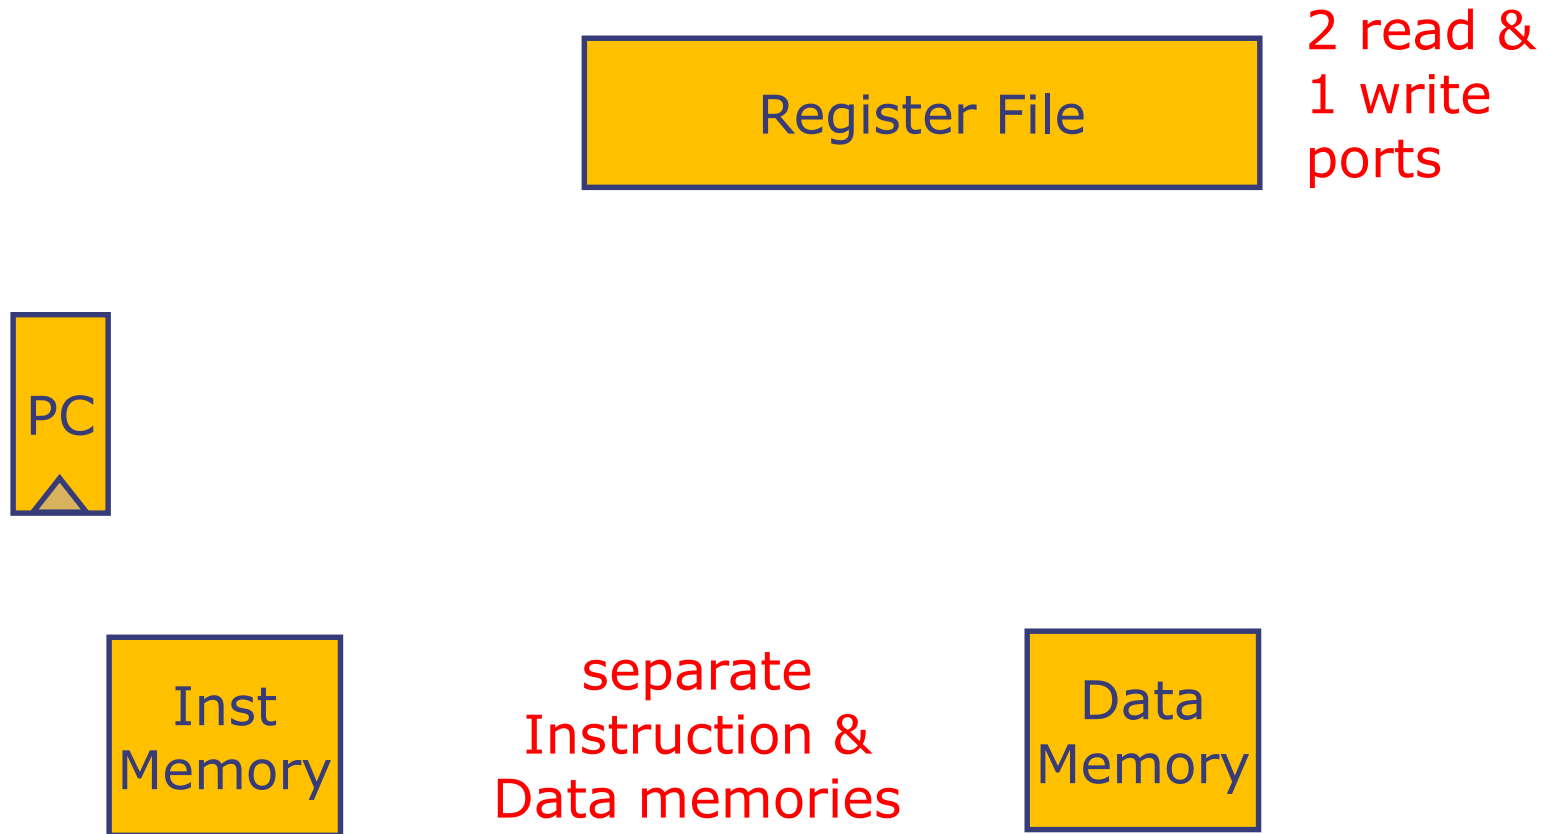
# Single-Cycle RISC-V Processor

---



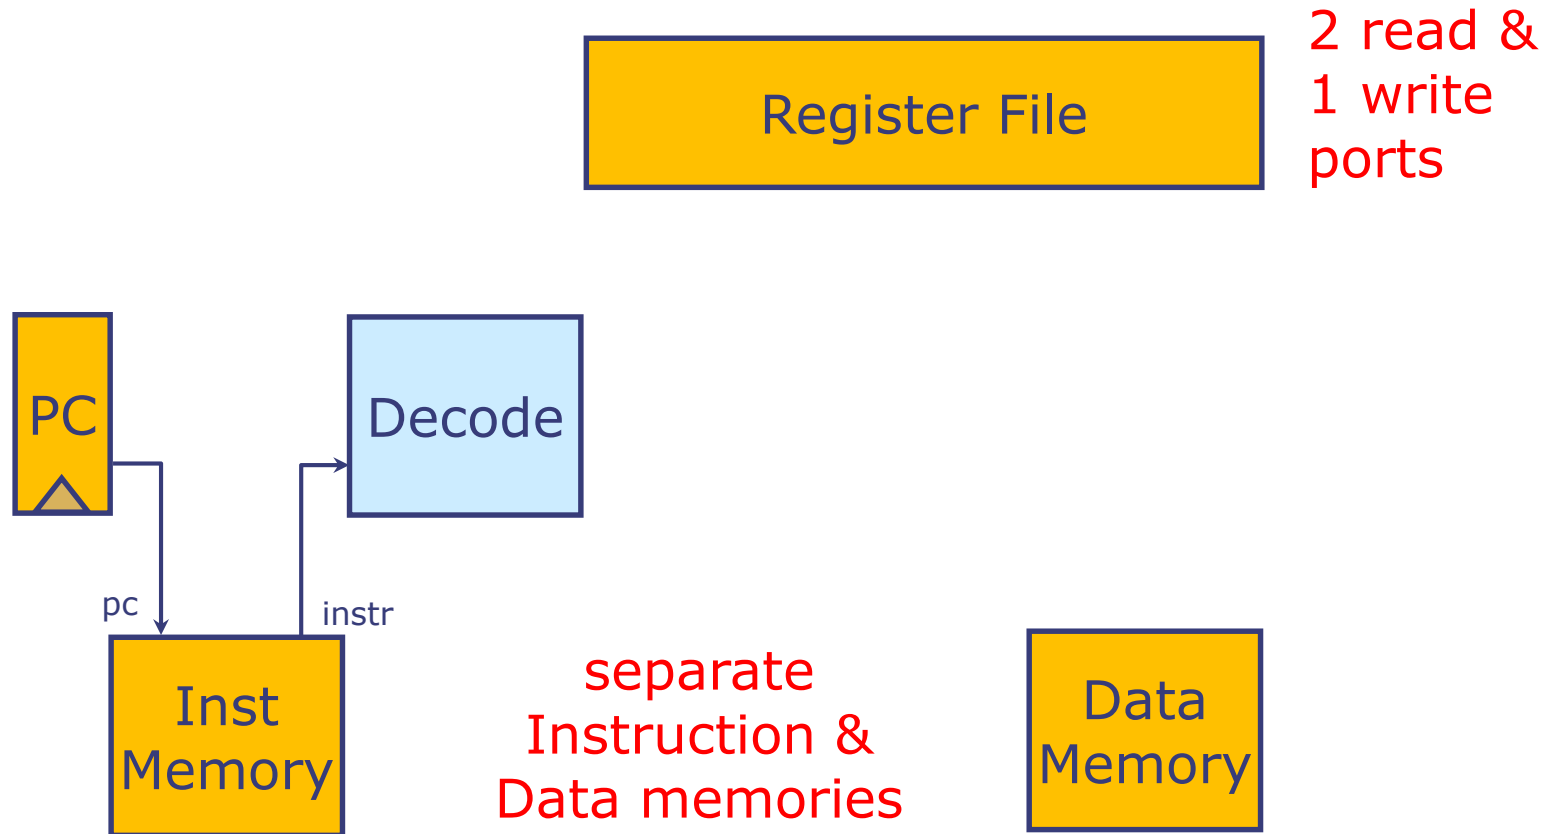
# Single-Cycle RISC-V Processor

---



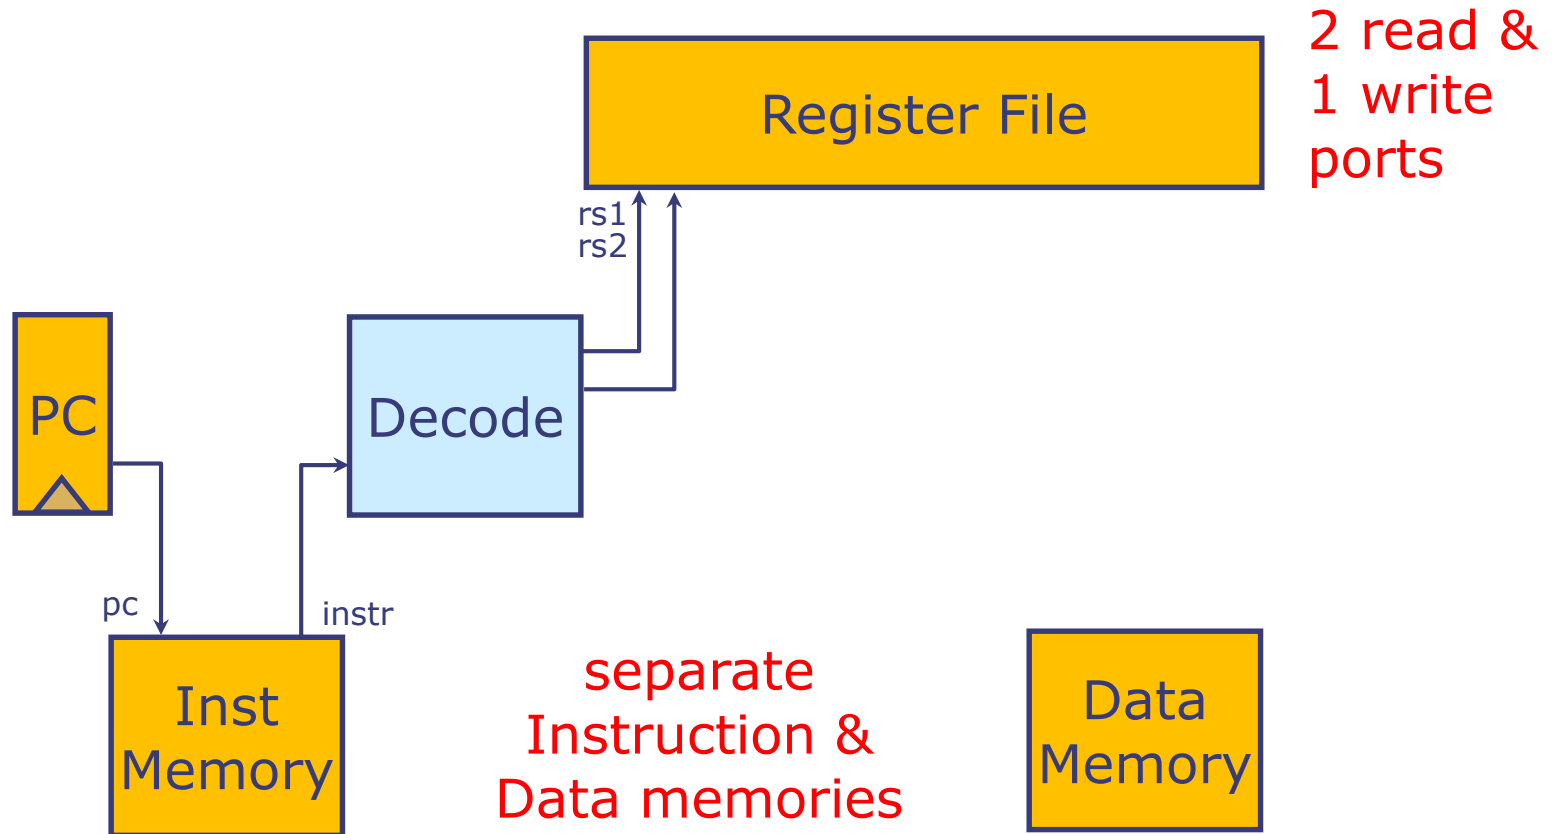
# Single-Cycle RISC-V Processor

---



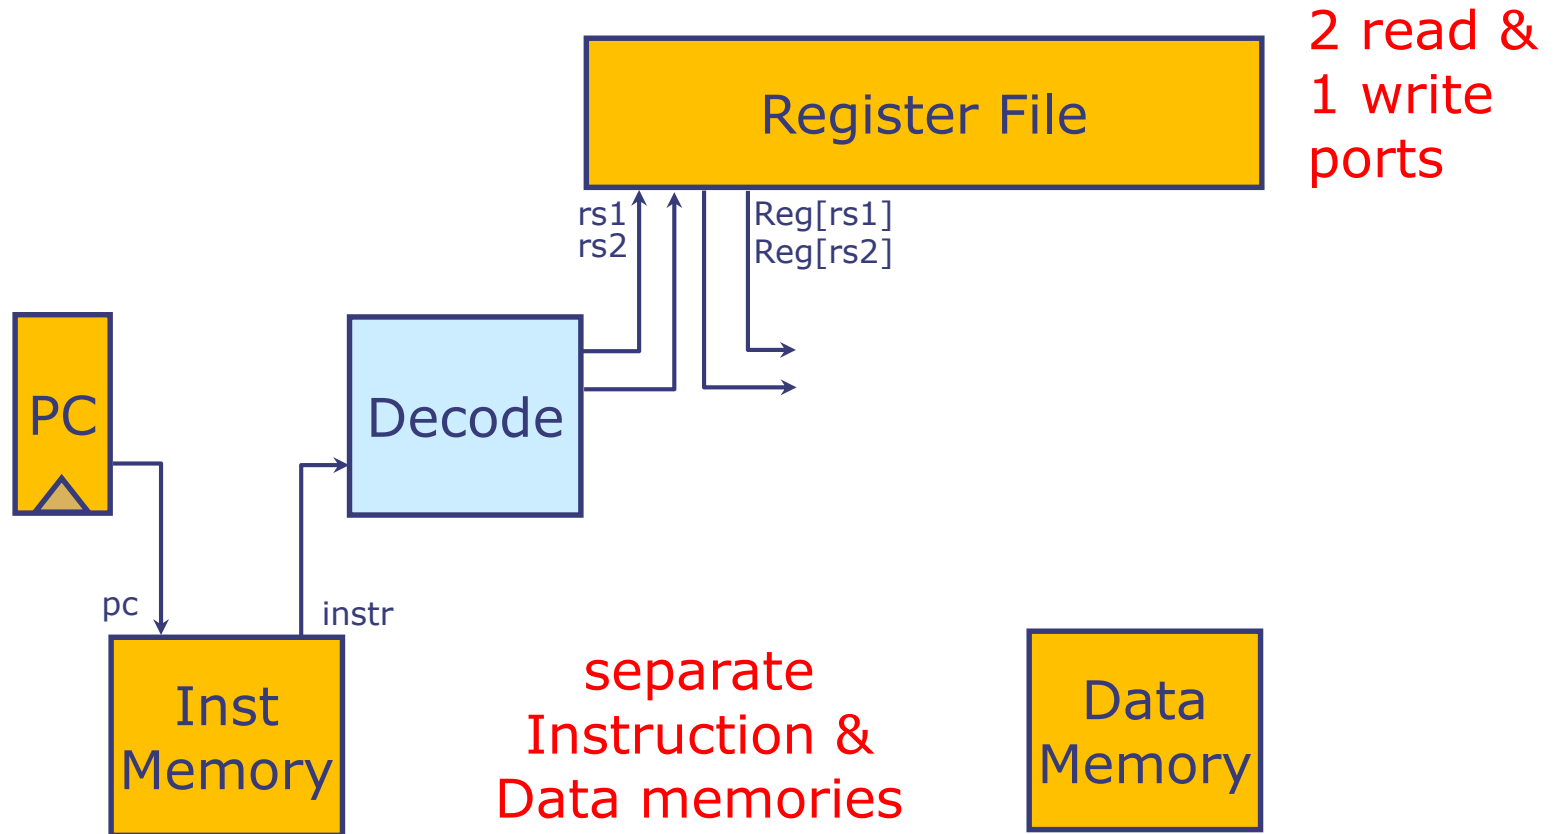
# Single-Cycle RISC-V Processor

---



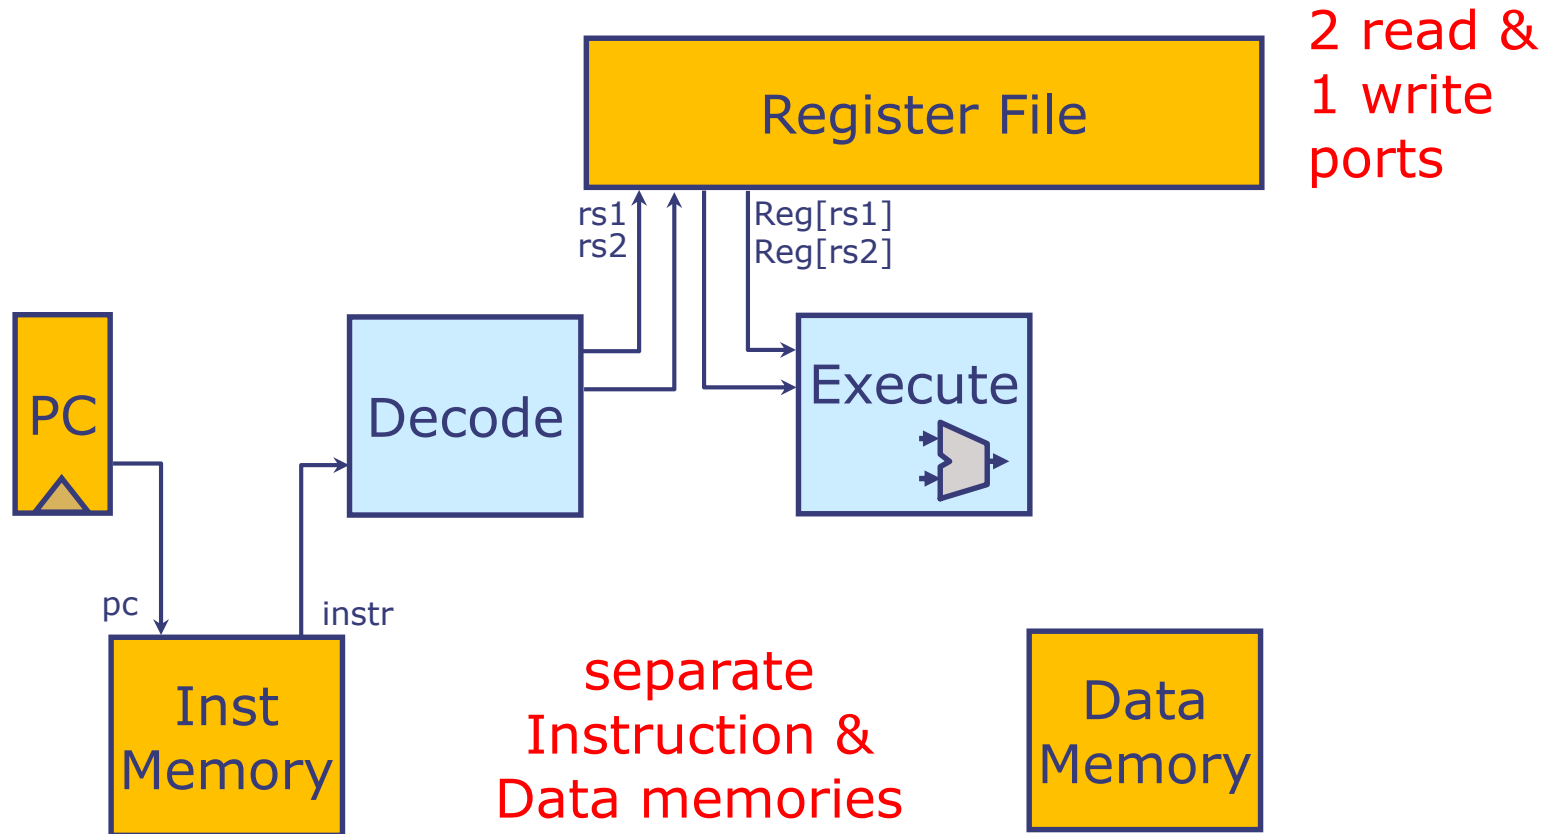
# Single-Cycle RISC-V Processor

---



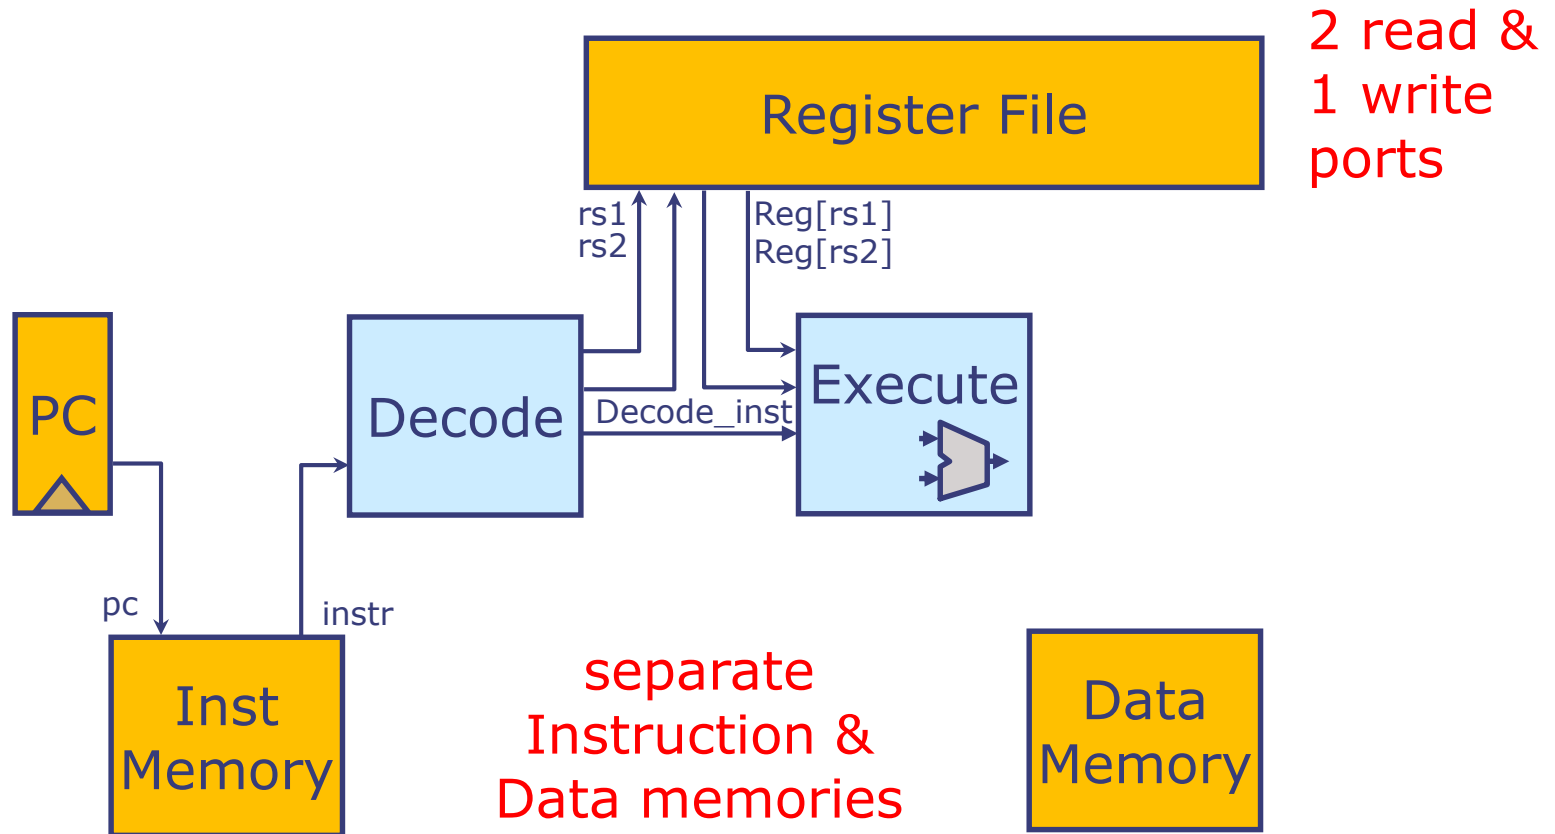
# Single-Cycle RISC-V Processor

---

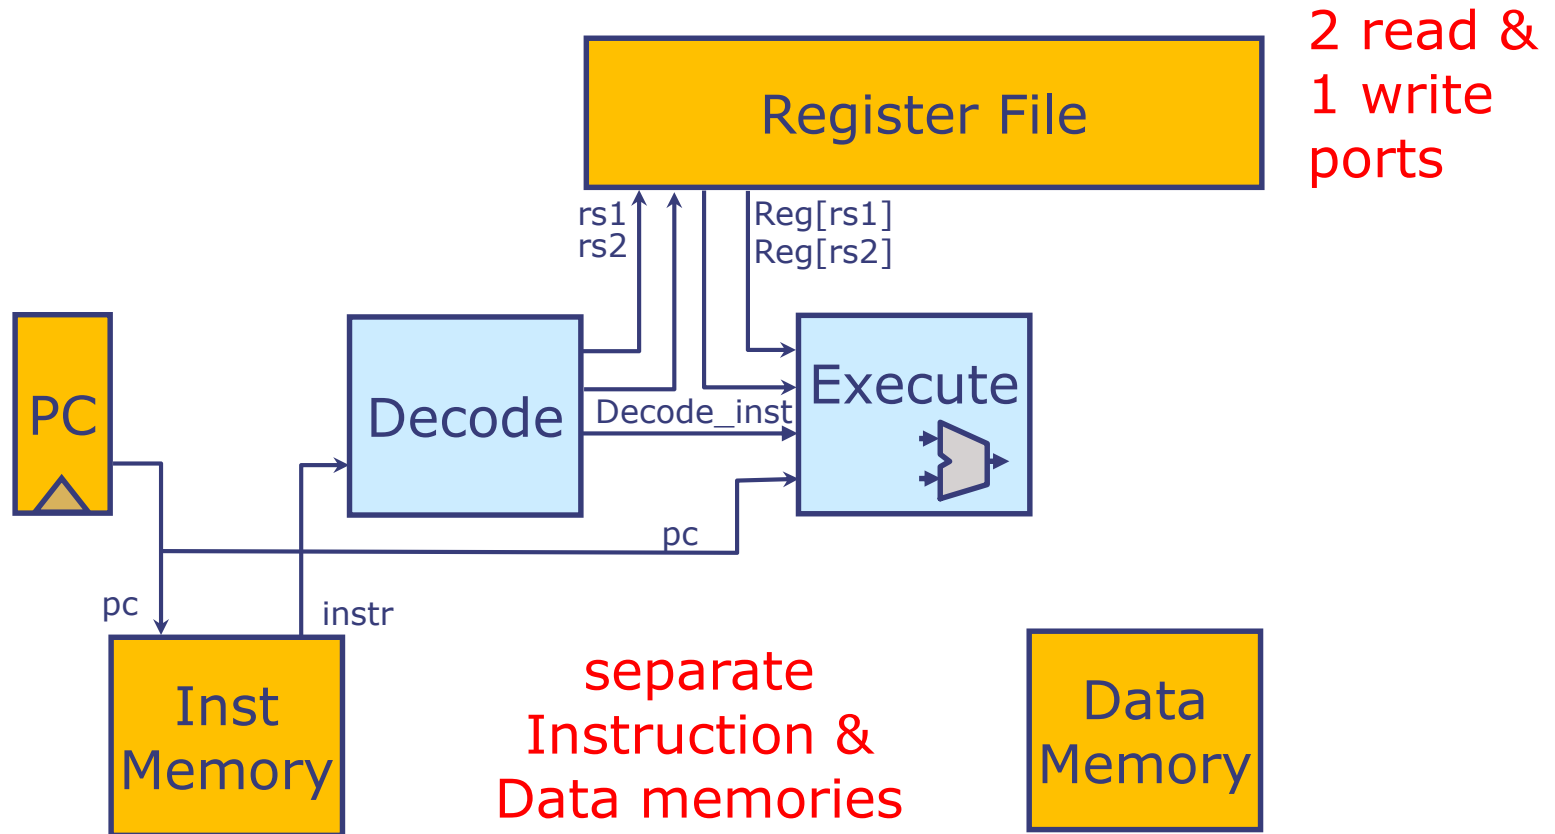


# Single-Cycle RISC-V Processor

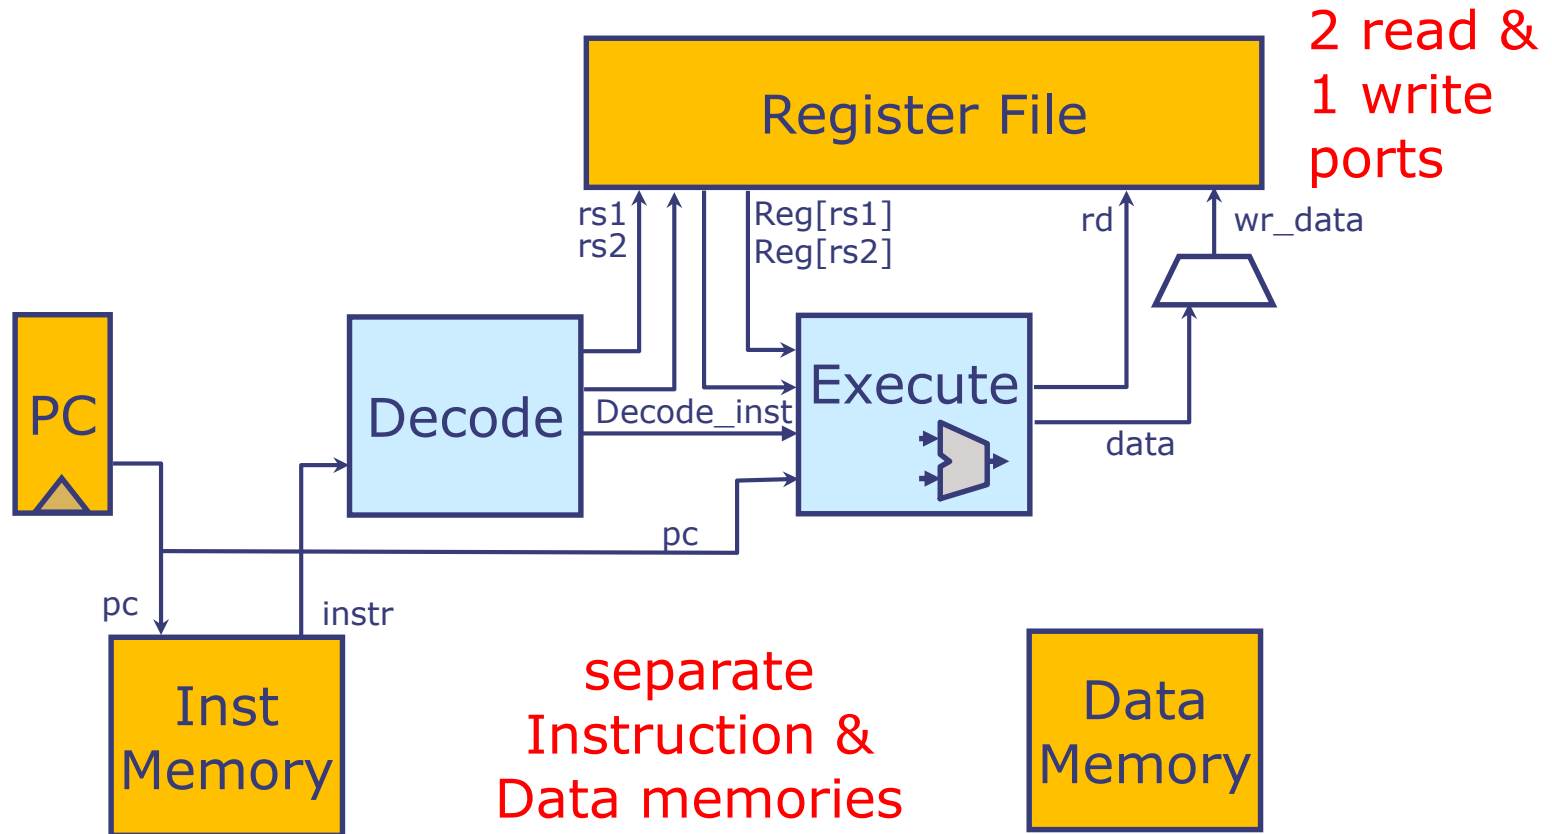
---



# Single-Cycle RISC-V Processor

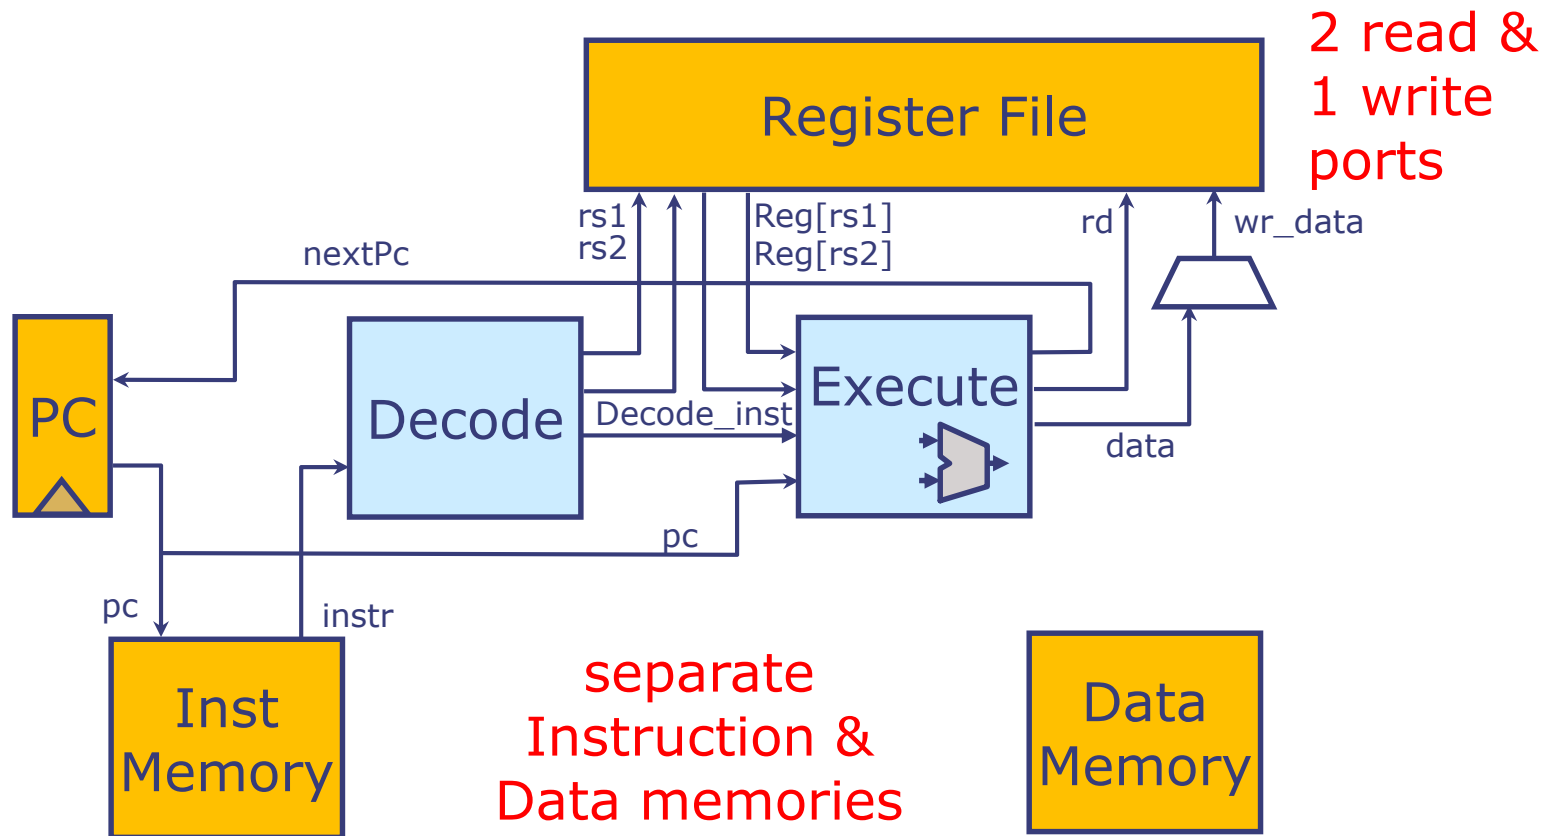


# Single-Cycle RISC-V Processor

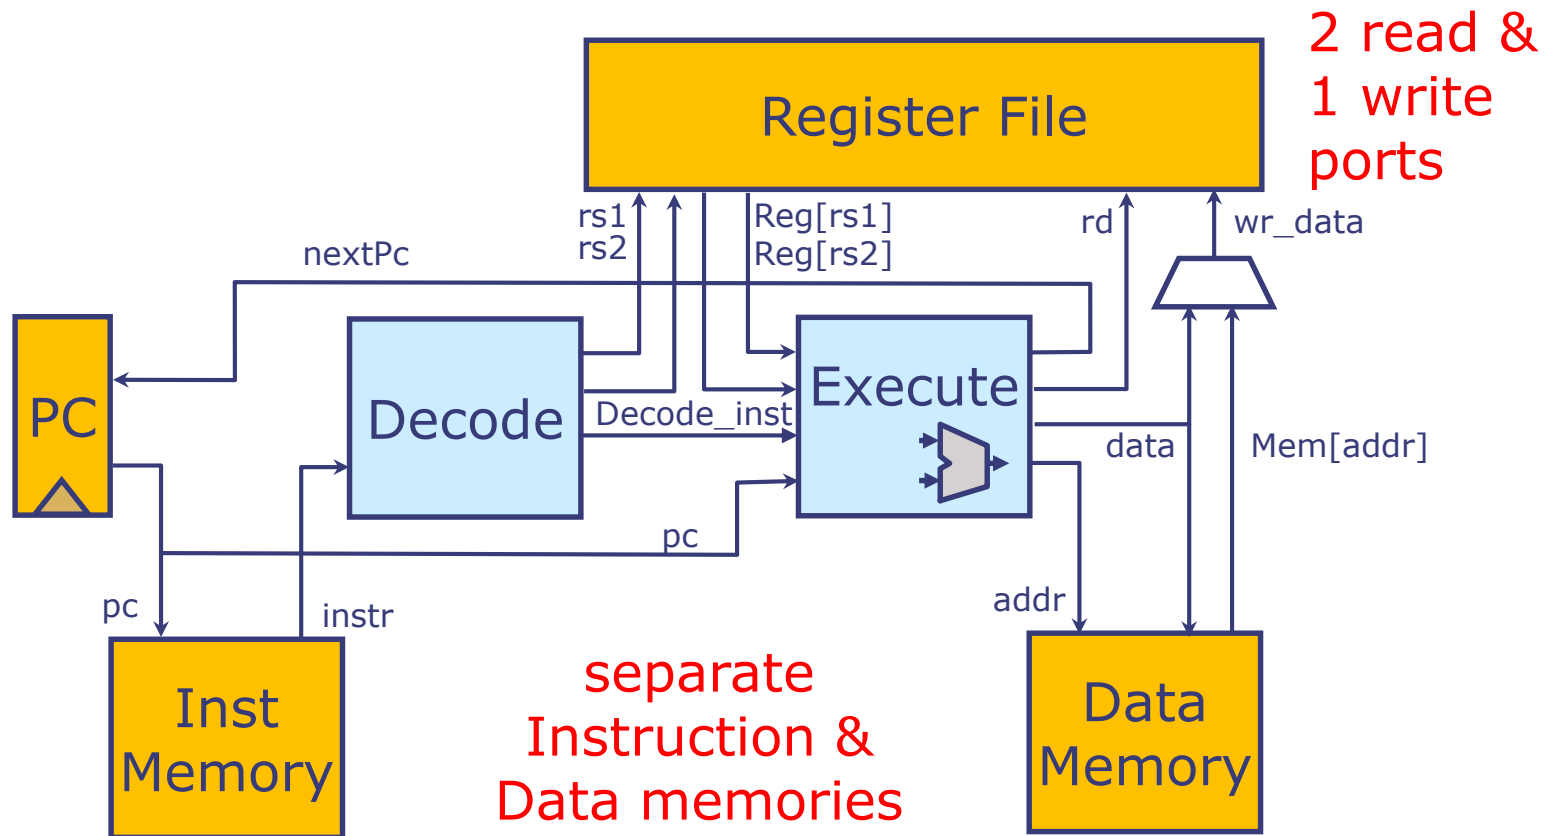





# Single-Cycle RISC-V Processor



# Single-Cycle RISC-V Processor



*Is **that** all  
there is to  
building a  
processor???*



*No.  
You've gotta print  
up all those little  
"RISC-V Inside"  
stickers.*

