

# Virtual Memory

Xuhao Chen

[cxh@mit.edu](mailto:cxh@mit.edu)


# Reminder: Operating Systems

- Goals of OS:

- **Protection and Privacy**: Process isolation
- **Abstraction**: Hide away details of underlying hardware
- **Resource Management**: Controls how processes share hardware resources (CPU, memory, disk, etc.)

- Key enabling technologies:


- User mode + supervisor mode
- Exceptions to safely transition into supervisor mode
- **Virtual Memory** to abstract the storage resources of the machine



# Virtual Memory (VM) Systems

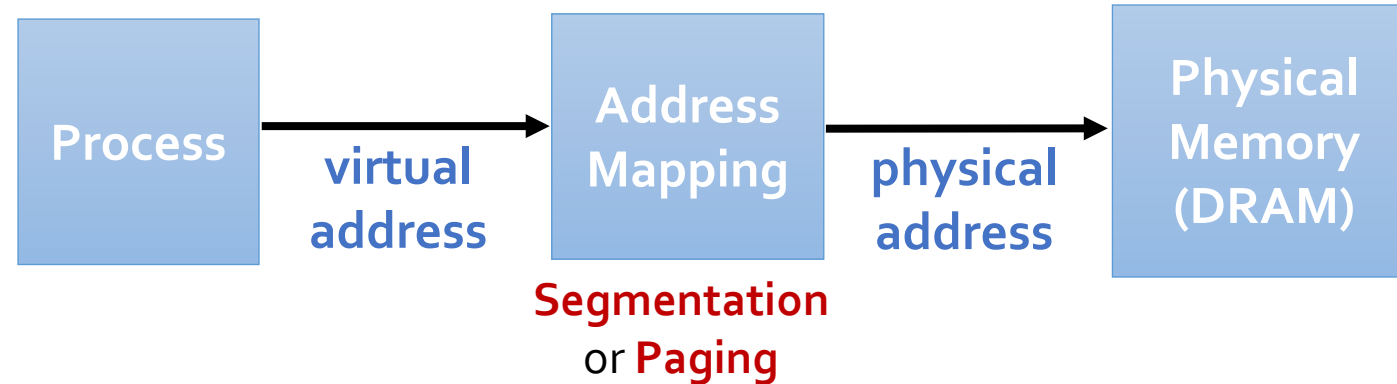
*Illusion of a large, private, uniform store*

- Protection & Privacy
  - Each process has a private address space
- Demand Paging
  - Use main memory as a cache of disk
  - Enables running programs larger than main memory
  - Hides differences in machine configuration



The price of VM is **address translation** on each memory reference

# Virtual Address vs. Physical Address



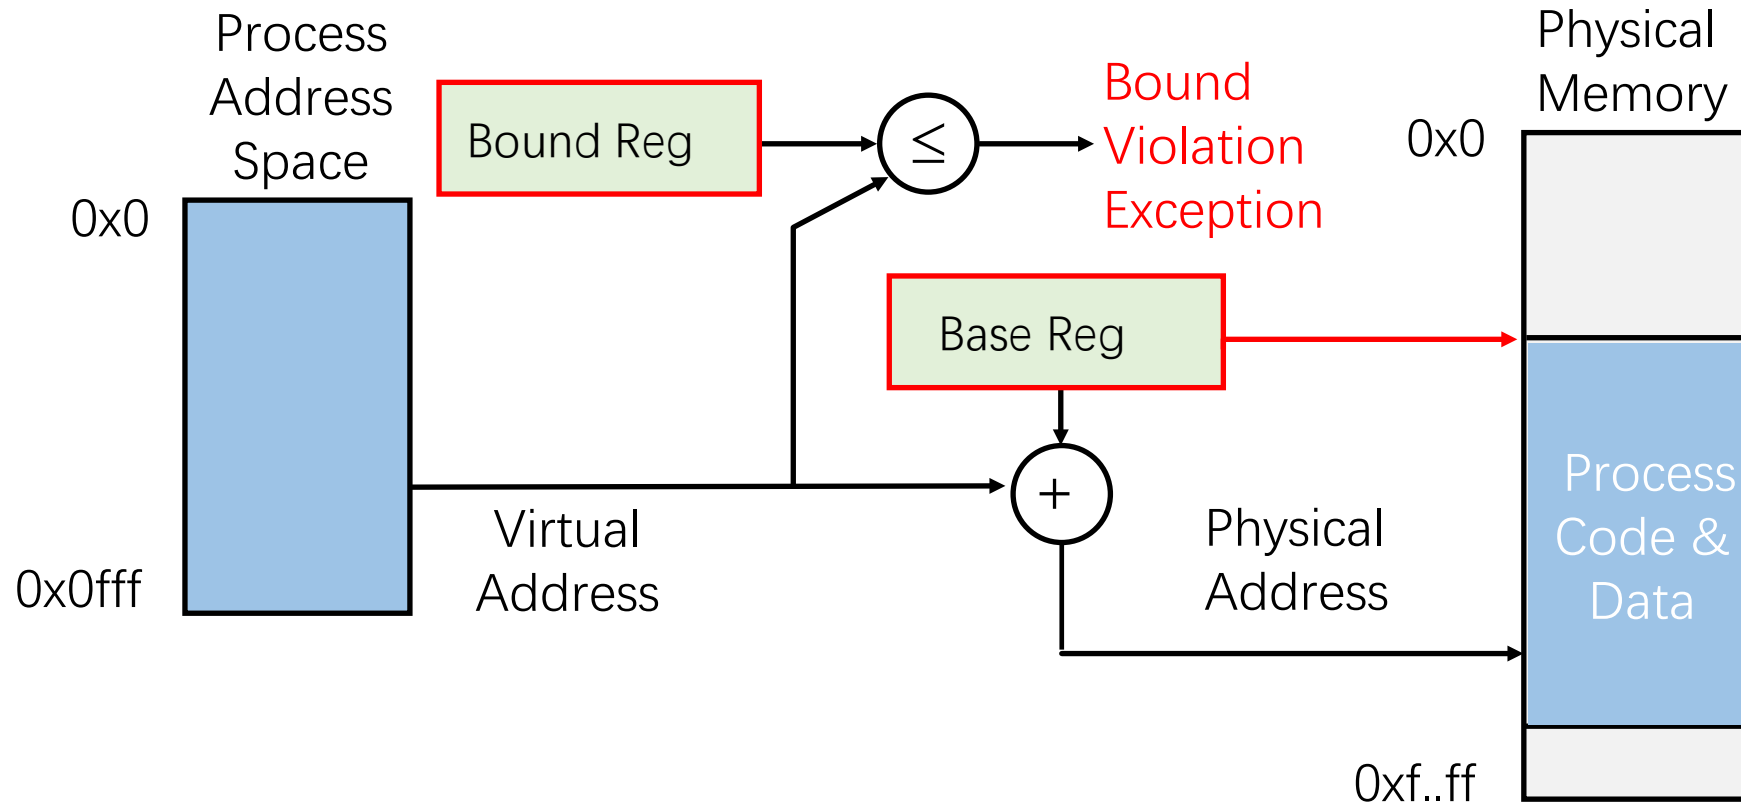
- **Virtual address**

- Address generated by the process
- Specific to the process's private address space

- **Physical address**

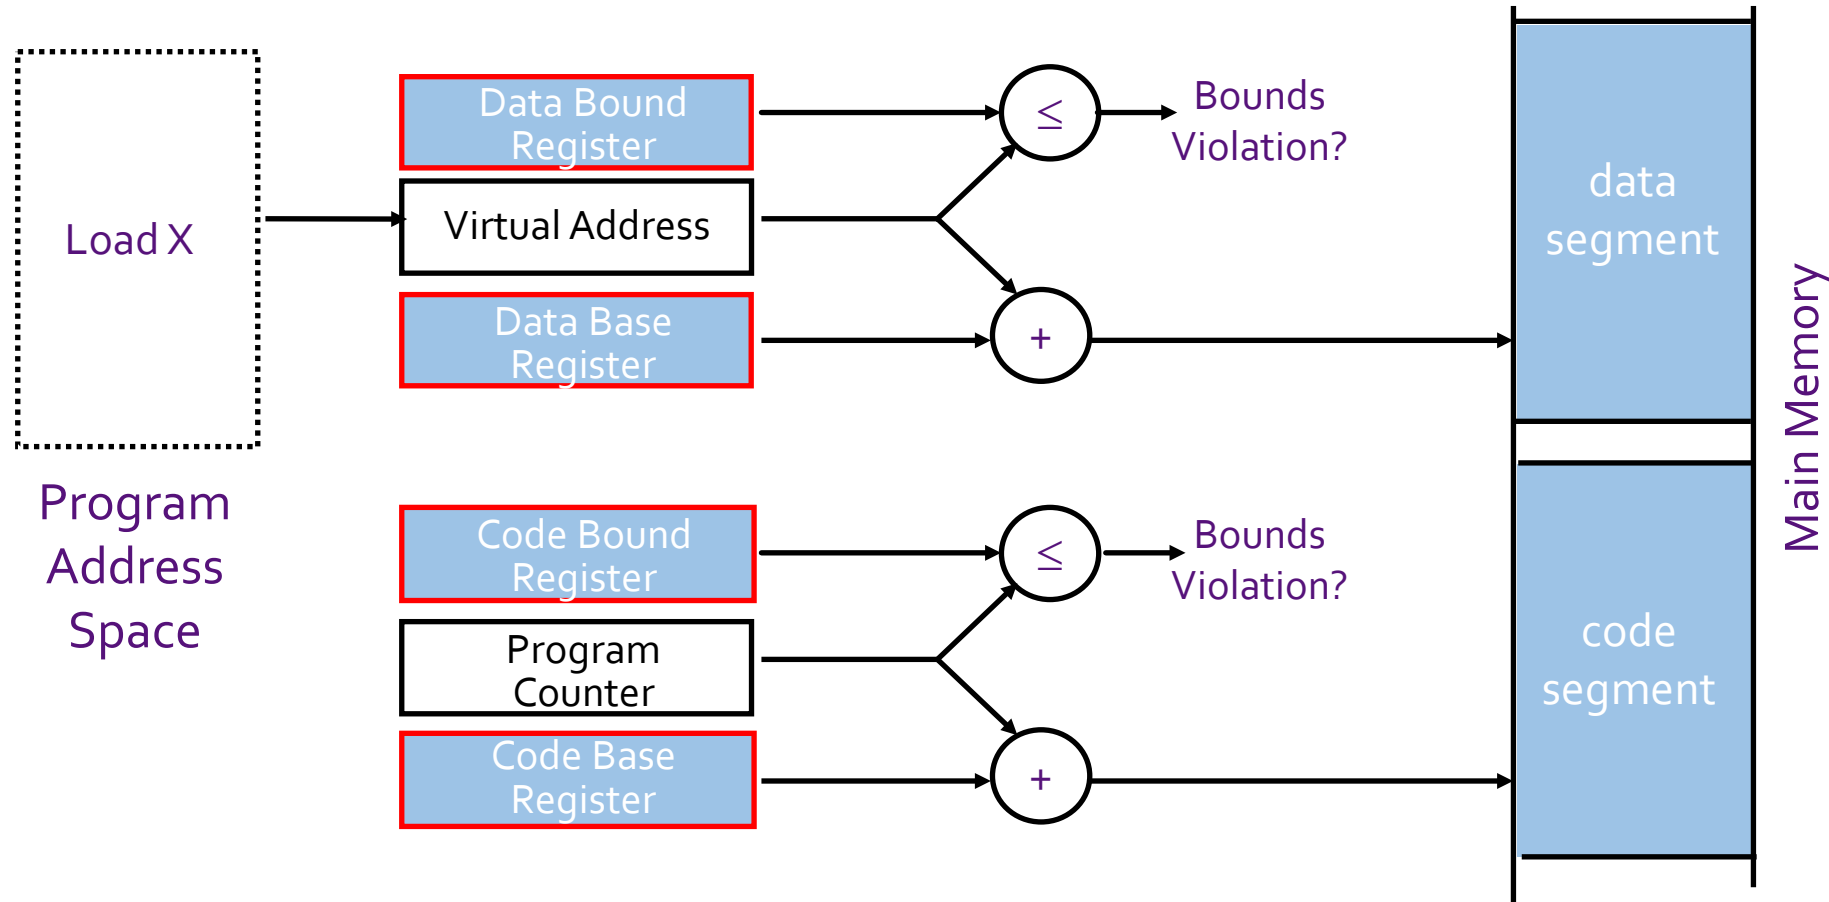
- Address used to access physical (hardware) memory
- **OS** specifies mapping of virtual addresses into physical addresses

# Segmentation: *Base-and-Bound* Address Translation



- Each program's data is allocated in a contiguous segment of physical memory
- Physical address = Virtual Address + Segment Base
- Bound register provides safety and isolation
- Base and Bound registers should not be accessed by user programs (only accessible in supervisor mode)

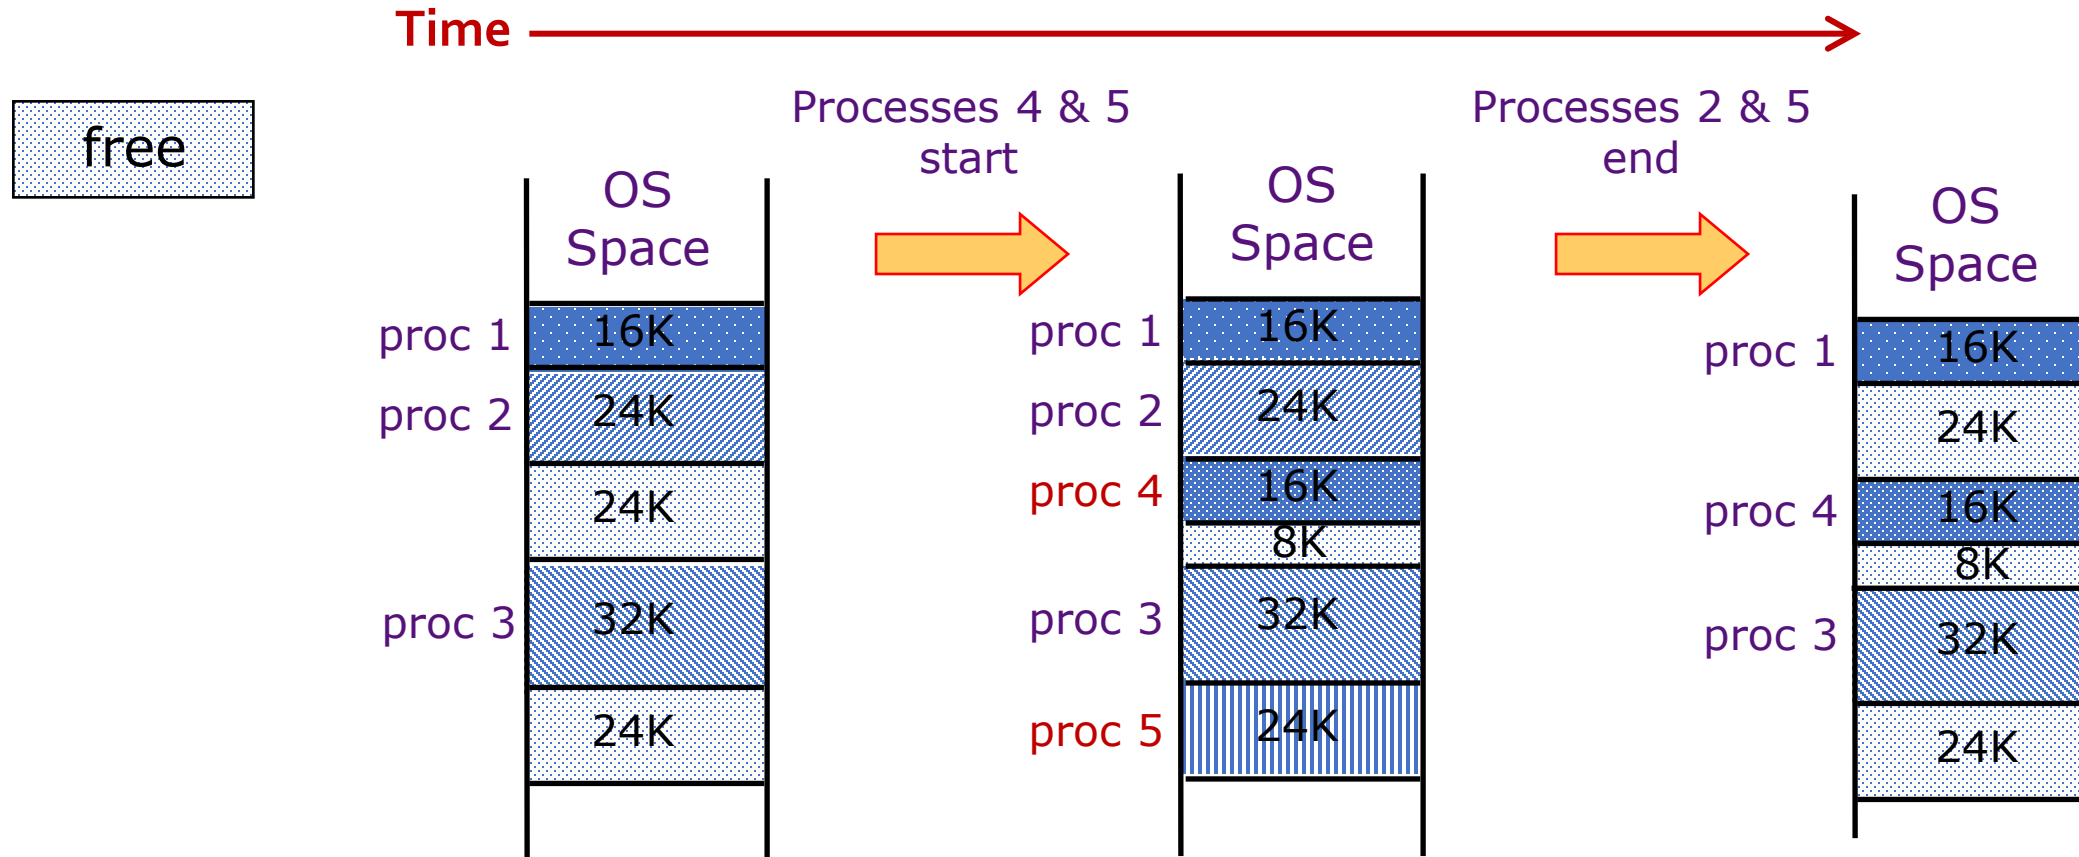
# Separate Segments for Code and Data



*Pros of this separation?*

- (1) Prevents buggy program from overwriting code
- (2) Multiple processes can share code segment


# Memory Fragmentation



- As processes start and end, storage is “**fragmented**”
- At some point, segments have to be moved around to compact the free space

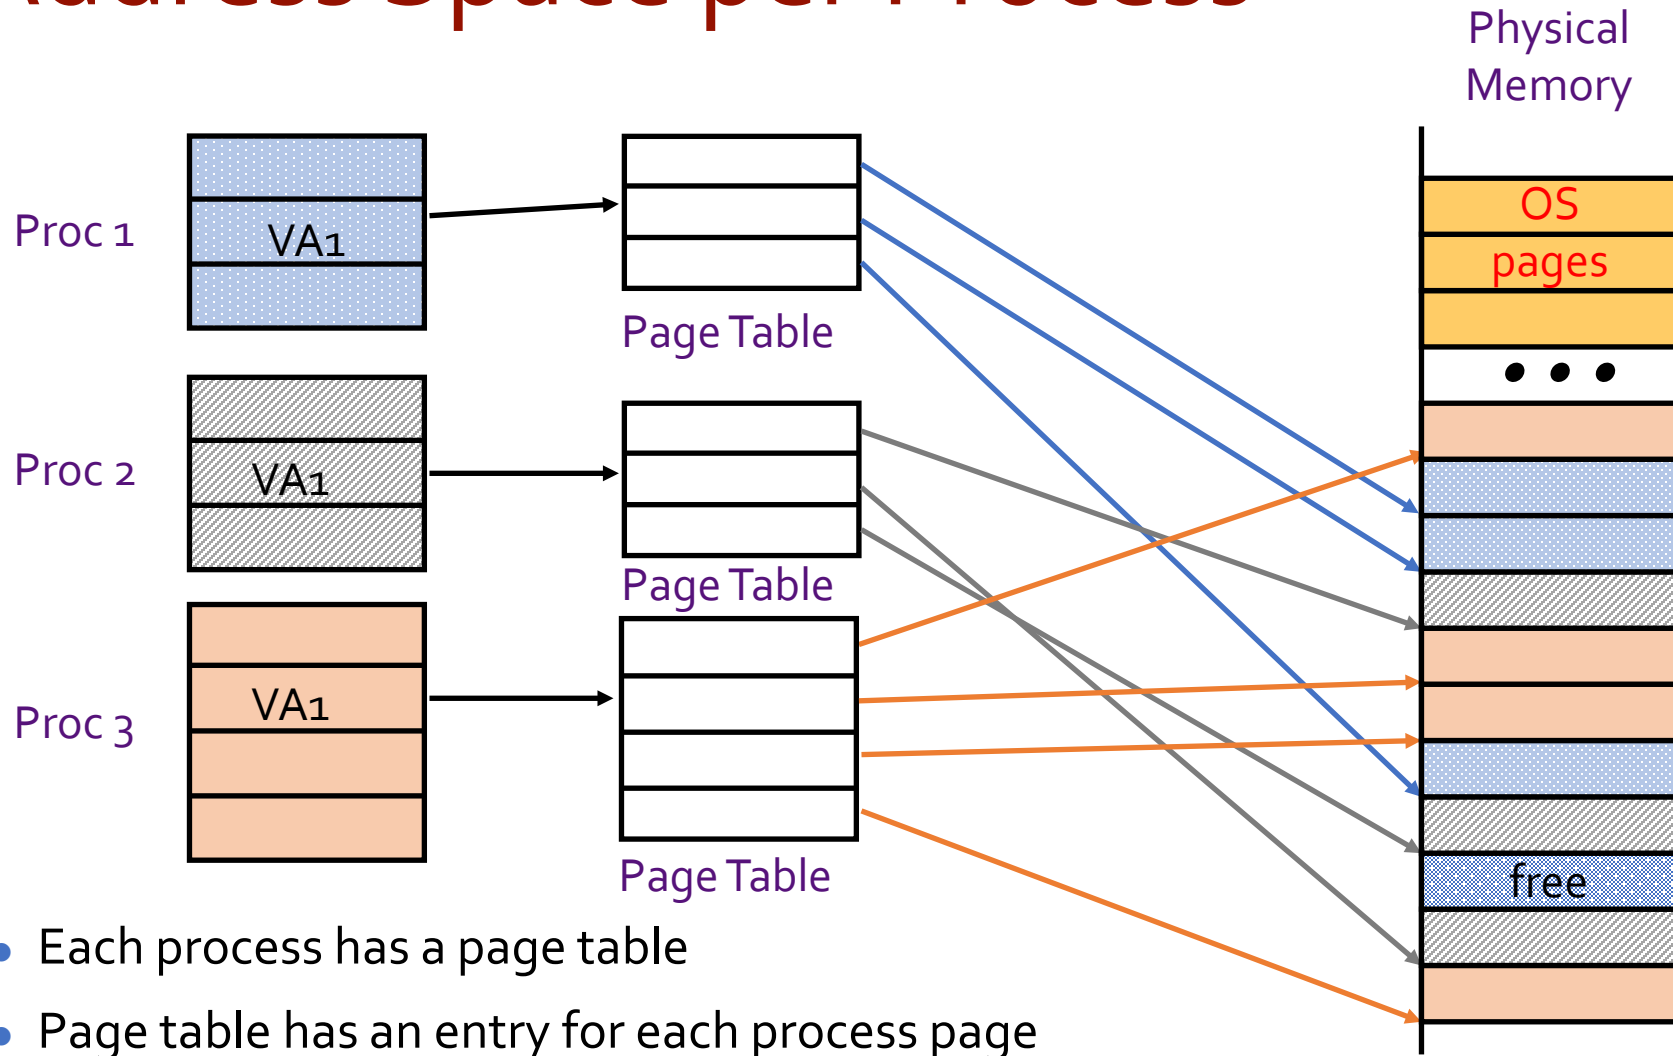
# Paged Memory Systems

- Divide physical memory in *fixed-size blocks* called **pages**
  - Typical page size: 4KB
- Interpret each virtual address as a pair **<virtual page number, offset>**
- Use a **Page Table** to translate from virtual to physical page numbers
  - Page table contains the physical page number (i.e., starting physical address) for each virtual page number





# Private Address Space per Process



Page tables make it possible to store the pages of a program **non-contiguously**

# Paging vs. Segmentation

- Pros of paging

- Paging avoids fragmentation
- Paging enables demand paging
  - Allows programs to use more VM than the machine's PM

- Cons of paging

- Page tables are much larger than base & bound registers

- What if all the pages can't fit in DRAM?

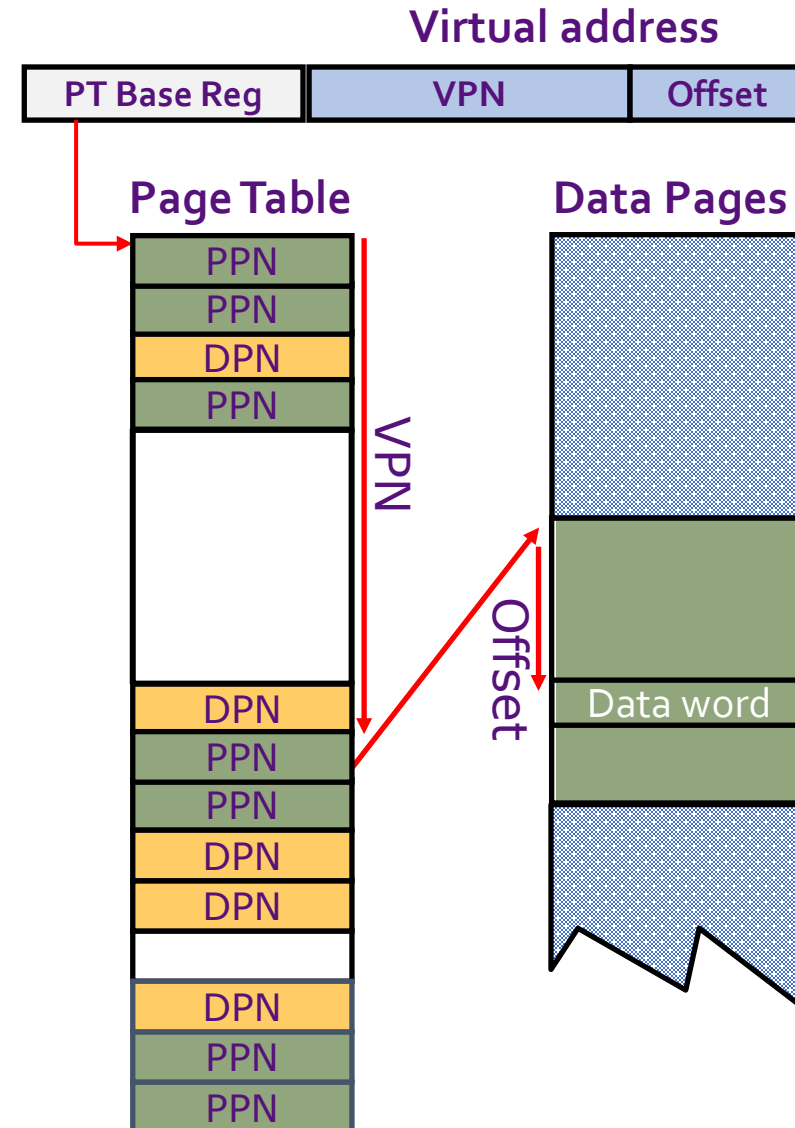
- Where do we store the page tables?

# Two Remaining Questions


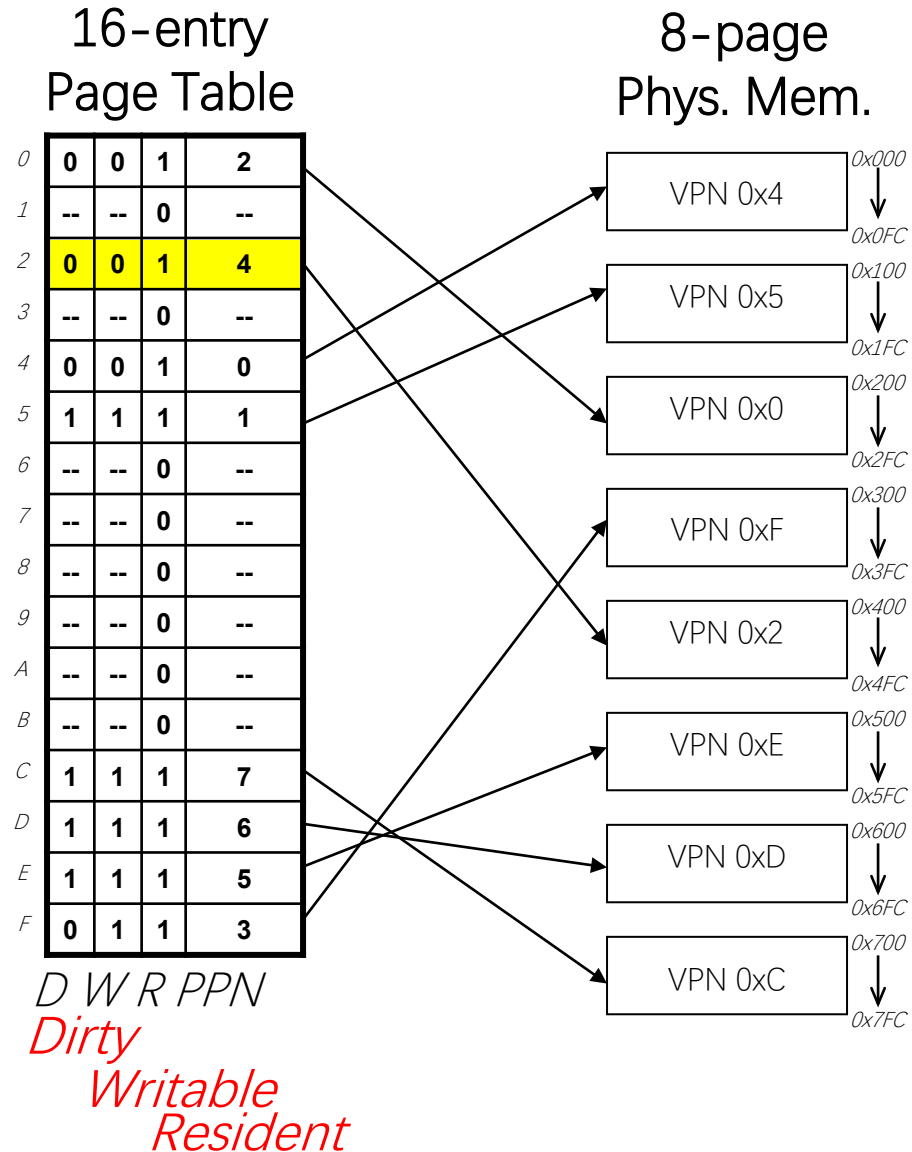
- What if all the pages can't fit in DRAM?
- Where do we store the page tables?

# Demand Paging: using main memory as a cache of disk

- All the pages of the processes may not fit in main memory → DRAM is backed up by *swap space* on disk.
- Page Table Entry (PTE) contains:
  - A **resident** bit to indicate if the page exists in main memory
  - **PPN** (physical page number) for a memory-resident page
  - **DPN** (disk page number) for a page on the disk
  - Protection and usage bits
- Even if all pages fit in memory, demand paging allows bringing only what is needed from disk
  - When a process starts, all code and data are on disk; bring pages in as they are accessed



# Example: Virtual → Physical Translation



Setup:

256 bytes/page ( $2^8$ )  
16 virtual pages ( $2^4$ )  
8 physical pages ( $2^3$ )  
12-bit VA (4 vpn, 8 offset)  
11-bit PA (3 ppn, 8 offset)

lw 0x2C8(x0)

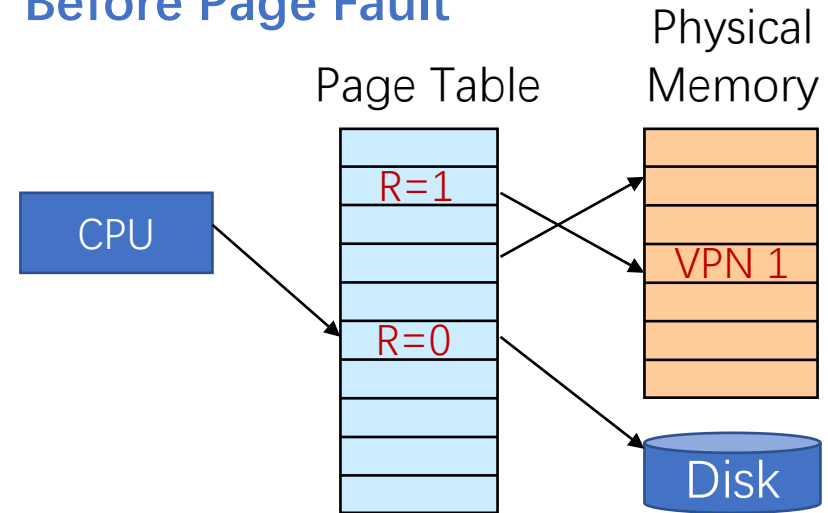
VA = 0x2C8, PA = 0x4C8

VPN = 0x2  
→ PPN = 0x4


# Page Faults

- An access to a page that does not have a valid translation causes a **page fault exception**
- OS page fault handler is invoked, handles miss:
  - Choose a page to replace, write it back if dirty. Mark page as no longer resident
  - Read page from disk into available physical page
  - Update page table to show new page is resident
  - Return control to program, which re-executes memory access

## Before Page Fault



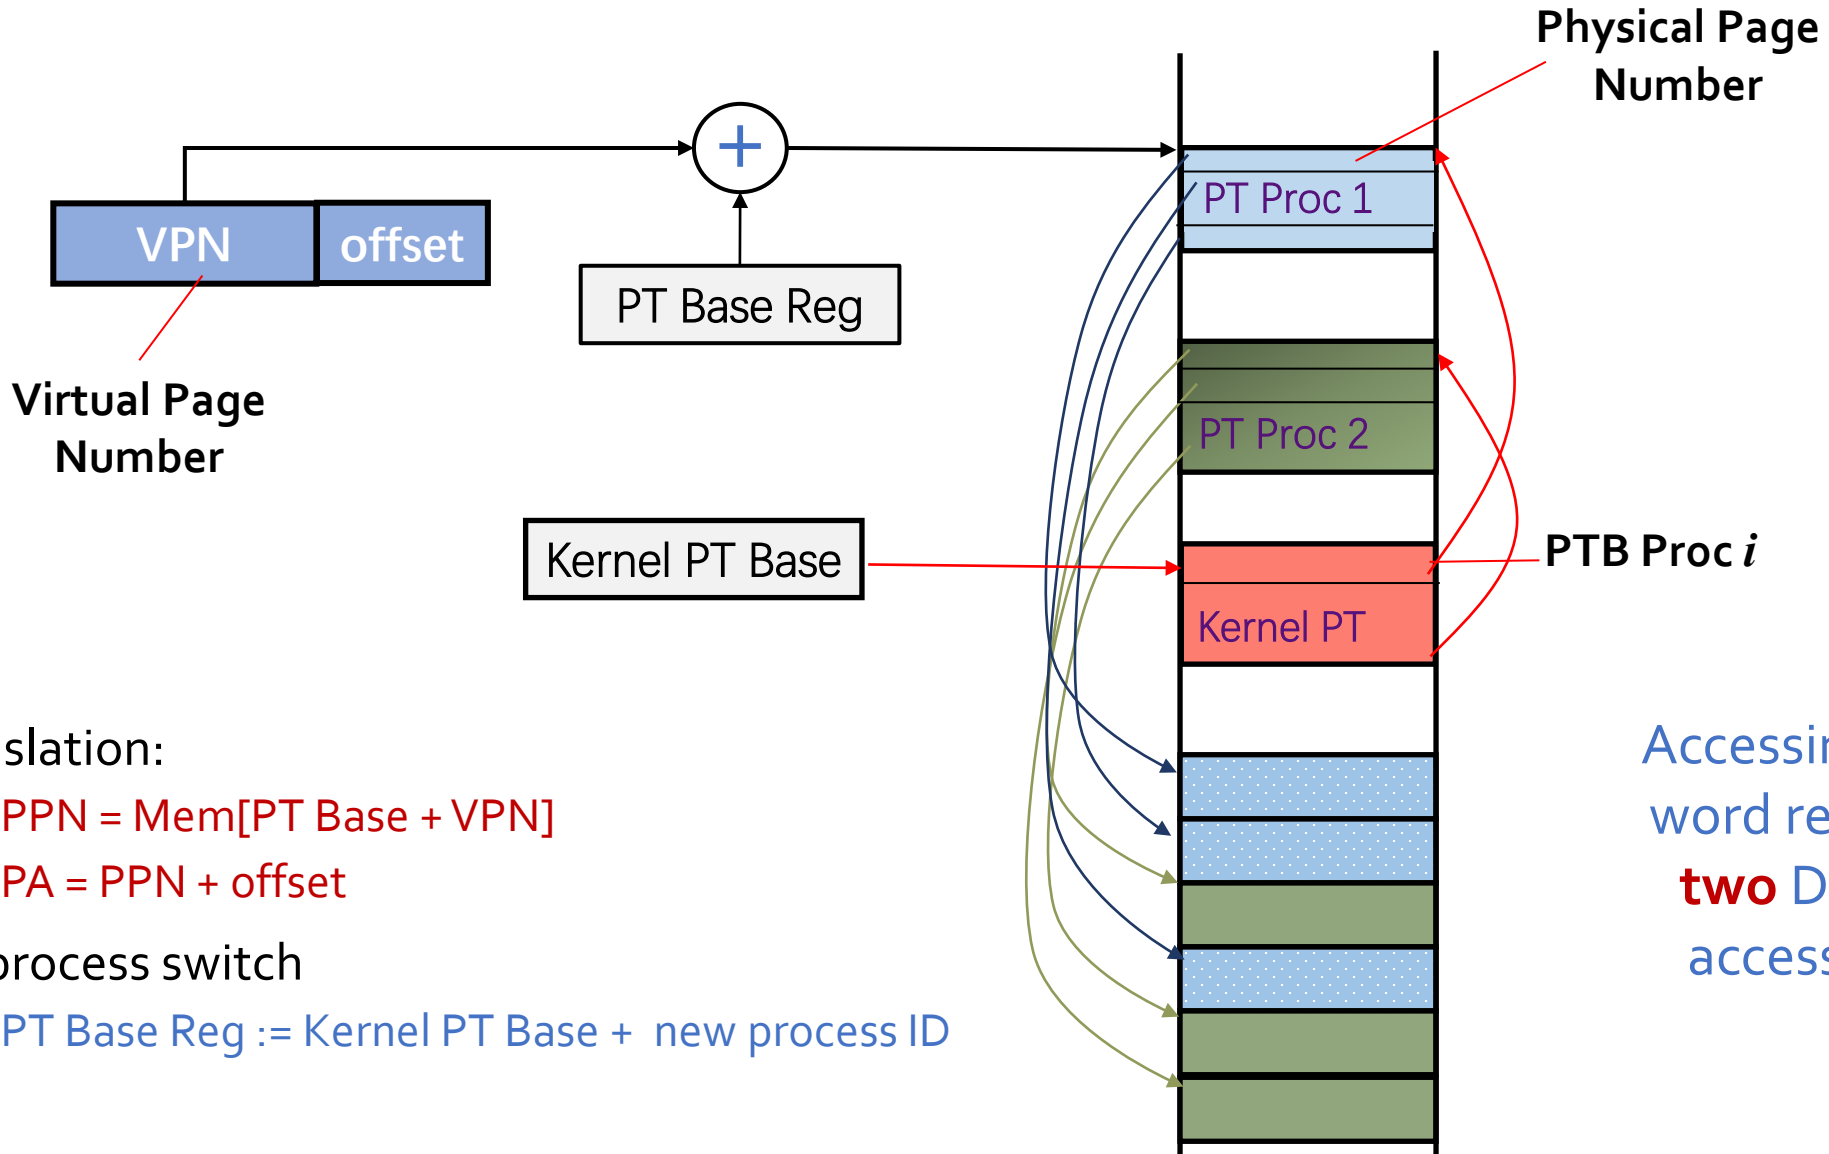
## After Page Fault



# Two Remaining Questions

- What if all the pages can't fit in DRAM?
- Where do we store the page tables?

# Suppose Page Tables reside in Memory



- Translation:
  - $PPN = \text{Mem}[\text{PT Base} + \text{VPN}]$
  - $PA = PPN + \text{offset}$
- On process switch
  - $\text{PT Base Reg} := \text{Kernel PT Base} + \text{new process ID}$




# Translation Lookaside Buffer (TLB)

**Problem:** Address translation is very expensive!  
Each reference requires accessing page table


**Solution:** **Cache translations in TLB**

TLB hit → *Single-cycle translation*

TLB miss → *Access page table to refill TLB*



# Address Translation: *Putting it all together*



# Summary

- Virtual memory benefits:
  - **Protection and Privacy**: Private address space per process
  - **Demand Paging**: use main memory as a cache of disk
- Segmentation: Each process address space is a contiguous block (a segment) in physical memory
  - **Simple: Base and bound registers**
  - **Suffers from fragmentation, no demand paging**
- Paging: Each process address space is stored on multiple fixed-size pages. A page table maps virtual to physical pages
  - **Avoids fragmentation**
  - **Enables demand paging**: pages can be in main memory or disk
  - **Requires a page table access on each memory reference**
- TLBs make paging efficient by caching the page table

Thank you!

Q & A