



GPU Teaching Kit  
Accelerated Computing



## Module 4.1 – Memory and Data Locality

CUDA Memories

# Objective

- To learn to effectively use the CUDA memory types in a parallel program
  - Importance of memory access efficiency
  - Registers, shared memory, global memory
  - Scope and lifetime

## Review: Image Blur Kernel.

```
// Get the average of the surrounding 2xBLUR_SIZE x 2xBLUR_SIZE box
for(int blurRow = -BLUR_SIZE; blurRow < BLUR_SIZE+1; ++blurRow) {
    for(int blurCol = -BLUR_SIZE; blurCol < BLUR_SIZE+1; ++blurCol) {


        int curRow = Row + blurRow;
        int curCol = Col + blurCol;
        // Verify we have a valid image pixel
        if(curRow > -1 && curRow < h && curCol > -1 && curCol < w) {
            ➡ pixVal += in[curRow * w + curCol];
            pixels++; // Keep track of number of pixels in the accumulated total
        }
    }
}

// Write our new pixel value out
out[Row * w + Col] = (unsigned char)(pixVal / pixels);
```

## How about performance on a GPU

- All threads access global memory for their input matrix elements
  - One memory accesses (4 bytes) per floating-point addition
  - 4B/s of memory bandwidth/FLOPS
- Assume a GPU with
  - Peak floating-point rate 1,600 GFLOPS with 600 GB/s DRAM bandwidth
  - $4 \times 1,600 = 6,400$  GB/s required to achieve peak FLOPS rating
  - The 600 GB/s memory bandwidth limits the execution at 150 GFLOPS
- This limits the execution rate to 9.3% (150/1600) of the peak floating-point execution rate of the device!
- Need to drastically cut down memory accesses to get close to the 1,600 GFLOPS

## Example – Matrix Multiplication




# A Basic Matrix Multiplication

```
__global__ void MatrixMulKernel(float* M, float* N, float* P, int Width) {  
    // Calculate the row index of the P element and M  
    int Row = blockIdx.y*blockDim.y+threadIdx.y;  
  
    // Calculate the column index of P and N  
    int Col = blockIdx.x*blockDim.x+threadIdx.x;  
  
    if ((Row < Width) && (Col < Width)) {  
        float Pvalue = 0;  
        // each thread computes one element of the block sub-matrix  
        for (int k = 0; k < Width; ++k) {  
            Pvalue += M[Row*Width+k]*N[k*Width+Col];  
        }  
        P[Row*Width+Col] = Pvalue;  
    }  
}
```

## Example – Matrix Multiplication


```
__global__ void MatrixMulKernel(float* M, float* N, float* P, int Width) {  
  
    // Calculate the row index of the P element and M  
    int Row = blockIdx.y*blockDim.y+threadIdx.y;  
  
    // Calculate the column index of P and N  
    int Col = blockIdx.x*blockDim.x+threadIdx.x;  
  
    if ((Row < Width) && (Col < Width)) {  
        float Pvalue = 0;  
        // each thread computes one element of the block sub-matrix  
        for (int k = 0; k < Width; ++k) {  
            Pvalue += M[Row*Width+k]*N[k*Width+Col];  
        }  
        P[Row*Width+Col] = Pvalue;  
    }  
}
```

## A Toy Example: Thread to P Data Mapping







## Calculation of $P_{0,0}$ and $P_{0,1}$



## Memory and Registers in the Von-Neumann Model



# Programmer View of CUDA Memories



# Declaring CUDA Variables


| Variable declaration                     | Memory   | Scope  | Lifetime    |
|--|----------|--------|-------------|
| int LocalVar;                            | register | thread | thread      |
| __device__ __shared__ int SharedVar;     | shared   | block  | block       |
| __device__ int GlobalVar;                | global   | grid   | application |
| __device__ __constant__ int ConstantVar; | constant | grid   | application |

- **\_\_device\_\_** is optional when used with **\_\_shared\_\_**, or **\_\_constant\_\_**
- Automatic variables reside in a **register**
  - Except **per-thread arrays** that reside in global memory

## Example: Shared Memory Variable Declaration

```
void blurKernel(unsigned char * in, unsigned char * out, int w, int h)
{
    __shared__ float ds_in[TILE_WIDTH][TILE_WIDTH];
    ...
}
```

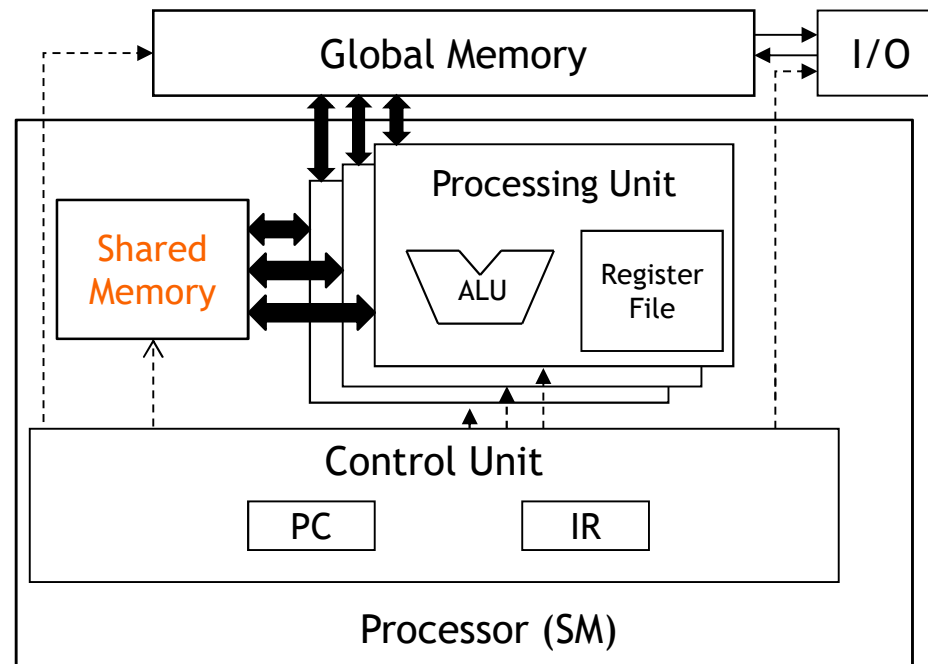
## Where to Declare Variables?



## Shared Memory in CUDA

- A special type of memory whose contents are explicitly defined and used in the kernel source code
  - One in each SM
  - Accessed at much higher speed (in both latency and throughput) than global memory
  - Scope of access and sharing - thread blocks
  - Lifetime – thread block, contents will disappear after the corresponding thread finishes terminates execution
  - Accessed by memory load/store instructions
  - A form of scratchpad memory in computer architecture

## Hardware View of CUDA Memories







## GPU Teaching Kit

Accelerated Computing



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the [Creative Commons Attribution-NonCommercial 4.0 International License](https://creativecommons.org/licenses/by-nc/4.0/).



GPU Teaching Kit  
Accelerated Computing



## Module 4.2 – Memory and Data Locality


Tiled Parallel Algorithms

# Objective

- To understand the motivation and ideas for tiled parallel algorithms
  - Reducing the limiting effect of memory bandwidth on parallel kernel performance
  - Tiled algorithms and barrier synchronization


# Global Memory Access Pattern of the Basic Matrix Multiplication Kernel

Global Memory



# Tiling/Blocking - Basic Idea

Global Memory



Divide the global memory content into tiles



Focus the computation of threads on one or a small number of tiles at each point in time

# Tiling/Blocking - Basic Idea

Global Memory



On-chip Memory



...

# Basic Concept of Tiling

- In a congested traffic system, significant reduction of vehicles can greatly improve the delay seen by all vehicles
  - Carpooling for commuters
  - Tiling for global memory accesses
    - drivers = threads accessing their memory data operands
    - cars = memory access requests



# Some Computations are More Challenging to Tile

- Some carpools may be easier than others
  - Car pool participants need to have similar work schedule
  - Some vehicles may be more suitable for carpooling
- Similar challenges exist in tiling






# Carpools need synchronization.

- Good: when people have similar schedule




# Carpools need synchronization.

- Bad: when people have very different schedule




# Same with Tiling

- Good: when threads have similar access timing



- Bad: when threads have very different timing

# Barrier Synchronization for Tiling



# Outline of Tiling Technique

- Identify a tile of global memory contents that are accessed by multiple threads
- Load the tile from global memory into on-chip memory
- Use barrier synchronization to make sure that all threads are ready to start the phase
- Have the multiple threads to access their data from the on-chip memory
- Use barrier synchronization to make sure that all threads have completed the current phase
- Move on to the next tile



# GPU Teaching Kit

Accelerated Computing



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the [Creative Commons Attribution-NonCommercial 4.0 International License](https://creativecommons.org/licenses/by-nc/4.0/).



GPU Teaching Kit  
Accelerated Computing



## Module 4.3 - Memory Model and Locality

Tiled Matrix Multiplication


# Objective

- To understand the design of a tiled parallel algorithm for matrix multiplication
  - Loading a tile
  - Phased execution
  - Barrier Synchronization



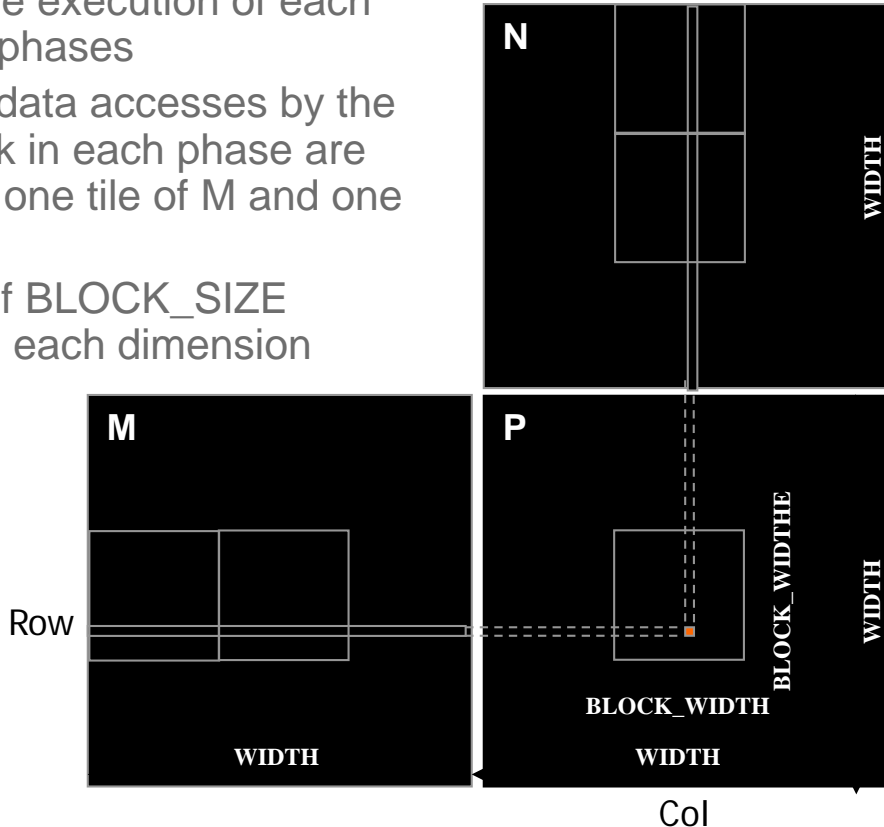
# Matrix Multiplication

- Data access pattern
  - Each thread - a row of M and a column of N
  - Each thread block – a strip of M and a strip of N



# Tiled Matrix Multiplication


- Break up the execution of each thread into phases
- so that the data accesses by the thread block in each phase are focused on one tile of M and one tile of N
- The tile is of BLOCK\_SIZE elements in each dimension



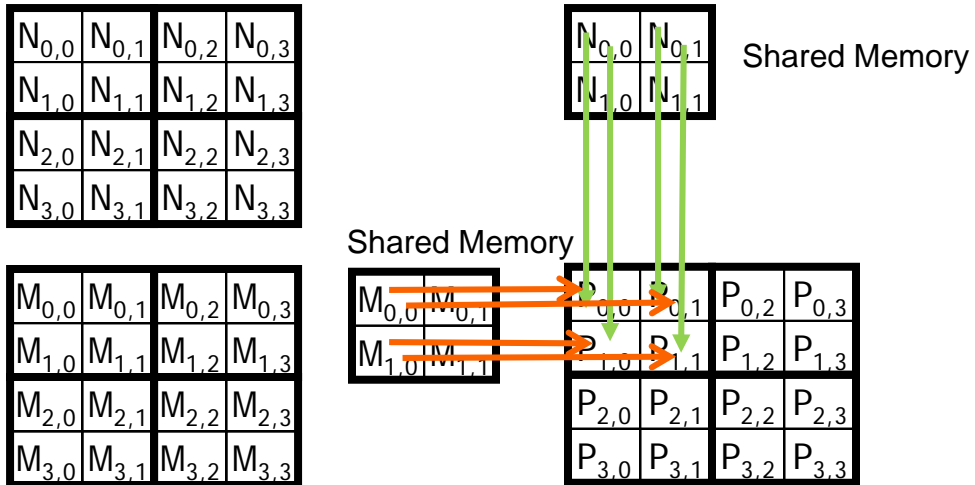
# Loading a Tile

- All threads in a block participate
  - Each thread loads one M element and one N element in tiled code

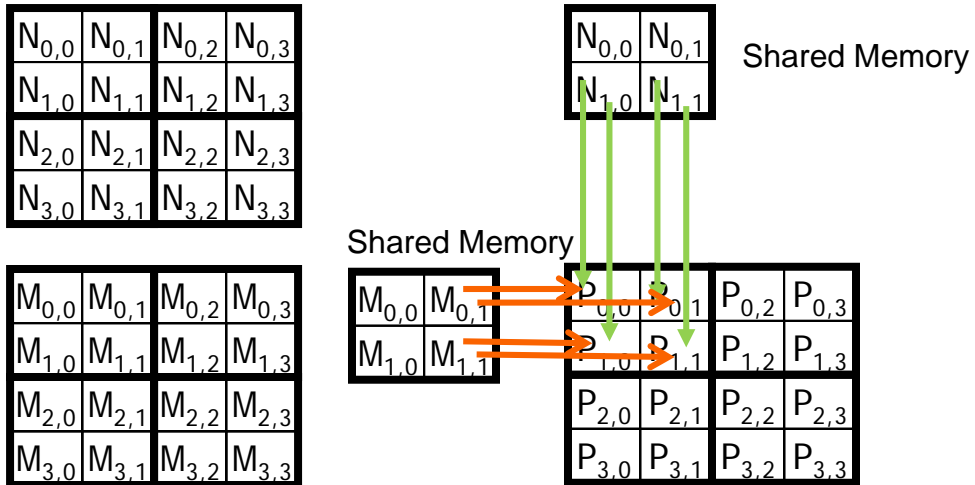
# Phase 0 Load for Block (0,0)



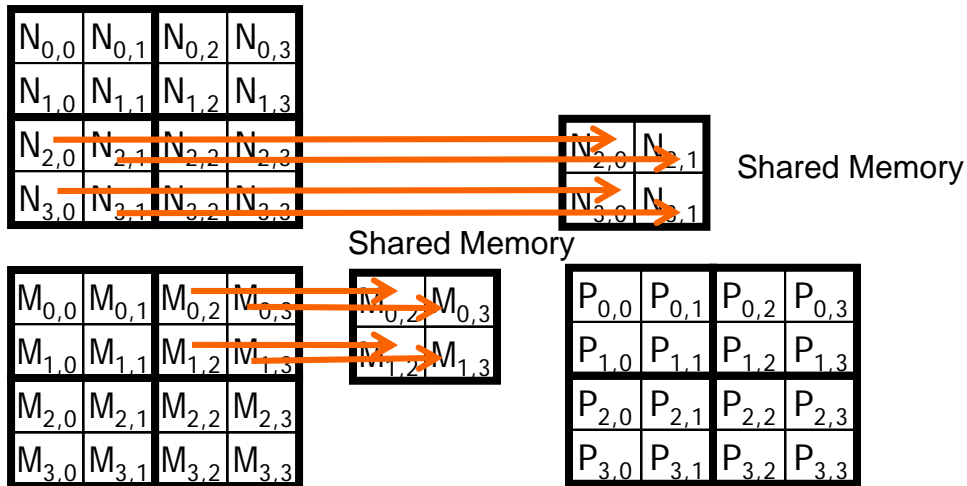
# Phase 0 Use for Block (0,0) (iteration 0)




# Phase 0 Use for Block (0,0) (iteration 1)



# Phase 1 Load for Block (0,0)




# Phase 1 Use for Block (0,0) (iteration 0)






# Phase 1 Use for Block (0,0) (iteration 1)



# Execution Phases of Toy Example

|                       | Phase 0   |   |  | Phase 1   |   |  |
|-----------------------|---|---|--|---|---|--|
| thread <sub>0,0</sub> | <b>M<sub>0,0</sub></b><br>↓<br>Mds <sub>0,0</sub> | <b>N<sub>0,0</sub></b><br>↓<br>Nds <sub>0,0</sub> | PValue <sub>0,0</sub> +=<br>Mds <sub>0,0</sub> *Nds <sub>0,0</sub> +<br>Mds <sub>0,1</sub> *Nds <sub>1,0</sub> | <b>M<sub>0,2</sub></b><br>↓<br>Mds <sub>0,0</sub> | <b>N<sub>2,0</sub></b><br>↓<br>Nds <sub>0,0</sub> | PValue <sub>0,0</sub> +=<br>Mds <sub>0,0</sub> *Nds <sub>0,0</sub> +<br>Mds <sub>0,1</sub> *Nds <sub>1,0</sub> |
| thread <sub>0,1</sub> | <b>M<sub>0,1</sub></b><br>↓<br>Mds <sub>0,1</sub> | <b>N<sub>0,1</sub></b><br>↓<br>Nds <sub>1,0</sub> | PValue <sub>0,1</sub> +=<br>Mds <sub>0,0</sub> *Nds <sub>0,1</sub> +<br>Mds <sub>0,1</sub> *Nds <sub>1,1</sub> | <b>M<sub>0,3</sub></b><br>↓<br>Mds <sub>0,1</sub> | <b>N<sub>2,1</sub></b><br>↓<br>Nds <sub>0,1</sub> | PValue <sub>0,1</sub> +=<br>Mds <sub>0,0</sub> *Nds <sub>0,1</sub> +<br>Mds <sub>0,1</sub> *Nds <sub>1,1</sub> |
| thread <sub>1,0</sub> | <b>M<sub>1,0</sub></b><br>↓<br>Mds <sub>1,0</sub> | <b>N<sub>1,0</sub></b><br>↓<br>Nds <sub>1,0</sub> | PValue <sub>1,0</sub> +=<br>Mds <sub>1,0</sub> *Nds <sub>0,0</sub> +<br>Mds <sub>1,1</sub> *Nds <sub>1,0</sub> | <b>M<sub>1,2</sub></b><br>↓<br>Mds <sub>1,0</sub> | <b>N<sub>3,0</sub></b><br>↓<br>Nds <sub>1,0</sub> | PValue <sub>1,0</sub> +=<br>Mds <sub>1,0</sub> *Nds <sub>0,0</sub> +<br>Mds <sub>1,1</sub> *Nds <sub>1,0</sub> |
| thread <sub>1,1</sub> | <b>M<sub>1,1</sub></b><br>↓<br>Mds <sub>1,1</sub> | <b>N<sub>1,1</sub></b><br>↓<br>Nds <sub>1,1</sub> | PValue <sub>1,1</sub> +=<br>Mds <sub>1,0</sub> *Nds <sub>0,1</sub> +<br>Mds <sub>1,1</sub> *Nds <sub>1,1</sub> | <b>M<sub>1,3</sub></b><br>↓<br>Mds <sub>1,1</sub> | <b>N<sub>3,1</sub></b><br>↓<br>Nds <sub>1,1</sub> | PValue <sub>1,1</sub> +=<br>Mds <sub>1,0</sub> *Nds <sub>0,1</sub> +<br>Mds <sub>1,1</sub> *Nds <sub>1,1</sub> |

time 

# Execution Phases of Toy Example (cont.)

|                       | Phase 0                                       |   |  | Phase 1                                       |   |  |
|-----------------------|---|---|--|---|---|--|
| thread <sub>0,0</sub> | $\mathbf{M}_{0,0}$<br>↓<br>Mds <sub>0,0</sub> | $\mathbf{N}_{0,0}$<br>↓<br>Nds <sub>0,0</sub> | PValue <sub>0,0</sub> +=<br><u>Mds<sub>0,0</sub></u> *Nds <sub>0,0</sub> +<br>Mds <sub>0,1</sub> *Nds <sub>1,0</sub> | $\mathbf{M}_{0,2}$<br>↓<br>Mds <sub>0,0</sub> | $\mathbf{N}_{2,0}$<br>↓<br>Nds <sub>0,0</sub> | PValue <sub>0,0</sub> +=<br>Mds <sub>0,0</sub> *Nds <sub>0,0</sub> +<br>Mds <sub>0,1</sub> *Nds <sub>1,0</sub> |
| thread <sub>0,1</sub> | $\mathbf{M}_{0,1}$<br>↓<br>Mds <sub>0,1</sub> | $\mathbf{N}_{0,1}$<br>↓<br>Nds <sub>1,0</sub> | PValue <sub>0,1</sub> +=<br><u>Mds<sub>0,0</sub></u> *Nds <sub>0,1</sub> +<br>Mds <sub>0,1</sub> *Nds <sub>1,1</sub> | $\mathbf{M}_{0,3}$<br>↓<br>Mds <sub>0,1</sub> | $\mathbf{N}_{2,1}$<br>↓<br>Nds <sub>0,1</sub> | PValue <sub>0,1</sub> +=<br>Mds <sub>0,0</sub> *Nds <sub>0,1</sub> +<br>Mds <sub>0,1</sub> *Nds <sub>1,1</sub> |
| thread <sub>1,0</sub> | $\mathbf{M}_{1,0}$<br>↓<br>Mds <sub>1,0</sub> | $\mathbf{N}_{1,0}$<br>↓<br>Nds <sub>1,0</sub> | PValue <sub>1,0</sub> +=<br>Mds <sub>1,0</sub> *Nds <sub>0,0</sub> +<br>Mds <sub>1,1</sub> *Nds <sub>1,0</sub>       | $\mathbf{M}_{1,2}$<br>↓<br>Mds <sub>1,0</sub> | $\mathbf{N}_{3,0}$<br>↓<br>Nds <sub>1,0</sub> | PValue <sub>1,0</sub> +=<br>Mds <sub>1,0</sub> *Nds <sub>0,0</sub> +<br>Mds <sub>1,1</sub> *Nds <sub>1,0</sub> |
| thread <sub>1,1</sub> | $\mathbf{M}_{1,1}$<br>↓<br>Mds <sub>1,1</sub> | $\mathbf{N}_{1,1}$<br>↓<br>Nds <sub>1,1</sub> | PValue <sub>1,1</sub> +=<br>Mds <sub>1,0</sub> *Nds <sub>0,1</sub> +<br>Mds <sub>1,1</sub> *Nds <sub>1,1</sub>       | $\mathbf{M}_{1,3}$<br>↓<br>Mds <sub>1,1</sub> | $\mathbf{N}_{3,1}$<br>↓<br>Nds <sub>1,1</sub> | PValue <sub>1,1</sub> +=<br>Mds <sub>1,0</sub> *Nds <sub>0,1</sub> +<br>Mds <sub>1,1</sub> *Nds <sub>1,1</sub> |

time →

Shared memory allows each value to be accessed by multiple threads

# Barrier Synchronization

- Synchronize all threads in a block
  - `__syncthreads()`
- All threads in the same block must reach the `__syncthreads()` before any of the them can move on
- Best used to coordinate the phased execution tiled algorithms
  - To ensure that all elements of a tile are loaded at the beginning of a phase
  - To ensure that all elements of a tile are consumed at the end of a phase



# GPU Teaching Kit

Accelerated Computing



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the [Creative Commons Attribution-NonCommercial 4.0 International License](https://creativecommons.org/licenses/by-nc/4.0/).



GPU Teaching Kit  
Accelerated Computing



# Module 4.4 - Memory and Data Locality

Tiled Matrix Multiplication Kernel

# Objective

- To learn to write a tiled matrix-multiplication kernel
  - Loading and using tiles for matrix multiplication
  - Barrier synchronization, shared memory
  - Resource Considerations
  - Assume that Width is a multiple of tile size for simplicity

# Loading Input Tile 0 of M (Phase 0)

- Have each thread load an M element and an N element at the same relative position as its P element.


```
int Row = by * blockDim.y + ty;
```

```
int Col = bx * blockDim.x + tx;
```

2D indexing for accessing Tile 0:

**M|Row||tx|**

N[ty][Col]





# Loading Input Tile 0 of N (Phase 0)

- Have each thread load an M element and an N element at the same relative position as its P element.


```
int Row = by * blockDim.y + ty;
```

```
int Col = bx * blockDim.x + tx;
```

2D indexing for accessing Tile 0:

```
M[Row][tx]
```

```
N[ty][Col]
```




# Loading Input Tile 1 of M (Phase 1)

2D indexing for accessing Tile 1:

$M[\text{Row}][1 * \text{TILE\_WIDTH} + \text{tx}]$

$N[1 * \text{TILE} * \text{WIDTH} + \text{ty}][\text{Col}]$




# Loading Input Tile 1 of N (Phase 1)

2D indexing for accessing Tile 1:

$M[\text{Row}][1 * \text{TILE\_WIDTH} + \text{tx}]$

$N[1 * \text{TILE} * \text{WIDTH} + \text{ty}][\text{Col}]$



## M and N are dynamically allocated - use 1D indexing

➡  $M[\text{Row}][p * \text{TILE\_WIDTH} + tx]$   
➡  $M[\text{Row} * \text{Width} + p * \text{TILE\_WIDTH} + tx]$

➡  $N[p * \text{TILE\_WIDTH} + ty][\text{Col}]$   
➡  $N[(p * \text{TILE\_WIDTH} + ty) * \text{Width} + \text{Col}]$

where p is the sequence number of the current phase

# Tiled Matrix Multiplication Kernel

```
__global__ void MatrixMulKernel(float* M, float* N, float* P, Int Width)
{
    __shared__ float ds_M[TILE_WIDTH][TILE_WIDTH];
    __shared__ float ds_N[TILE_WIDTH][TILE_WIDTH];

    int bx = blockIdx.x;  int by = blockIdx.y;
    int tx = threadIdx.x; int ty = threadIdx.y;

    int Row = by * blockDim.y + ty;
    int Col = bx * blockDim.x + tx;
    float Pvalue = 0;

    // Loop over the M and N tiles required to compute the P element
    for (int p = 0; p < n/TILE_WIDTH; ++p) {
        // Collaborative loading of M and N tiles into shared memory
        ds_M[ty][tx] = M[Row*Width + p*TILE_WIDTH+tx];
        ds_N[ty][tx] = N[(t*TILE_WIDTH+ty)*Width + Col];
        __syncthreads();

        for (int i = 0; i < TILE_WIDTH; ++i) Pvalue += ds_M[ty][i] * ds_N[i][tx];
        __syncthreads();
    }
    P[Row*Width+Col] = Pvalue;
}
```

# Tiled Matrix Multiplication Kernel

```
__global__ void MatrixMulKernel(float* M, float* N, float* P, Int Width)
{
    __shared__ float ds_M[TILE_WIDTH][TILE_WIDTH];
    __shared__ float ds_N[TILE_WIDTH][TILE_WIDTH];

    int bx = blockIdx.x;  int by = blockIdx.y;
    int tx = threadIdx.x; int ty = threadIdx.y;

    int Row = by * blockDim.y + ty;
    int Col = bx * blockDim.x + tx;
    float Pvalue = 0;

    // Loop over the M and N tiles required to compute the P element
    for (int p = 0; p < n/TILE_WIDTH; ++p) {
        // Collaborative loading of M and N tiles into shared memory
        ds_M[ty][tx] = M[Row*Width + p*TILE_WIDTH+tx];
        ds_N[ty][tx] = N[(t*TILE_WIDTH+ty)*Width + Col];
        __syncthreads();

        for (int i = 0; i < TILE_WIDTH; ++i) Pvalue += ds_M[ty][i] * ds_N[i][tx];
        __syncthreads();
    }
    P[Row*Width+Col] = Pvalue;
}
```

# Tiled Matrix Multiplication Kernel

```
__global__ void MatrixMulKernel(float* M, float* N, float* P, Int Width)
{
    __shared__ float ds_M[TILE_WIDTH][TILE_WIDTH];
    __shared__ float ds_N[TILE_WIDTH][TILE_WIDTH];

    int bx = blockIdx.x;  int by = blockIdx.y;
    int tx = threadIdx.x; int ty = threadIdx.y;

    int Row = by * blockDim.y + ty;
    int Col = bx * blockDim.x + tx;
    float Pvalue = 0;

    // Loop over the M and N tiles required to compute the P element
    for (int p = 0; p < n/TILE_WIDTH; ++p) {
        // Collaborative loading of M and N tiles into shared memory
        ds_M[ty][tx] = M[Row*Width + p*TILE_WIDTH+tx];
        ds_N[ty][tx] = N[(t*TILE_WIDTH+ty)*Width + Col];
        __syncthreads();

        for (int i = 0; i < TILE_WIDTH; ++i) Pvalue += ds_M[ty][i] * ds_N[i][tx];
        __syncthreads();
    }
    P[Row*Width+Col] = Pvalue;
}
```

# Tile (Thread Block) Size Considerations

- Each **thread block** should have many threads
  - TILE\_WIDTH of 16 gives  $16*16 = 256$  threads
  - TILE\_WIDTH of 32 gives  $32*32 = 1024$  threads
- For 16, in each phase, each block performs  $2*256 = 512$  float loads from global memory for  $256 * (2*16) = 8,192$  mul/add operations. (16 floating-point operations for each memory load)
- For 32, in each phase, each block performs  $2*1024 = 2048$  float loads from global memory for  $1024 * (2*32) = 65,536$  mul/add operations. (32 floating-point operation for each memory load)



# Shared Memory and Threading

- For an SM with 16KB shared memory
  - Shared memory size is implementation dependent!
  - For `TILE_WIDTH = 16`, each thread block uses  $2 \times 256 \times 4\text{B} = 2\text{KB}$  of shared memory.
  - For 16KB shared memory, one can potentially have up to 8 thread blocks executing
    - This allows up to  $8 \times 512 = 4,096$  pending loads. (2 per thread, 256 threads per block)
  - The next `TILE_WIDTH 32` would lead to  $2 \times 32 \times 32 \times 4\text{Byte} = 8\text{K Byte}$  shared memory usage per thread block, allowing 2 thread blocks active at the same time
    - However, in a GPU where the thread count is limited to 1536 threads per SM, the number of blocks per SM is reduced to one!
- Each `__syncthread()` can reduce the number of active threads for a block
  - More thread blocks can be advantageous



# GPU Teaching Kit

Accelerated Computing



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the [Creative Commons Attribution-NonCommercial 4.0 International License](https://creativecommons.org/licenses/by-nc/4.0/).



GPU Teaching Kit  
Accelerated Computing



## Module 4.5 - Memory and Data Locality

Handling Arbitrary Matrix Sizes in Tiled Algorithms

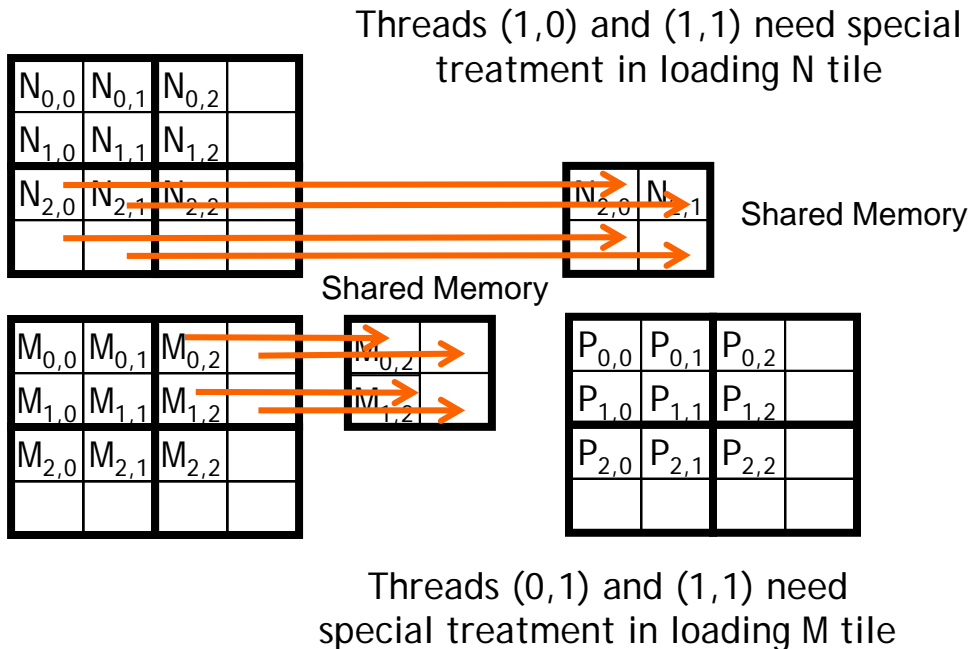
# Objective

- To learn to handle arbitrary matrix sizes in tiled matrix multiplication
  - Boundary condition checking
  - Regularizing tile contents
  - Rectangular matrices


# Handling Matrix of Arbitrary Size

- The tiled matrix multiplication kernel we presented so far can handle only square matrices whose dimensions (Width) are multiples of the tile width (TILE\_WIDTH)
  - However, real applications need to handle arbitrary sized matrices.
  - One could pad (add elements to) the rows and columns into multiples of the tile size, but would have significant space and data transfer time overhead.
- We will take a different approach.

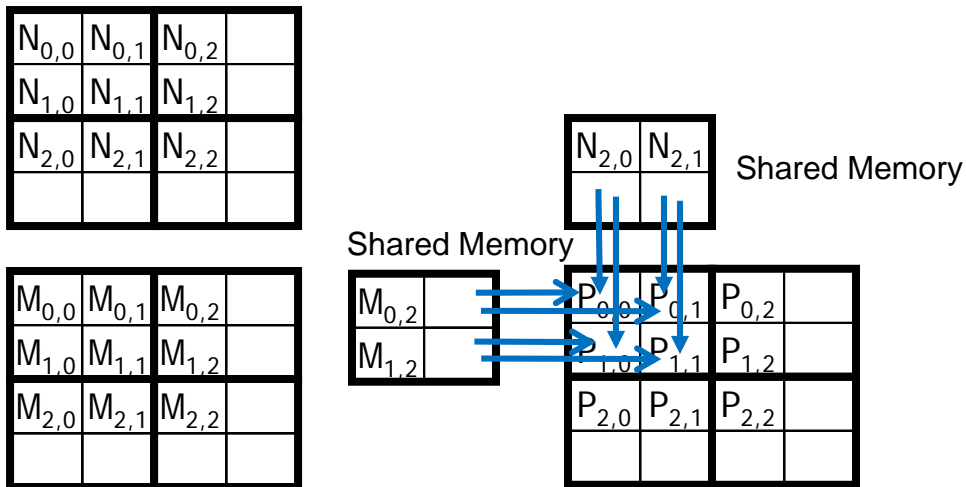
# Phase 1 Loads for Block (0,0) for a 3x3 Example



## Phase 1 Use for Block (0,0) (iteration 0)



## Phase 1 Use for Block (0,0) (iteration 1)





All Threads need special treatment. None of them should introduce invalidate contributions to their P elements.



## Phase 0 Loads for Block (1,1) for a 3x3 Example

Threads (0,1) and (1,1) need special treatment in loading N tile



Threads (1,0) and (1,1) need special treatment in loading M tile

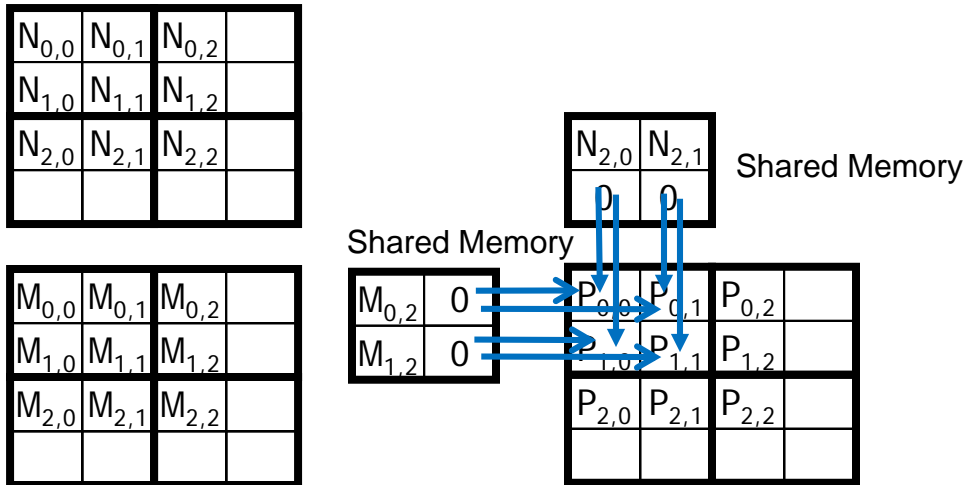
# Major Cases in Toy Example

- Threads that do not calculate valid P elements but still need to participate in loading the input tiles
  - Phase 0 of Block(1,1), Thread(1,0), assigned to calculate non-existent  $P[3,2]$  but need to participate in loading tile element  $N[1,2]$
- Threads that calculate valid P elements may attempt to load non-existing input elements when loading input tiles
  - Phase 0 of Block(0,0), Thread(1,0), assigned to calculate valid  $P[1,0]$  but attempts to load non-existing  $N[3,0]$

# A “Simple” Solution


- When a thread is to load any input element, test if it is in the valid index range
  - If valid, proceed to load
  - Else, do not load, just write a 0
- Rationale: a 0 value will ensure that that the multiply-add step does not affect the final value of the output element
- The condition tested for loading input elements is different from the test for calculating output P element
  - A thread that does not calculate valid P element can still participate in loading input tile elements

# Phase 1 Use for Block (0,0) (iteration 1)




# Boundary Condition for Input M Tile

- Each thread loads
  - $M[\text{Row}][p * \text{TILE\_WIDTH} + tx]$
  - $M[\text{Row} * \text{Width} + p * \text{TILE\_WIDTH} + tx]$
- Need to test
  - $(\text{Row} < \text{Width}) \ \&\& \ (p * \text{TILE\_WIDTH} + tx < \text{Width})$
  - If true, load M element
  - Else , load 0



# Boundary Condition for Input N Tile

- Each thread loads
  - $N[p * \text{TILE\_WIDTH} + ty][\text{Col}]$
  - $N[(p * \text{TILE\_WIDTH} + ty) * \text{Width} + \text{Col}]$
- Need to test
  - $(p * \text{TILE\_WIDTH} + ty < \text{Width}) \ \&\& \ (\text{Col} < \text{Width})$
  - If true, load N element
  - Else , load 0



# Loading Elements – with boundary check

```
- 8  for (int p = 0; p < (Width-1) / TILE_WIDTH + 1; ++p) {  
-  
-  ++    if (Row < Width && t * TILE_WIDTH + tx < Width) {  
- 9      ds_M[ty][tx] = M[Row * Width + p * TILE_WIDTH + tx];  
-  ++    } else {  
-  ++      ds_M[ty][tx] = 0.0;  
-  ++    }  
-  ++    if (p * TILE_WIDTH + ty < Width && Col < Width) {  
- 10     ds_N[ty][tx] = N[(p * TILE_WIDTH + ty) * Width + Col];  
-  ++    } else {  
-  ++      ds_N[ty][tx] = 0.0;  
-  ++    }  
- 11    __syncthreads();  
-  
-
```

# Inner Product – Before and After

```
-  ++  if(Row < Width && Col < Width) {  
-  12    for (int i = 0; i < TILE_WIDTH; ++i) {  
-  13        Pvalue += ds_M[ty][i] * ds_N[i][tx];  
-      }  
-  14    __syncthreads();  
-  15  } /* end of outer for loop */  
-  ++  if (Row < Width && Col < Width)  
-  16      P[Row*Width + Col] = Pvalue;  
-  } /* end of kernel */
```



# Some Important Points

- For each thread the conditions are different for
  - Loading M element
  - Loading N element
  - Calculating and storing output elements
- The effect of control divergence should be small for large matrices

# Handling General Rectangular Matrices

- In general, the matrix multiplication is defined in terms of rectangular matrices
  - A  $j \times k$   $M$  matrix multiplied with a  $k \times l$   $N$  matrix results in a  $j \times l$   $P$  matrix
- We have presented square matrix multiplication, a special case
- The kernel function needs to be generalized to handle general rectangular matrices
  - The Width argument is replaced by three arguments:  $j$ ,  $k$ ,  $l$
  - When Width is used to refer to the height of  $M$  or height of  $P$ , replace it with  $j$
  - When Width is used to refer to the width of  $M$  or height of  $N$ , replace it with  $k$
  - When Width is used to refer to the width of  $N$  or width of  $P$ , replace it with  $l$



# GPU Teaching Kit

Accelerated Computing



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the [Creative Commons Attribution-NonCommercial 4.0 International License](https://creativecommons.org/licenses/by-nc/4.0/).