GPU Teaching Kit

Accelerated Computing

ILLINOIS
UNIVERSITY OF ILLINOIS AT URBANA-CHAMPAIGN

Module 5.1 – Thread Execution Efficiency

Warps and SIMD Hardware

Objective

– To understand how CUDA threads execute on SIMD Hardware
 – Warp partitioning
  – SIMD Hardware
  – Control divergence

Warps as Scheduling Units

Block 1 Warps — t0 t1 t2 ... t31

Block 2 Warps — t0 t1 t2 ... t31

Block 3 Warps — t0 t1 t2 ... t31

– Each block is divided into 32-thread warps

- An implementation technique, not part of the CUDA programming model
- Warps are scheduling units in SM
- Threads in a warp execute in Single Instruction Multiple Data (SIMD) manner
- The number of threads in a warp may vary in future generations

Warps in Multi-dimensional Thread Blocks

– The thread blocks are first linearized into 1D in row major order
 – In x-dimension first, y-dimension next, and z-dimension last

$$T_{row,col} = T_{y,x}$$

logical 2-D organization

linear order

Blocks are partitioned after linearization

– Linearized thread blocks are partitioned
  – Thread indices within a warp are consecutive and increasing
  – Warp 0 starts with Thread 0

– Partitioning scheme is consistent across devices
  – Thus you can use this knowledge in control flow
  – However, the exact size of warps may change from generation to generation

– DO NOT rely on any ordering within or between warps
  – If there are any dependencies between threads, you must __syncthreads() to get correct results (more later).

SMs are SIMD Processors

– Control unit for instruction fetch, decode, and control is shared among multiple processing units

    – Control overhead is minimized (Module 1)

SIMD Execution Among Threads in a Warp

– All threads in a warp must execute the same instruction at any point in time

– This works efficiently if all threads follow the same control flow path
  – All if-then-else statements make the same decision
  – All loops iterate the same number of times

Control Divergence

– Control divergence occurs when threads in a warp take different control flow paths by making different control decisions
  – Some take the then-path and others take the else-path of an if-statement
  – Some threads take different number of loop iterations than others

– The execution of threads taking different paths are serialized in current GPUs
  – The control paths taken by the threads in a warp are traversed one at a time until there is no more.
  – During the execution of each path, all threads taking that path will be executed in parallel
  – The number of different paths can be large when considering nested control flow statements

Control Divergence Examples

– Divergence can arise when branch or loop condition is a function of thread indices

– Example kernel statement with divergence:

 – if (threadIdx.x > 2) { }

 – This creates two different control paths for threads in a block

 – Decision granularity < warp size; threads 0, 1 and 2 follow different path than the rest of the threads in the first warp

– Example without divergence:

 – If (blockIdx.x > 2) { }

 – Decision granularity is a multiple of blocks size; all threads in any given warp follow the same path

Example: Vector Addition Kernel

Device Code

```
// Compute vector sum C = A + B
// Each thread performs one pair-wise addition

__global__
void vecAddKernel(float* A, float* B, float* C,
  int n)
{
  int i = threadIdx.x + blockDim.x * blockIdx.x;
    if(i<n) C[i] = A[i] + B[i];
}
```

Analysis for vector size of 1,000 elements

– Assume that block size is 256 threads
  – 8 warps in each block

– All threads in Blocks 0, 1, and 2 are within valid range
  – i values from 0 to 767
  – There are 24 warps in these three blocks, none will have control divergence

– Most warps in Block 3 will not control divergence
  – Threads in the warps 0-6 are all within valid range, thus no control divergence

– One warp in Block 3 will have control divergence
  – Threads with i values 992-999 will all be within valid range
  – Threads with i values of 1000-1023 will be outside valid range

– Effect of serialization on control divergence will be small
  – 1 out of 32 warps has control divergence
  – The impact on performance will likely be less than 3%

GPU Teaching Kit

Accelerated Computing

ILLINOIS
UNIVERSITY OF ILLINOIS AT URBANA-CHAMPAIGN

GPU Teaching Kit

Accelerated Computing

Module 5.2 – Thread Execution Efficiency

Performance Impact of Control Divergence

Objective

- To learn to analyze the performance impact of control divergence
  - Boundary condition checking
  - Control divergence is data-dependent

Performance Impact of Control Divergence

– Boundary condition checks are vital for complete functionality and robustness of parallel code

- The tiled matrix multiplication kernel has many boundary condition checks
- The concern is that these checks may cause significant performance degradation
- For example, see the tile loading code below:

```
if(Row < Width && t * TILE_WIDTH+tx < Width) {
    ds_M[ty][tx] = M[Row * Width + p * TILE_WIDTH + tx];
} else {
    ds_M[ty][tx] = 0.0;
}

if (p*TILE_WIDTH+ty < Width && Col < Width) {
    ds_N[ty][tx] = N[(p*TILE_WIDTH + ty) * Width + Col];
} else {
    ds_N[ty][tx] = 0.0;
}
```

Two types of blocks in loading M Tiles

– 1. Blocks whose tiles are all within valid range until the last phase.
– 2. Blocks whose tiles are partially outside the valid range all the way

Analysis of Control Divergence Impact

– Assume 16x16 tiles and thread blocks
– Each thread block has 8 warps (256/32)
– Assume square matrices of 100x100
– Each thread will go through 7 phases (ceiling of 100/16)

– There are 49 thread blocks (7 in each dimension)

Control Divergence in Loading M Tiles

– Assume 16x16 tiles and thread blocks
– Each thread block has 8 warps (256/32)
– Assume square matrices of 100x100
– Each warp will go through 7 phases (ceiling of 100/16)

– There are 42 (6*7) Type 1 blocks, with a total of 336 (8*42) warps
– They all have 7 phases, so there are 2,352 (336*7) warp-phases
– The warps have control divergence only in their last phase
– 336 warp-phases have control divergence

Type 1

M

TILE_WIDTH

Control Divergence in Loading M Tiles (Type 2)

- Type 2: the 7 block assigned to load the bottom tiles, with a total of 56 (8*7) warps
- They all have 7 phases, so there are 392 (56*7) warp-phases
- The first 2 warps in each Type 2 block will stay within the valid range until the last phase
- The 6 remaining warps stay outside the valid range
- So, only 14 (2*7) warp-phases have control divergence

M

TILE_WIDTH

Type 2

Overall Impact of Control Divergence

– Type 1 Blocks: 336 out of 2,352 warp-phases have control divergence

– Type 2 Blocks: 14 out of 392 warp-phases have control divergence

– The performance impact is expected to be less than 12% (350/2,944 or (336+14)/(2352+14))

Additional Comments

– The calculation of impact of control divergence in loading N tiles is somewhat different and is left as an exercise

– The estimated performance impact is data dependent.
  – For larger matrices, the impact will be significantly smaller

– In general, the impact of control divergence for boundary condition checking for large input data sets should be insignificant
  – One should not hesitate to use boundary checks to ensure full functionality

– The fact that a kernel is full of control flow constructs does not mean that there will be heavy occurrence of control divergence

– We will cover some algorithm patterns that naturally incur control divergence (such as parallel reduction) in the Parallel Algorithm Patterns modules

NVIDIA   ILLINOIS

GPU Teaching Kit

Accelerated Computing

ILLINOIS
UNIVERSITY OF ILLINOIS AT URBANA-CHAMPAIGN